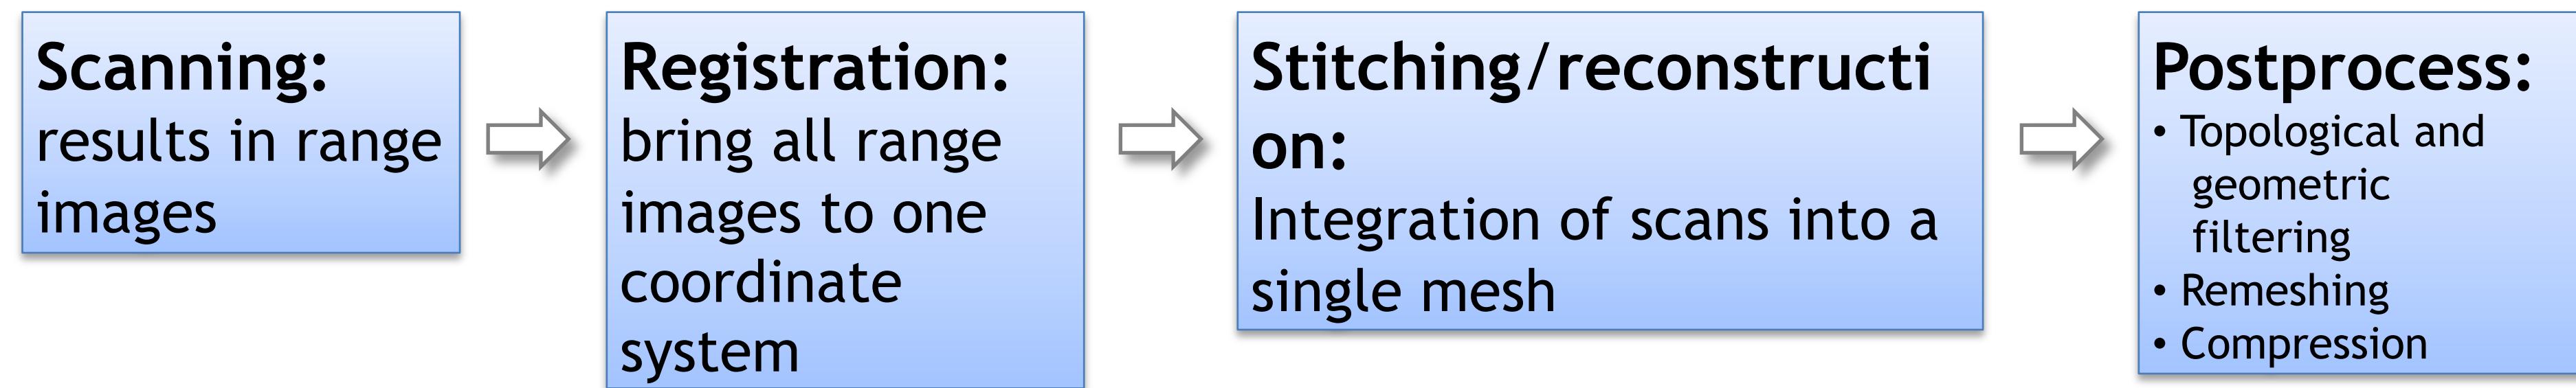
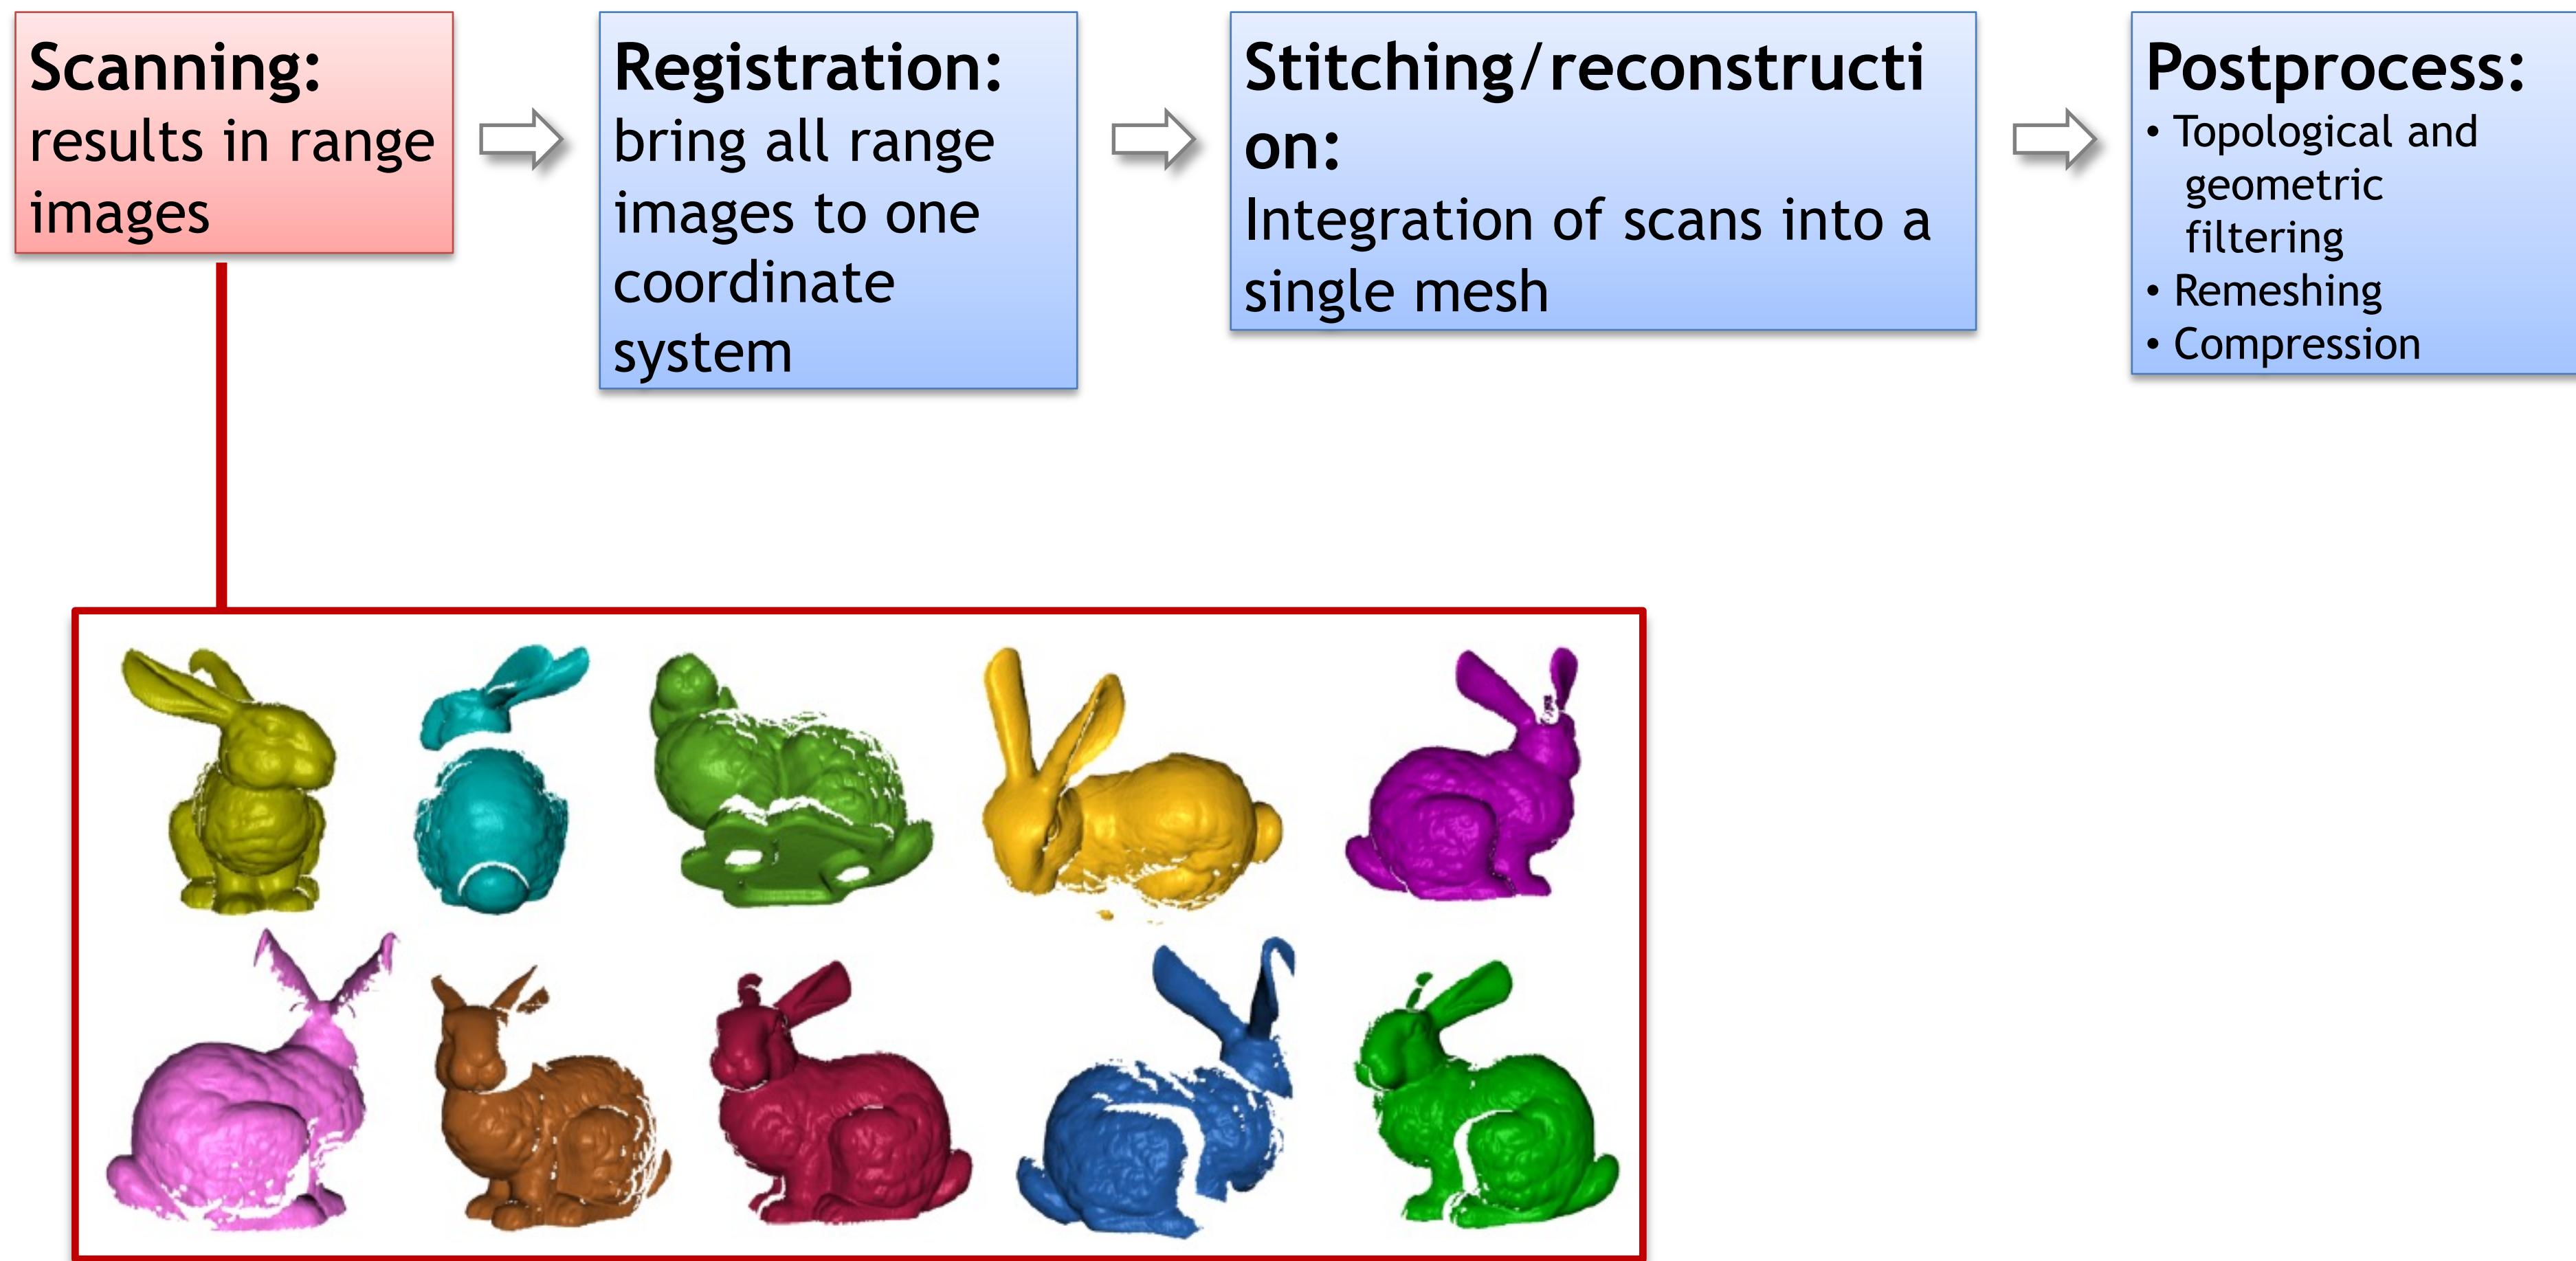


03 – Acquisition and Mesh

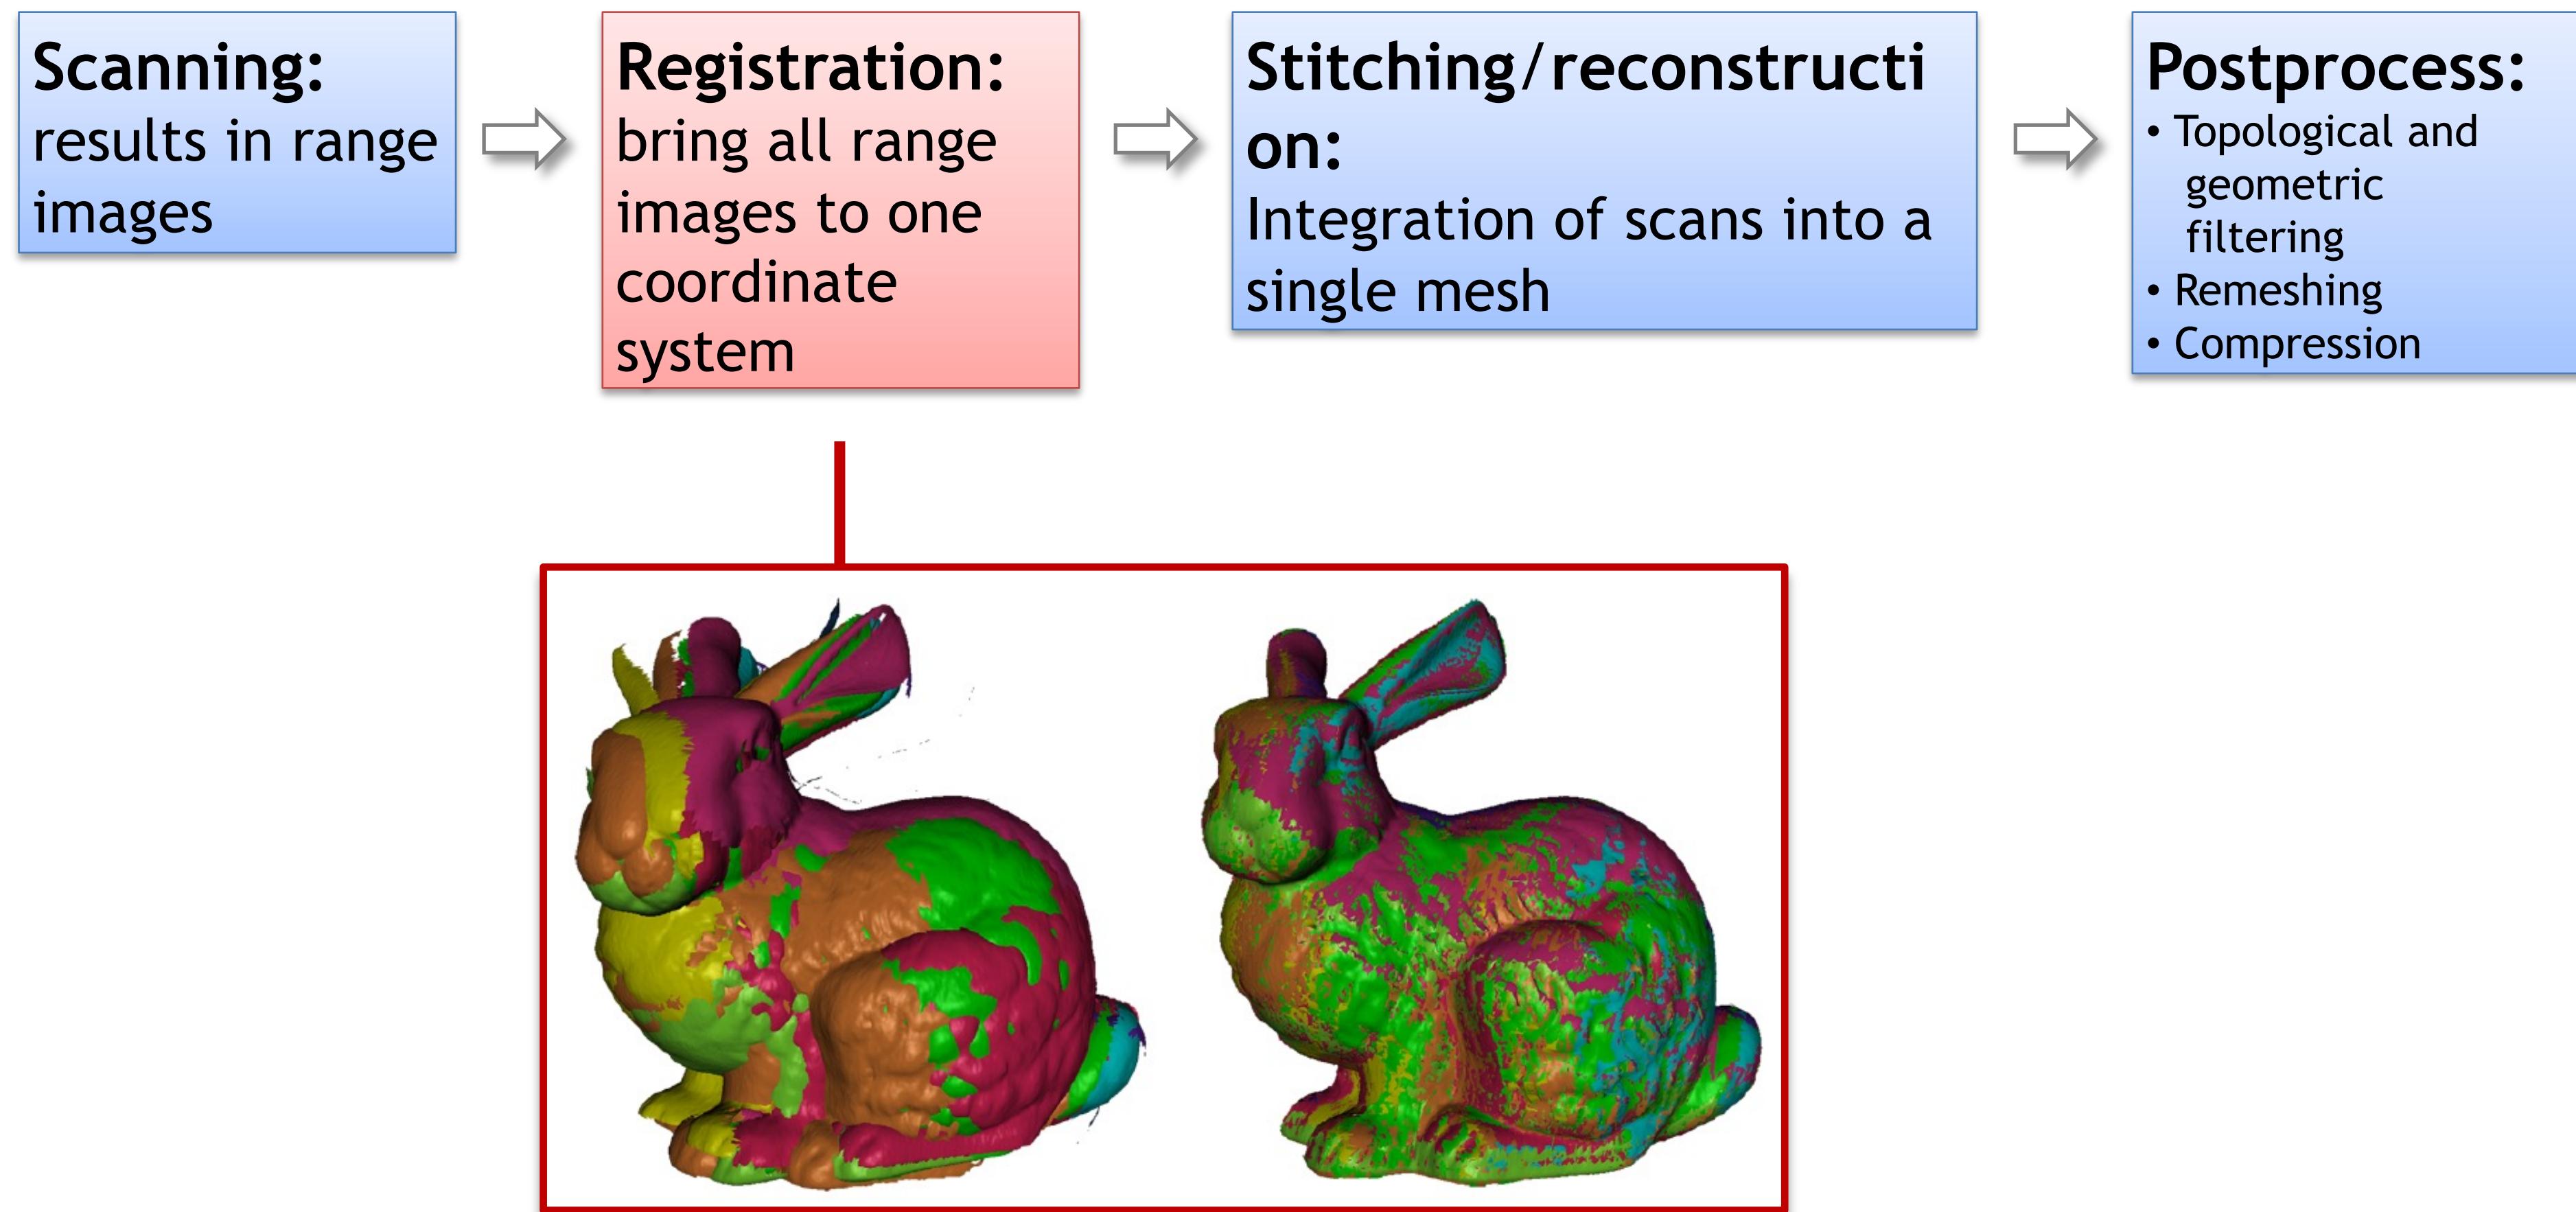
Geometry Acquisition Pipeline



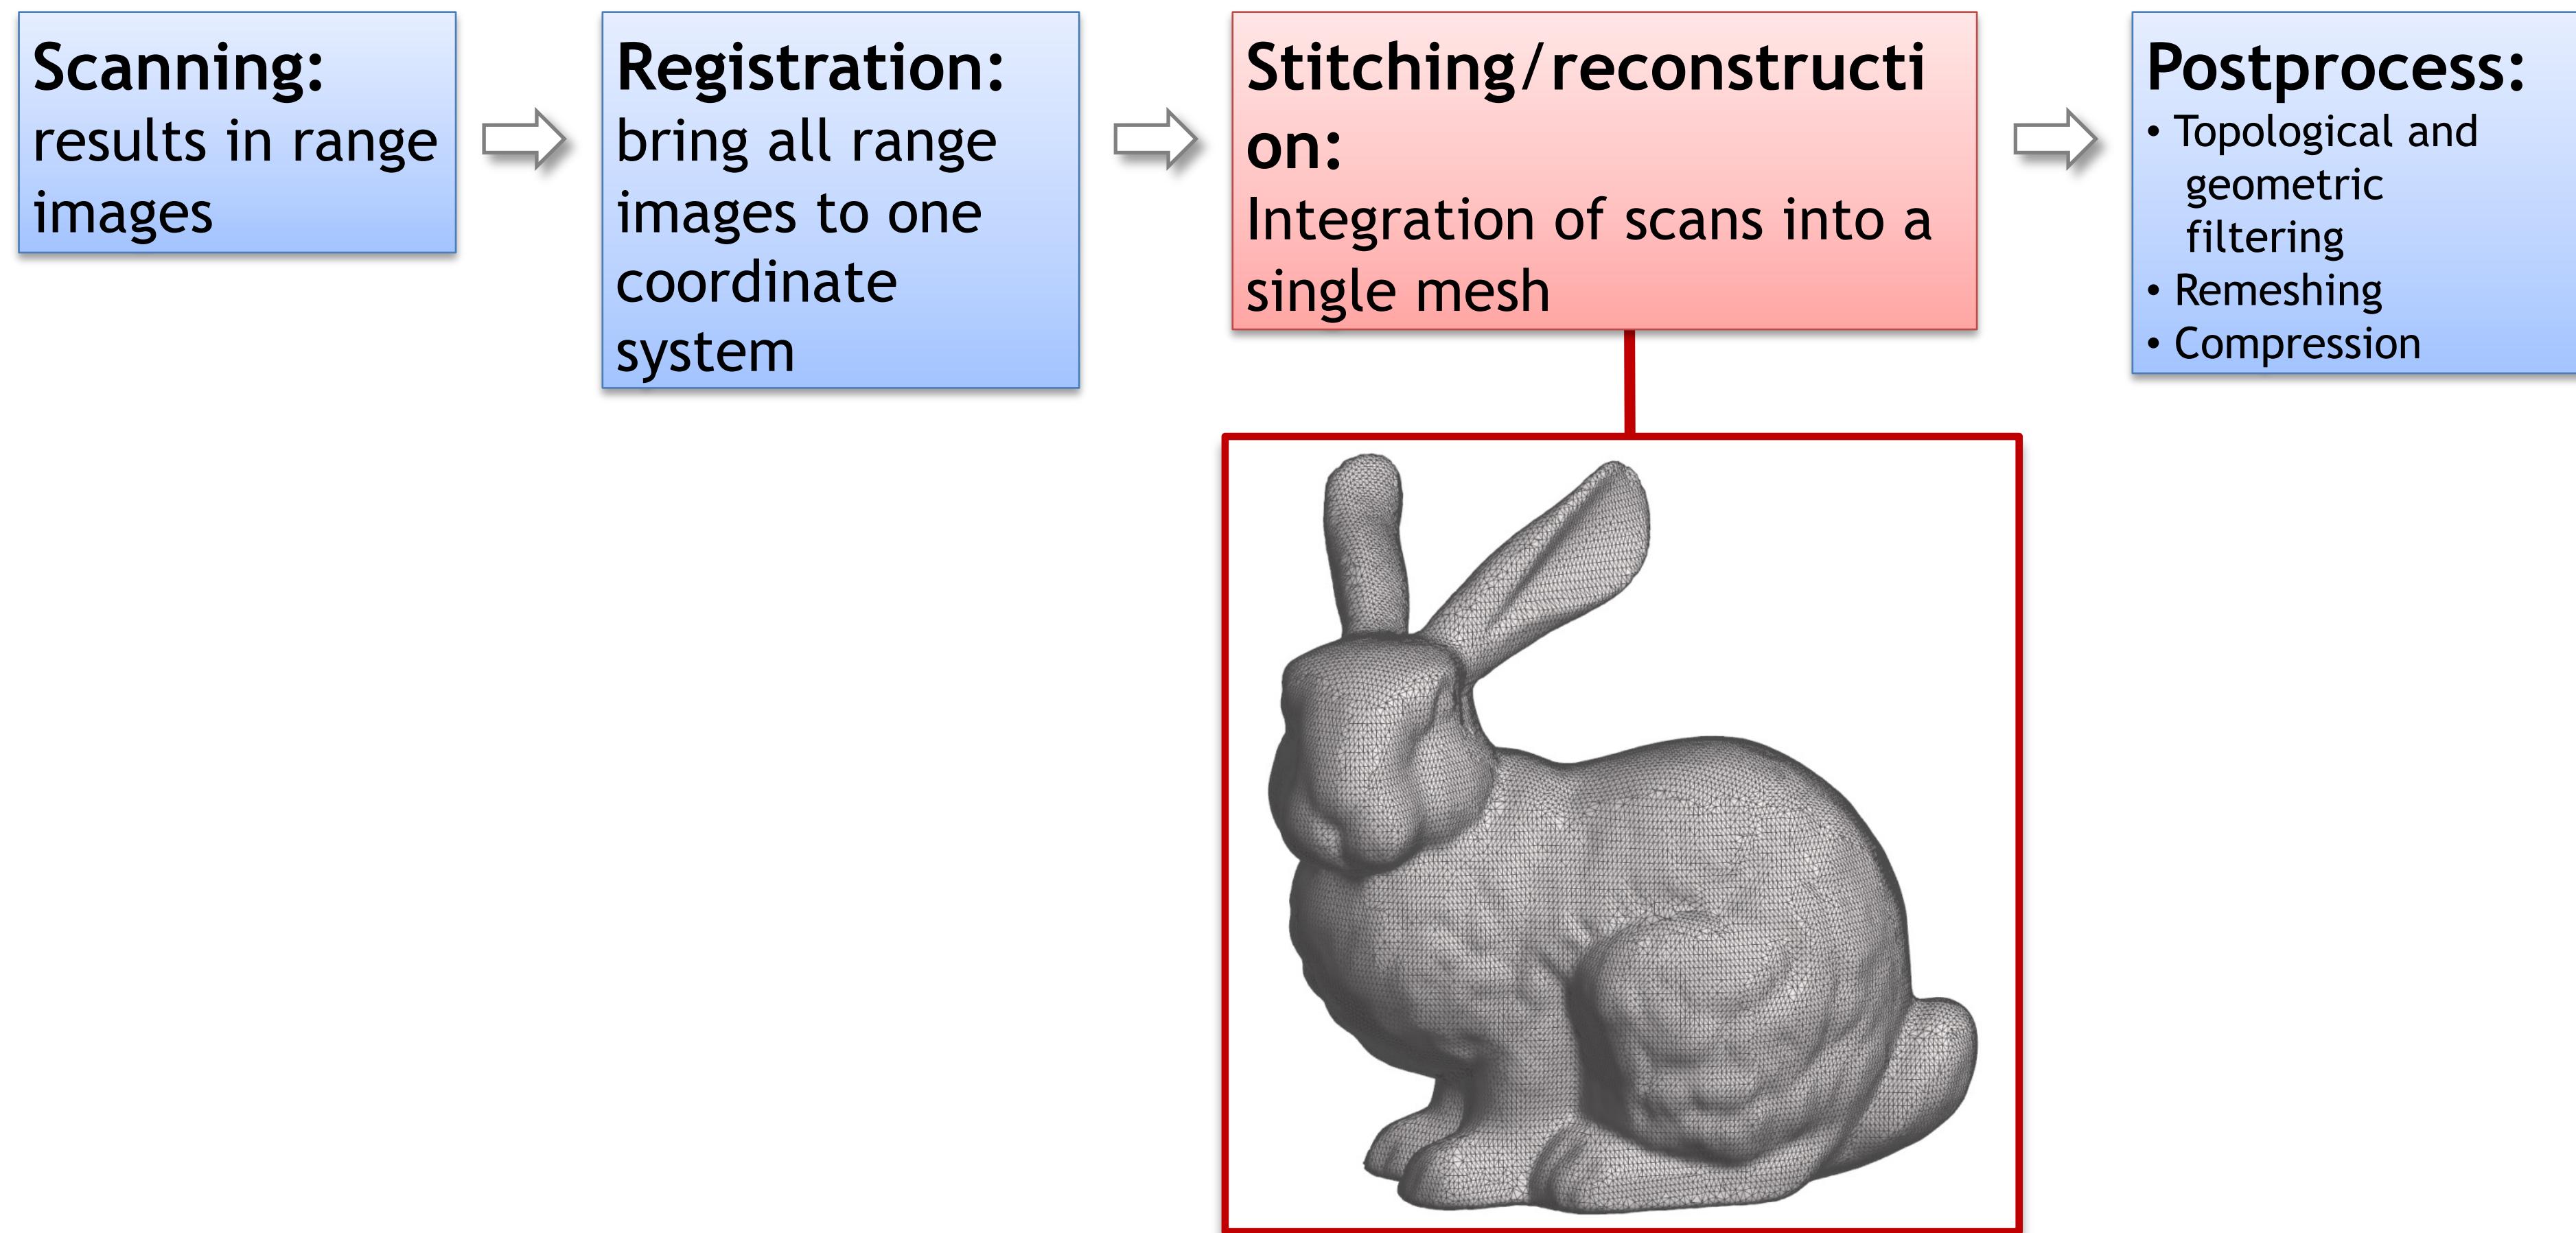
Geometry Acquisition Pipeline



Geometry Acquisition Pipeline



Geometry Acquisition Pipeline

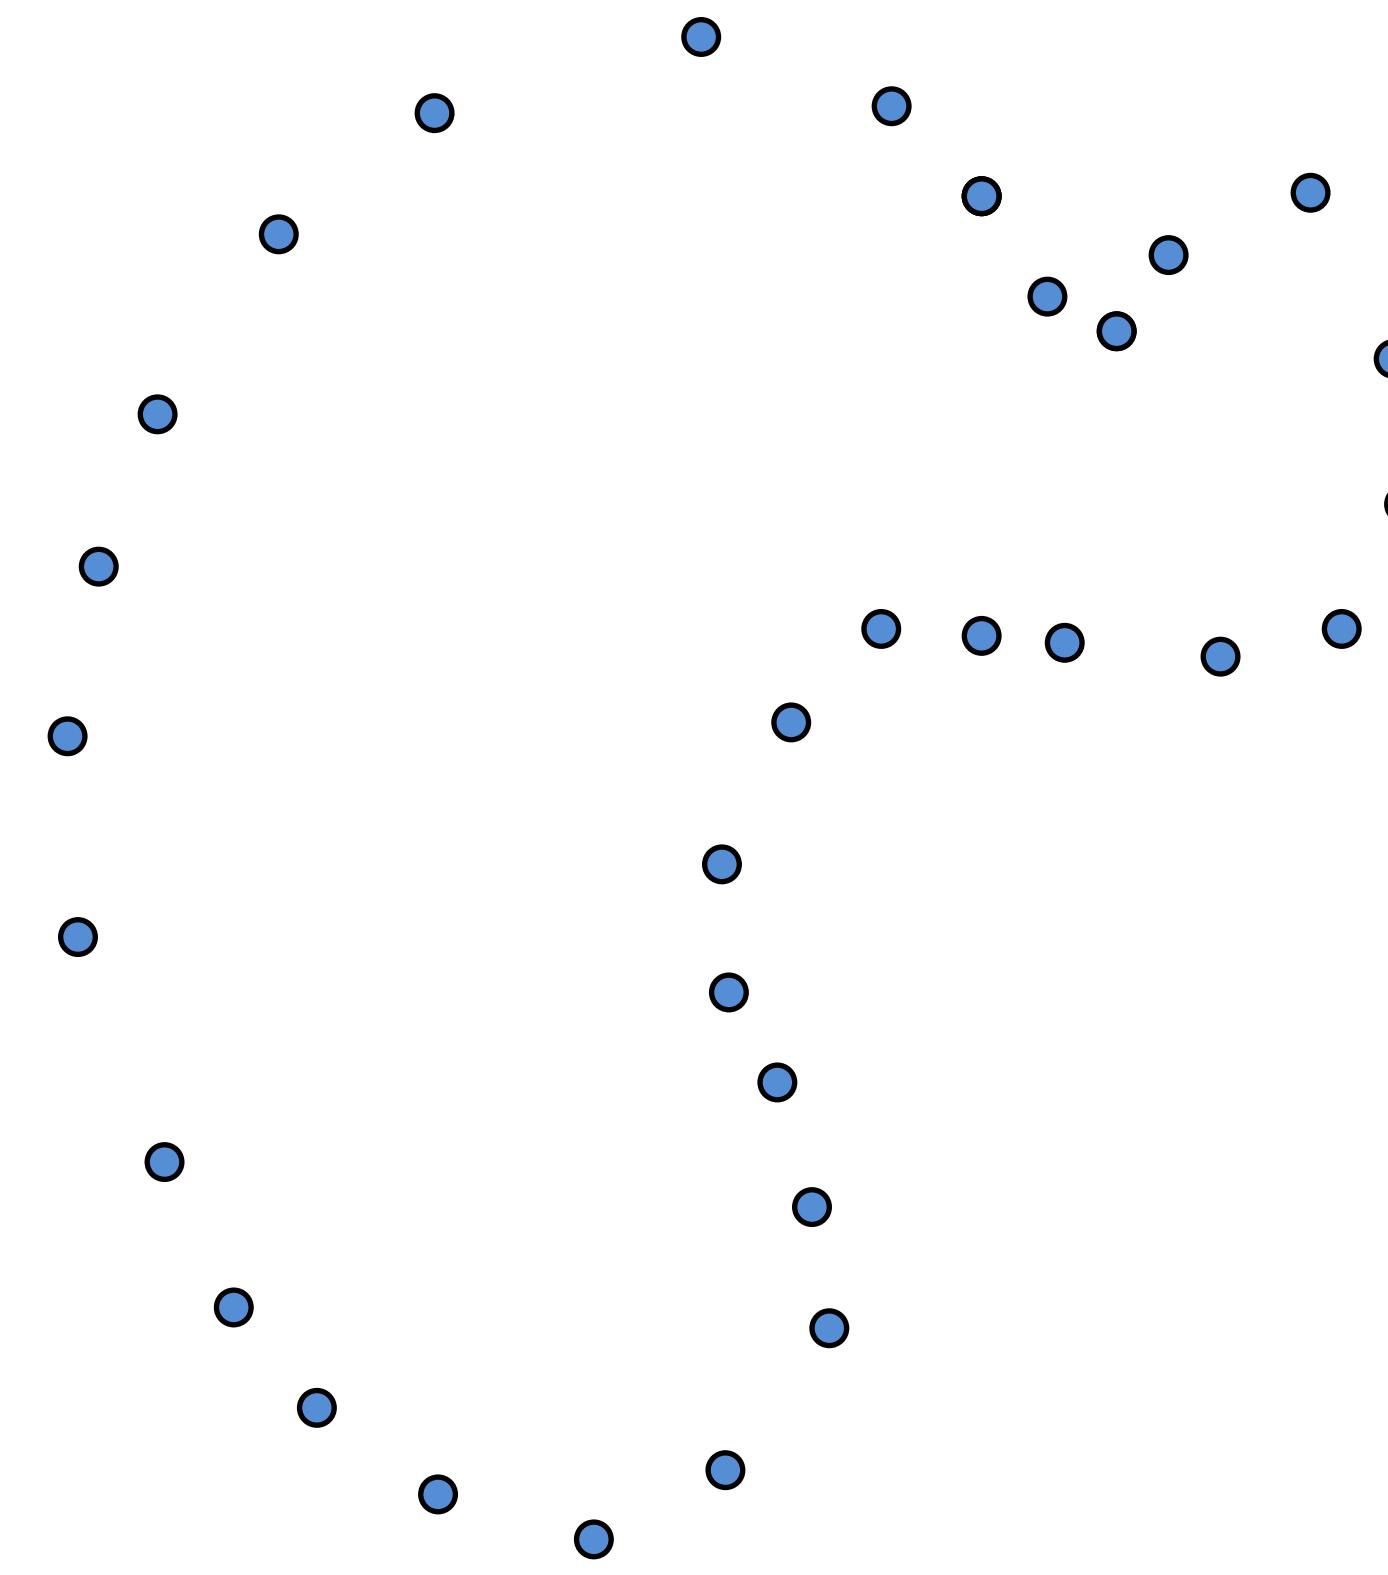


Surface Reconstruction

- Generate a mesh from a set of surface samples



Implicit Function Approach

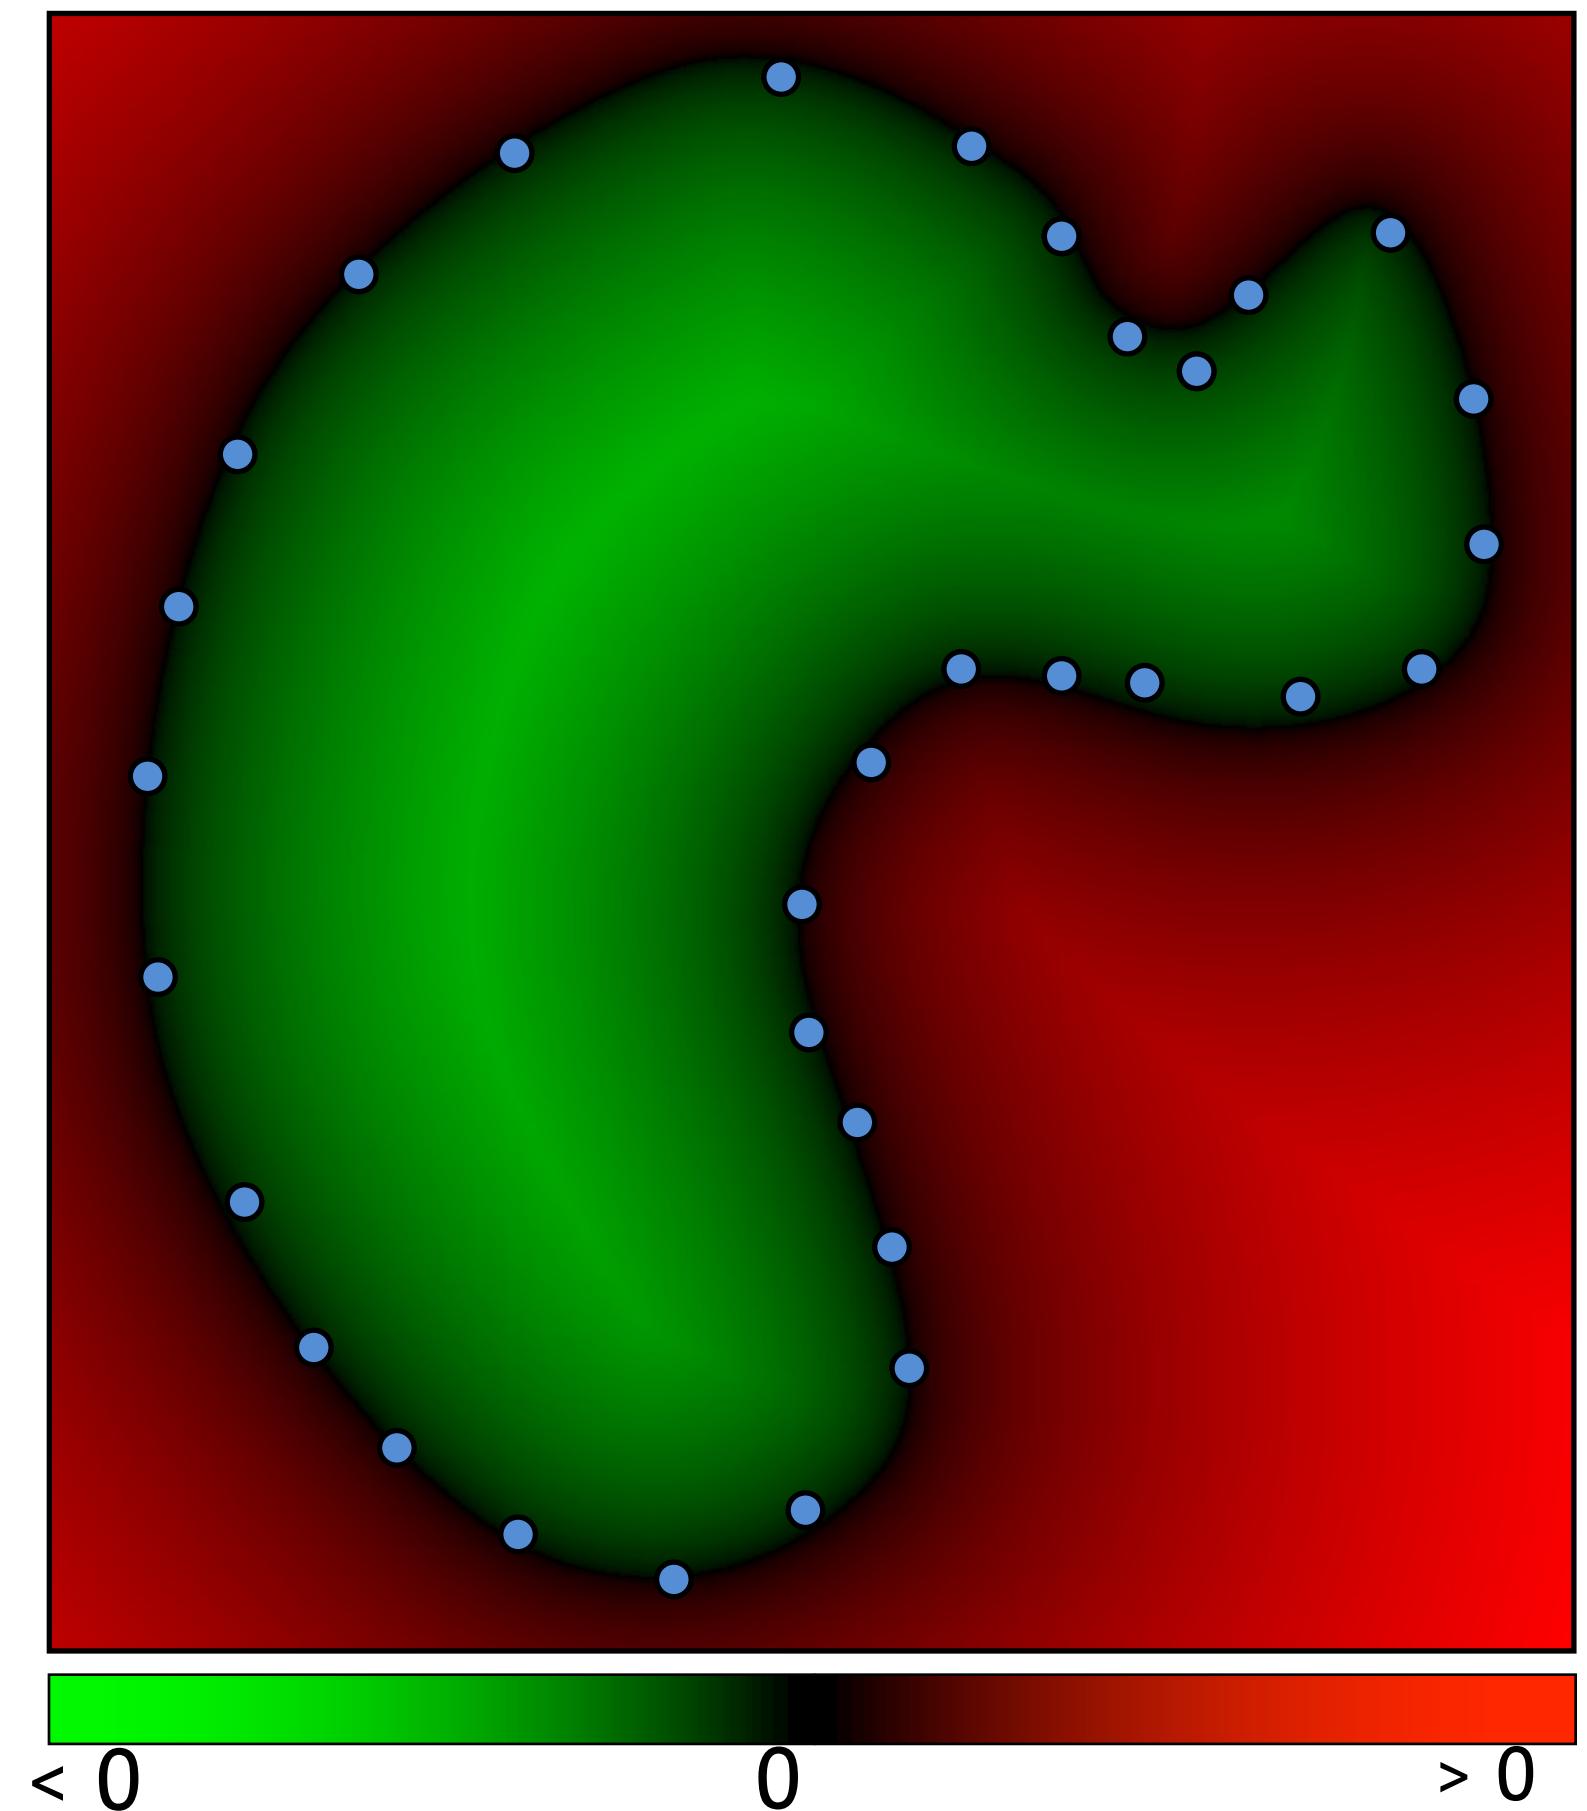


Implicit Function Approach

- Define a function

$$f : \mathbb{R}^3 \rightarrow \mathbb{R}$$

with value > 0 outside the shape and < 0 inside



Implicit Function Approach

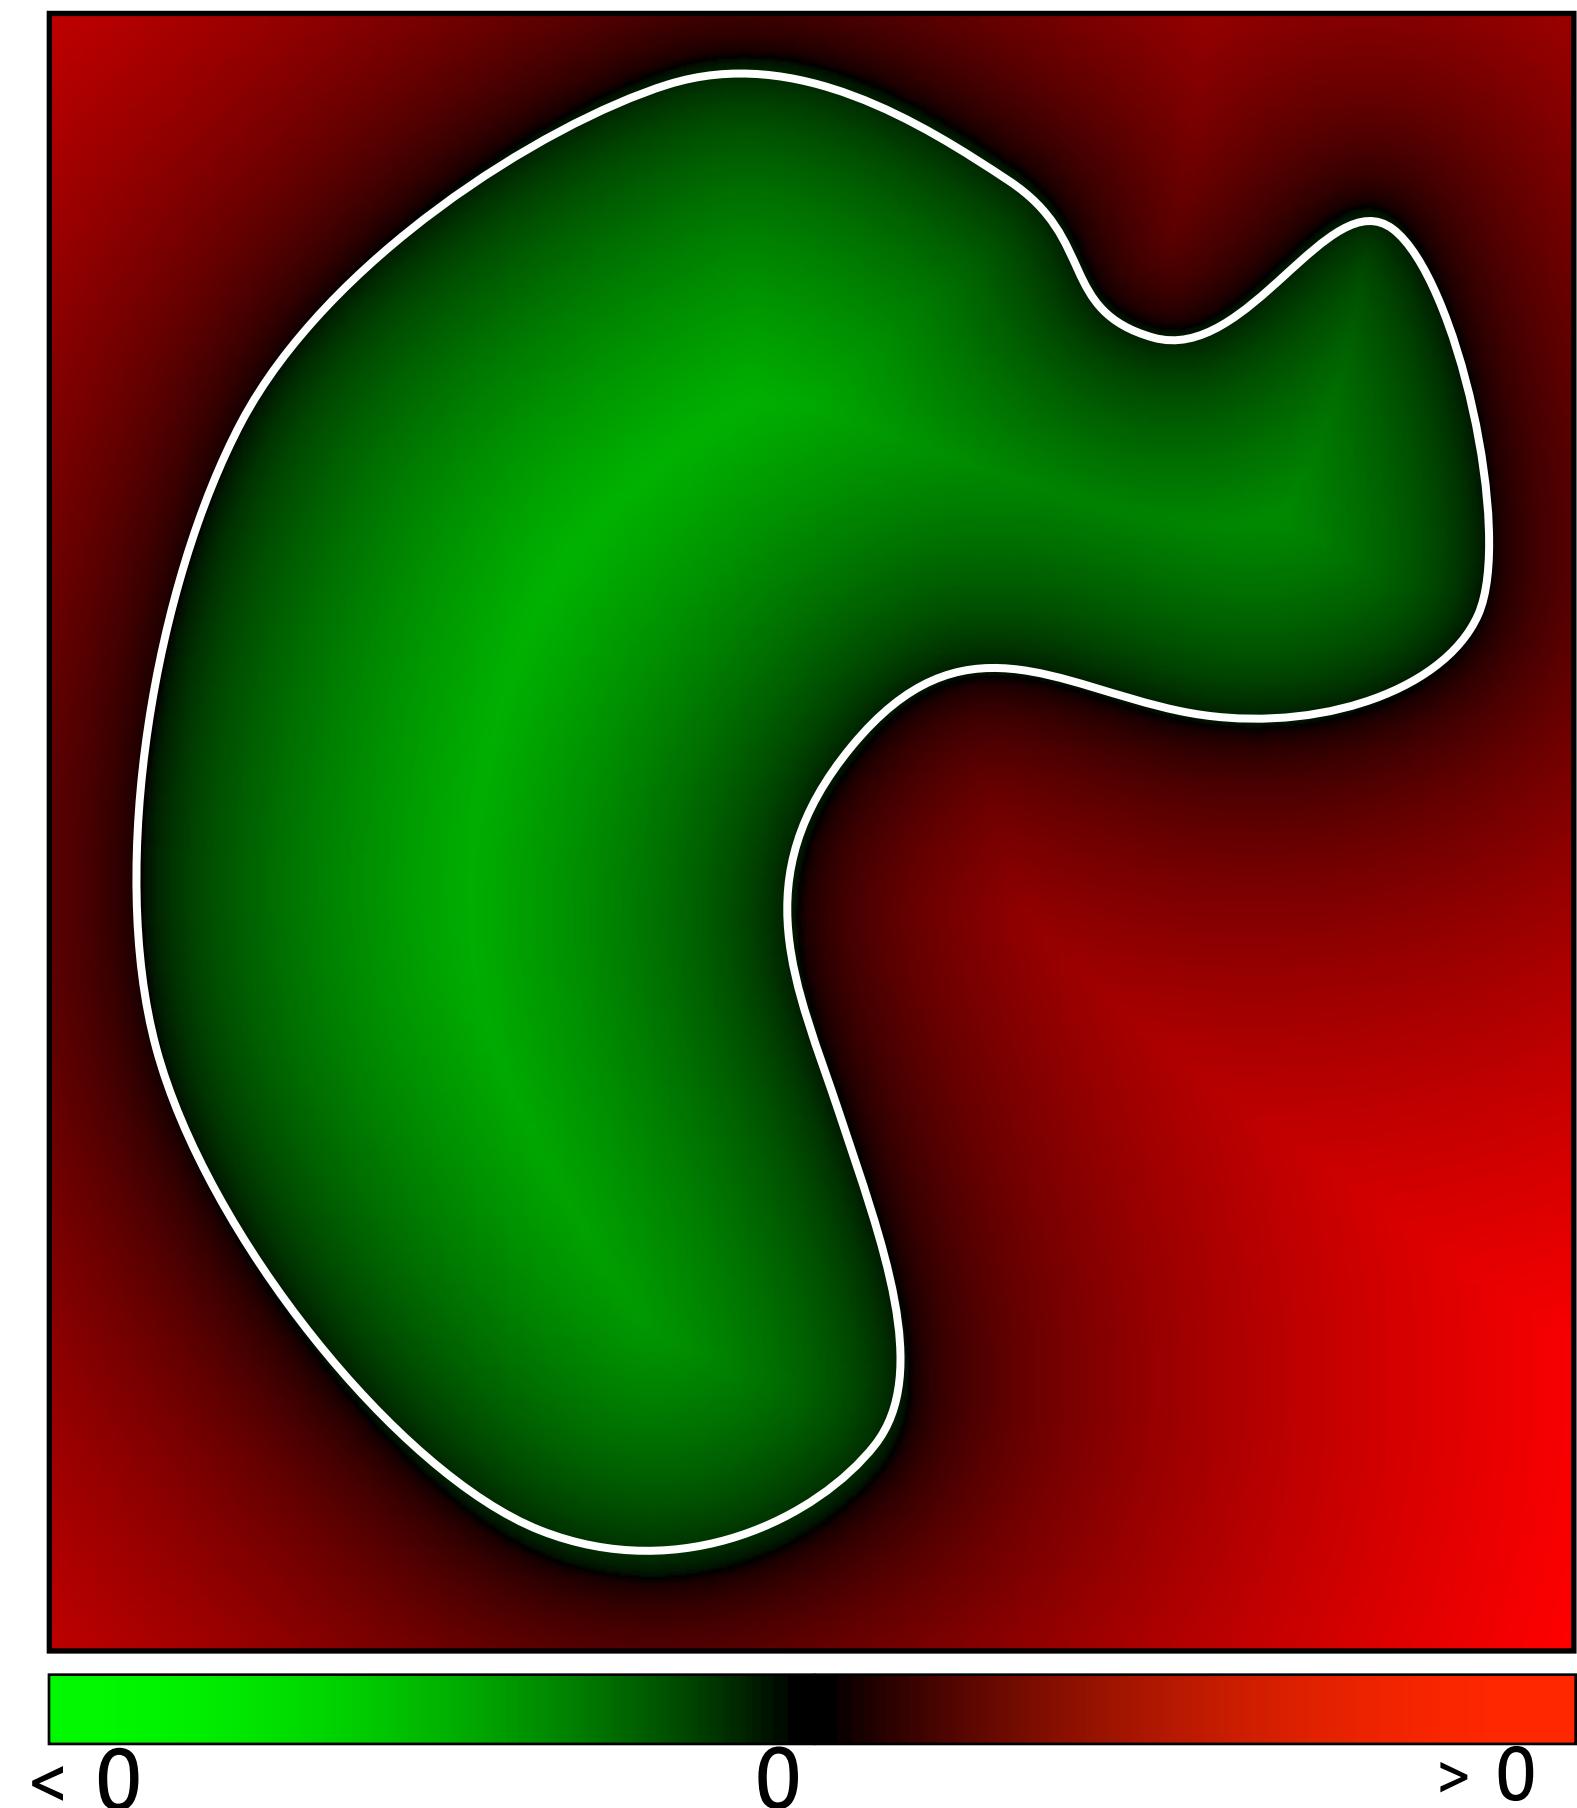
- Define a function

$$f : \mathbb{R}^3 \rightarrow \mathbb{R}$$

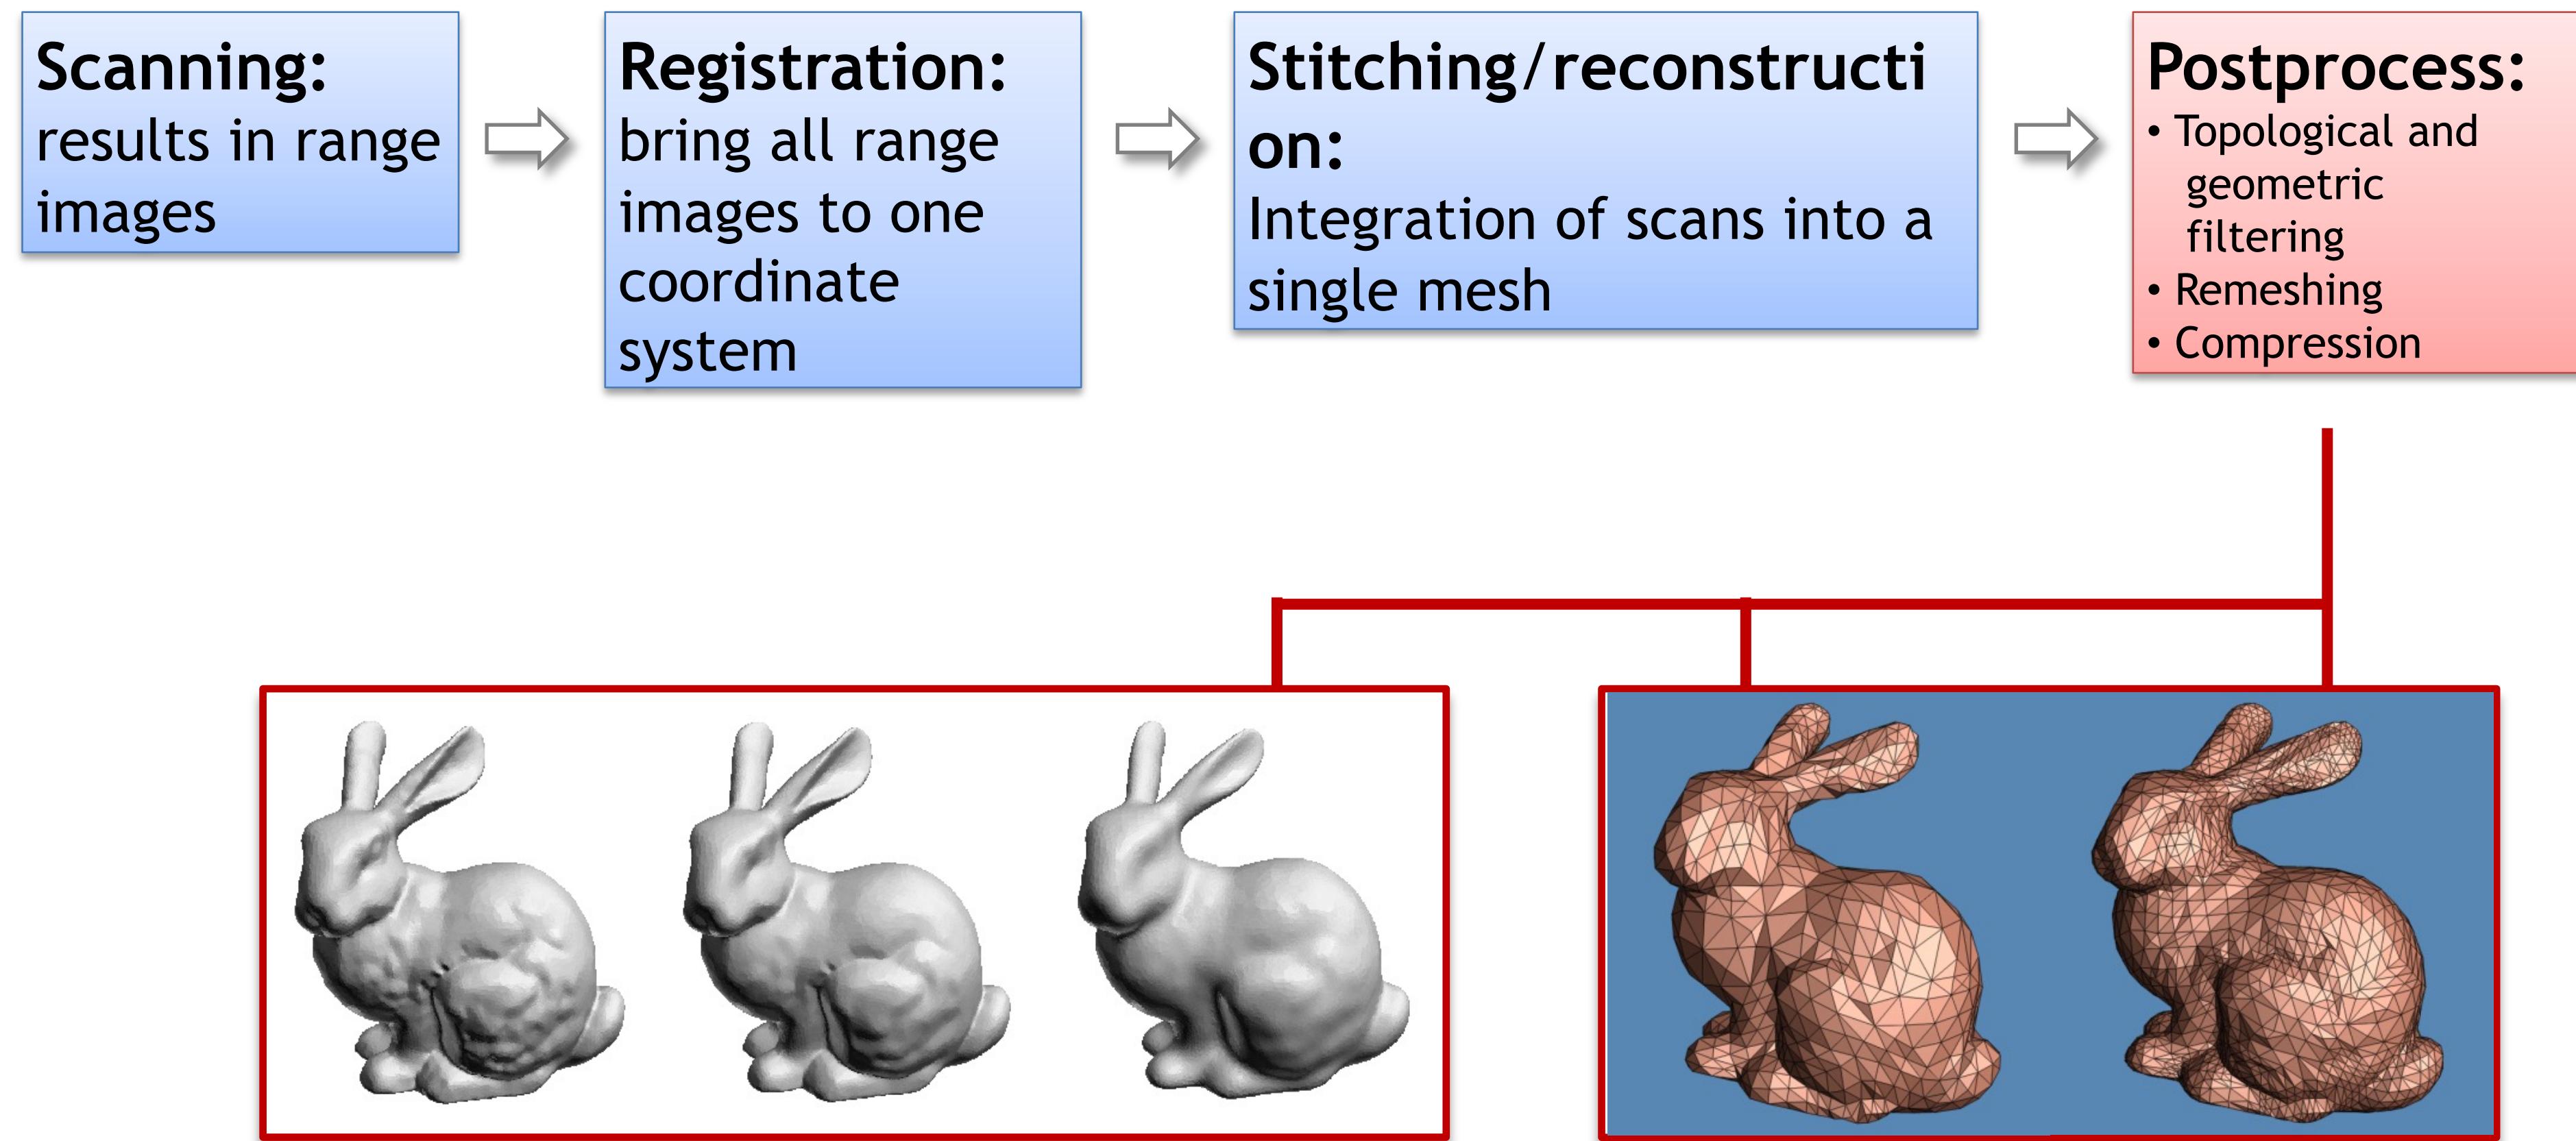
with value > 0 outside the shape and < 0 inside

- Extract the zero-set

$$\{\mathbf{x} : f(\mathbf{x}) = 0\}$$



Geometry Acquisition Pipeline



Scanning

Touch Probes



Touch Probes

- Physical contact with the object
- Manual or computer-guided
- Advantages:
 - Can be very precise
 - Can scan **any** solid surface
- Disadvantages:
 - Slow, small scale
 - Can't use on fragile objects



Optical Scanning

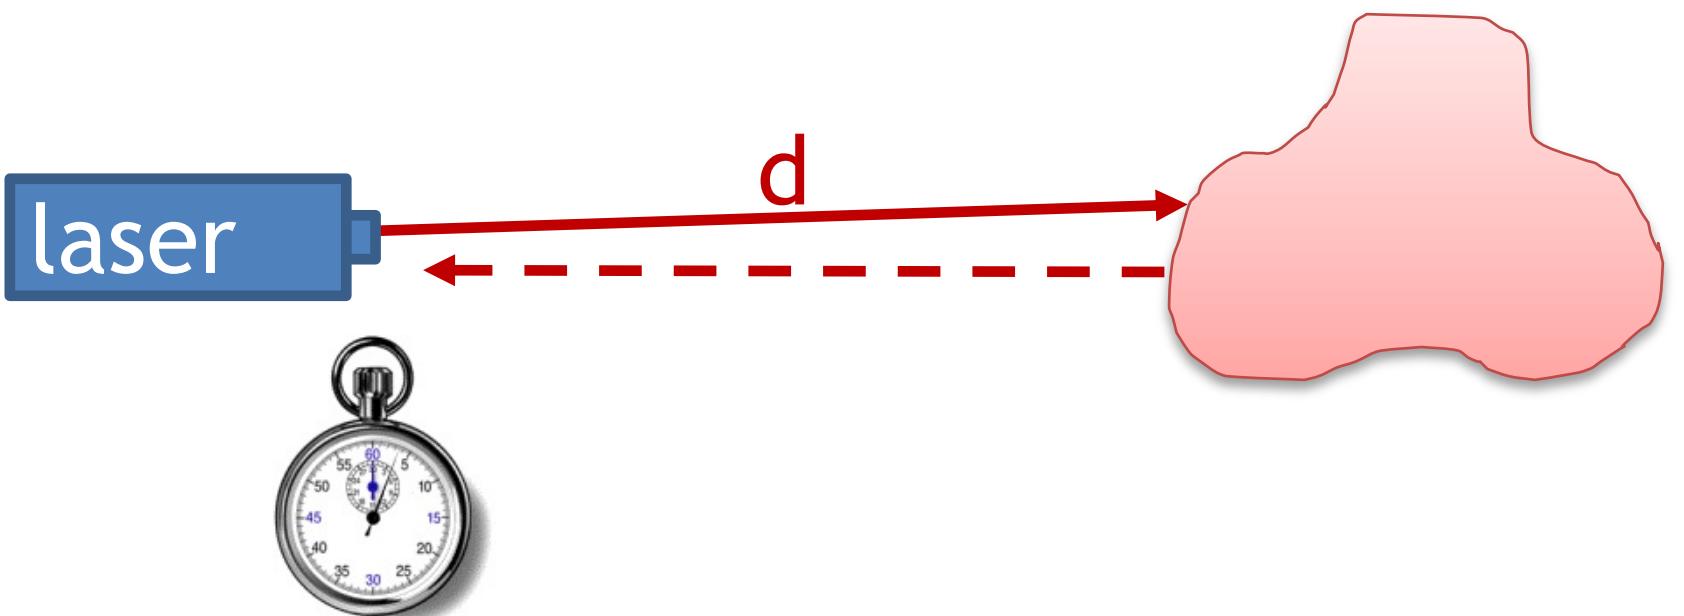
- Infer the geometry from light reflectance
- Advantages:
 - Less invasive than touch
 - Fast, large scale possible
- Disadvantages:
 - Difficulty with transparent and shiny objects



Optical scanning – active lighting

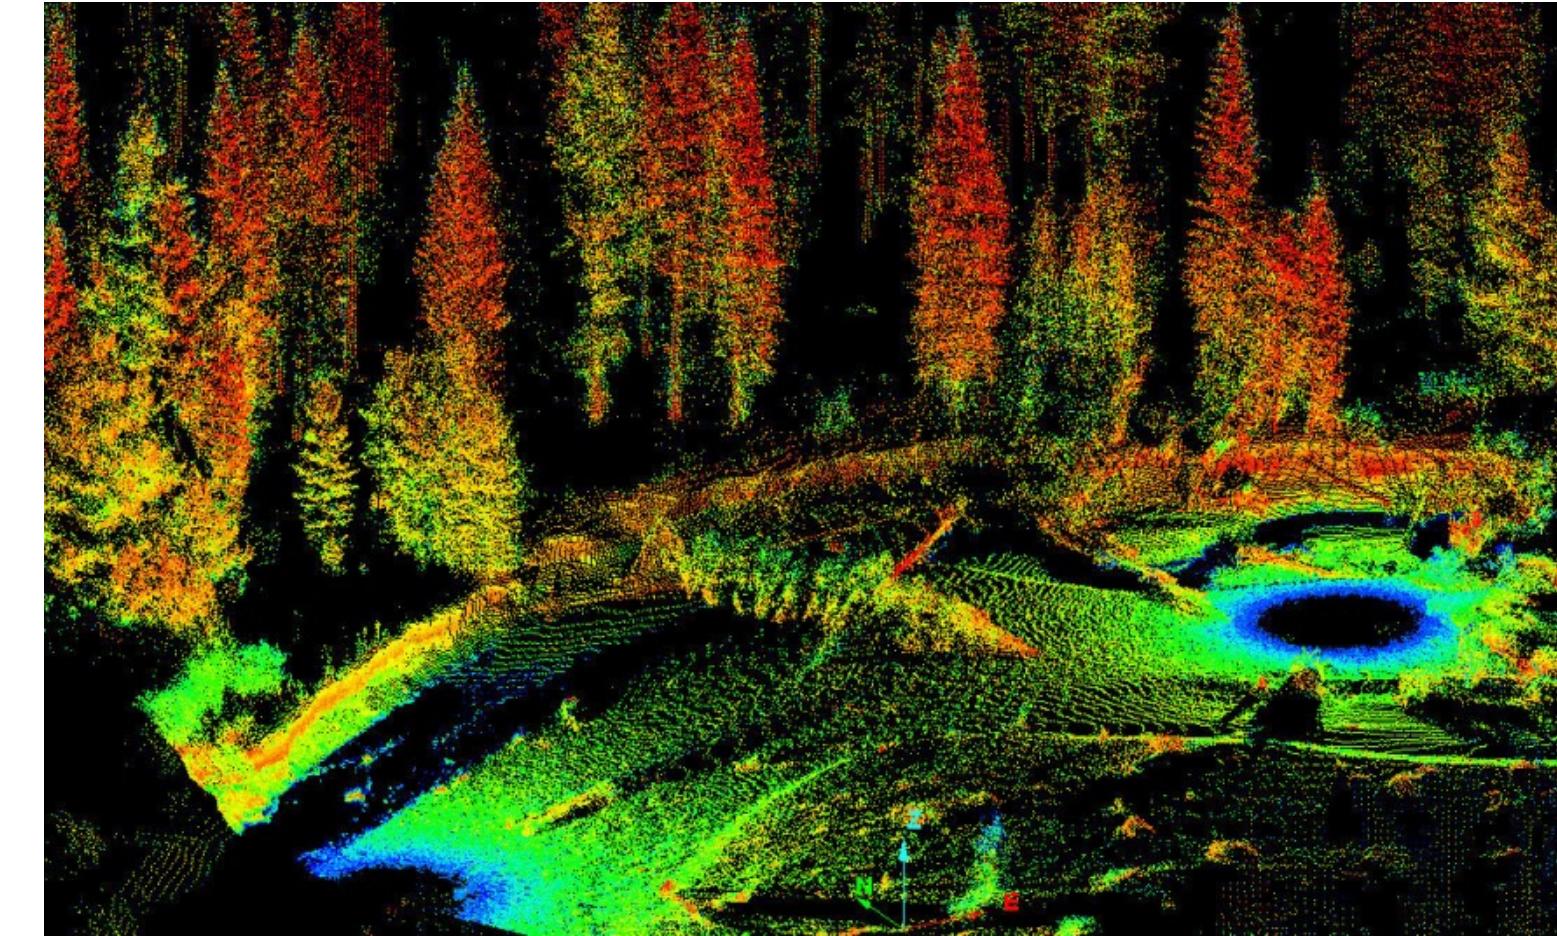
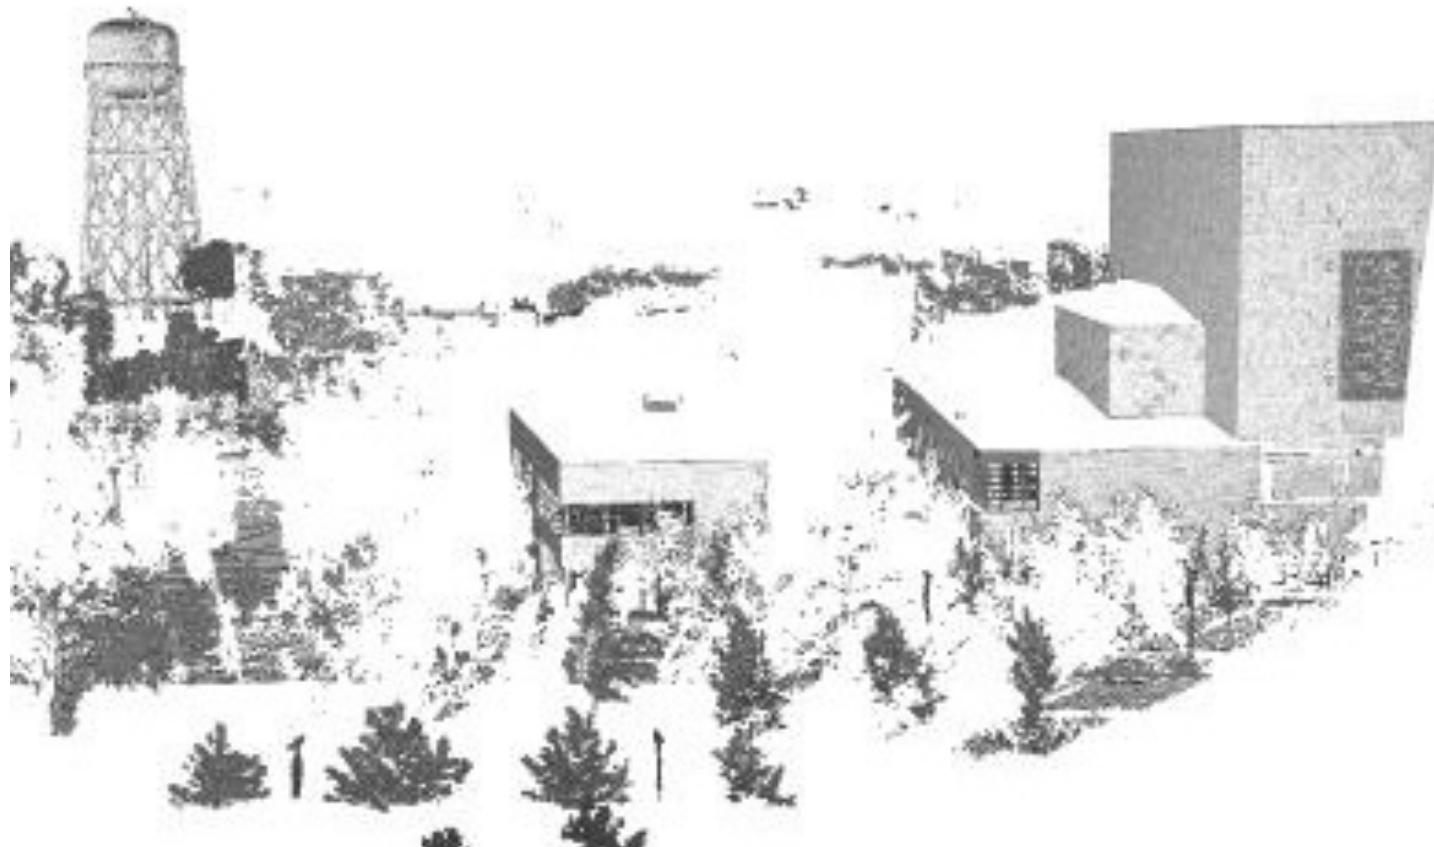
Time of flight laser

- A type of laser rangefinder (LIDAR)
- Measures the time it takes the laser beam to hit the object and come back



Optical scanning – active lighting

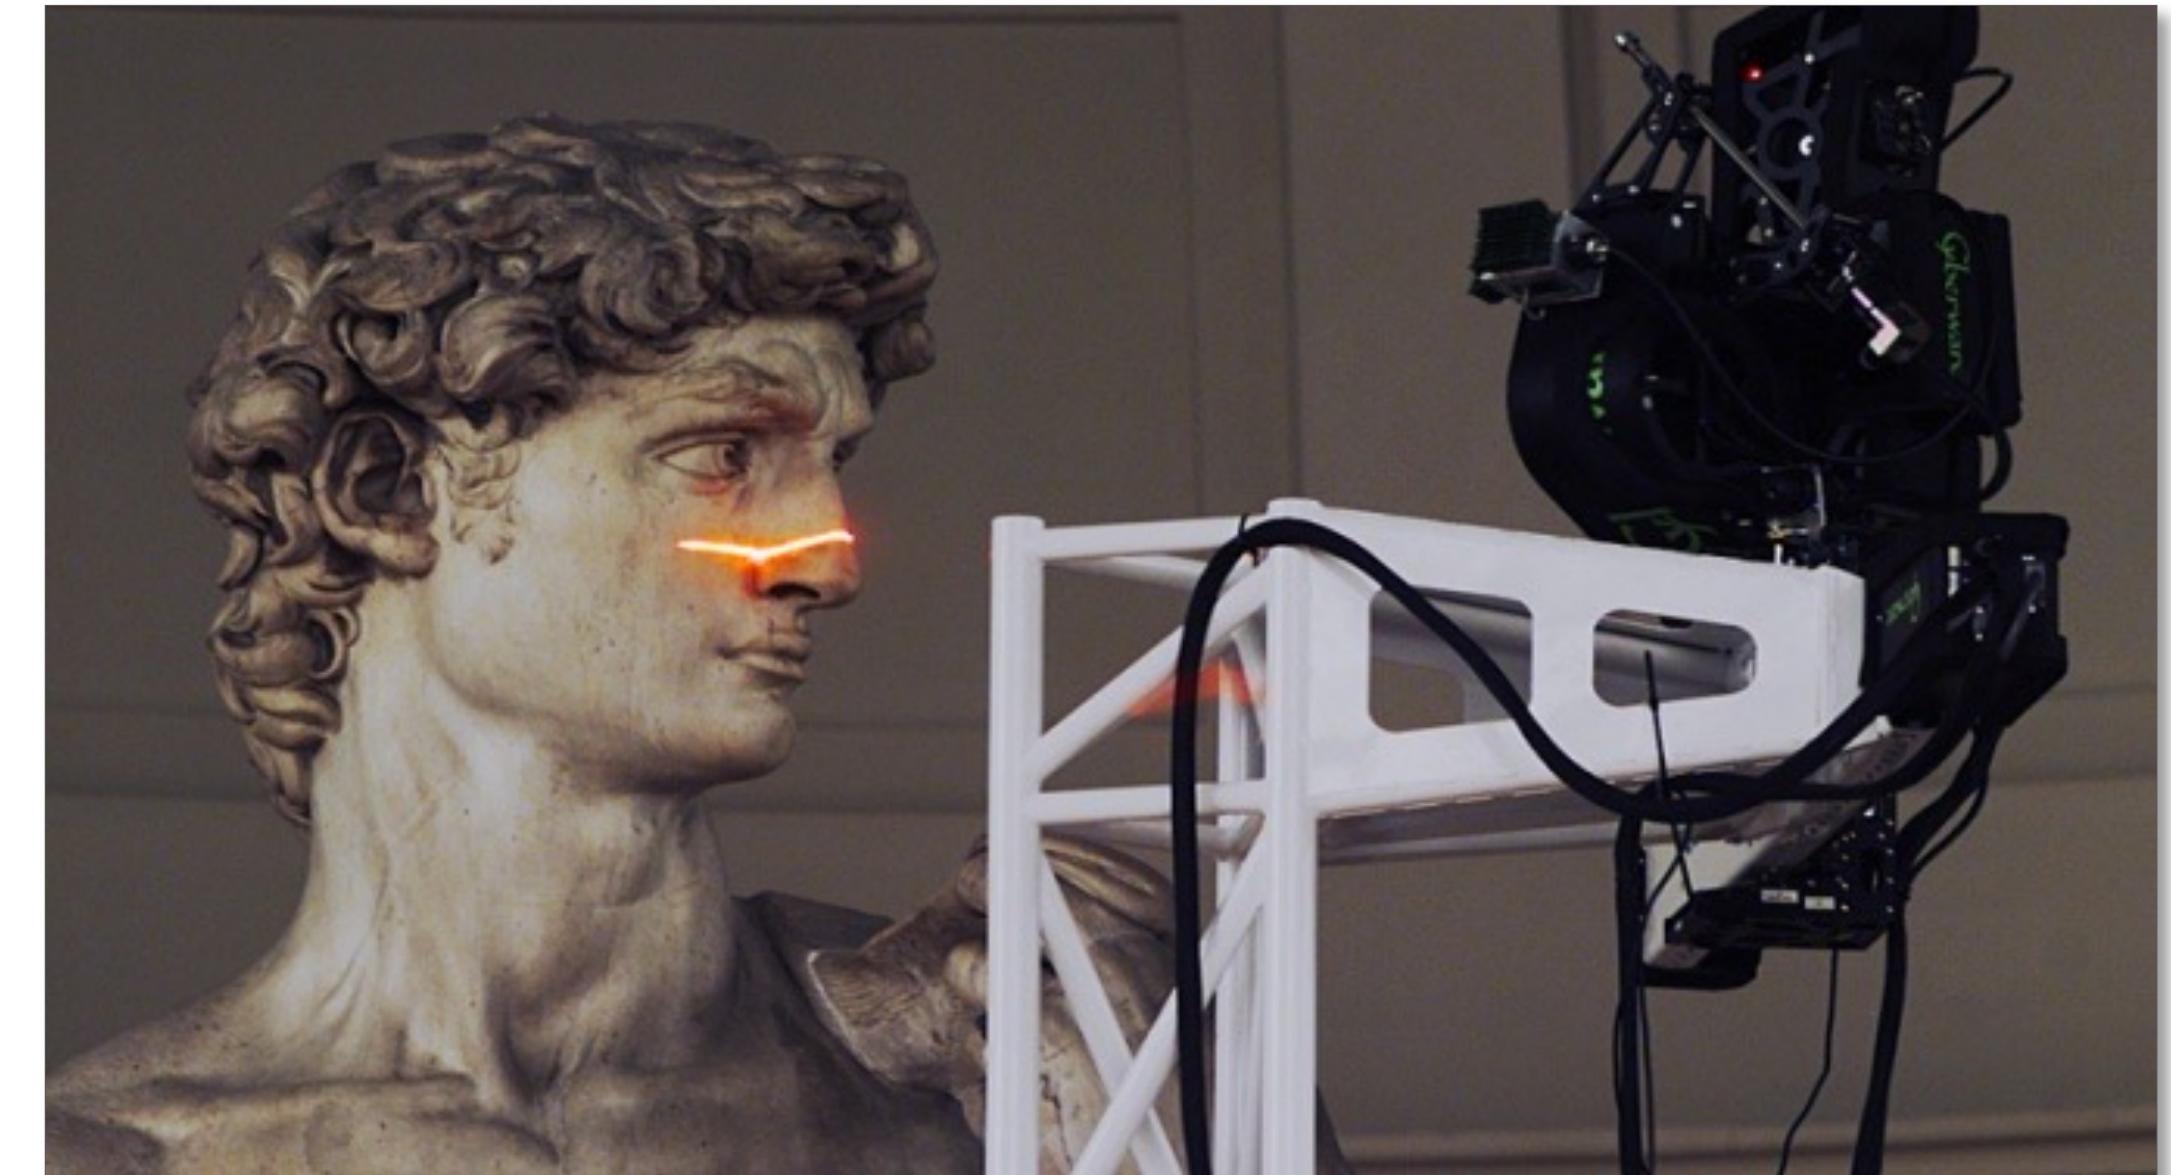
- Accommodates large range – up to several miles (suitable for buildings, rocks)
- Only for static scenes, object motion introduces noise



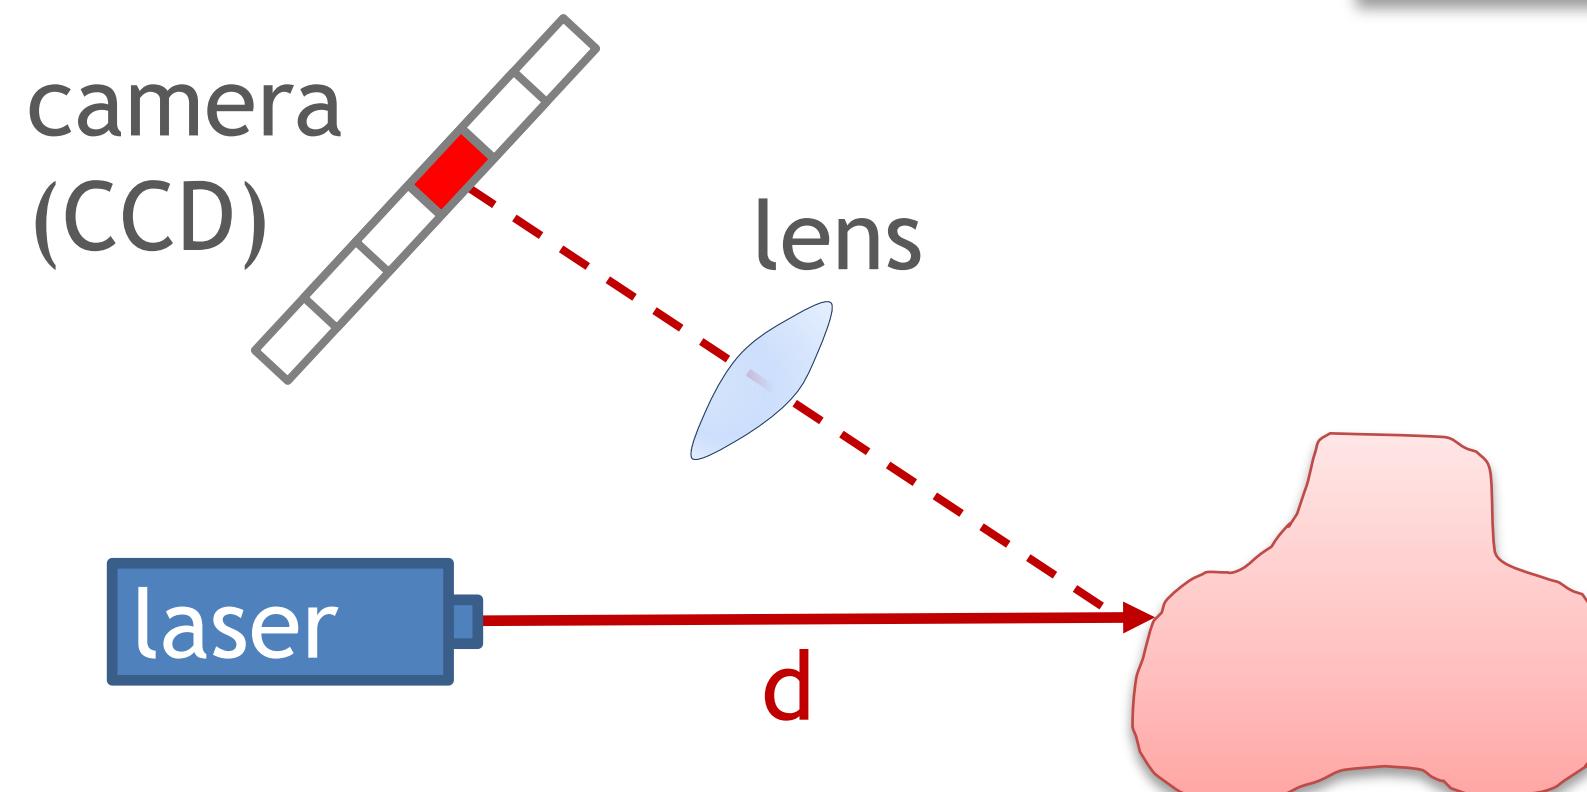
Optical scanning – active lighting

Triangulation laser

- Laser beam and camera
- Laser dot is photographed
- The location of the dot in the image allows triangulation: we get the distance to the object



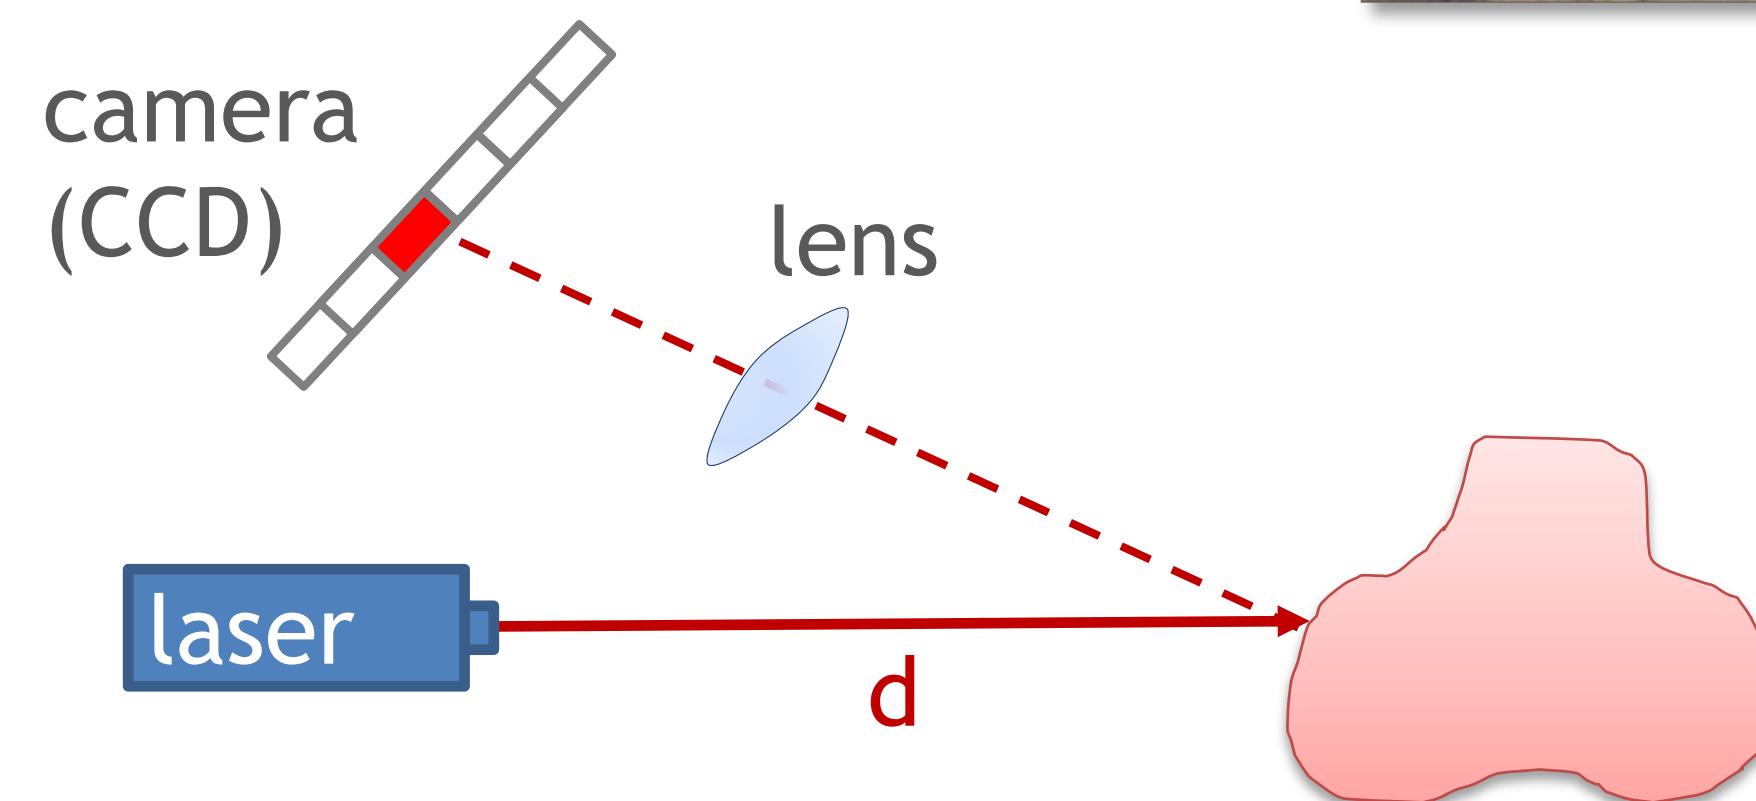
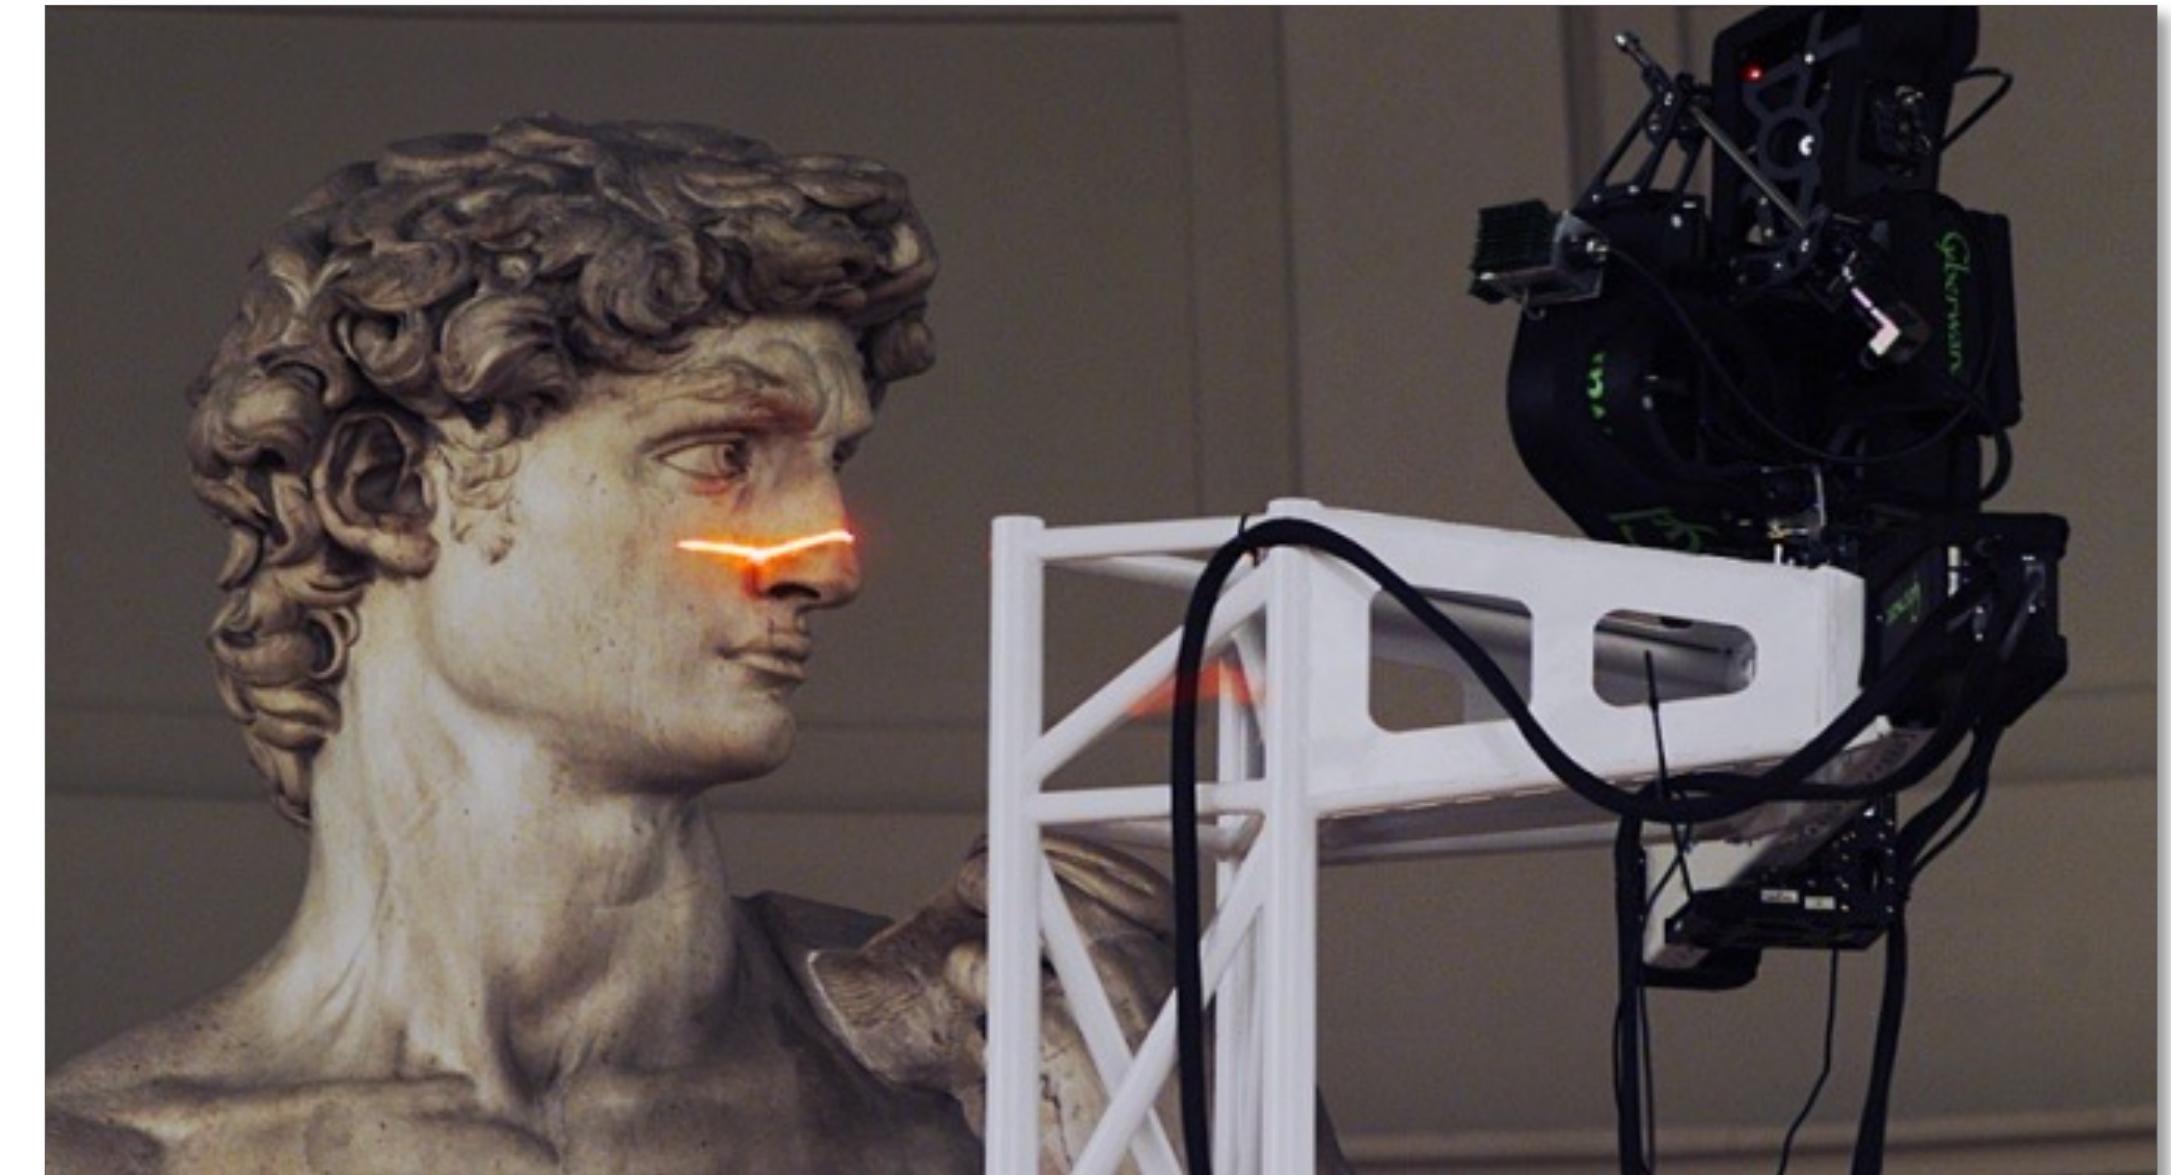
Digital Michelangelo Project



Optical scanning – active lighting

Triangulation laser

- Laser beam and camera
- Laser dot is photographed
- The location of the dot in the image allows triangulation: we get the distance to the object

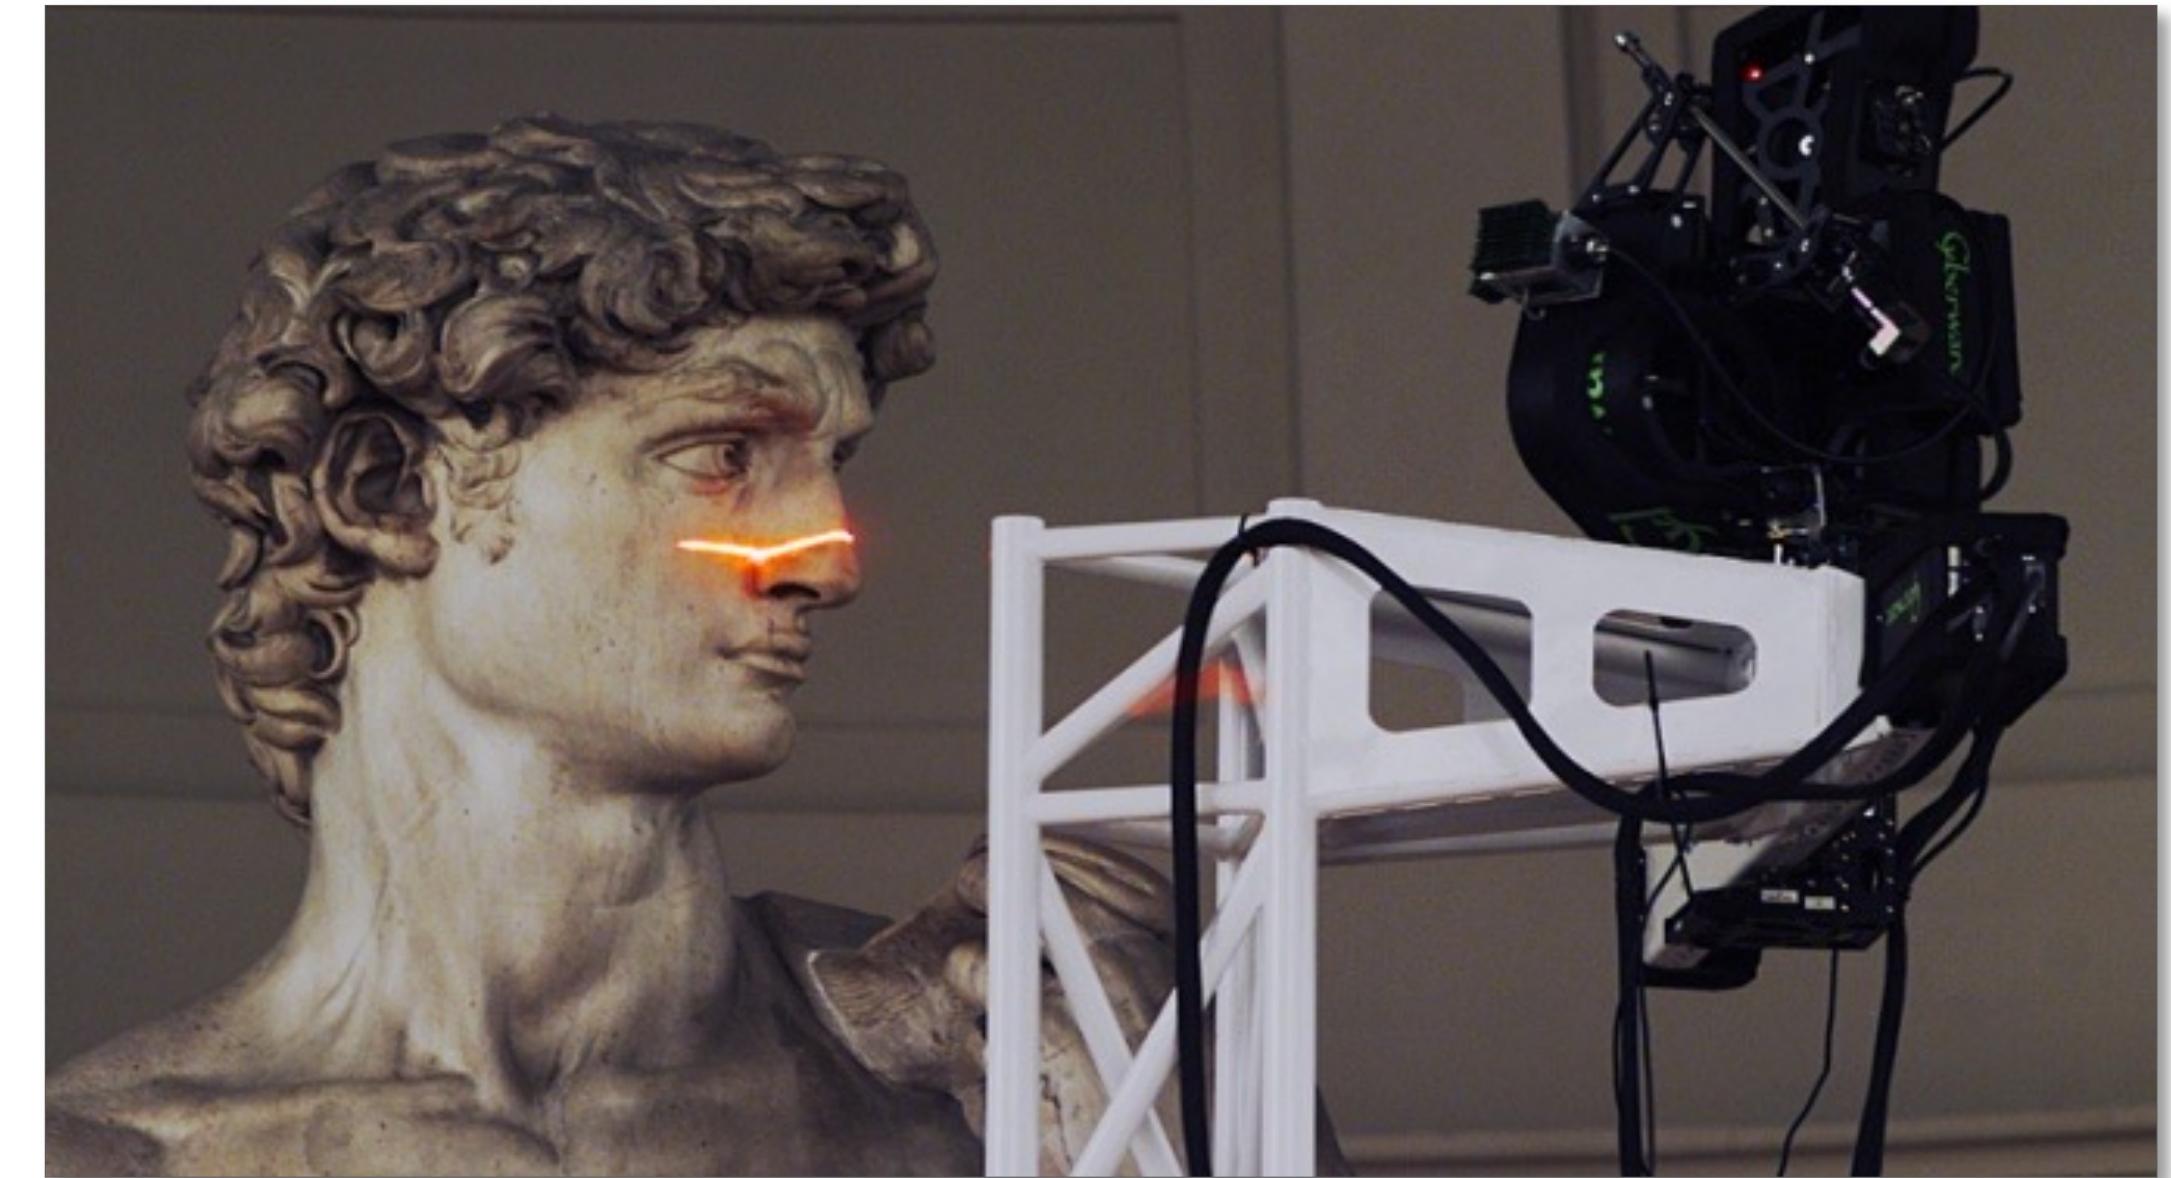


Digital Michelangelo Project

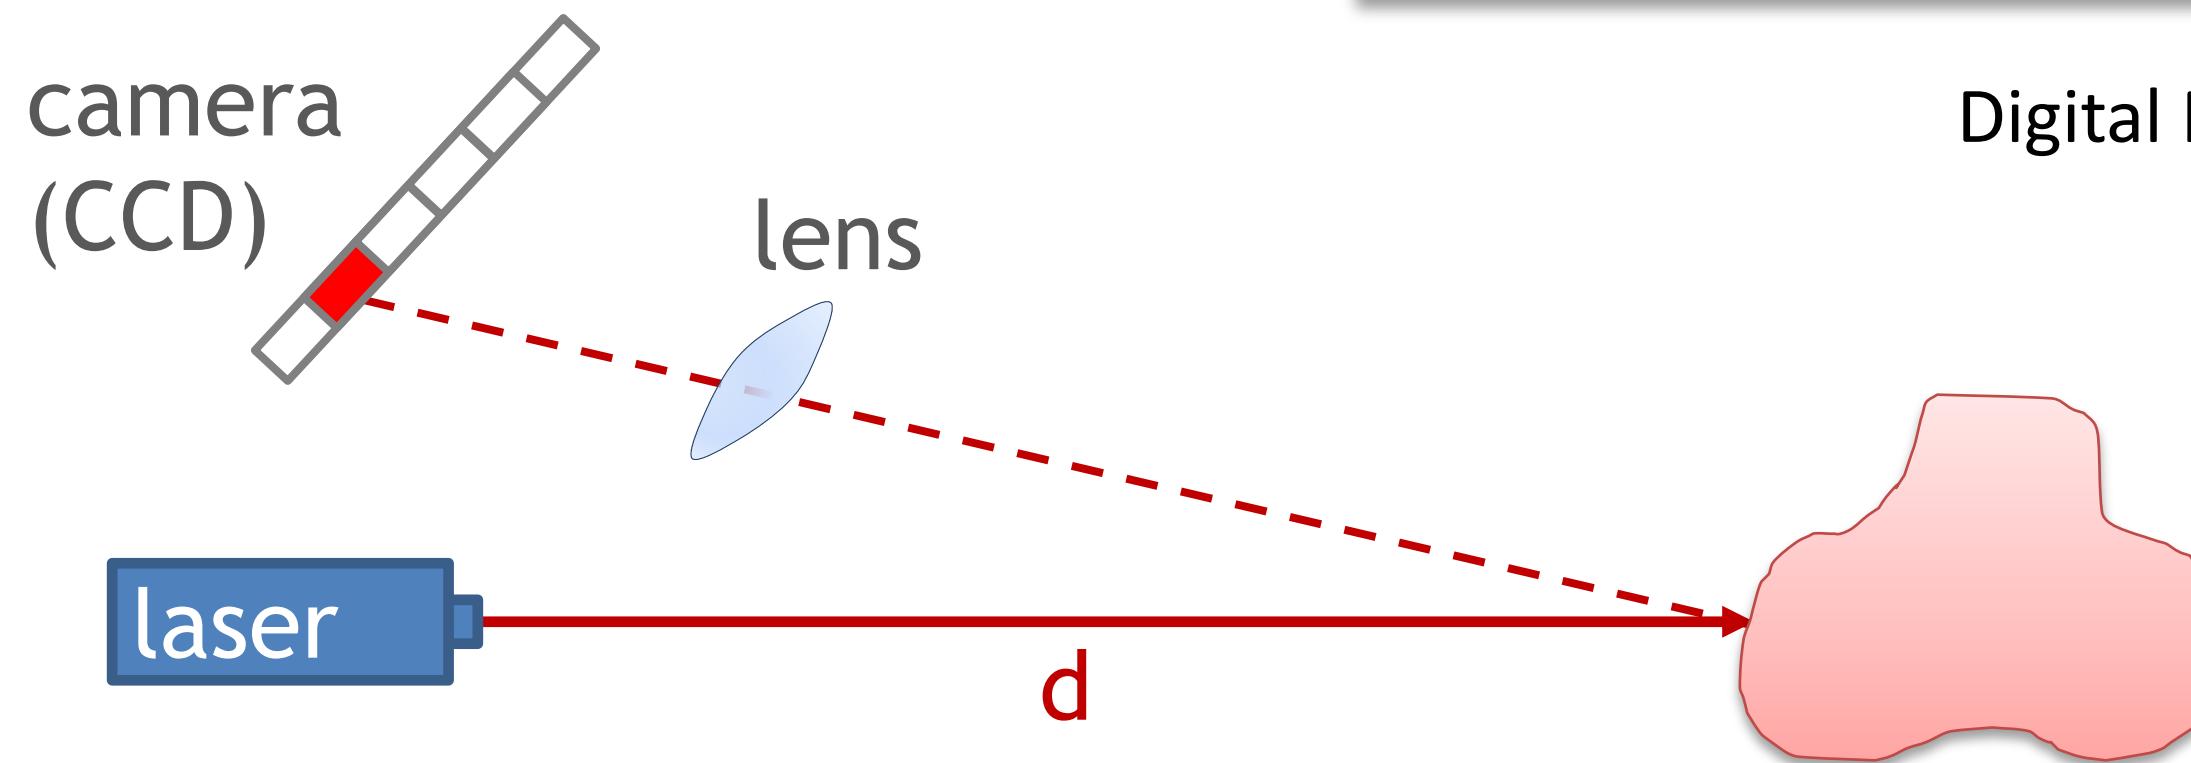
Optical scanning – active lighting

Triangulation laser

- Laser beam and camera
- Laser dot is photographed
- The location of the dot in the image allows triangulation: we get the distance to the object



Digital Michelangelo Project



Optical scanning – active lighting

Triangulation laser

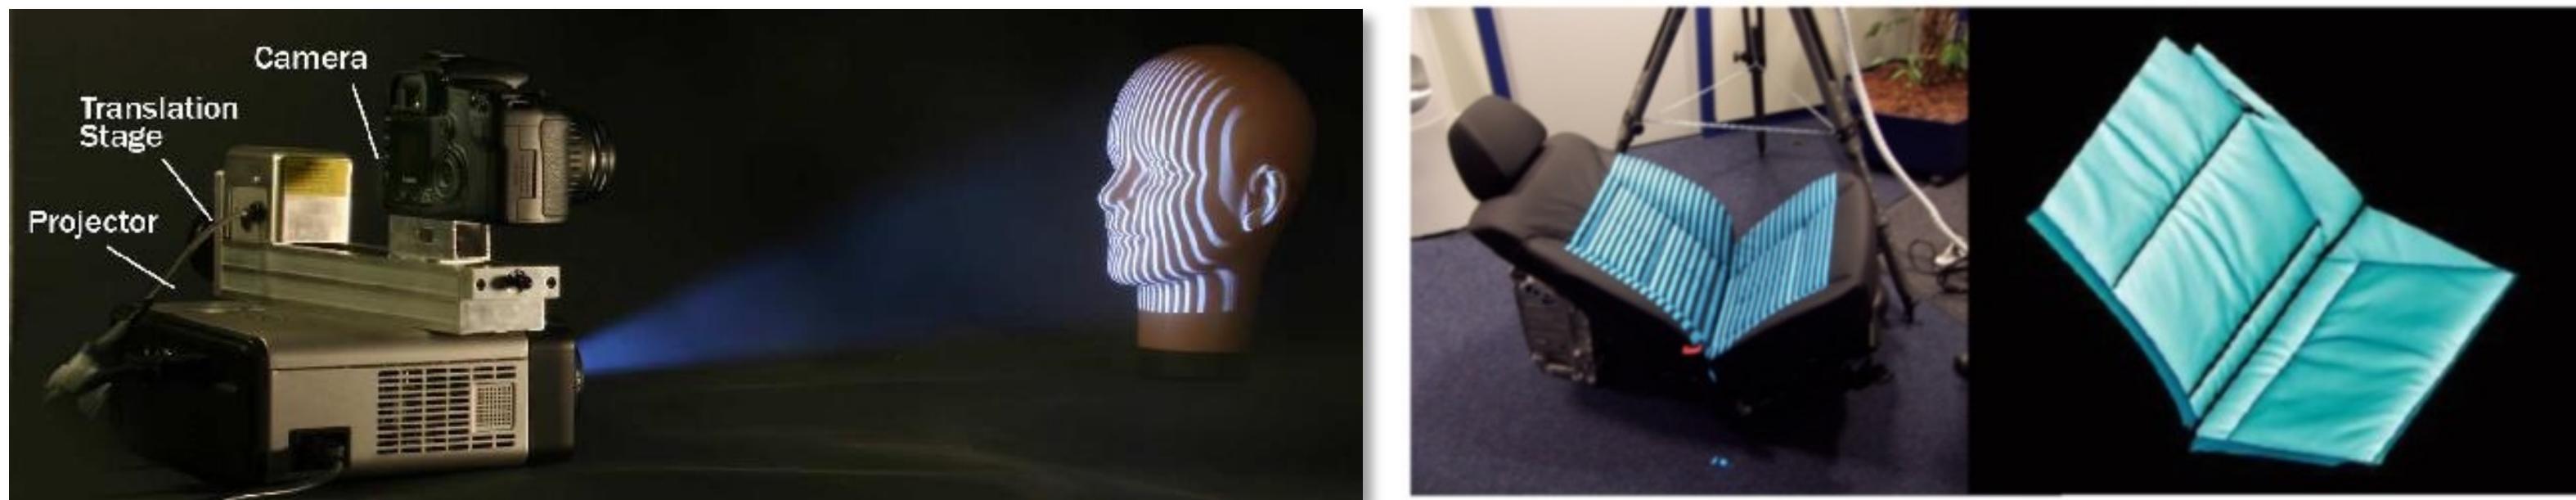
- Very precise (tens of microns)
- Small distances (meters)



Optical scanning – active lighting

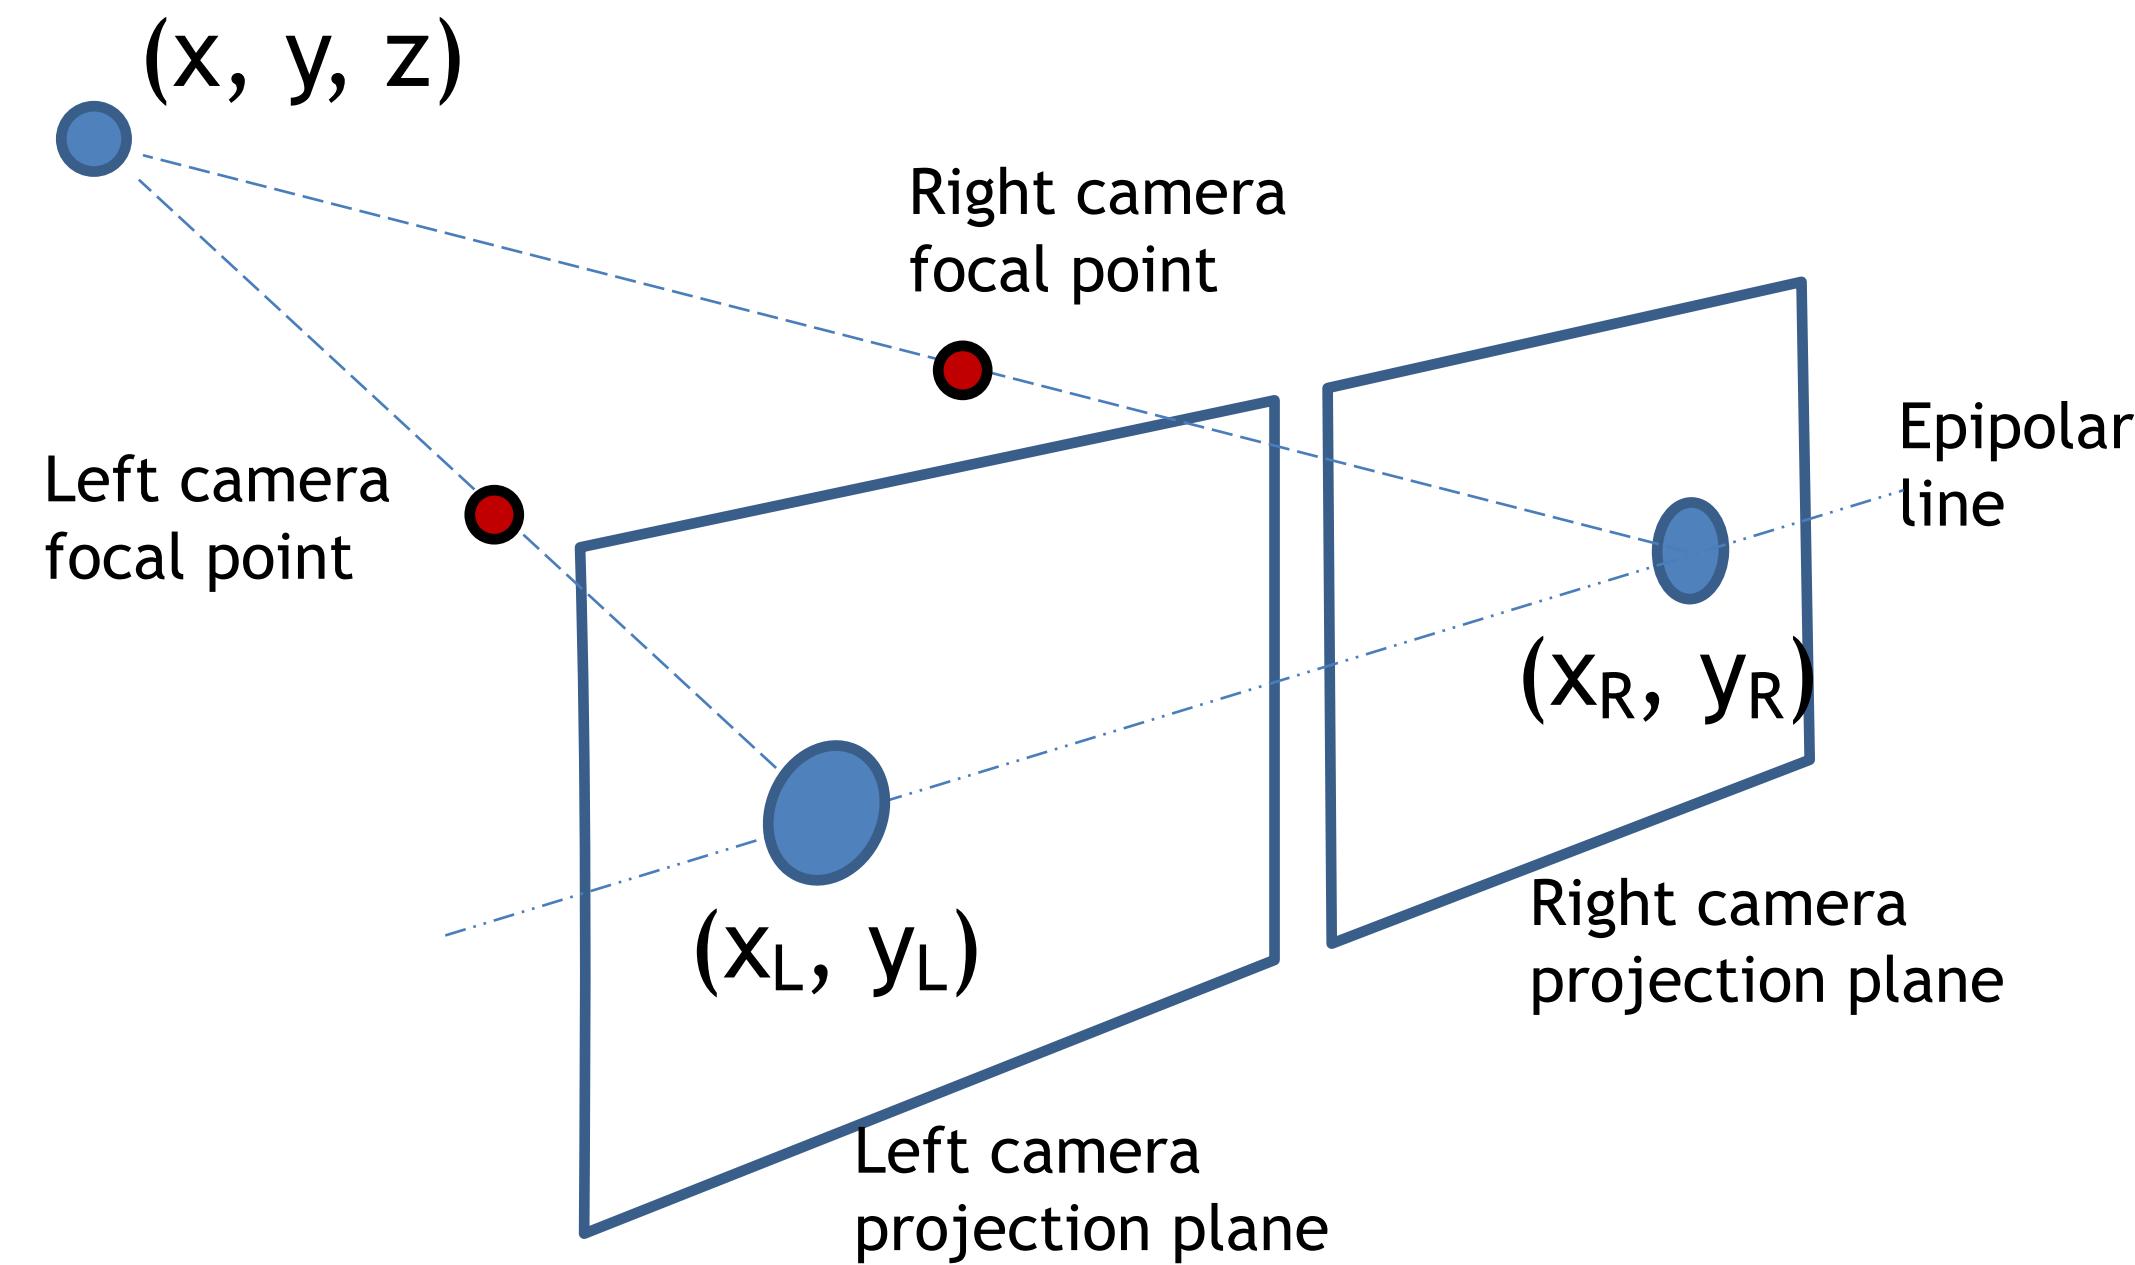
Structured light

- Pattern of visible or **infrared** light is projected onto the object
- The distortion of the pattern, recorded by the camera, provides geometric information
- Very fast – 2D pattern at once
 - Even in real time, like Kinect 1.0
- Complex distance calculation, prone to noise



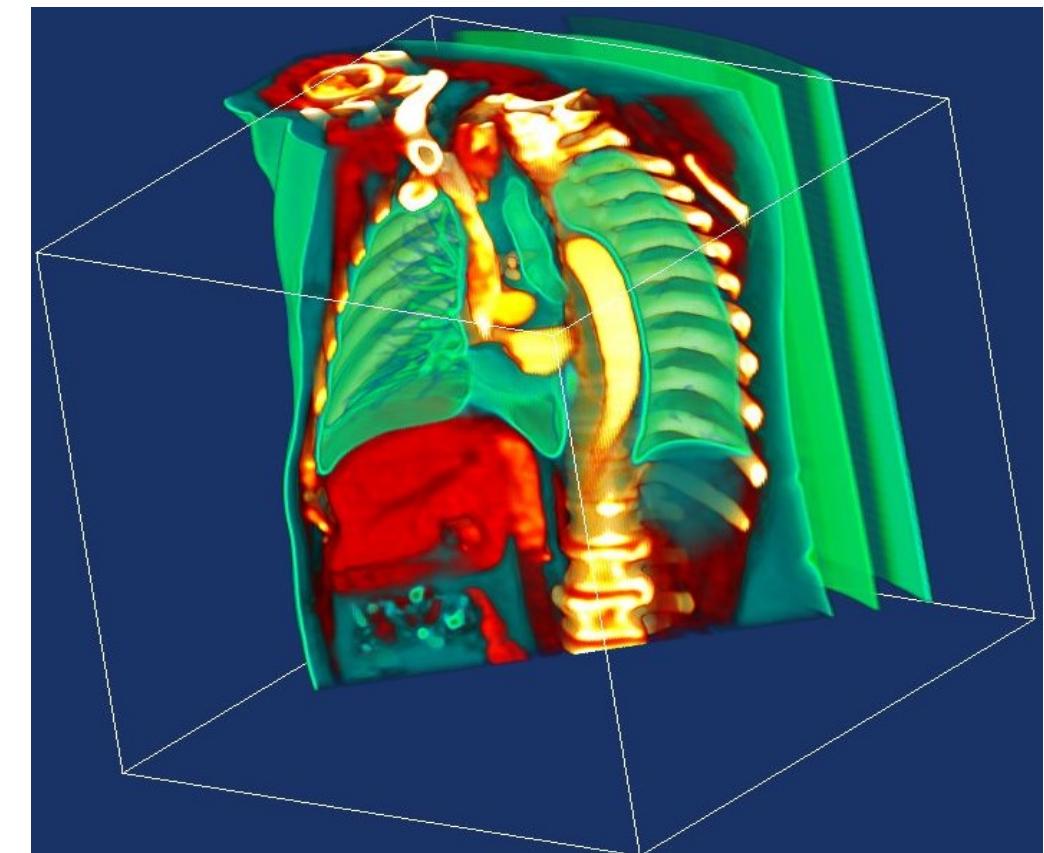
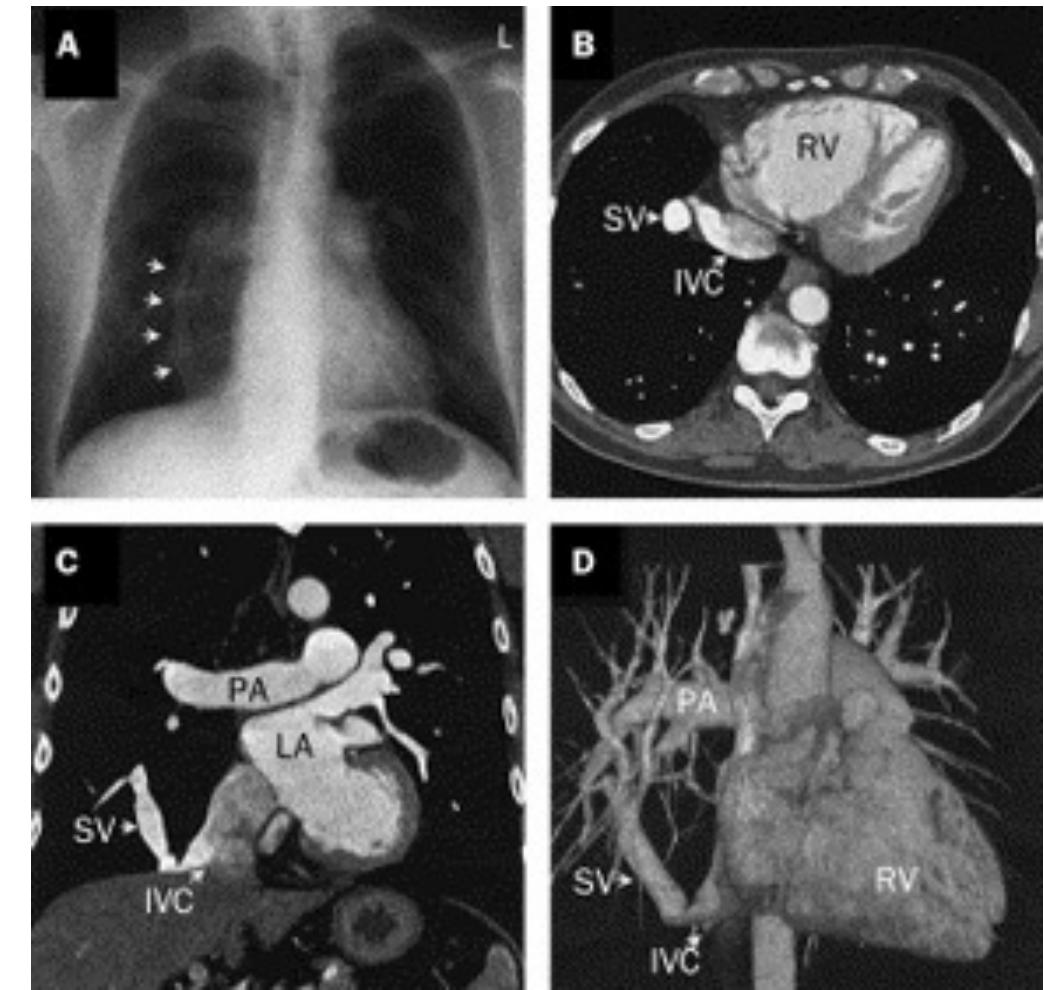
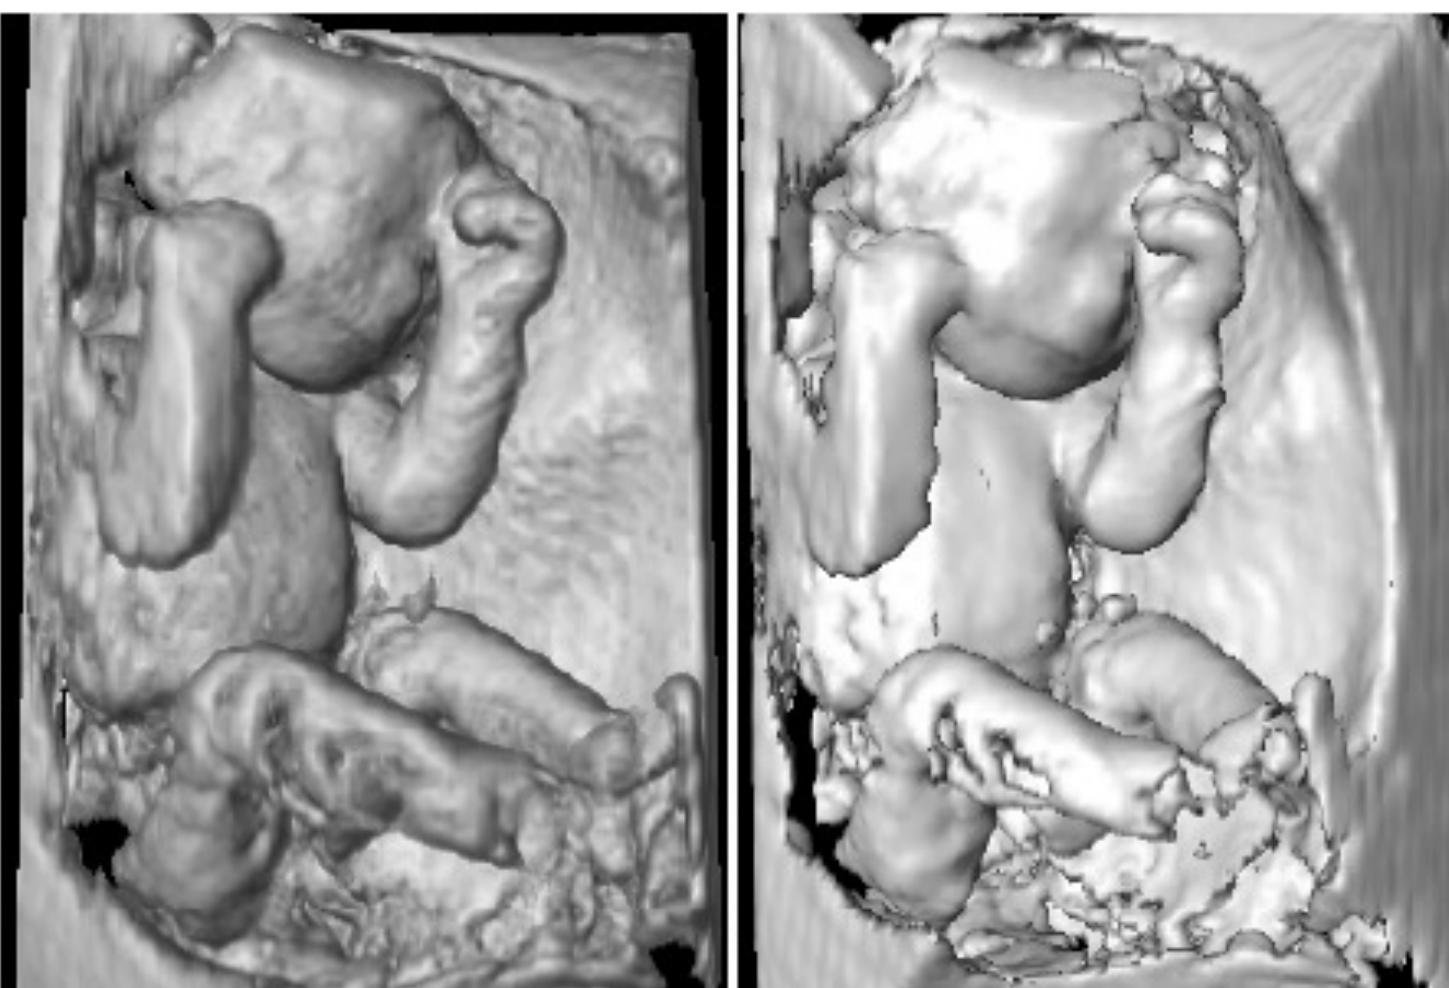
Optical scanning – passive stereo

- No need for special lighting/radiation
- Two (or more) cameras
- Feature matching and triangulation



Imaging

- Ultrasound, CT, MRI
- Discrete volume of density data
- First need to segment the desired object (contouring)

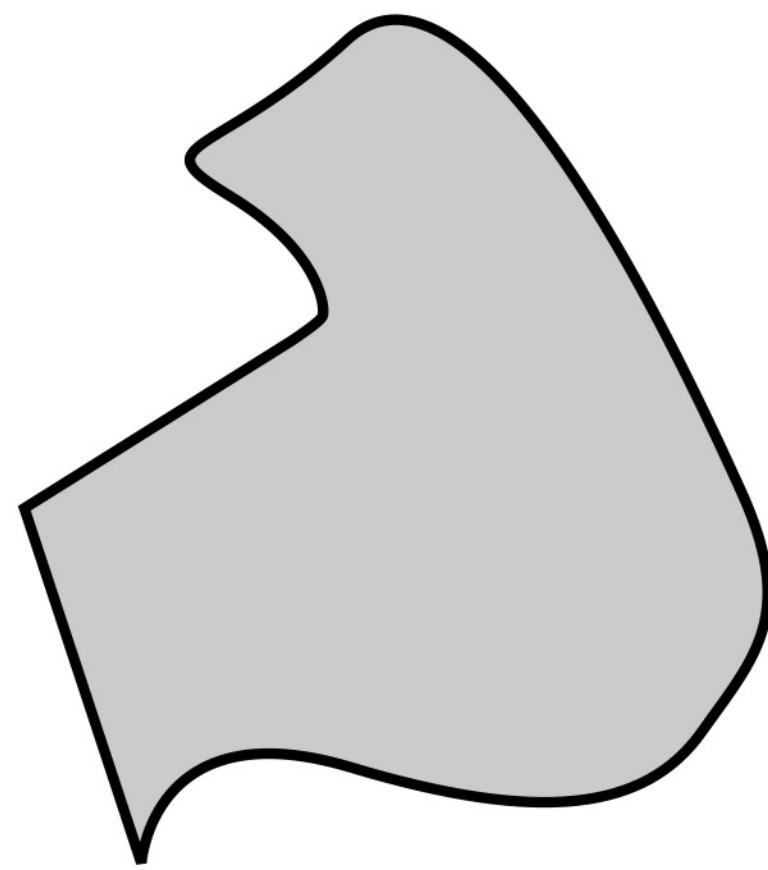


Registration

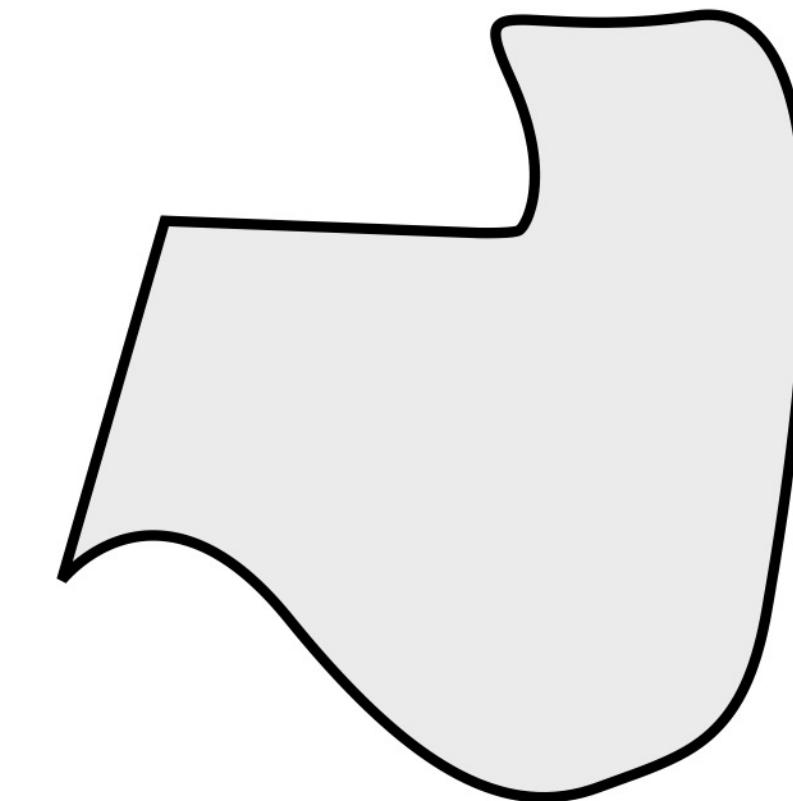
Acknowledgement: Niloy Mitra
http://resources.mpi-inf.mpg.de/deformableShapeMatching/EG2012_Tutorial/

Problem Statement

M_1



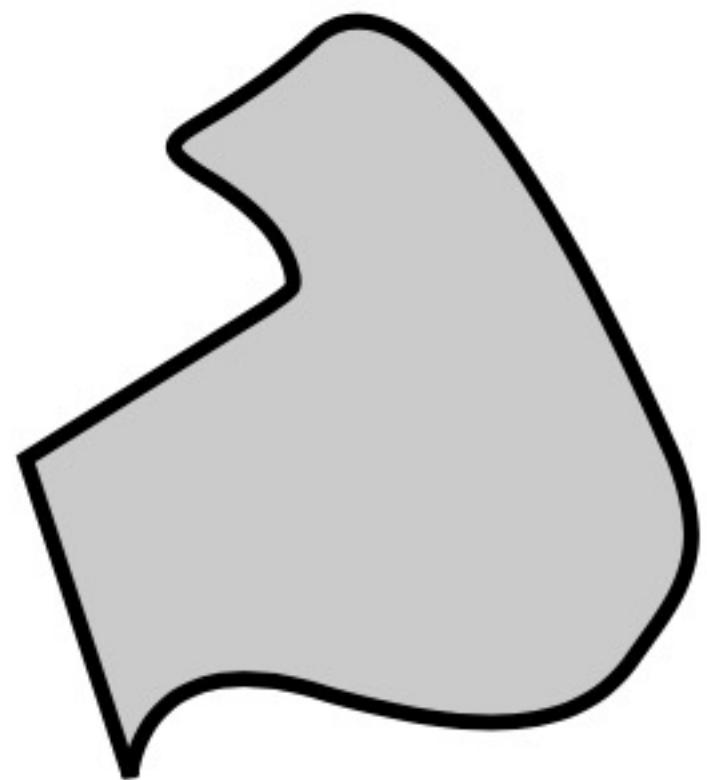
M_2



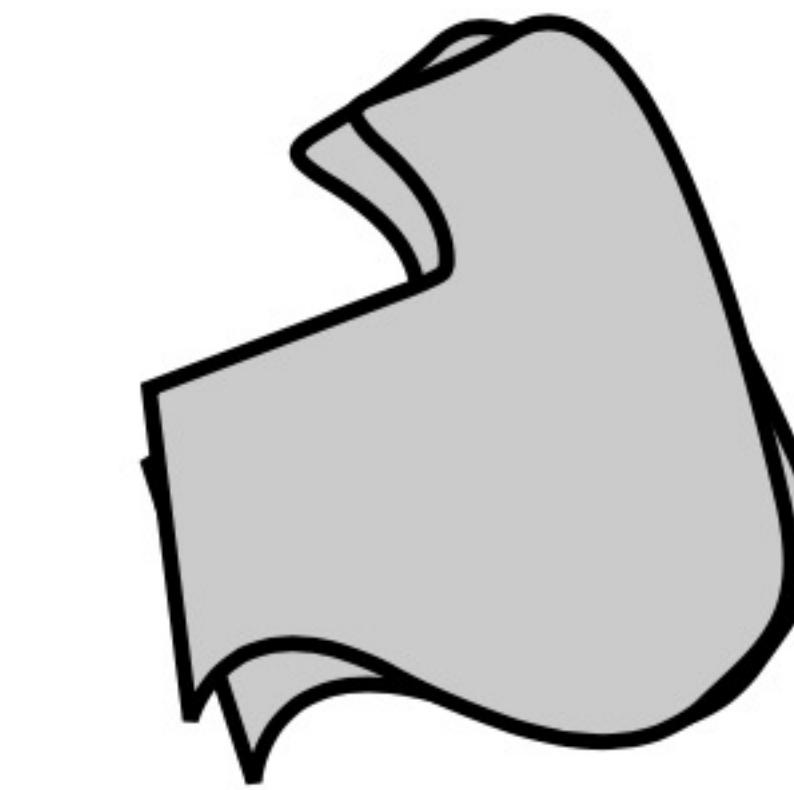
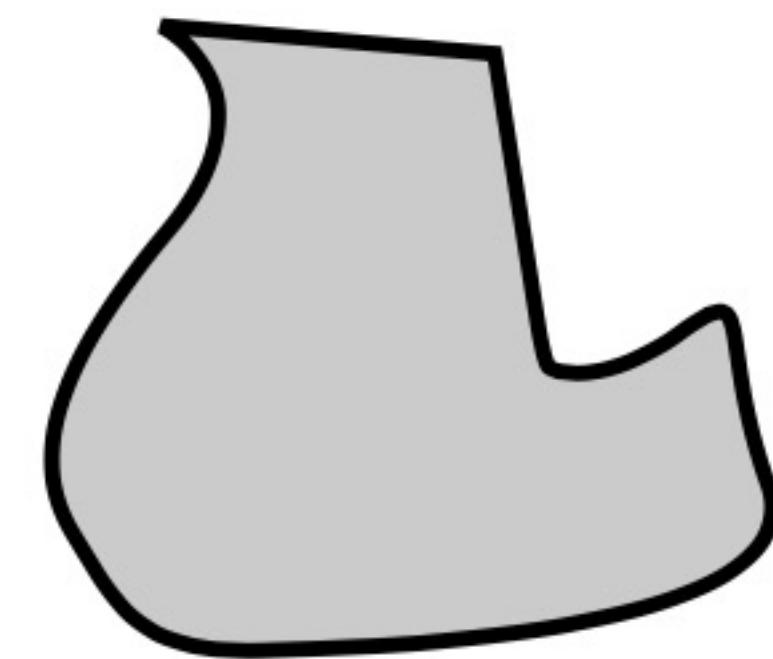
$$M_1 \approx T(M_2)$$

T: Translation + Rotation

Local vs Global



Global Registration
Arbitrary Transformation



Local Registration
“Small” Transformation

Given M_1, \dots, M_n , find T_2, \dots, T_n such that

$$M_1 \approx T_2(M_2) \cdots \approx T_n(M_n)$$

Correspondences

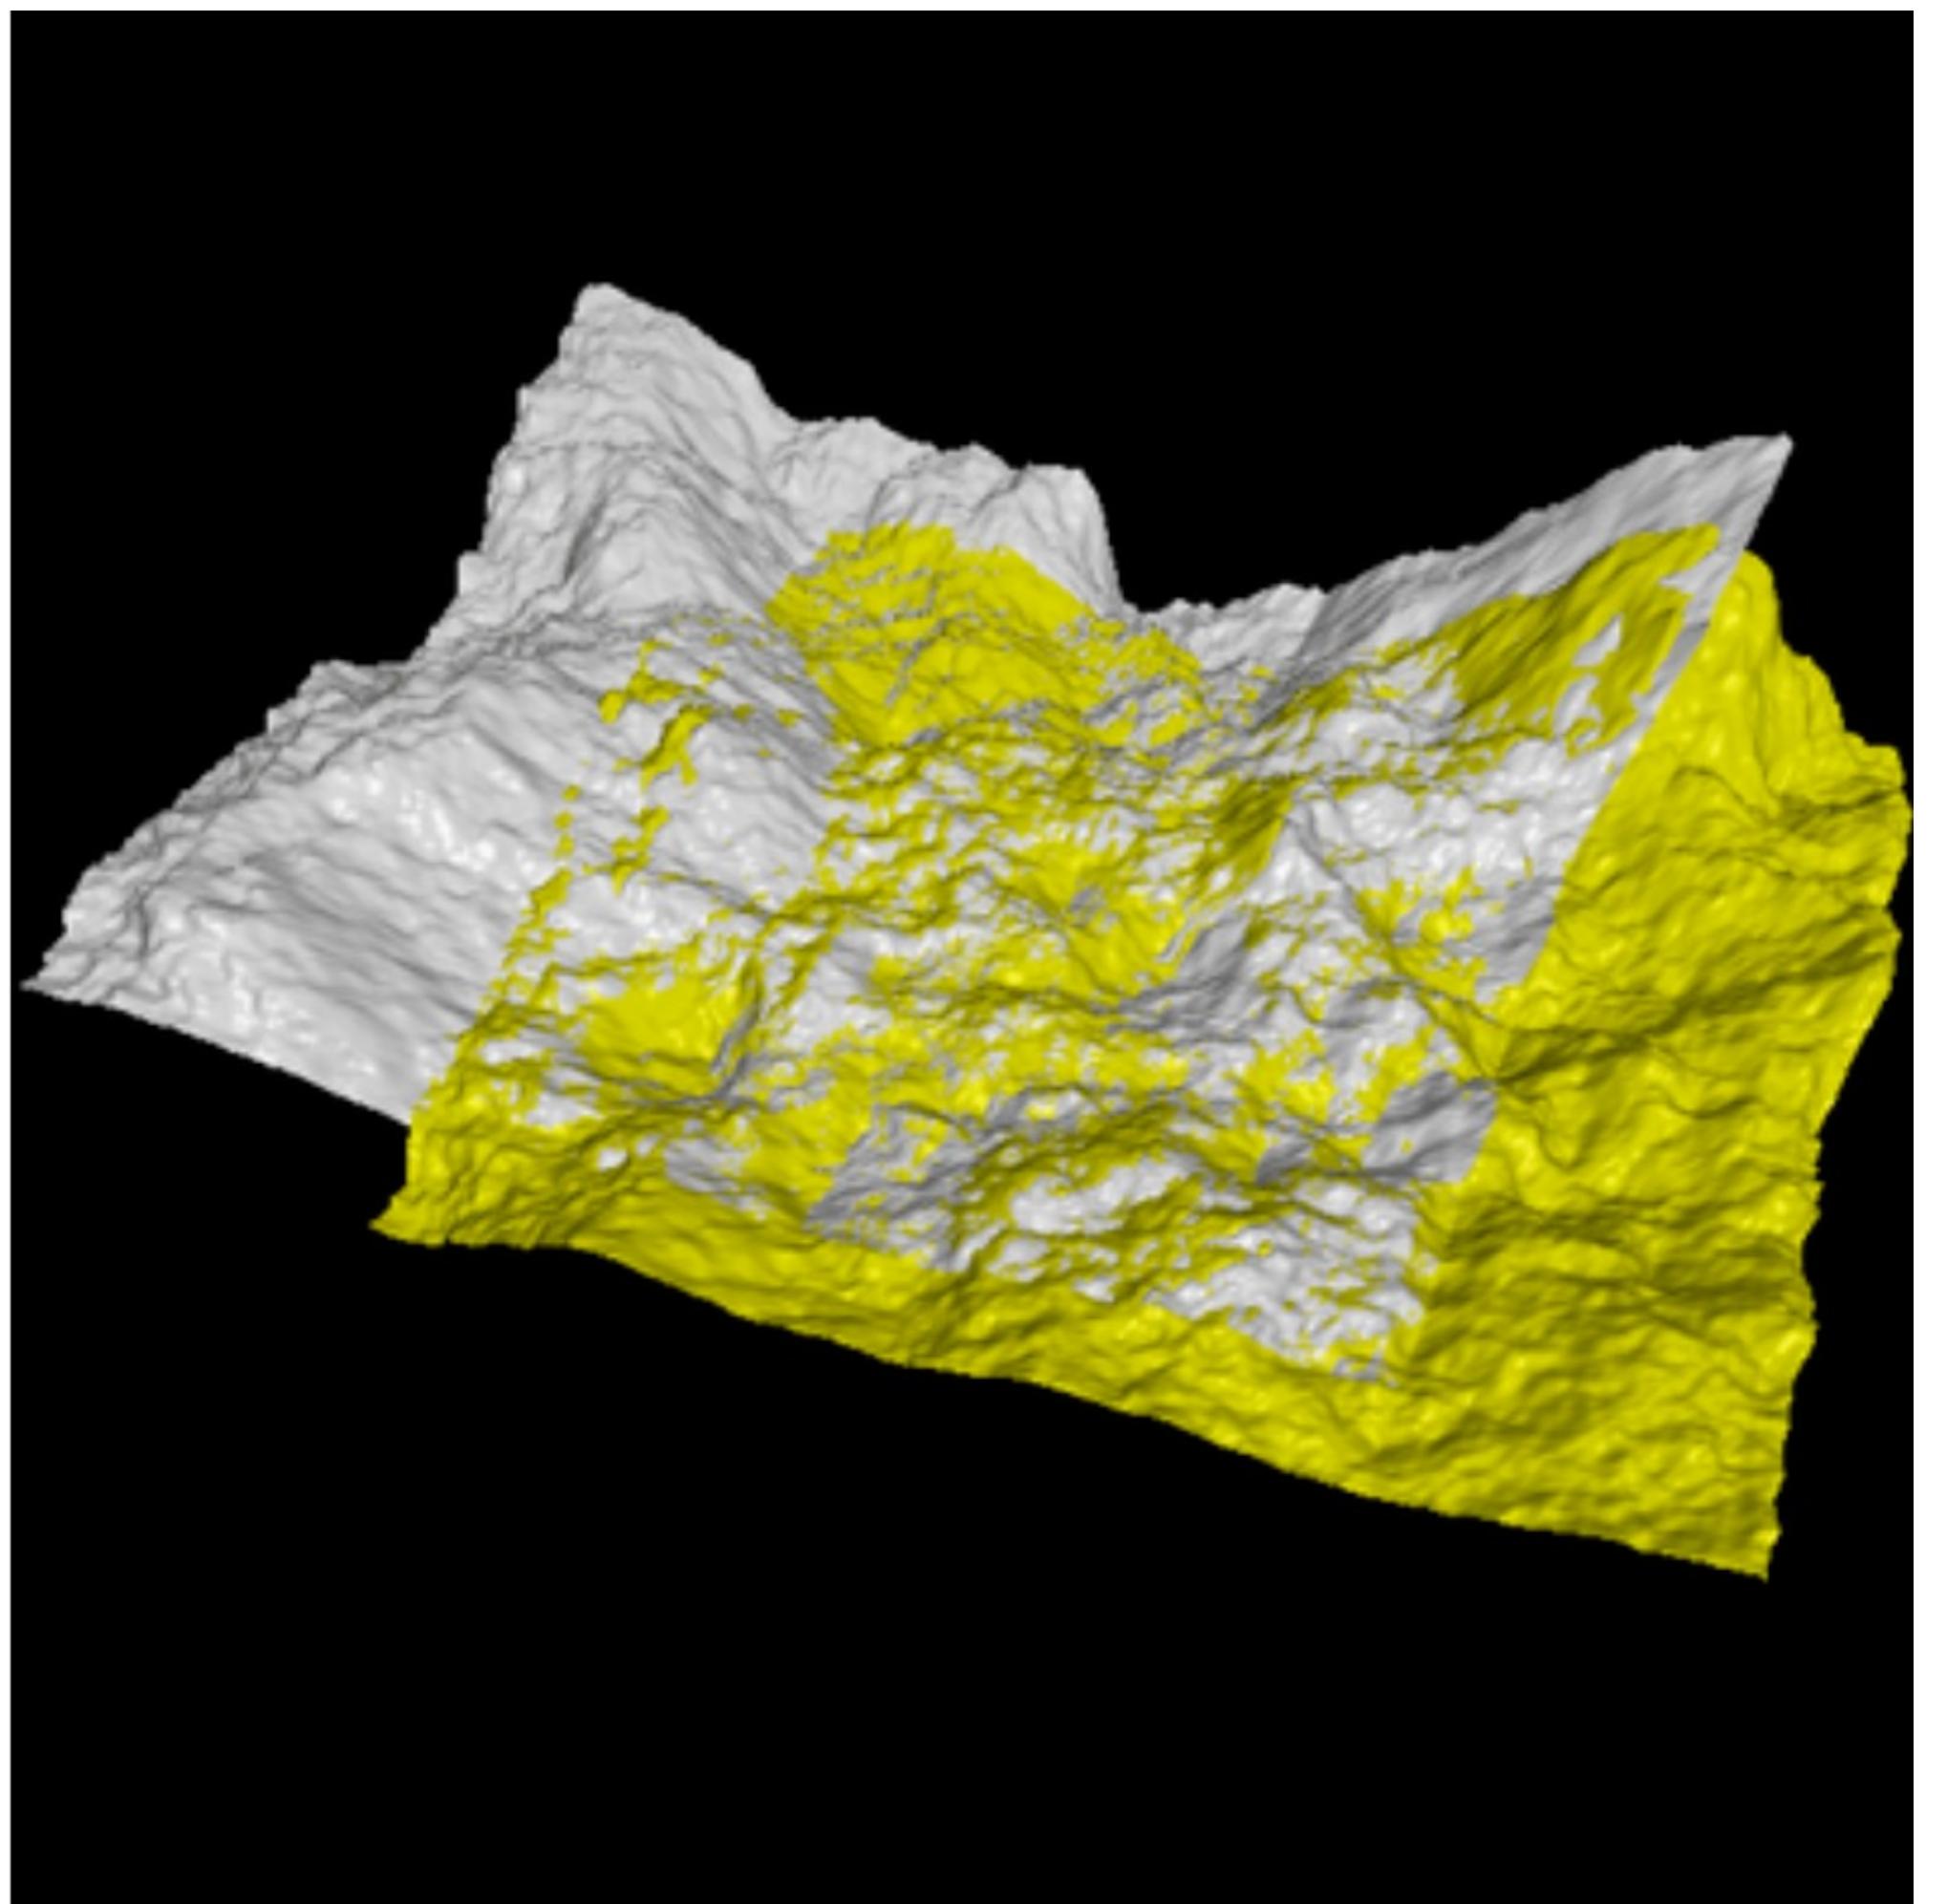
- How many points are needed to define a unique rigid transformation?
- The first problem is finding corresponding pairs!

$$\mathbf{p}_1 \rightarrow \mathbf{q}_1$$

$$\mathbf{p}_2 \rightarrow \mathbf{q}_2$$

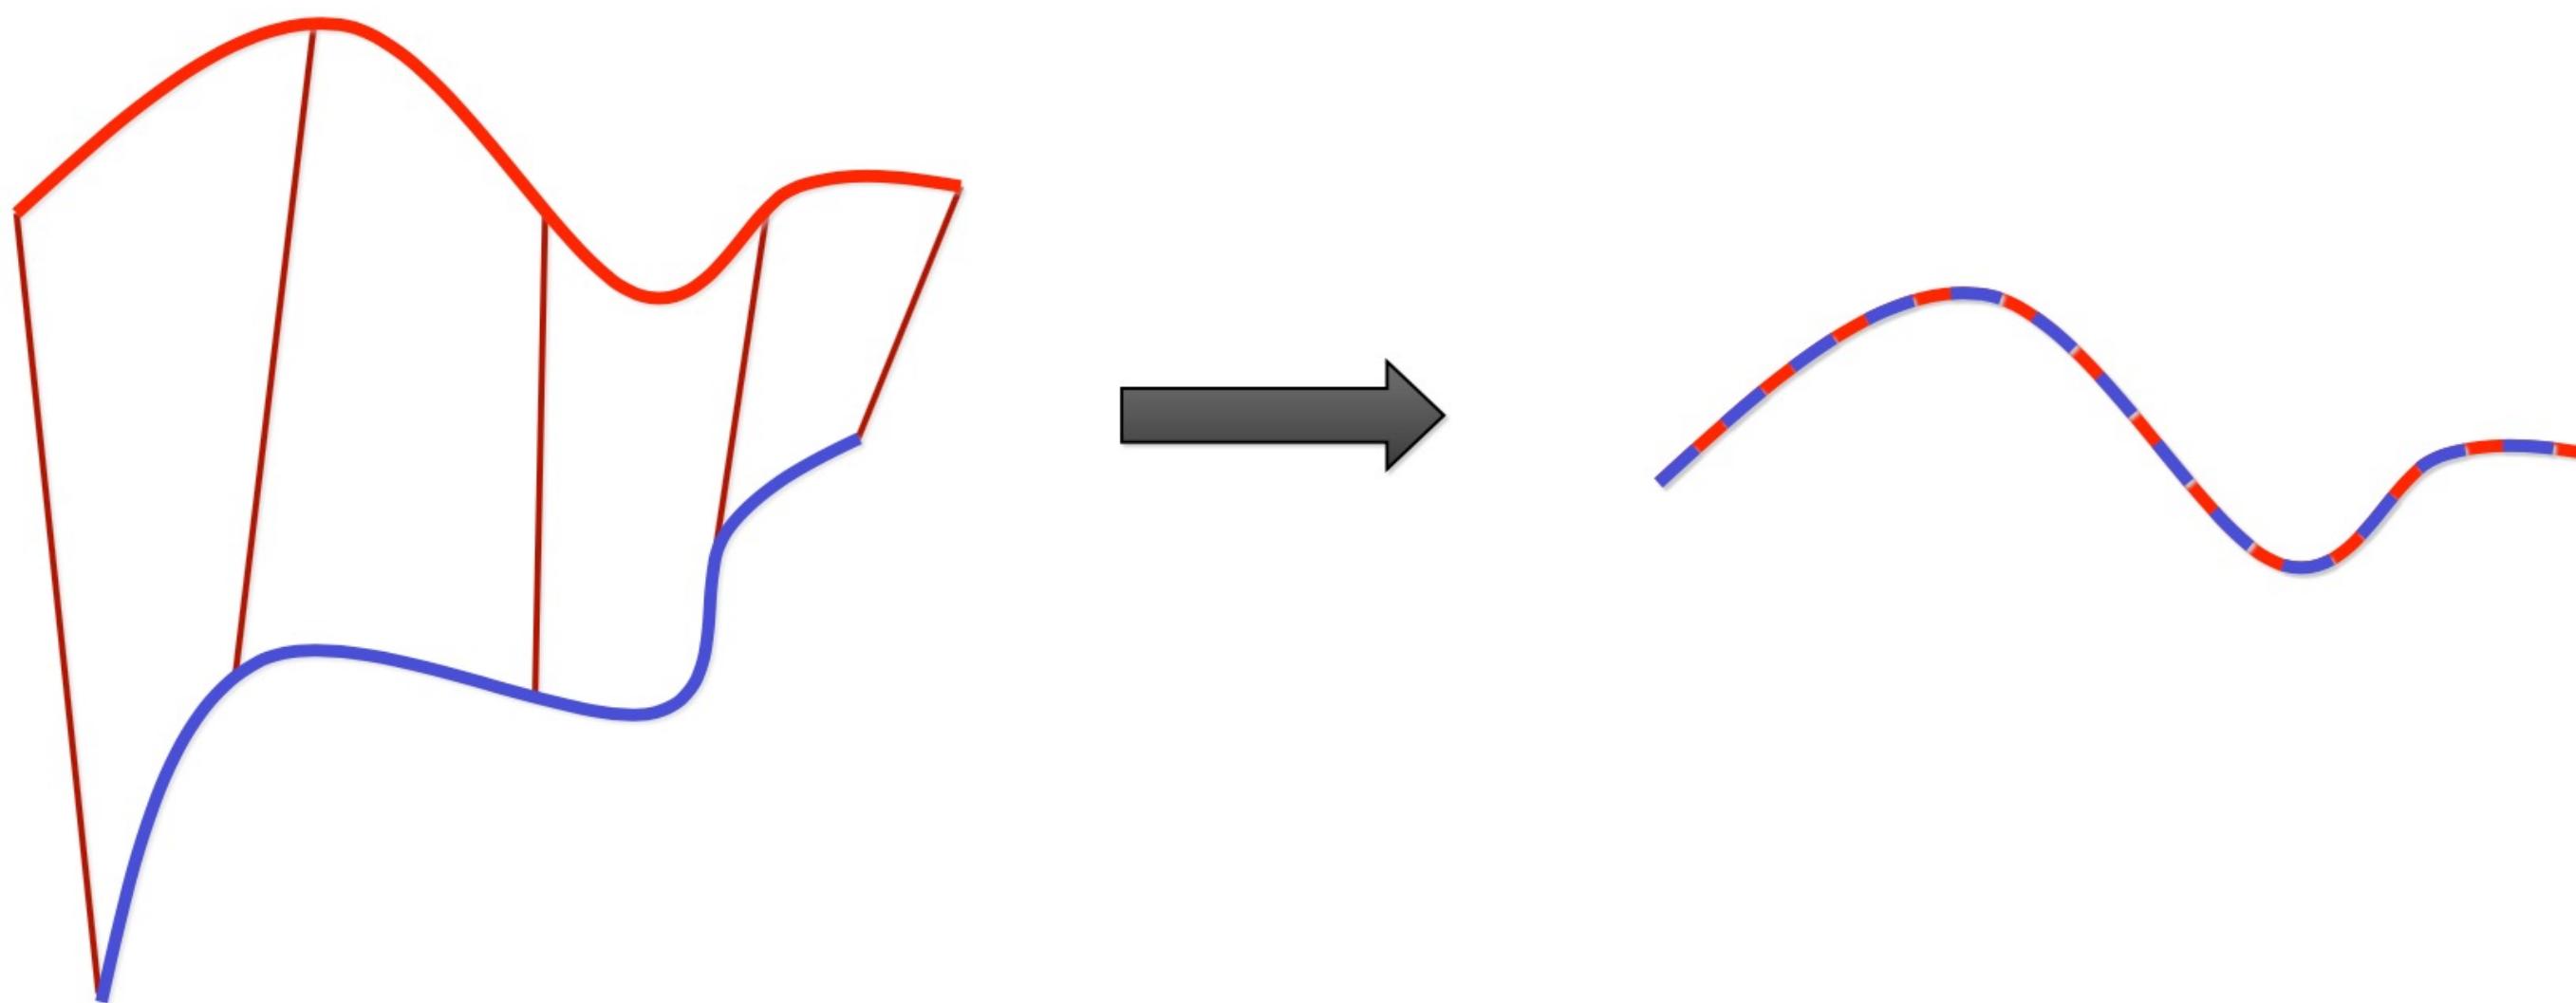
$$\mathbf{p}_3 \rightarrow \mathbf{q}_3$$

$$R\mathbf{p}_i + t \approx \mathbf{q}_i$$



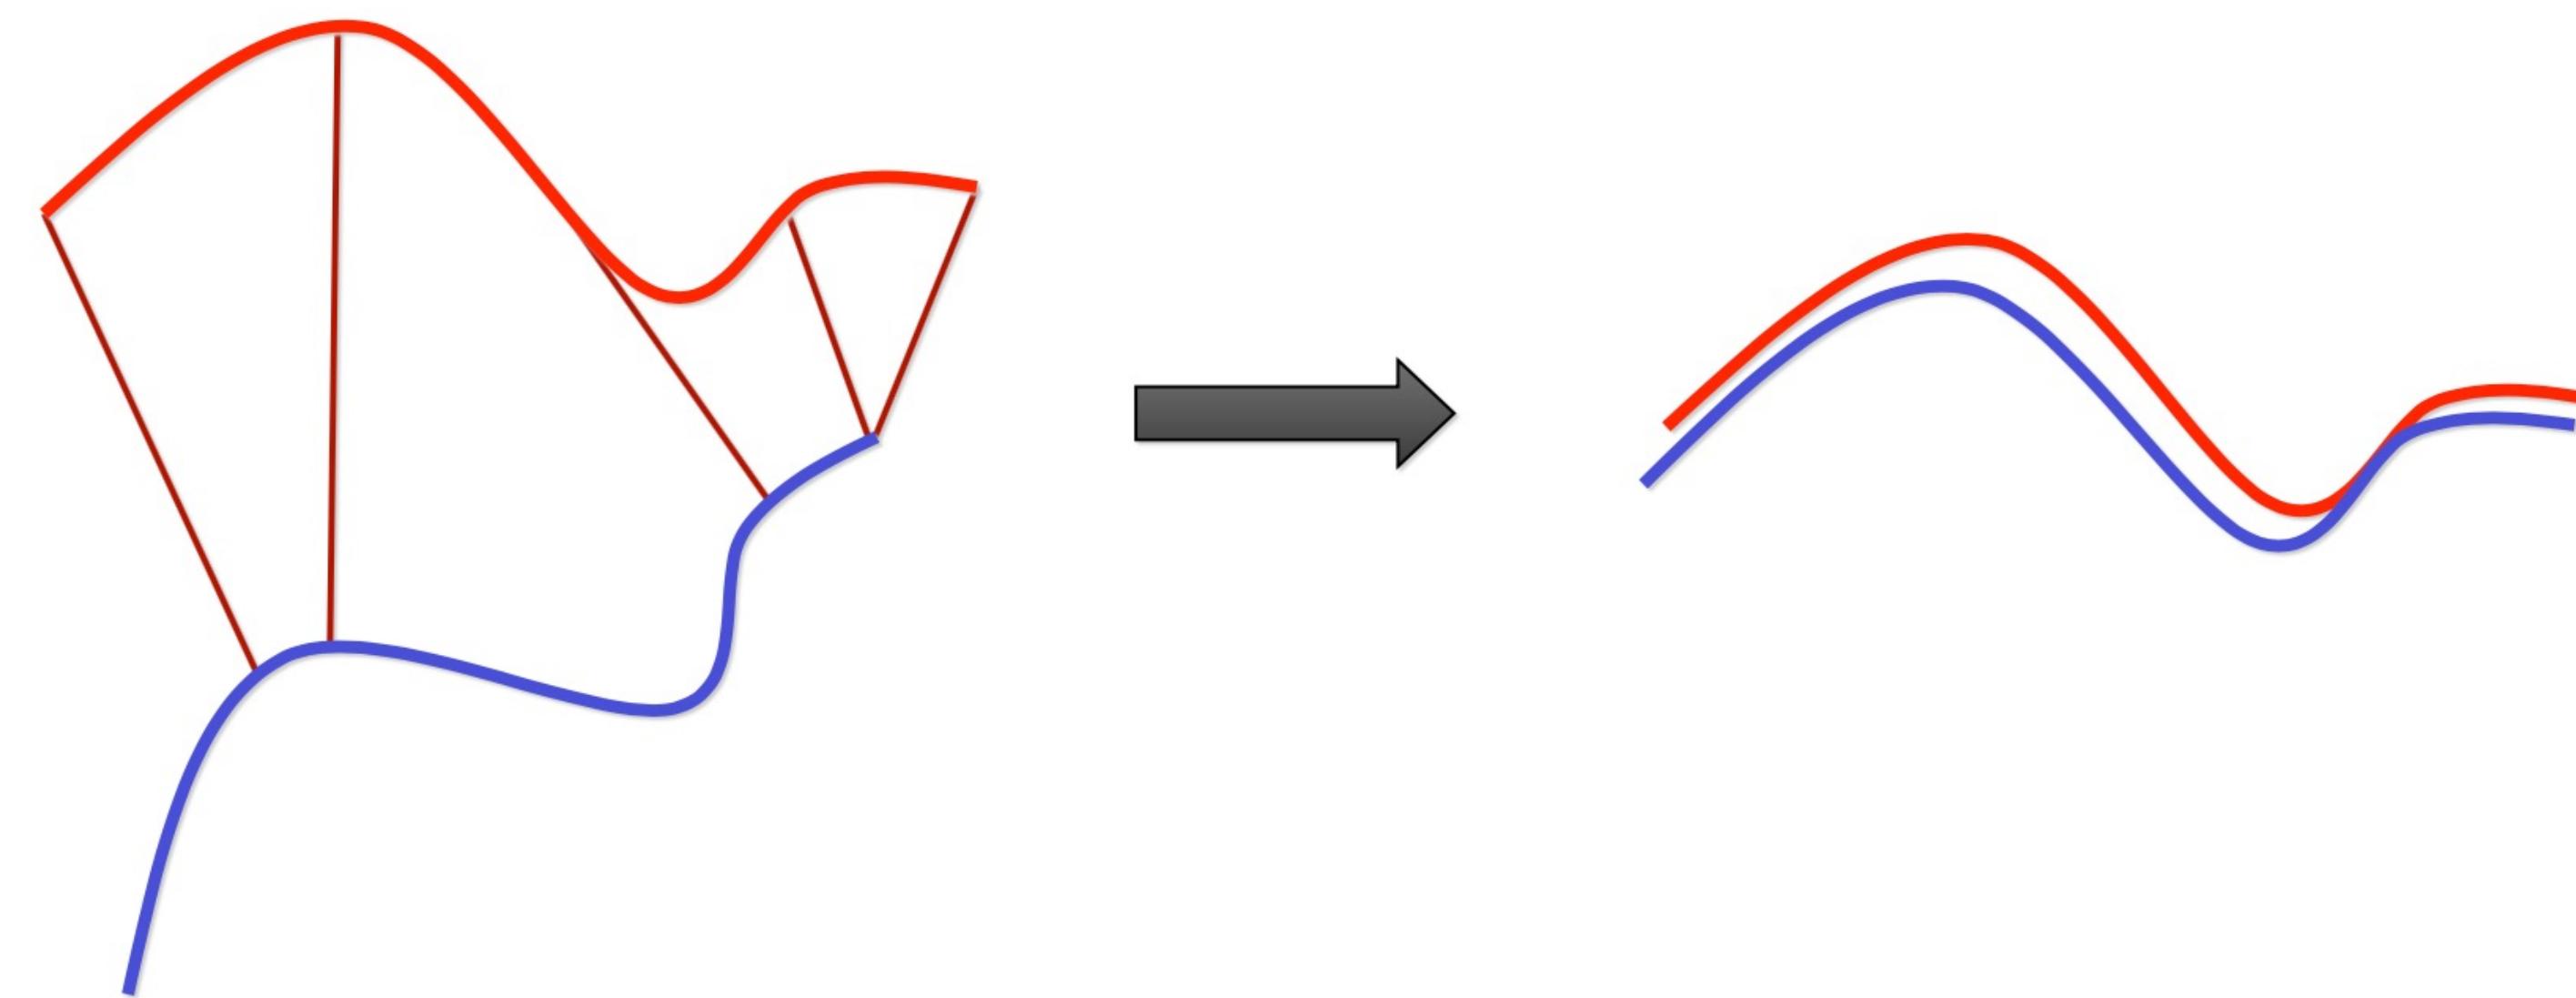
ICP: Iterative Closest Point

- Idea: Iteratively (1) find correspondences and (2) use them to find a transformation
- Intuition: If you have the right correspondences, then the problem is easy

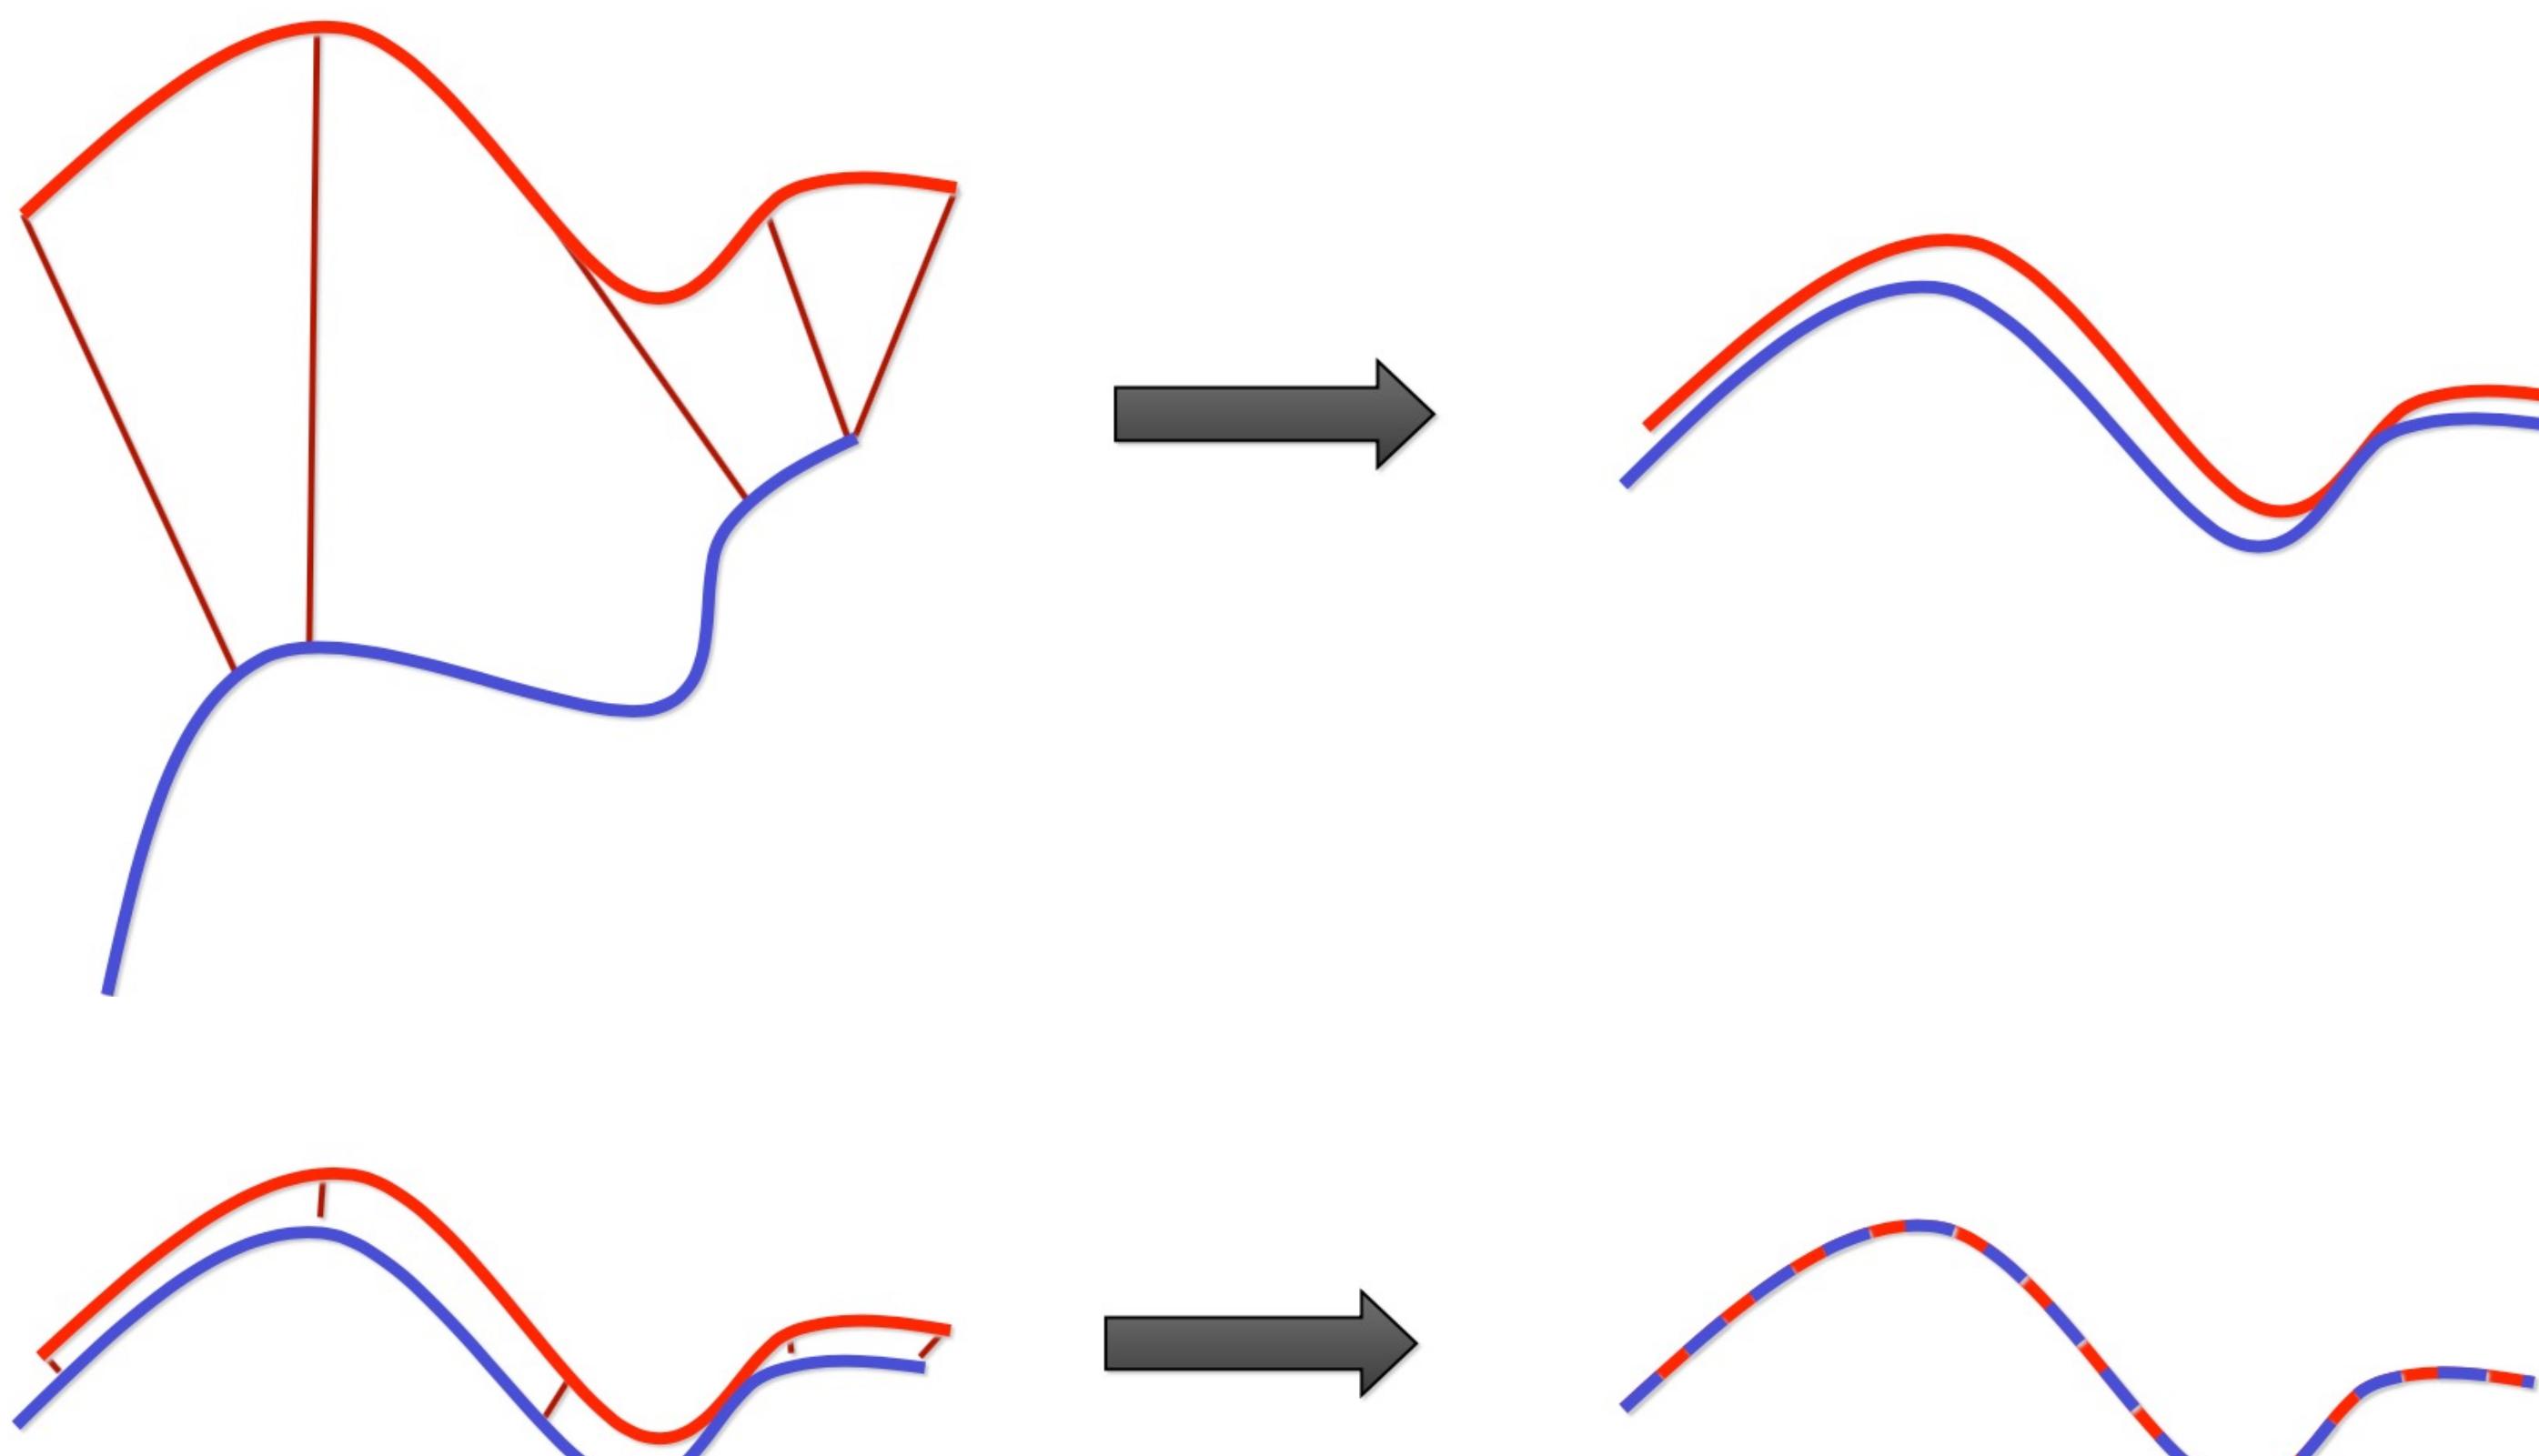


ICP: Iterative Closest Point

- Idea: Iteratively (1) find correspondences and (2) use them to find a transformation
- Intuition: If you don't have the right correspondences, you still can make progress



ICL: Iterative Closest Point

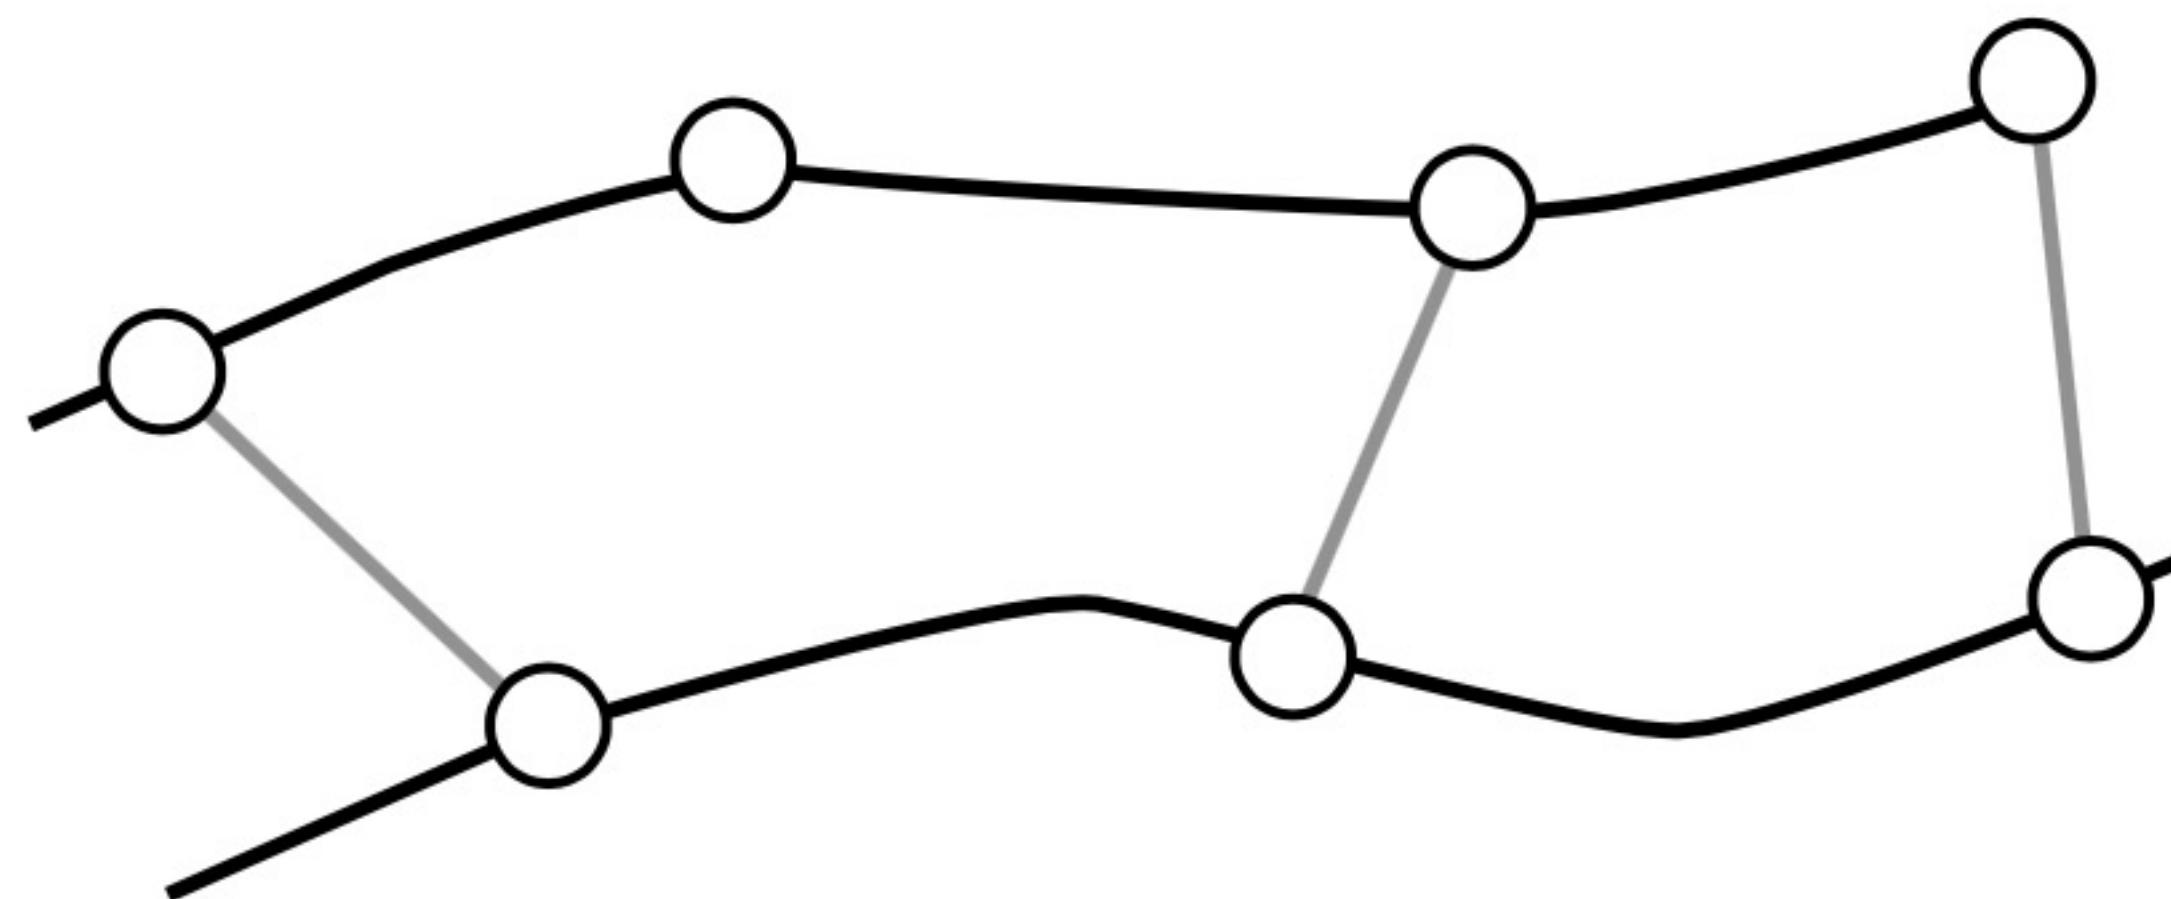


This algorithm converges to the correct solution only
if the starting scans are “close enough”

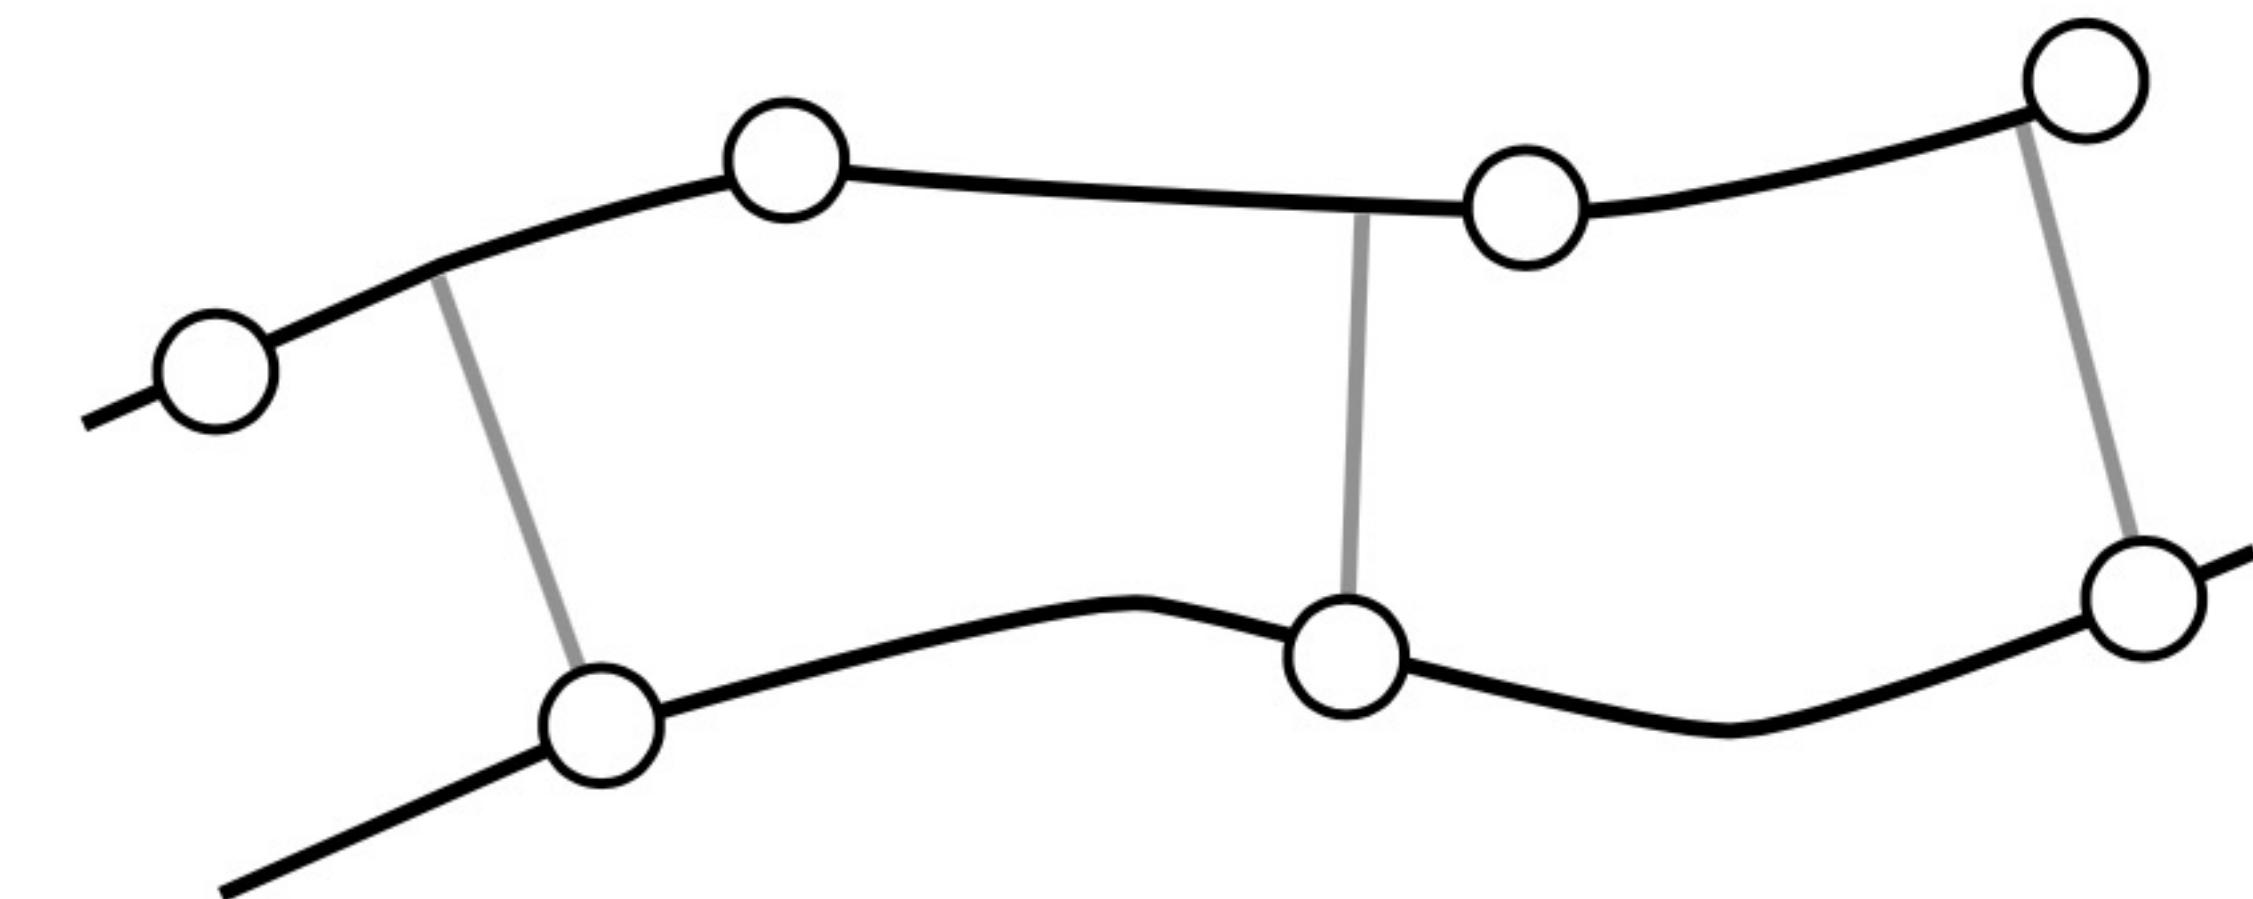
Basic Algorithm

- Select (e.g., 1000) random points
- Match each to closest point on other scan, using data structure such as k -d tree
- Reject pairs with distance $> k$ times median
- Construct error function:
$$E := \sum_i (R\mathbf{p}_i + t - \mathbf{q}_i)^2$$
- Minimize (closed form solution in “Estimating 3-D rigid body transformations: a comparison of four major algorithms”, <http://dl.acm.org/citation.cfm?id=250160>)

Important Variant



Point-to-Point



Point-to-Plane

See http://resources.mpi-inf.mpg.de/deformableShapeMatching/EG2012_Tutorial/ for details

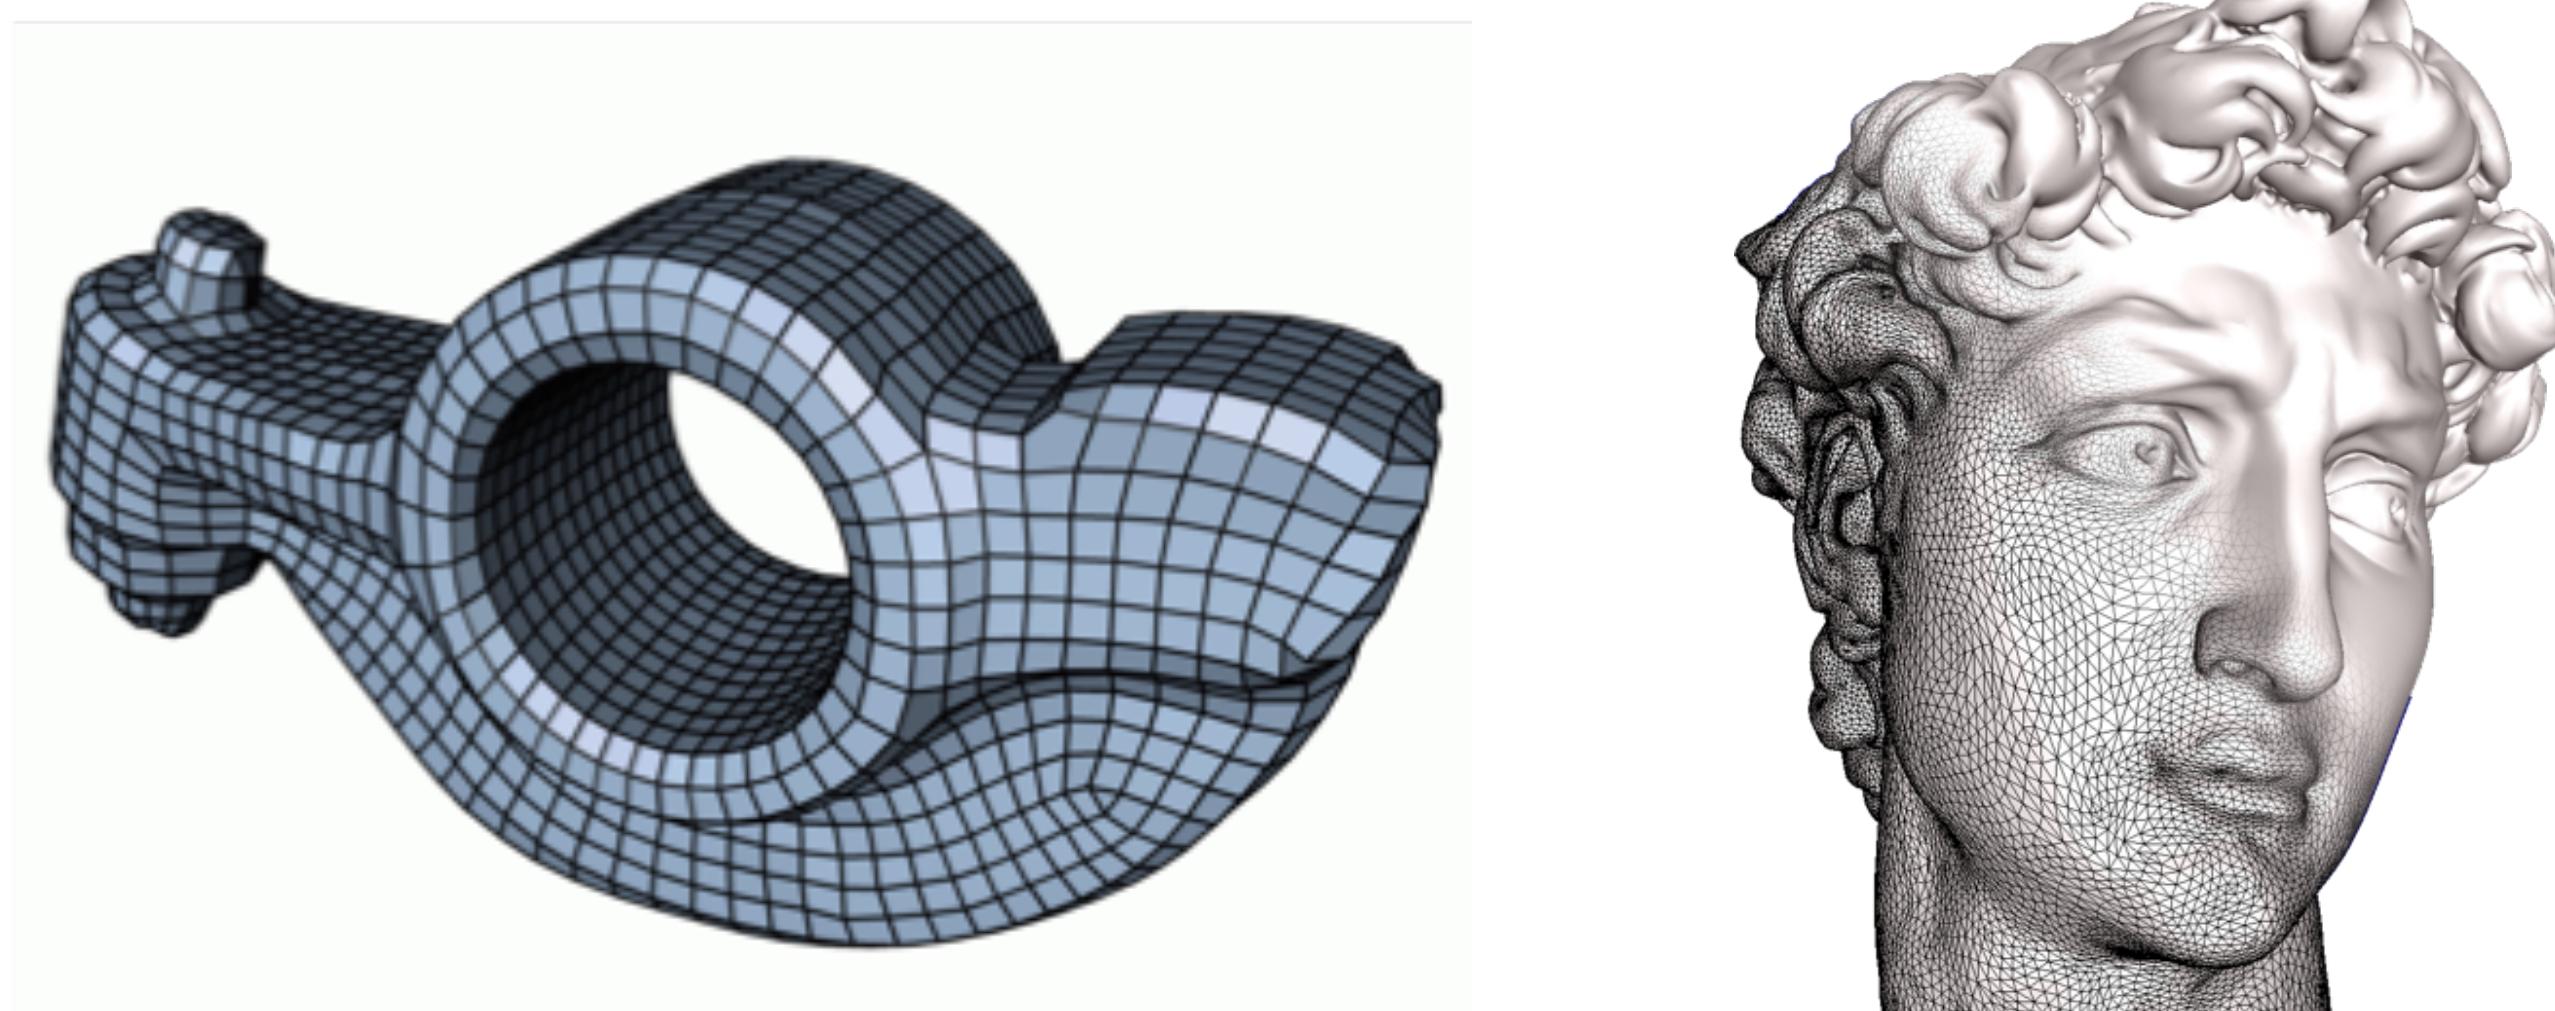
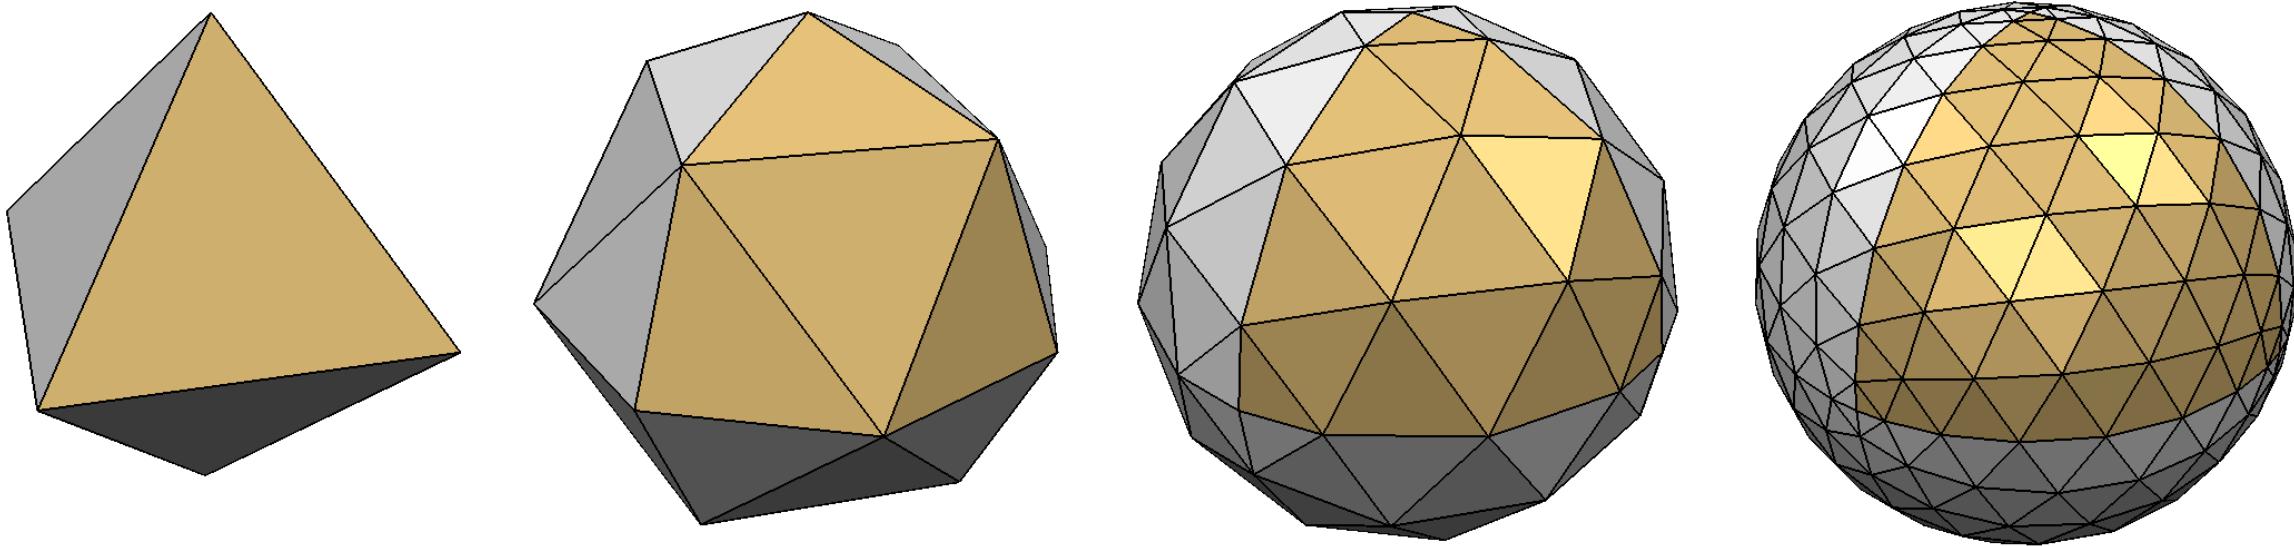
Mesh Representation

Libigl Tutorial

[102_DrawMesh]

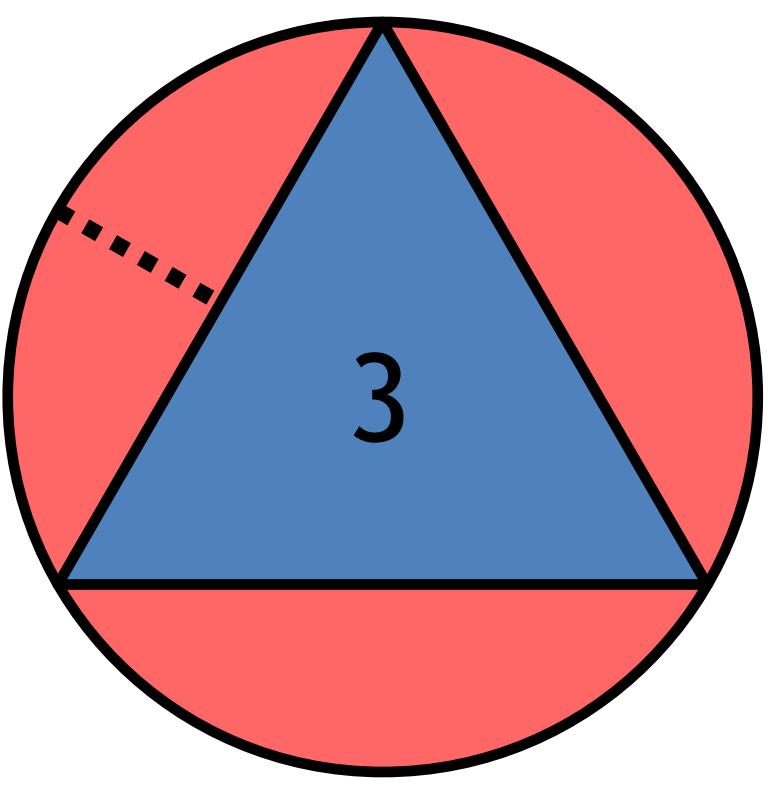
Polygonal Meshes

- Boundary representations of objects

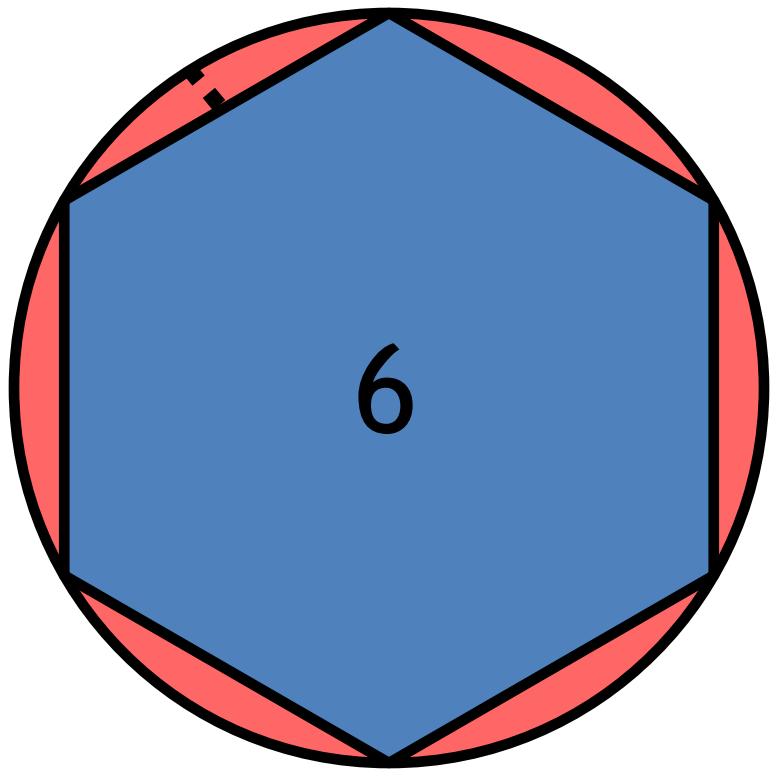


Meshes as Approximations of Smooth Surfaces

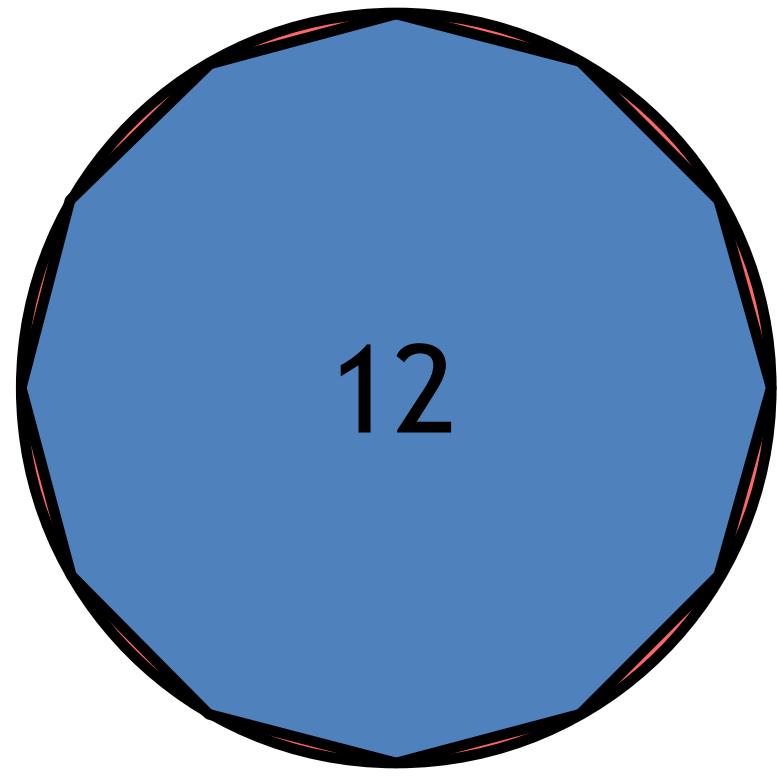
- Piecewise linear approximation



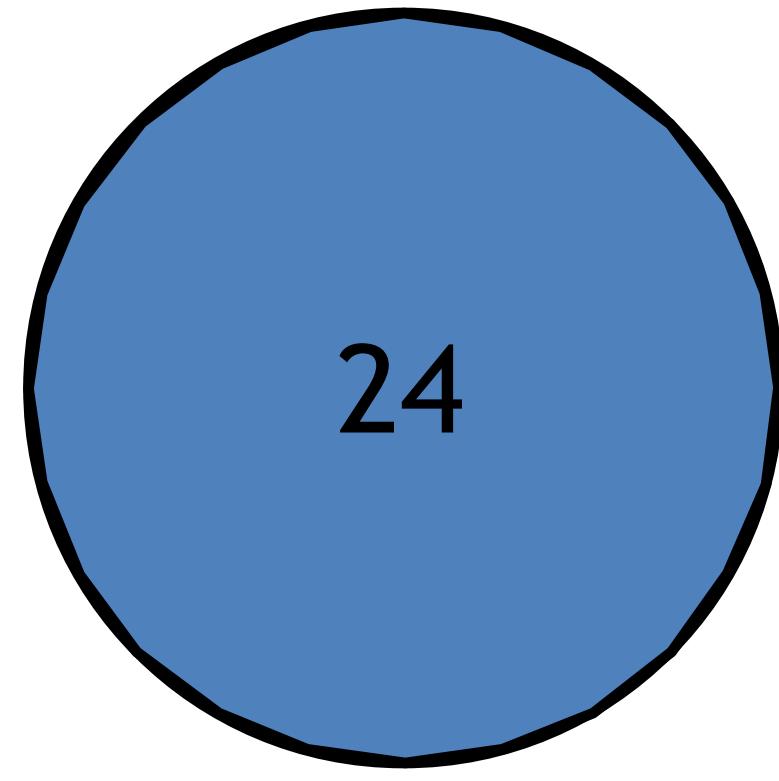
25%



6.5%



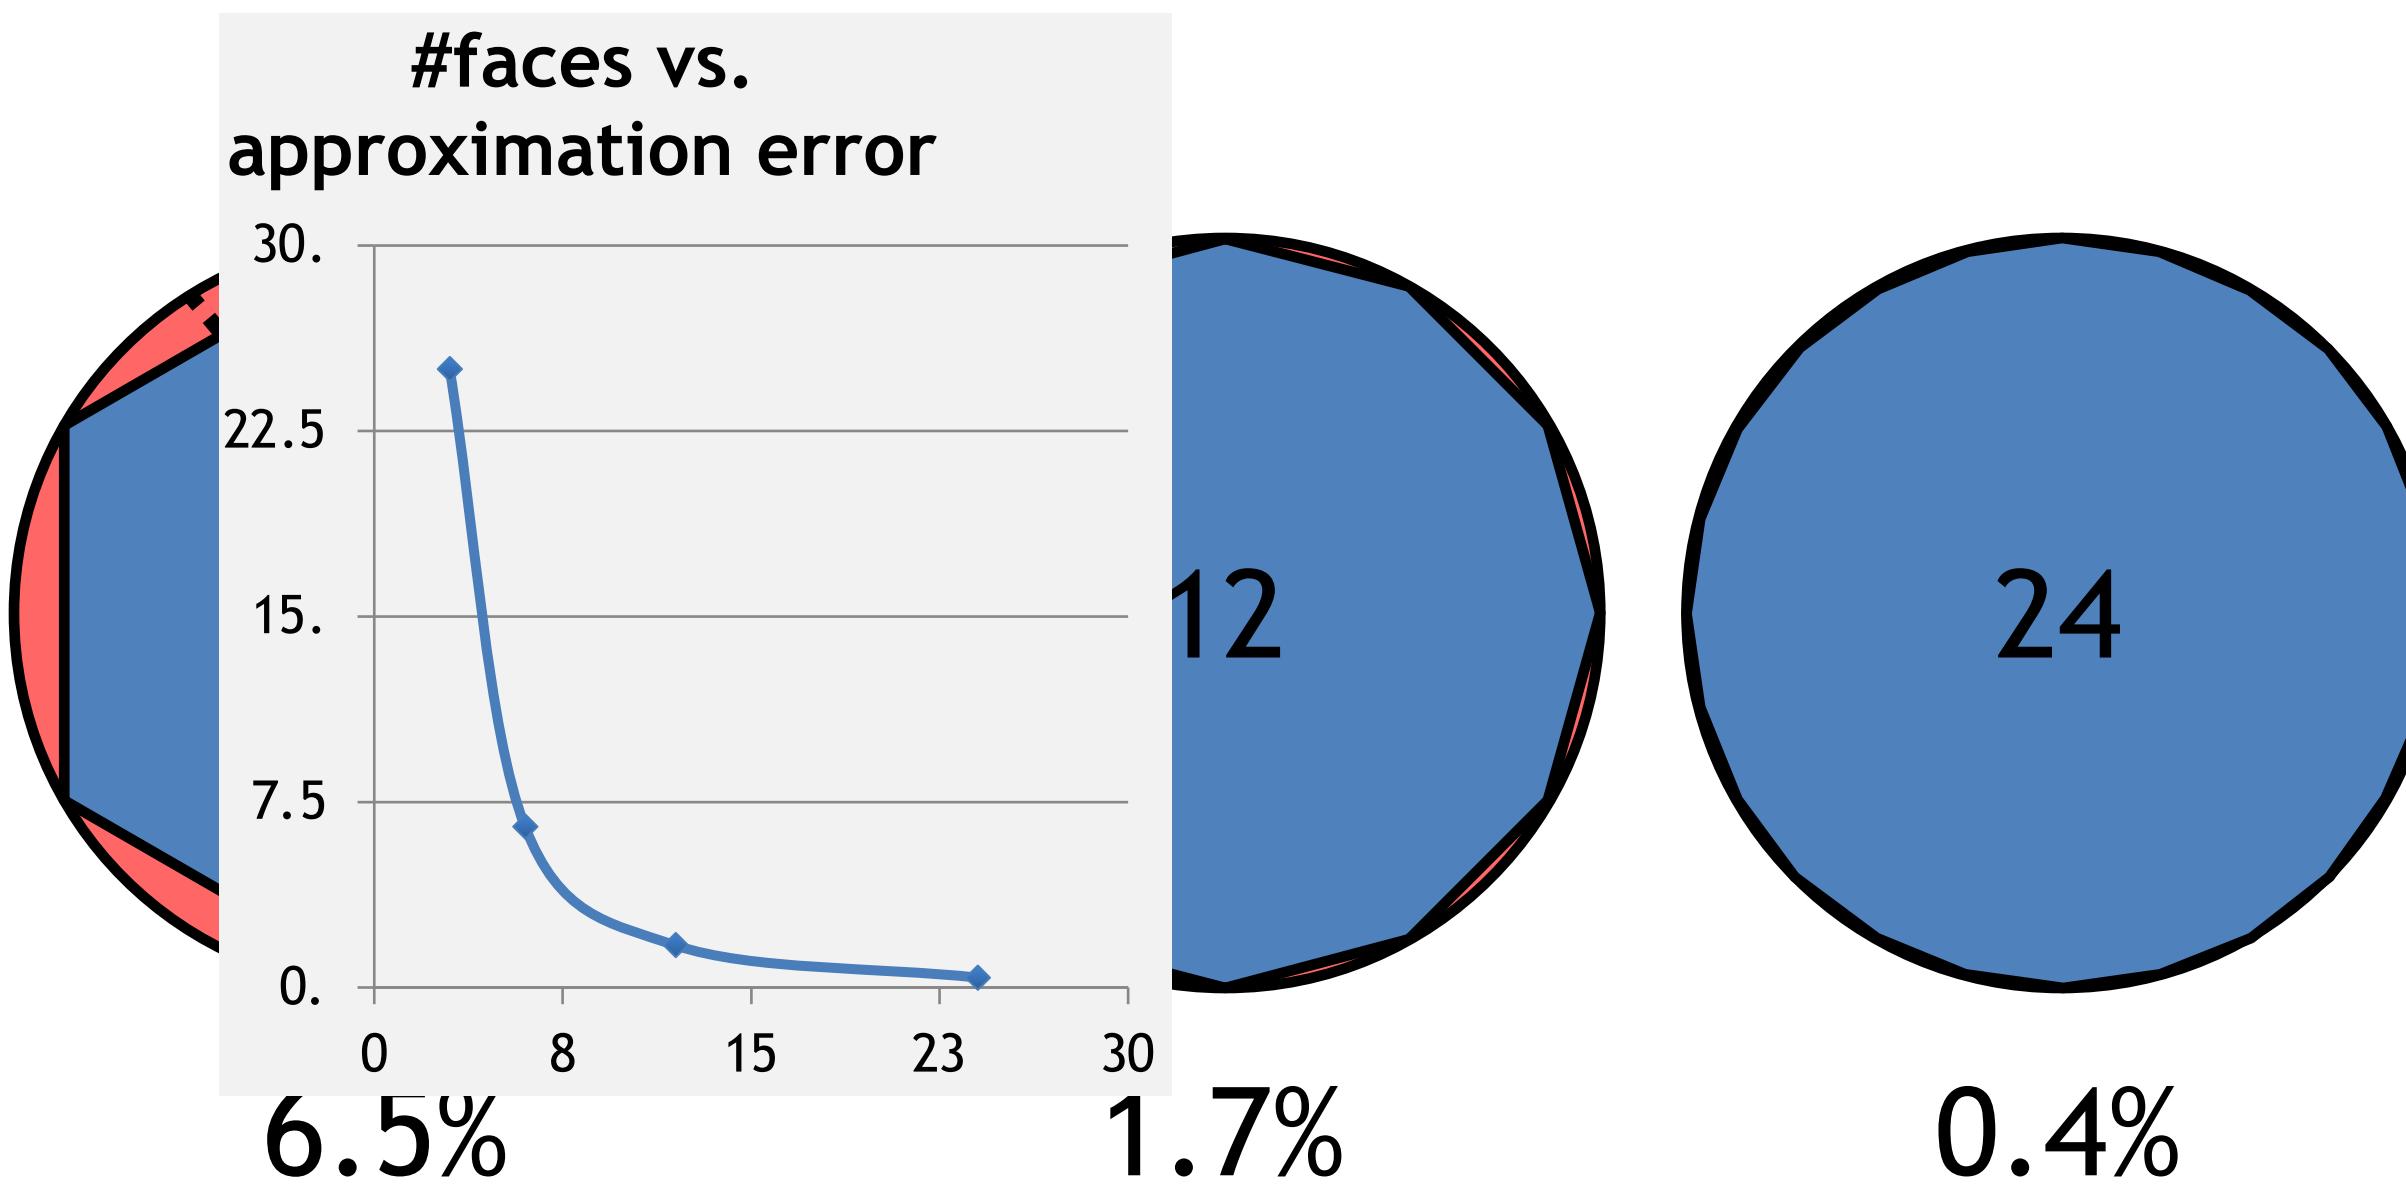
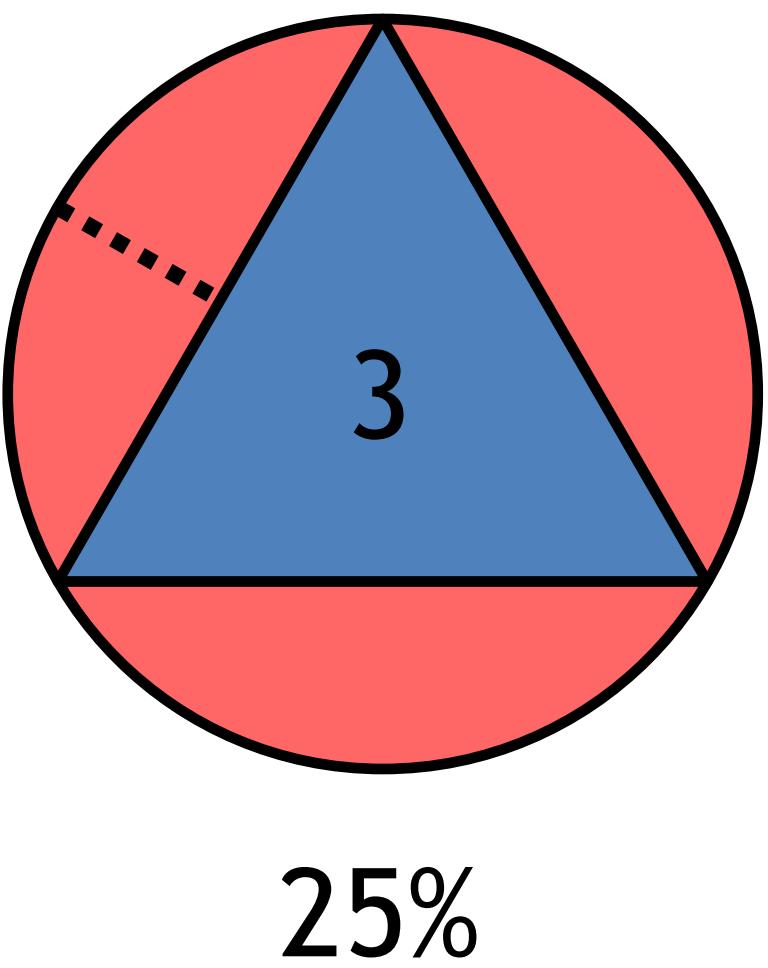
1.7%



0.4%

Meshes as Approximations of Smooth Surfaces

- Piecewise linear approximation



Polygonal Meshes

- Polygonal meshes are a good representation
 - Approximation error drops fast
 - arbitrary topology
 - piecewise smooth surfaces
 - adaptive refinement
 - efficient rendering

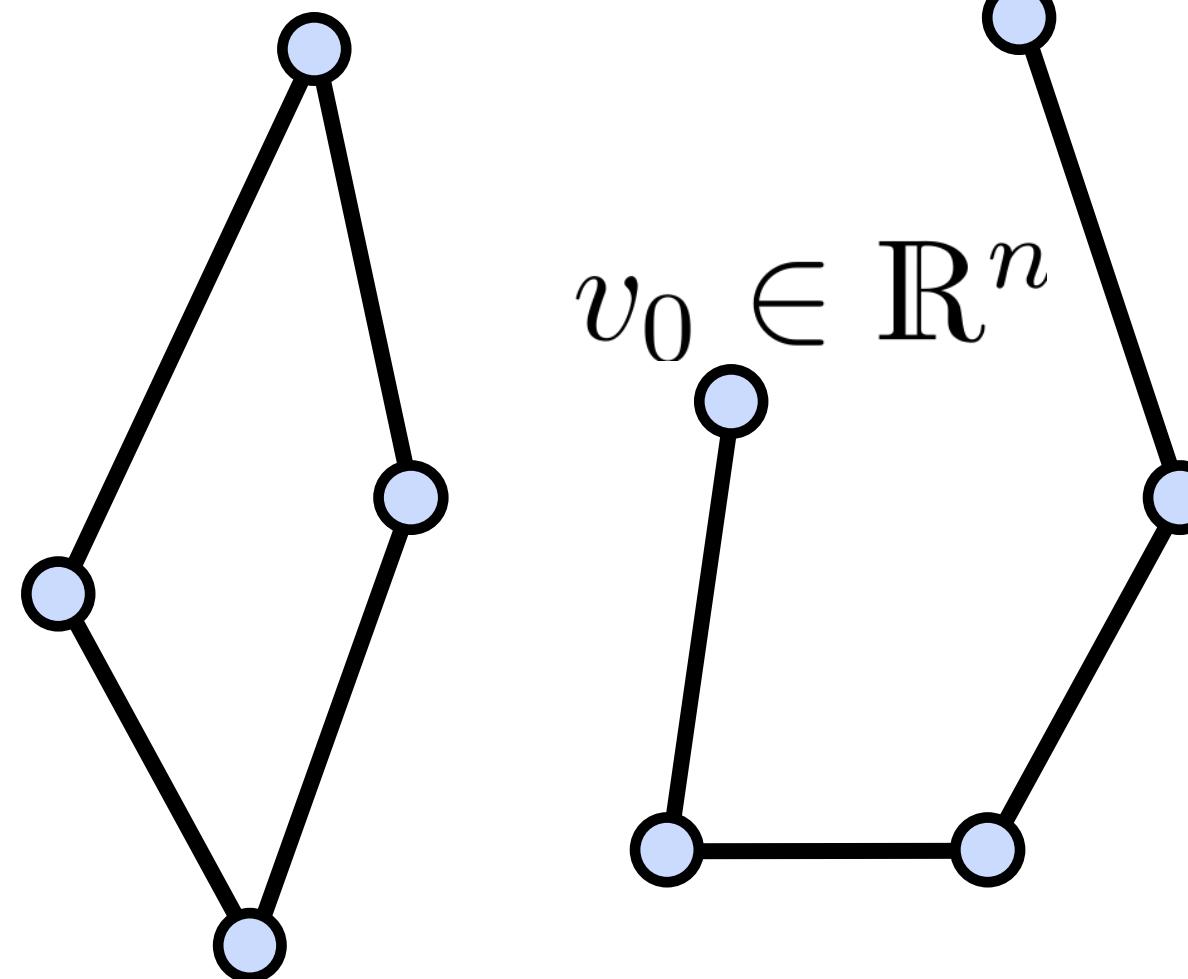


Polygon

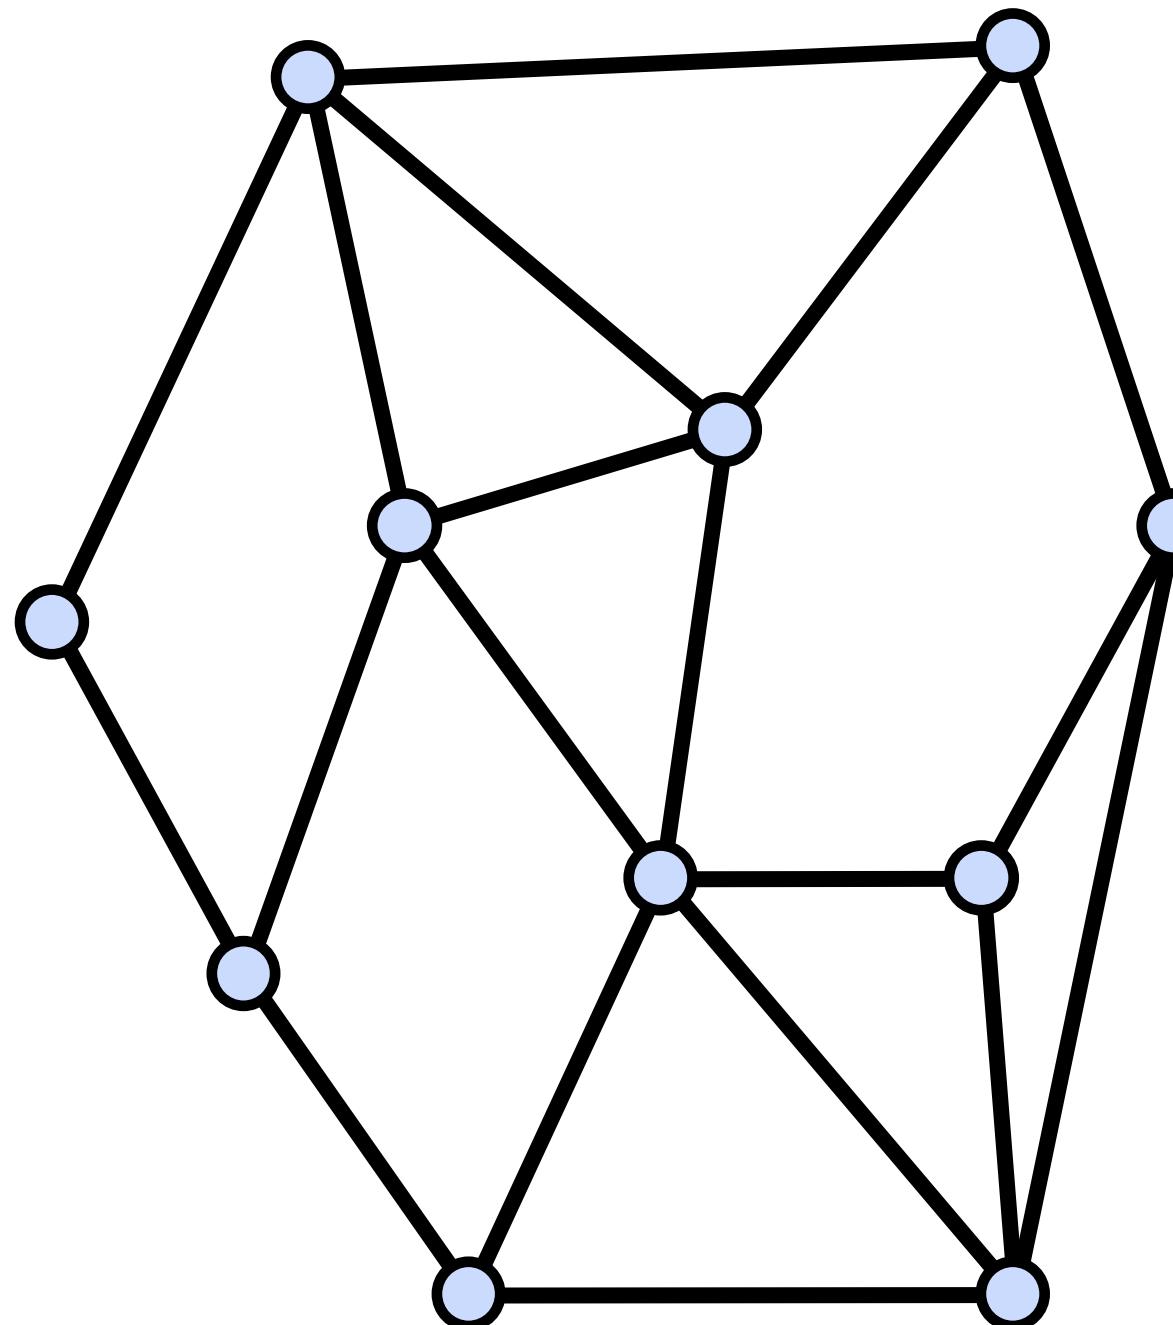
- Vertices:
- Edges:

$$v_0, v_1, \dots, v_{n-1}$$
$$\{(v_0, v_1), \dots, (v_{n-2}, v_{n-1})\}$$

- Closed:
- Planar: all vertices on a plane
- Simple: not self-intersecting

$$v_0 = v_{n-1}$$

$$v_0 \in \mathbb{R}^n$$

Polygonal Mesh

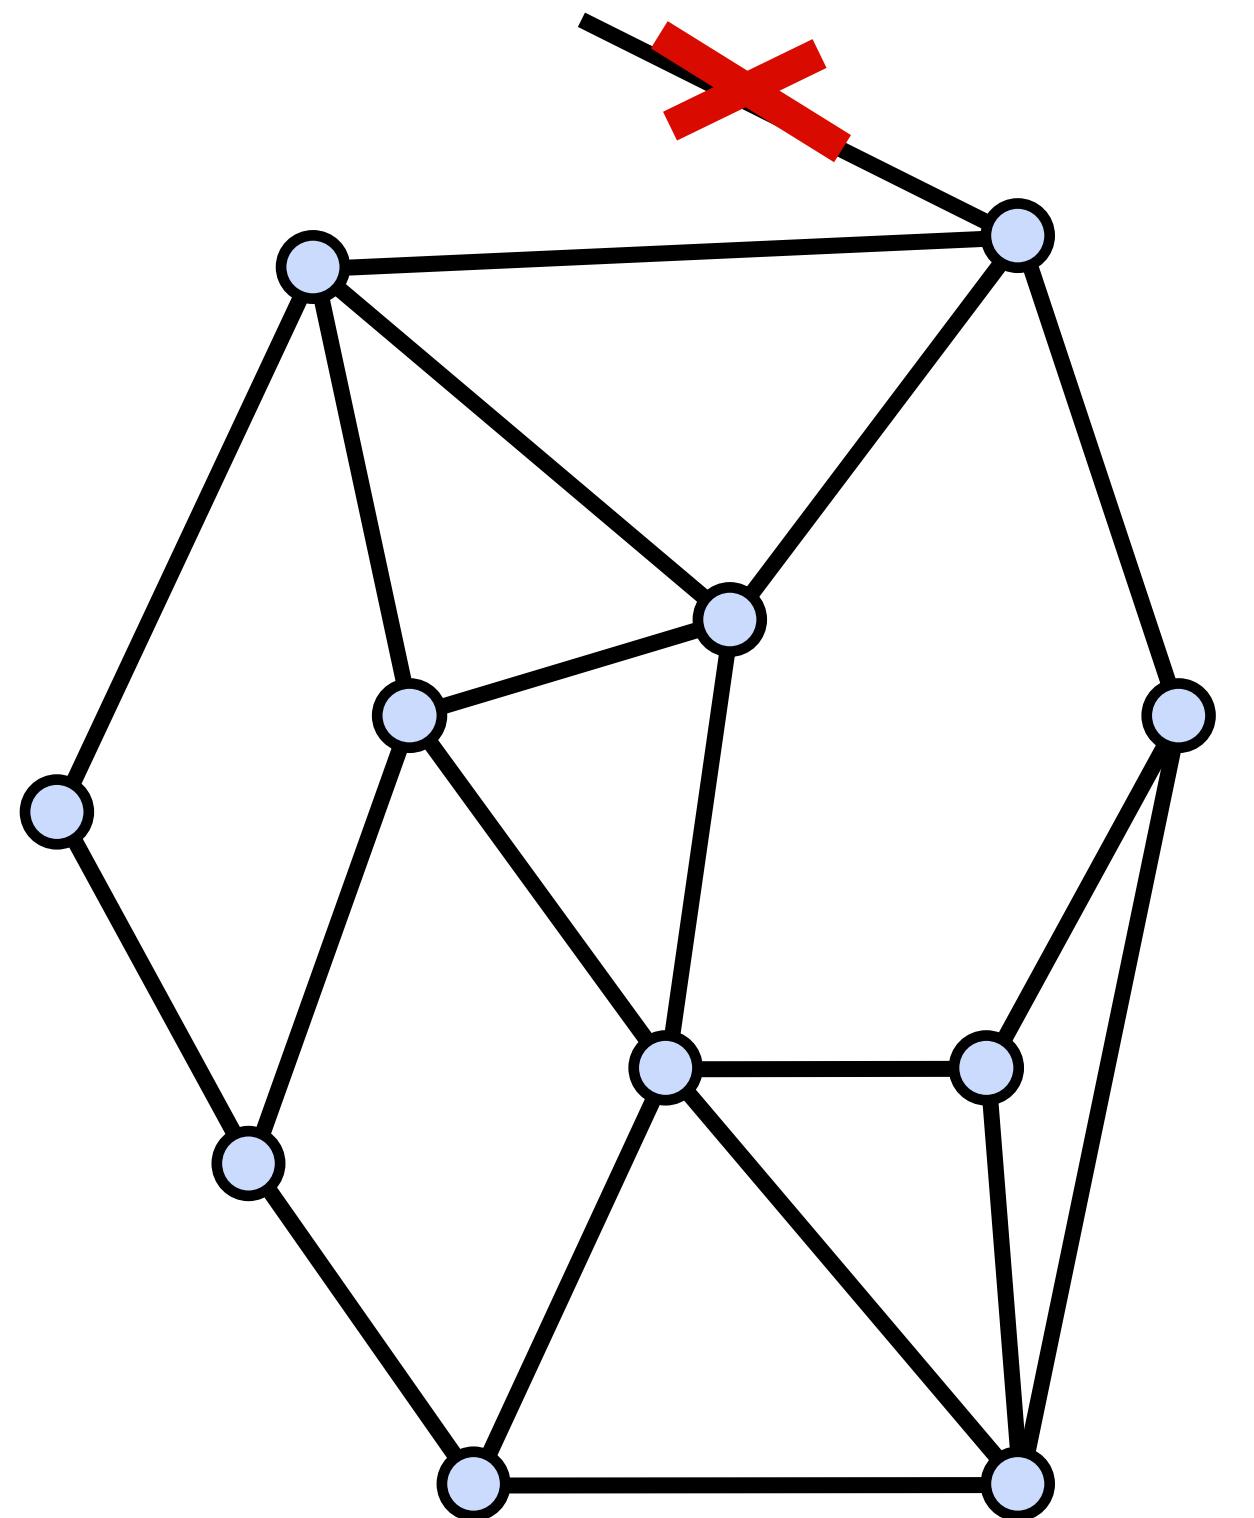


- A finite set M of closed, simple polygons Q_i is a polygonal mesh
- The intersection of two polygons in M is either empty, a vertex, or an edge

$$M = \langle V, E, F \rangle$$

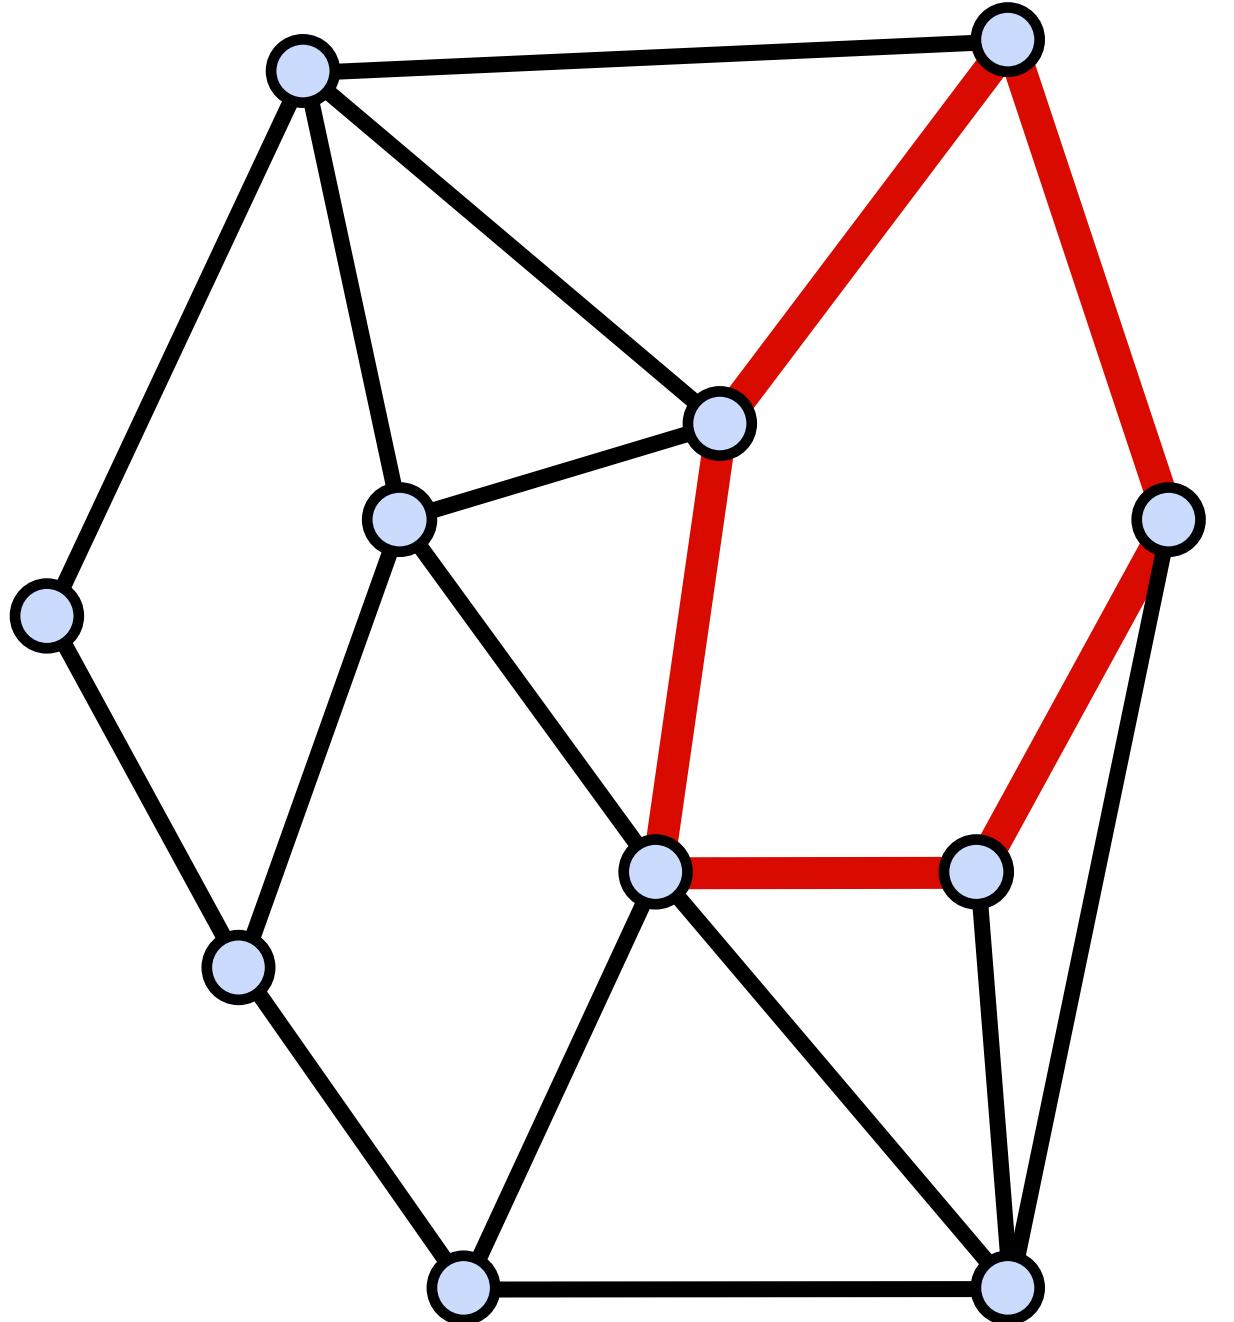
vertices edges faces

Polygonal Mesh



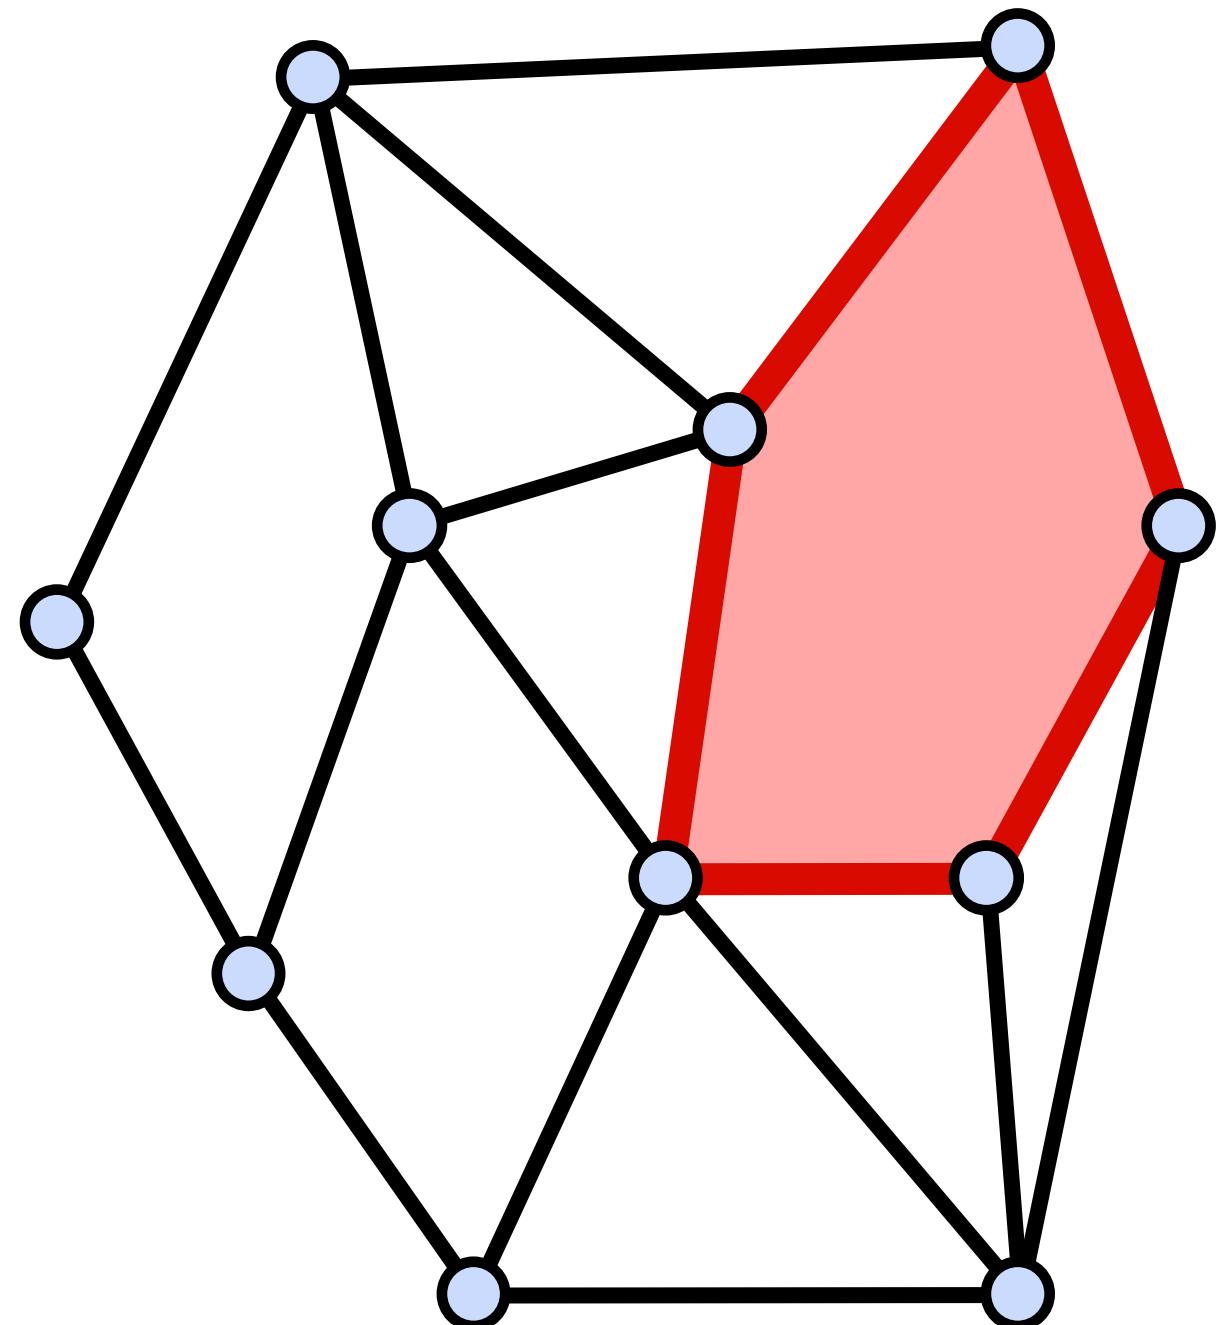
- A finite set M of closed, simple polygons Q_i is a polygonal mesh
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon

Polygonal Mesh



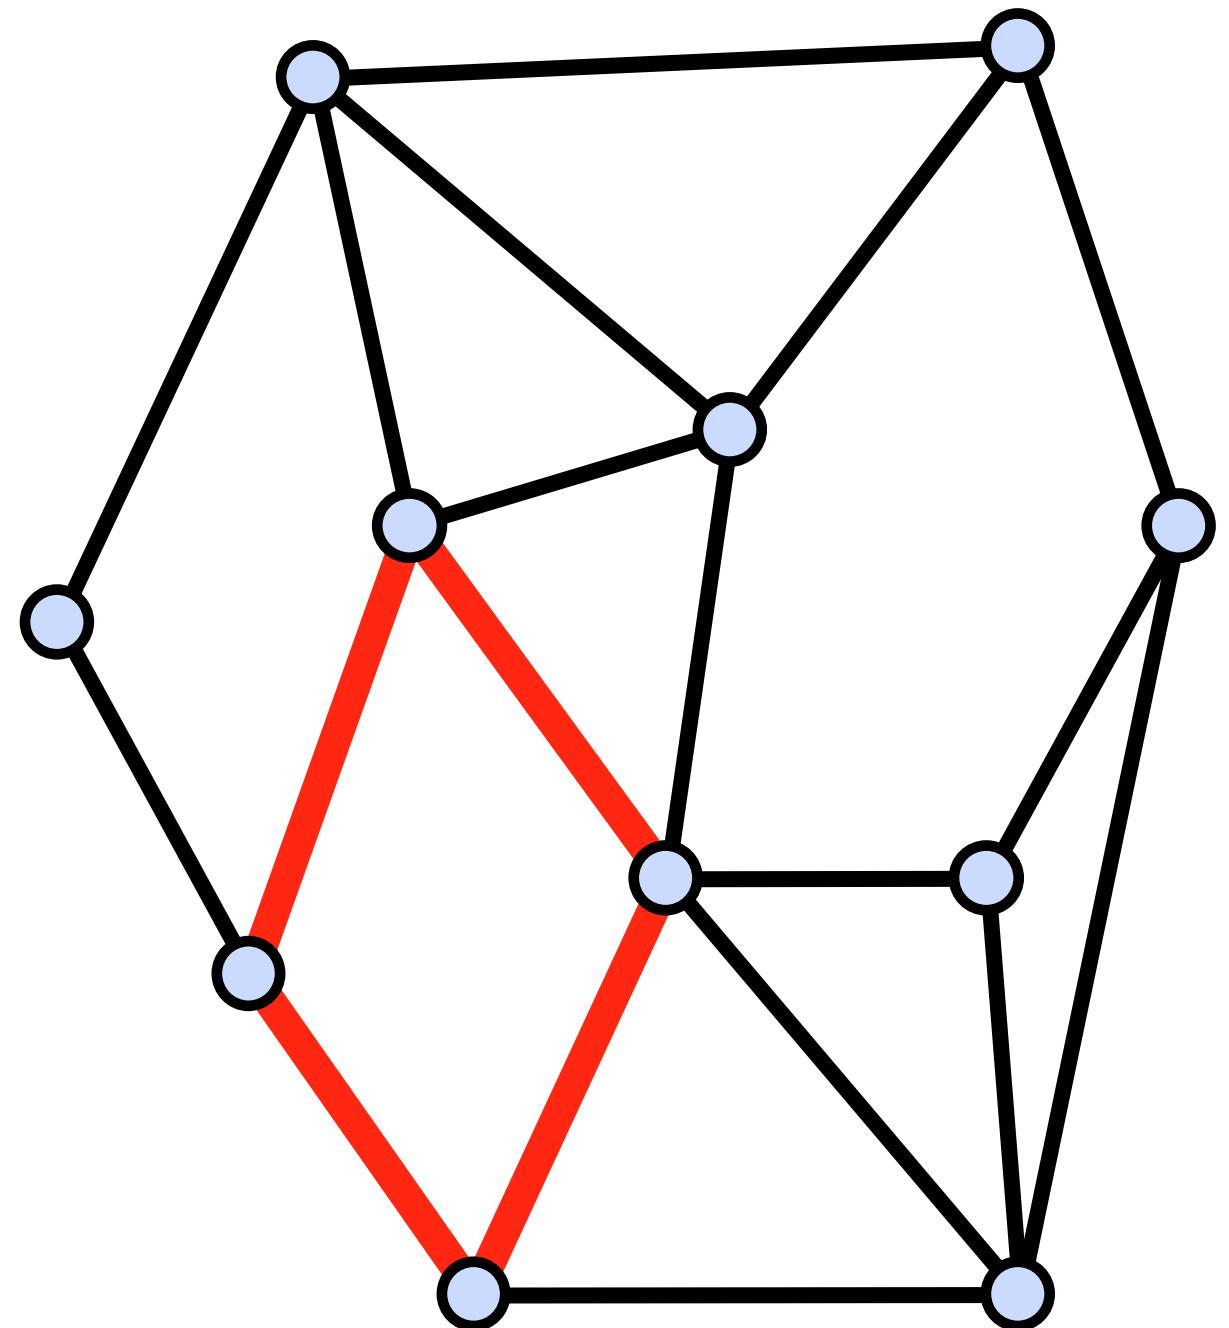
- A finite set M of closed, simple polygons Q_i is a **polygonal mesh**
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each Q_i defines a **face** of the polygonal mesh

Polygonal Mesh



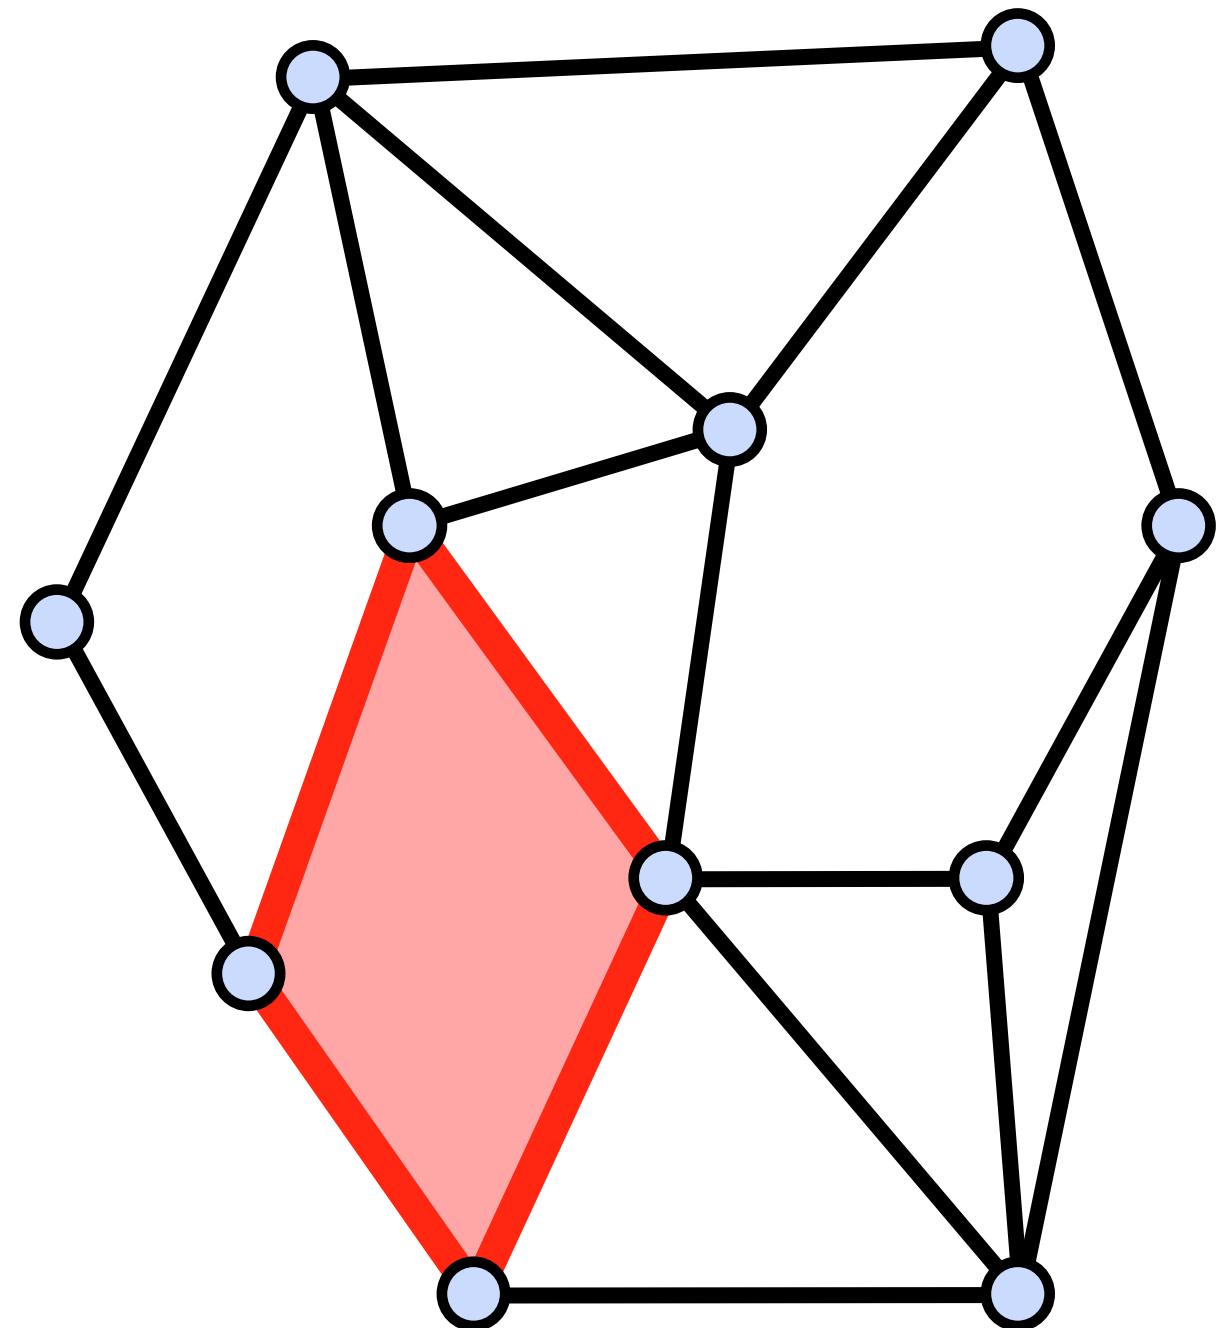
- A finite set M of closed, simple polygons Q_i is a **polygonal mesh**
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each Q_i defines a **face** of the polygonal mesh

Polygonal Mesh



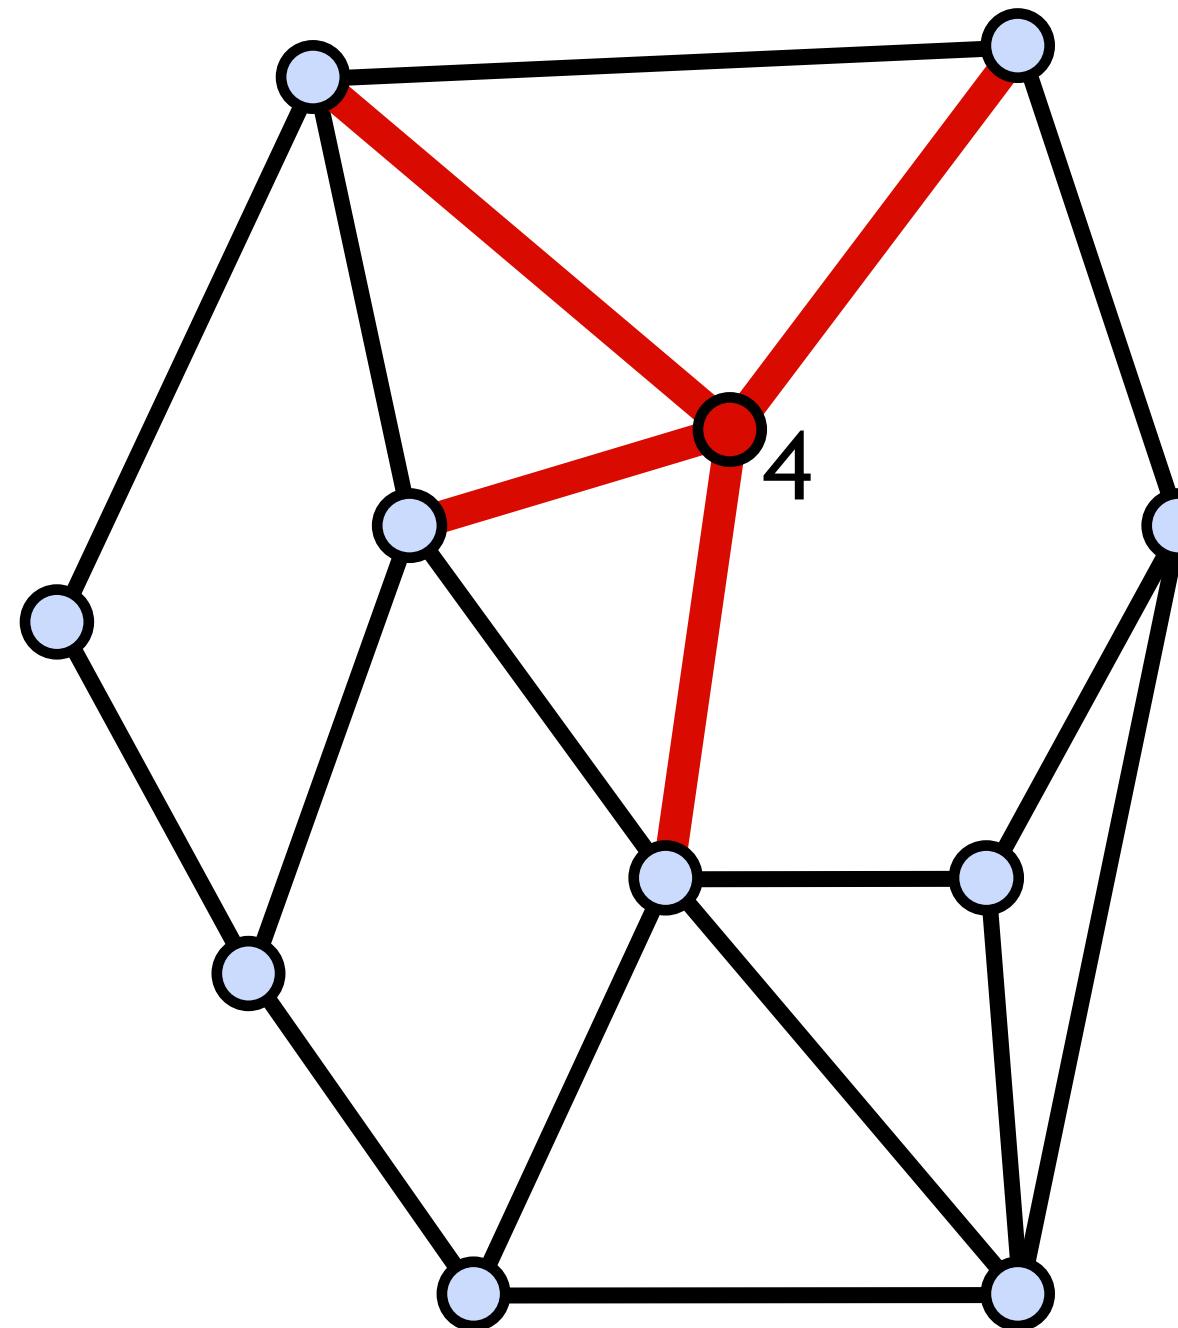
- A finite set M of closed, simple polygons Q_i is a **polygonal mesh**
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each Q_i defines a **face** of the polygonal mesh

Polygonal Mesh



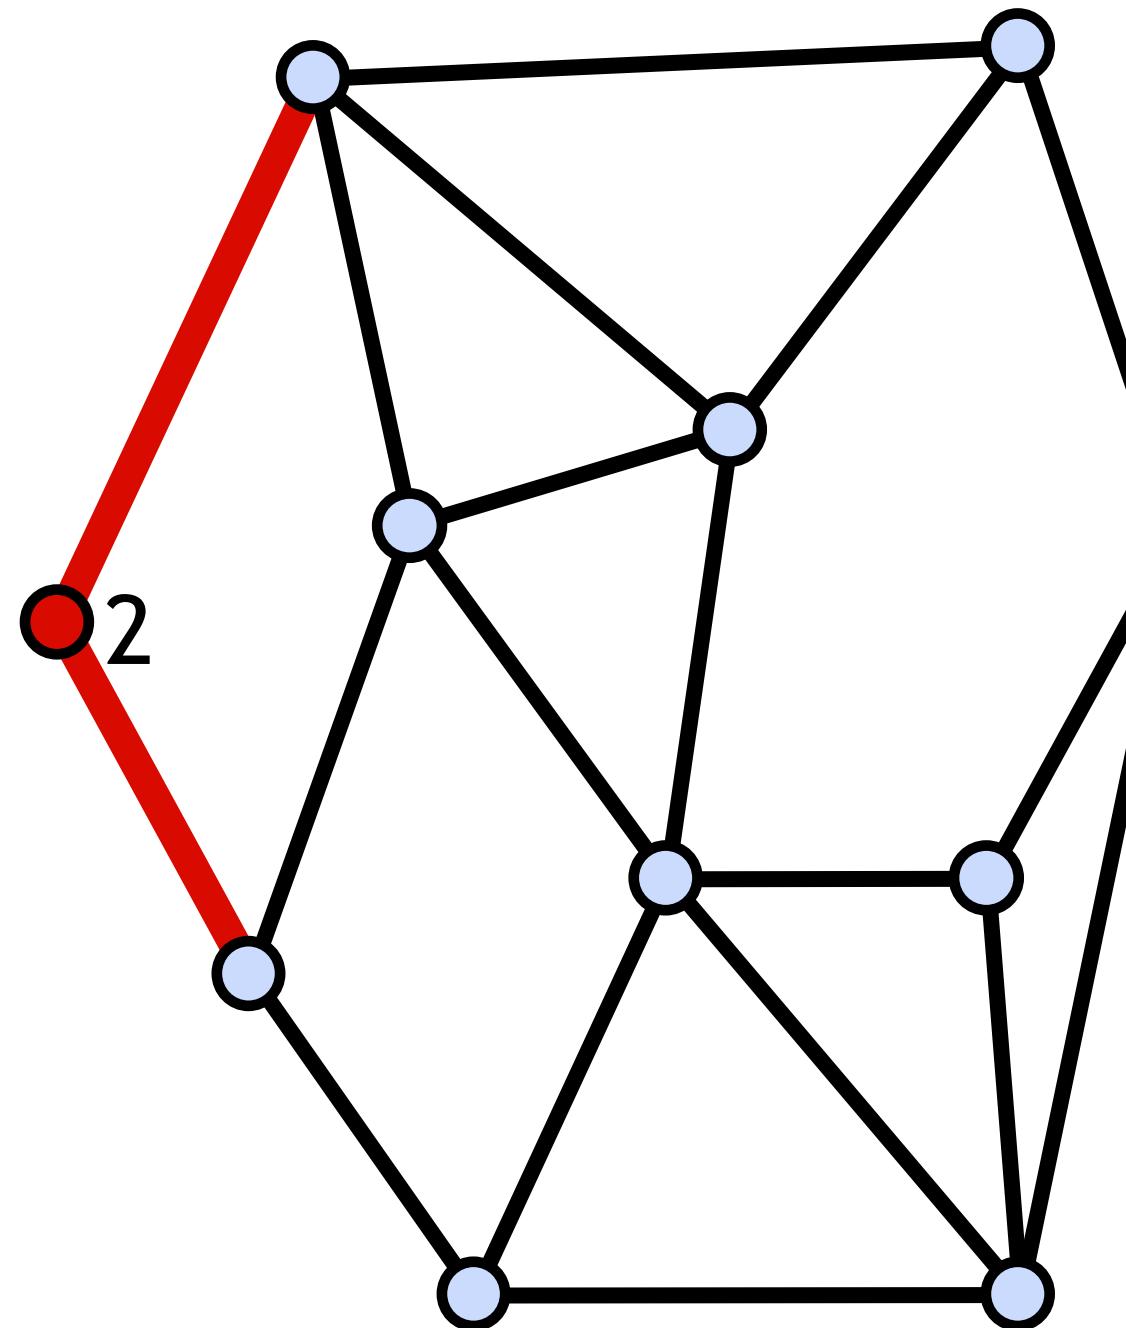
- A finite set M of closed, simple polygons Q_i is a **polygonal mesh**
- The intersection of two polygons in M is either empty, a vertex, or an edge
- Every edge belongs to at least one polygon
- Each Q_i defines a **face** of the polygonal mesh

Polygonal Mesh



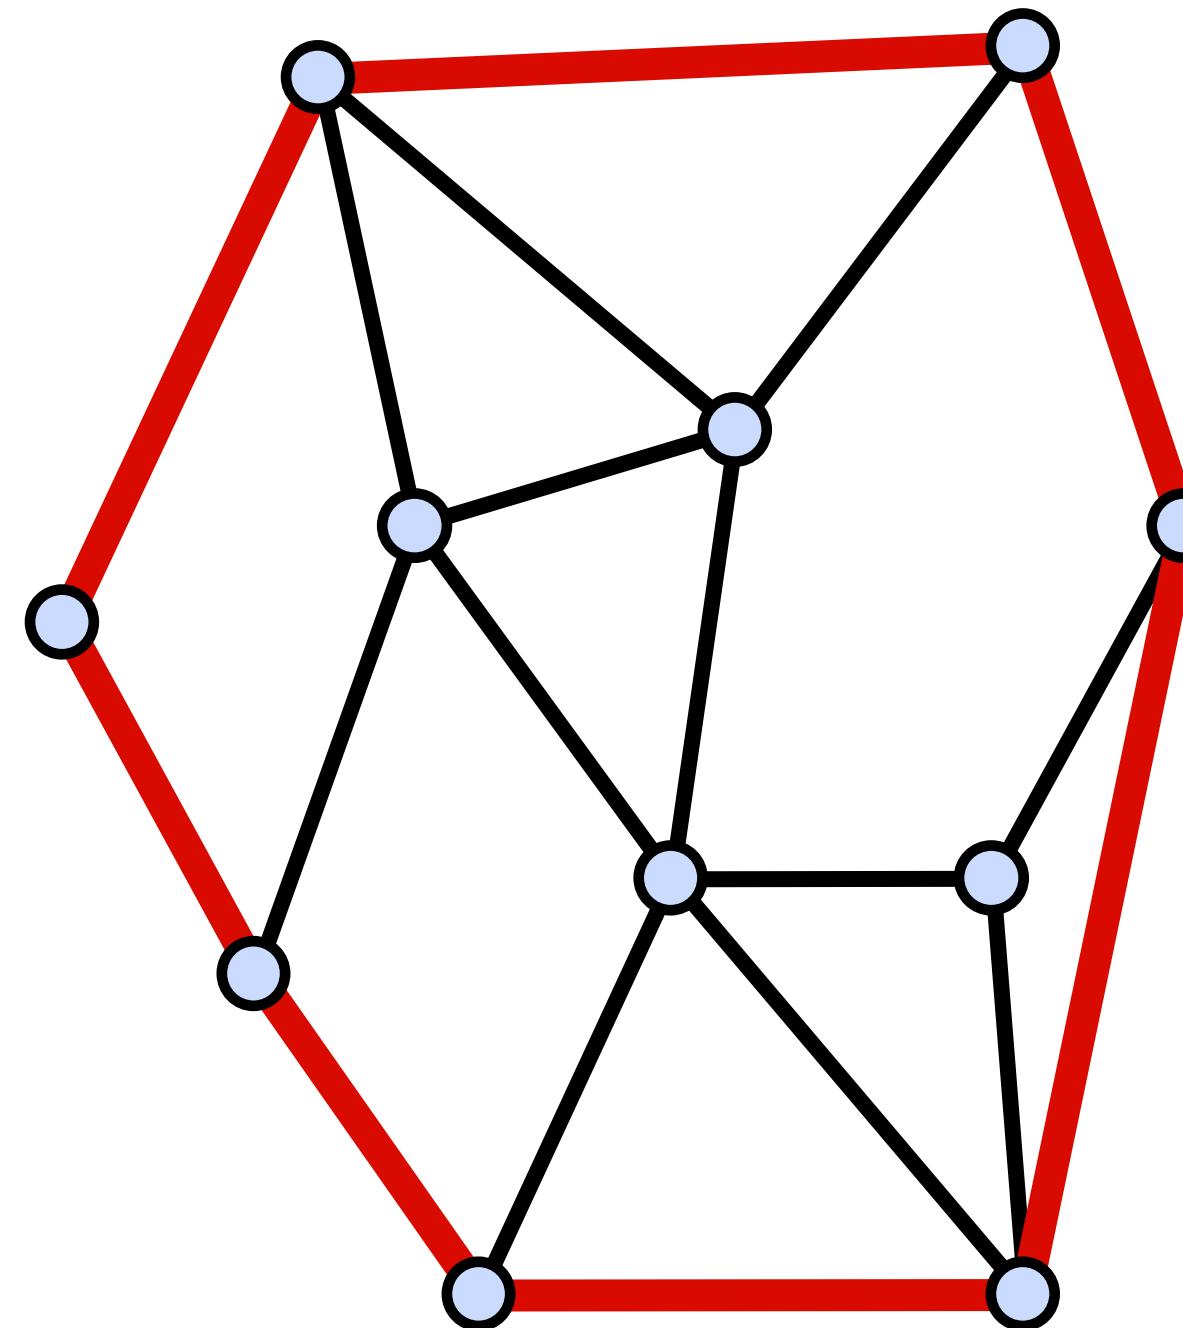
- Vertex **degree** or **valence** = number of incident edges

Polygonal Mesh

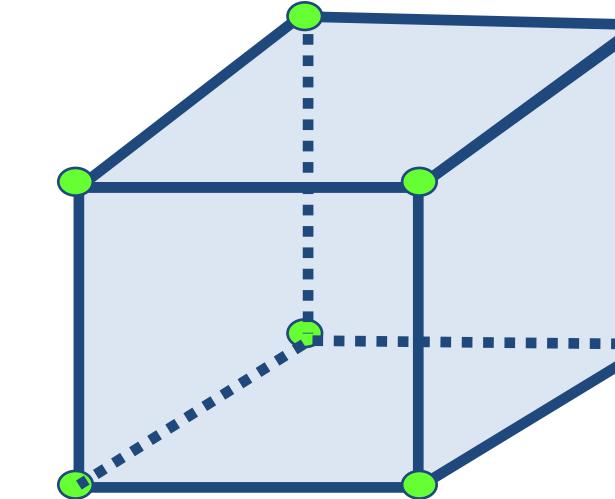


- Vertex **degree** or **valence** = number of incident edges

Polygonal Mesh



- **Boundary:** the set of all edges that belong to only one polygon
 - Either empty or forms closed loops
 - If empty, then the polygonal mesh is closed



Triangle Meshes

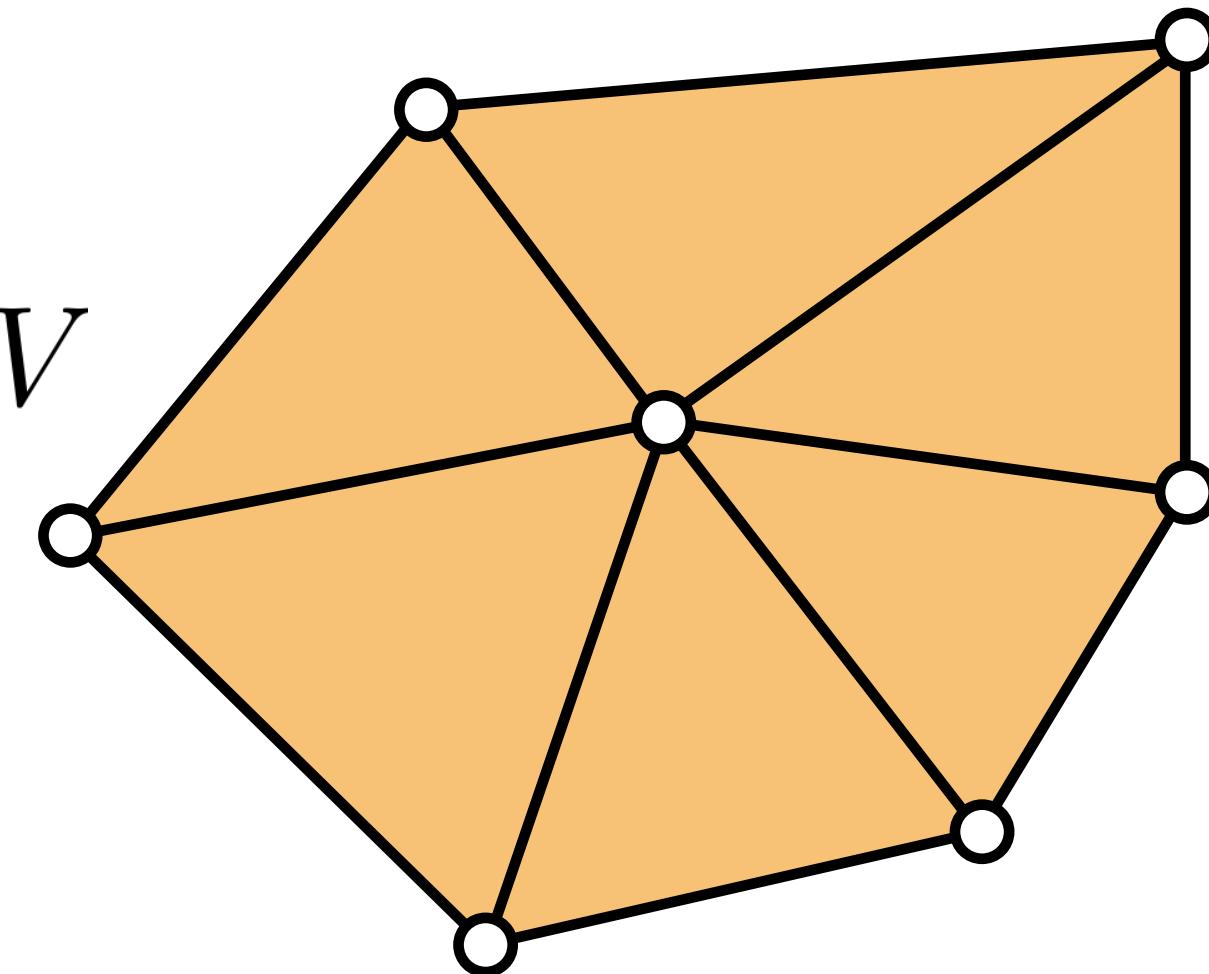
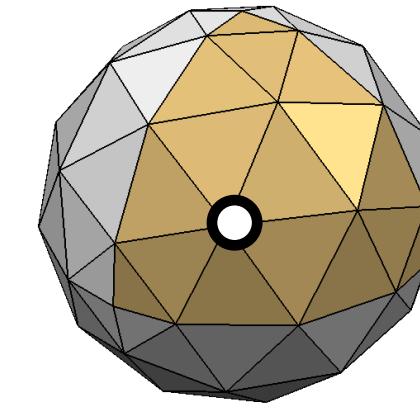
- Connectivity: vertices, edges, triangles
- Geometry: vertex positions

$$V = \{v_1, \dots, v_n\}$$

$$E = \{e_1, \dots, e_k\}, \quad e_i \in V \times V$$

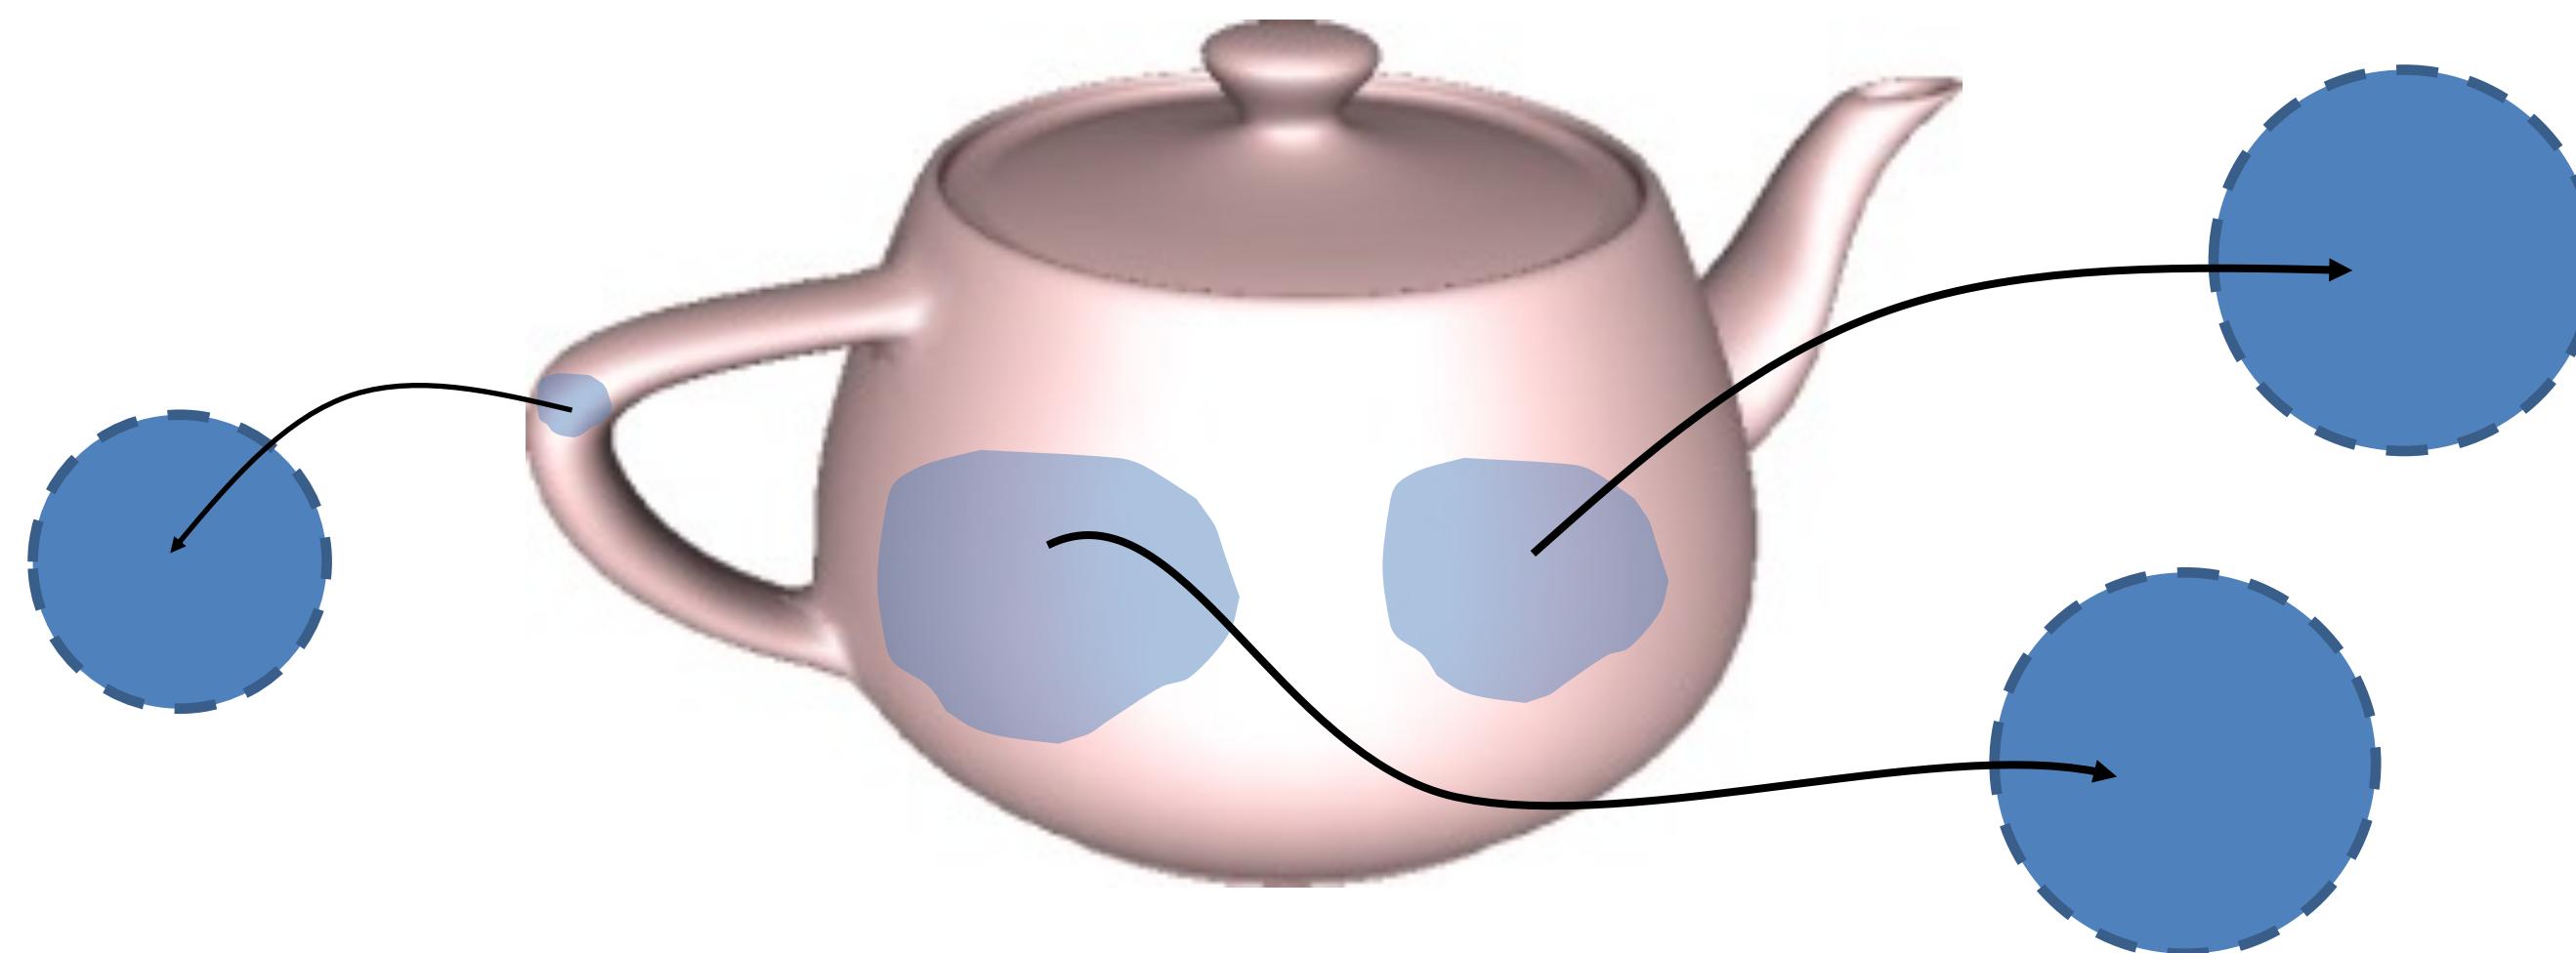
$$F = \{f_1, \dots, f_m\}, \quad f_i \in V \times V \times V$$

$$P = \{\mathbf{p}_1, \dots, \mathbf{p}_n\}, \quad \mathbf{p}_i \in \mathbb{R}^3$$



Manifolds

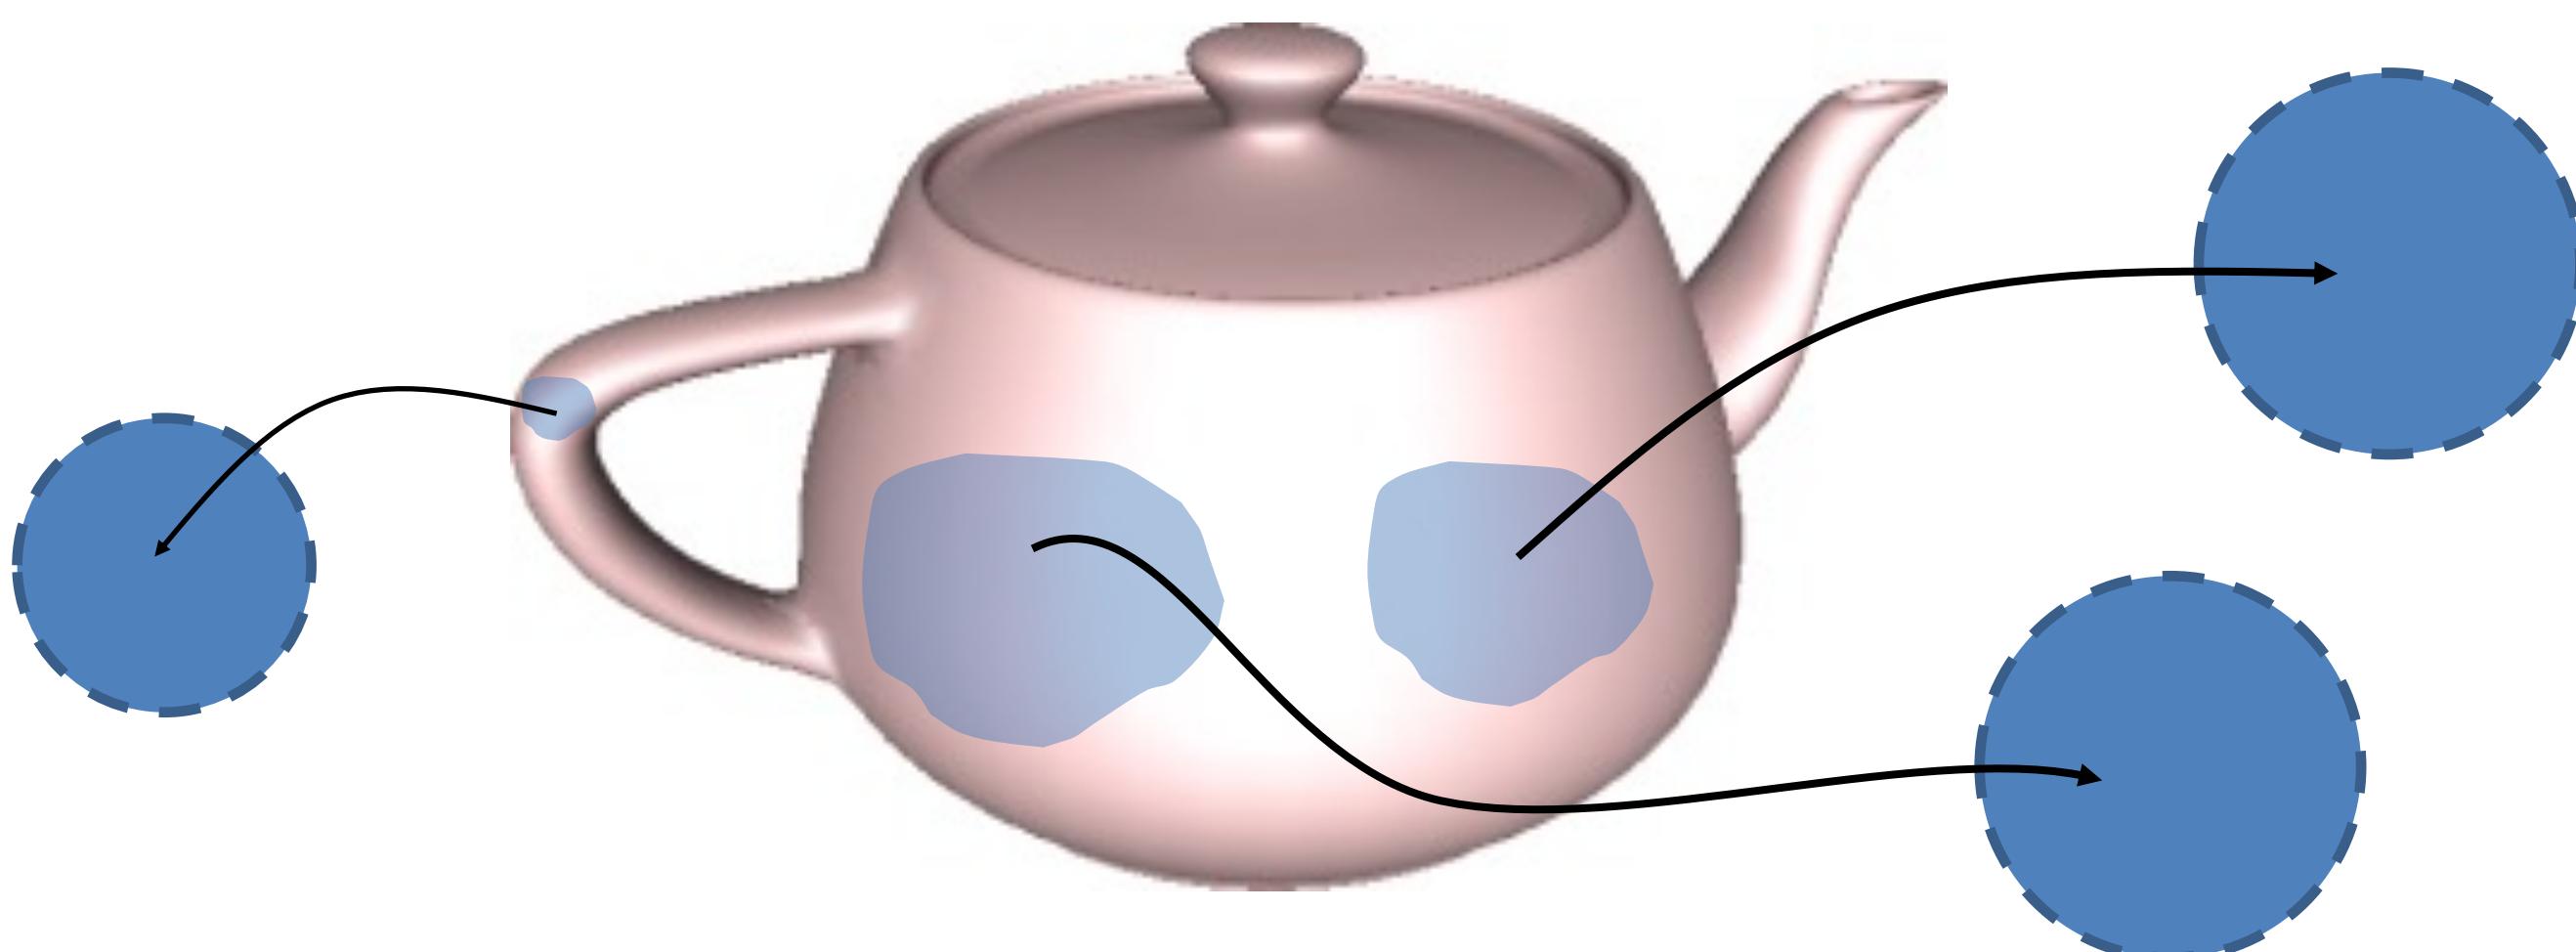
- A surface is a closed **2-manifold** if it is everywhere locally homeomorphic to a disk



Manifolds

- For every point \mathbf{x} in M , there is an **open** ball $B_{\mathbf{x}}(r)$ of radius $r > 0$ centered at \mathbf{x} such that $M \cap B_{\mathbf{x}}$ is homeomorphic to an open disk

$$B_{\mathbf{x}}(r) = \{\mathbf{y} \in \mathbb{R}^3 \text{ s.t. } \|\mathbf{y} - \mathbf{x}\| < r\}$$



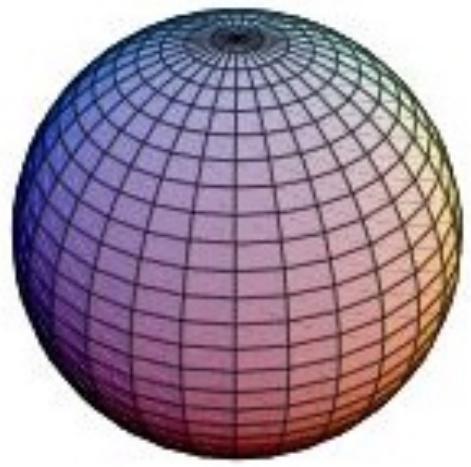
Manifolds

- Manifold with boundary: a vicinity of each boundary point is homeomorphic to a half-disk

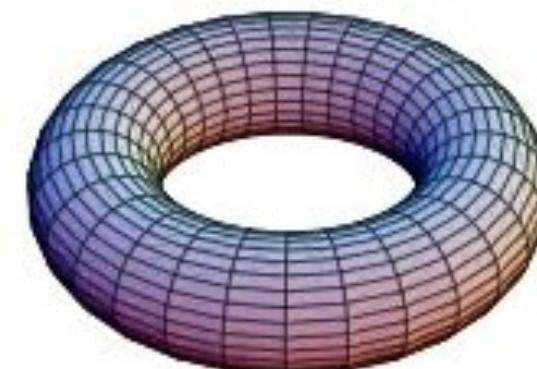


Examples

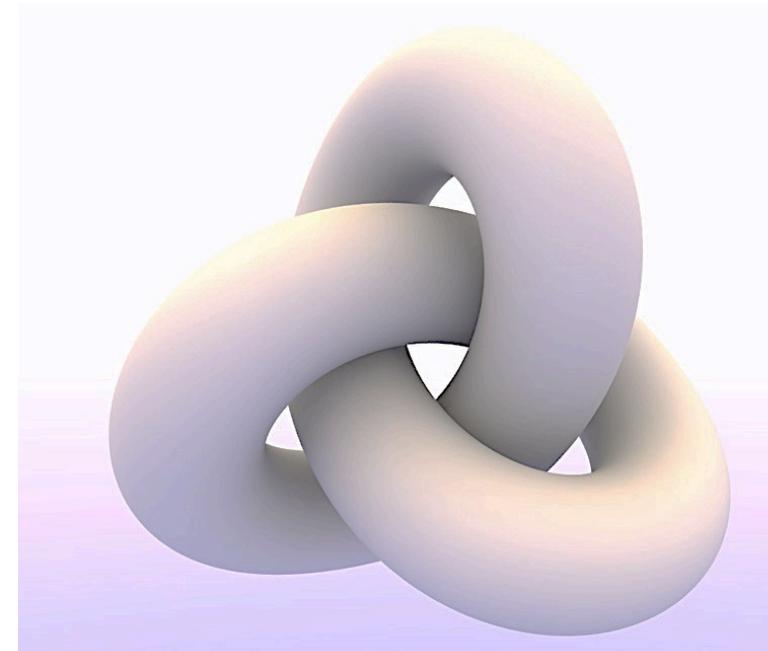
- For each case, decide if it is a 2-manifold (possibly with boundary) or not. If not, explain why not.



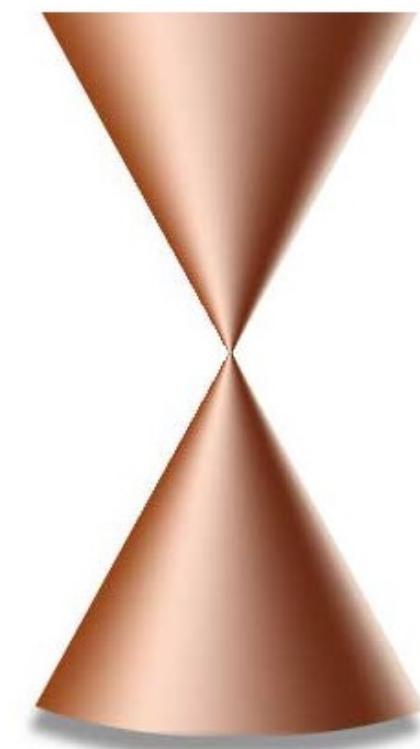
Case 1



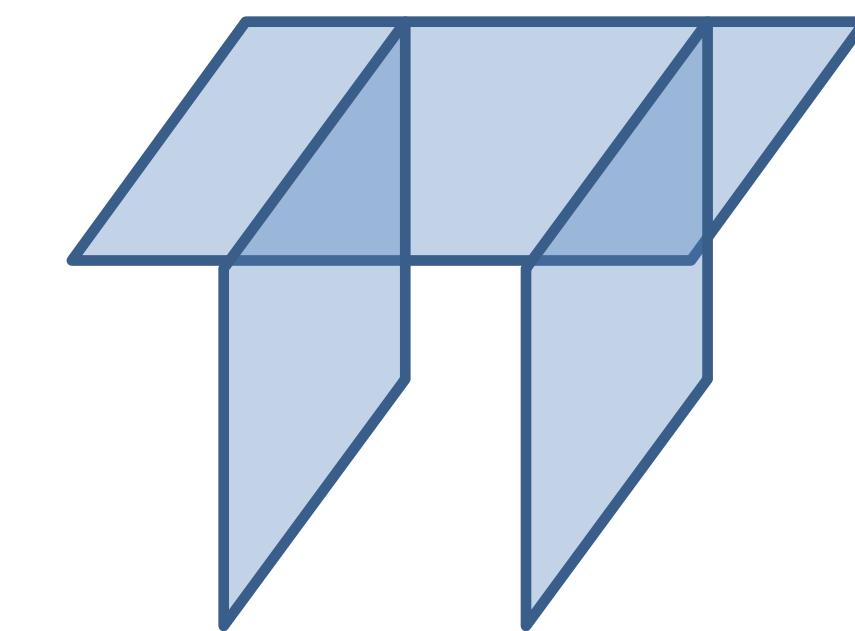
Case 2



Case 3



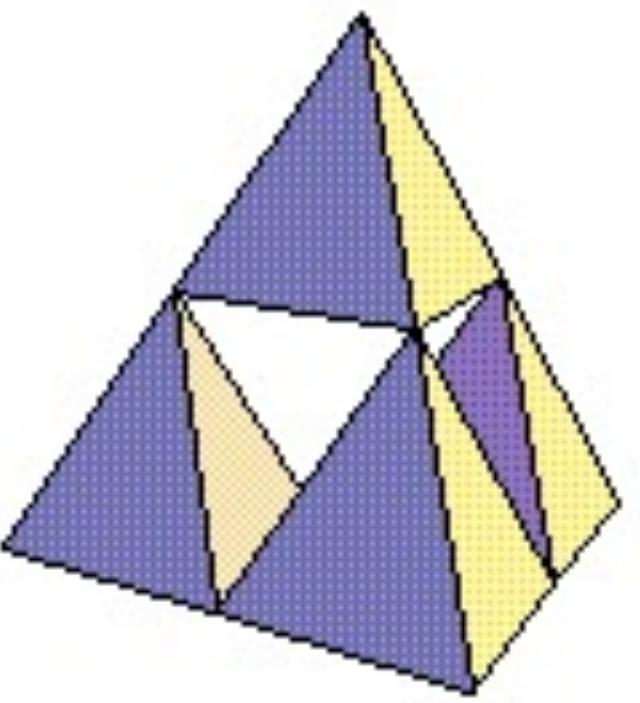
Case 4



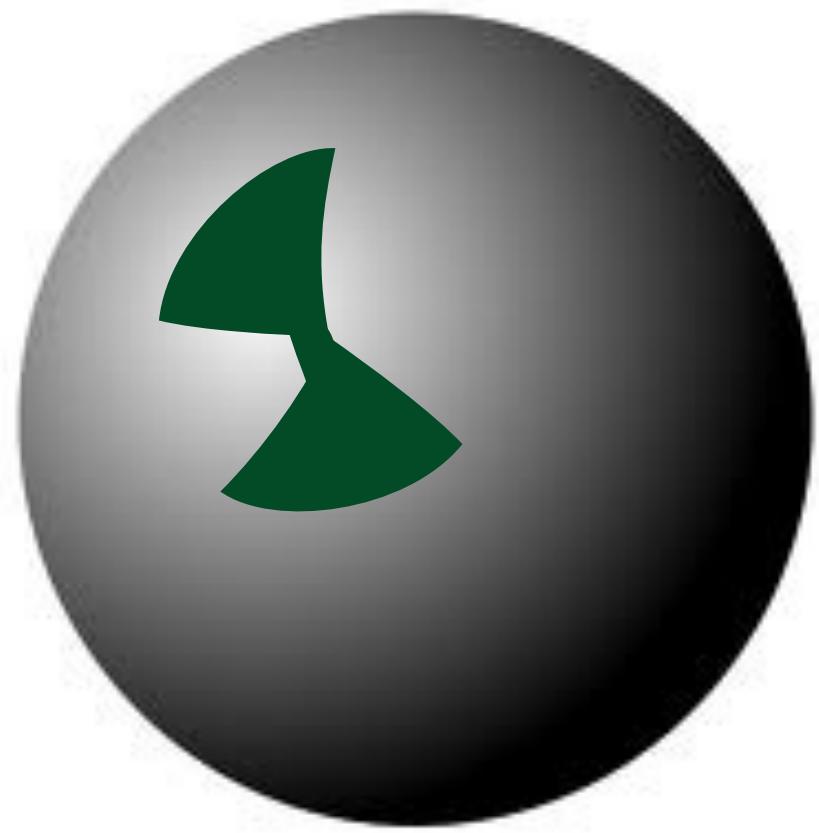
Case 5

Examples

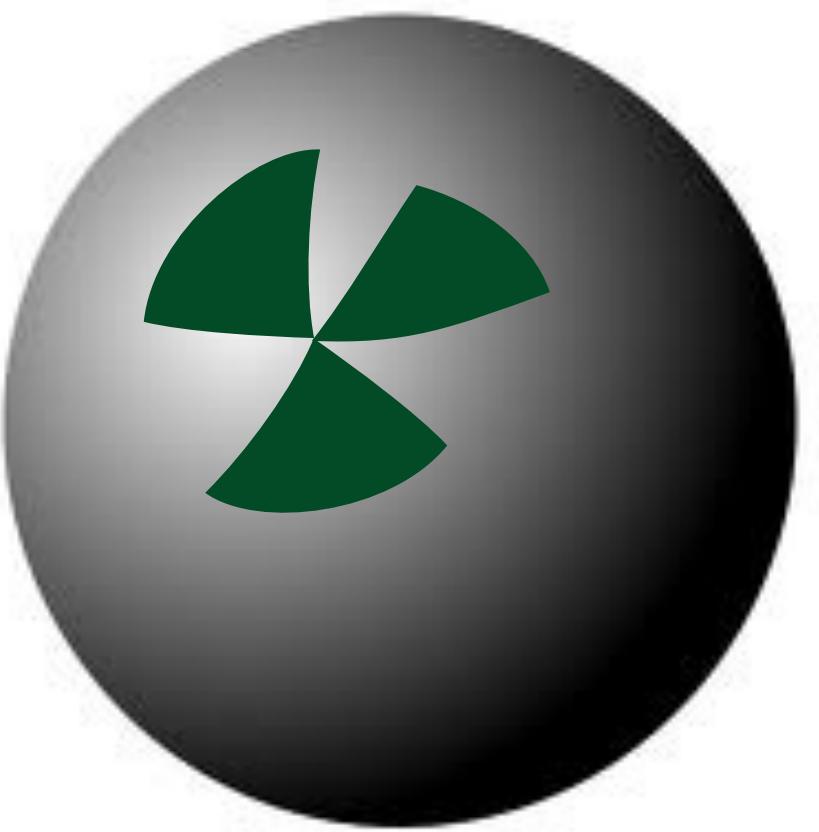
- Bonus cases



Case 6



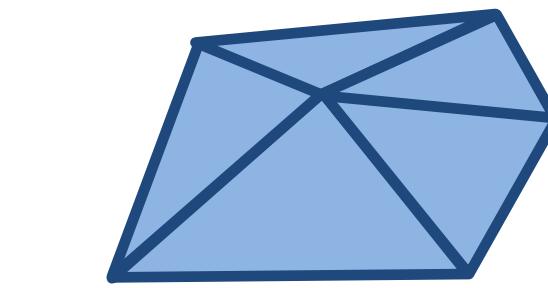
Case 7



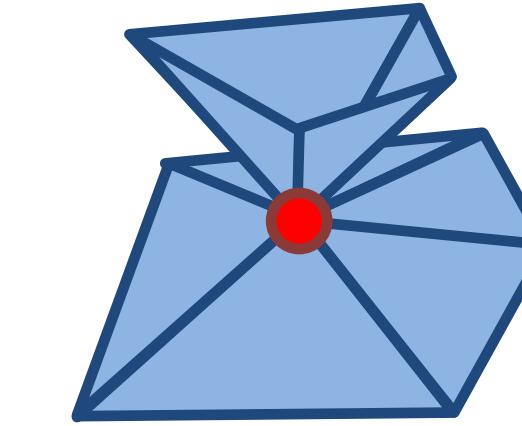
Case 8

Manifolds

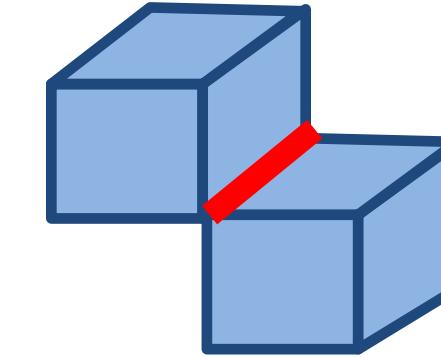
- In a manifold mesh, there are at most 2 faces sharing an edge
 - Boundary edges: have one incident face
 - Inner edges have two incident faces
- A manifold vertex has 1 connected ring of faces around it, or 1 connected half-ring (boundary)



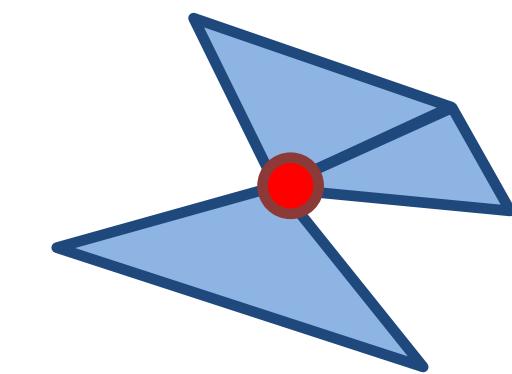
✓ manifold



✗ non-manifold vertex



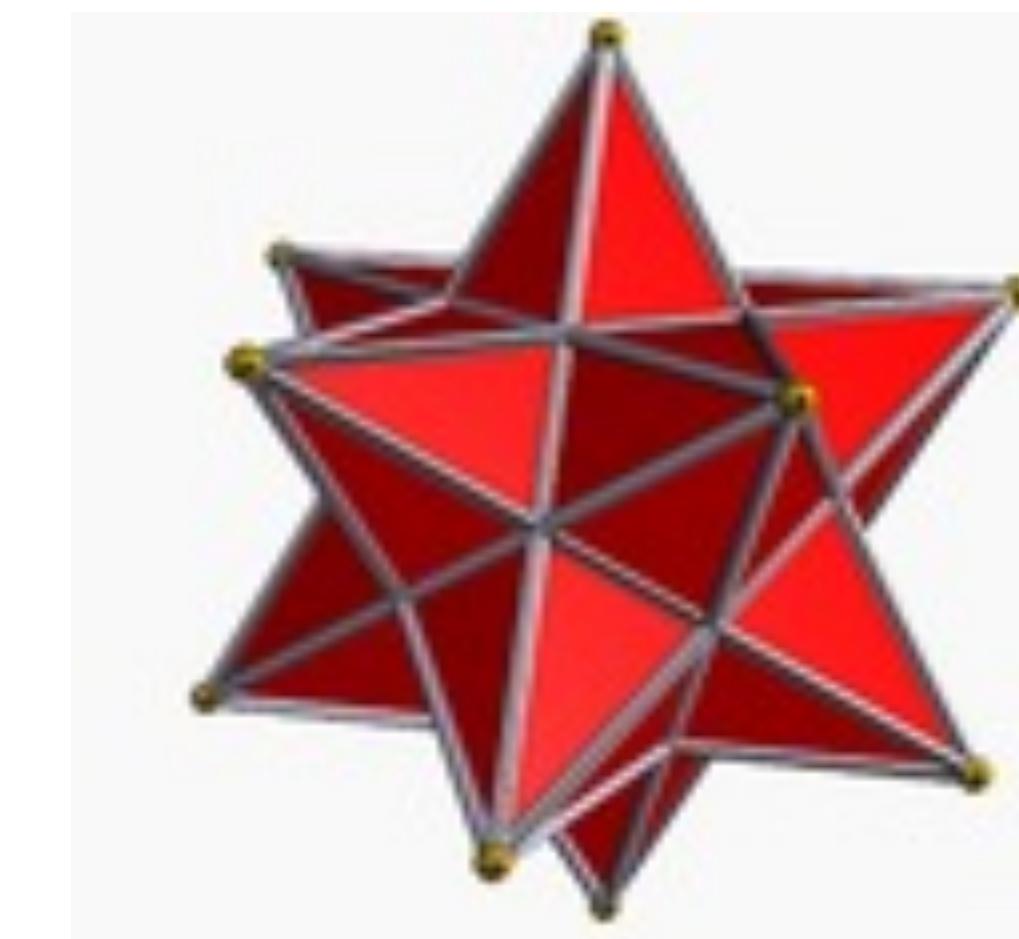
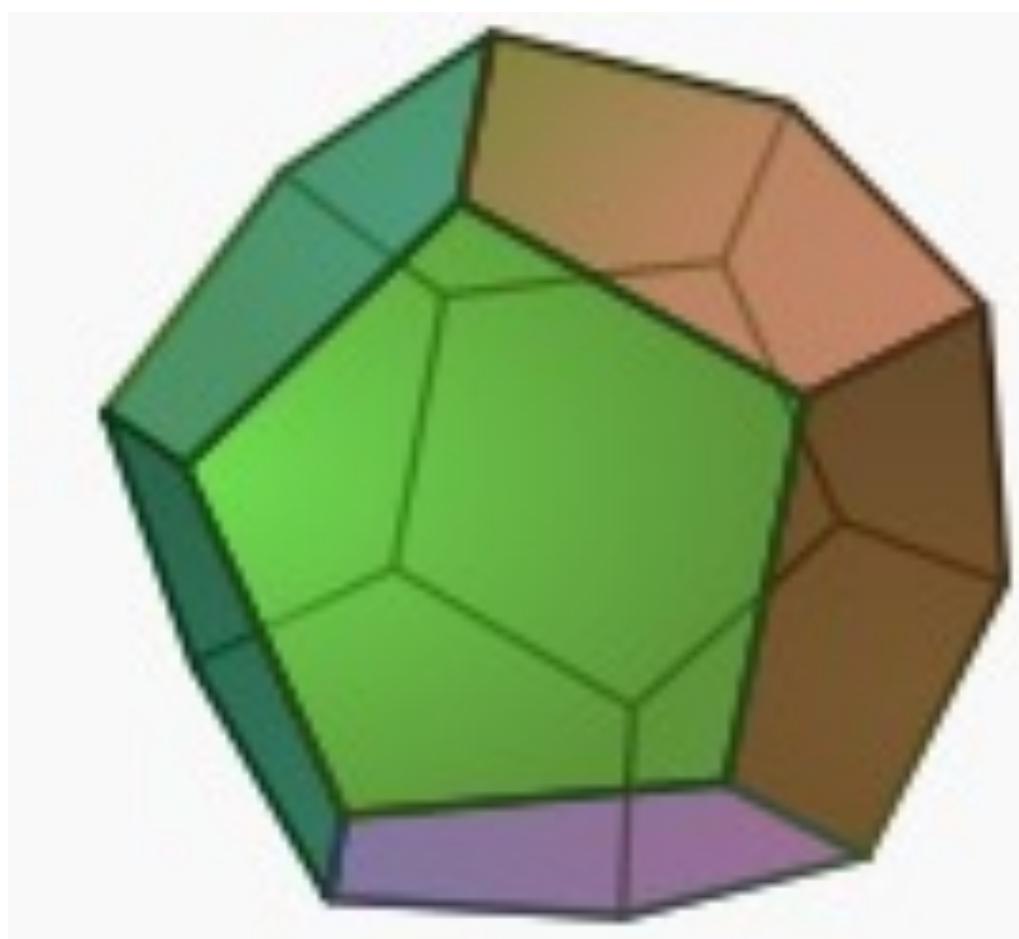
✗ non-manifold edge



✗ non-manifold vertex

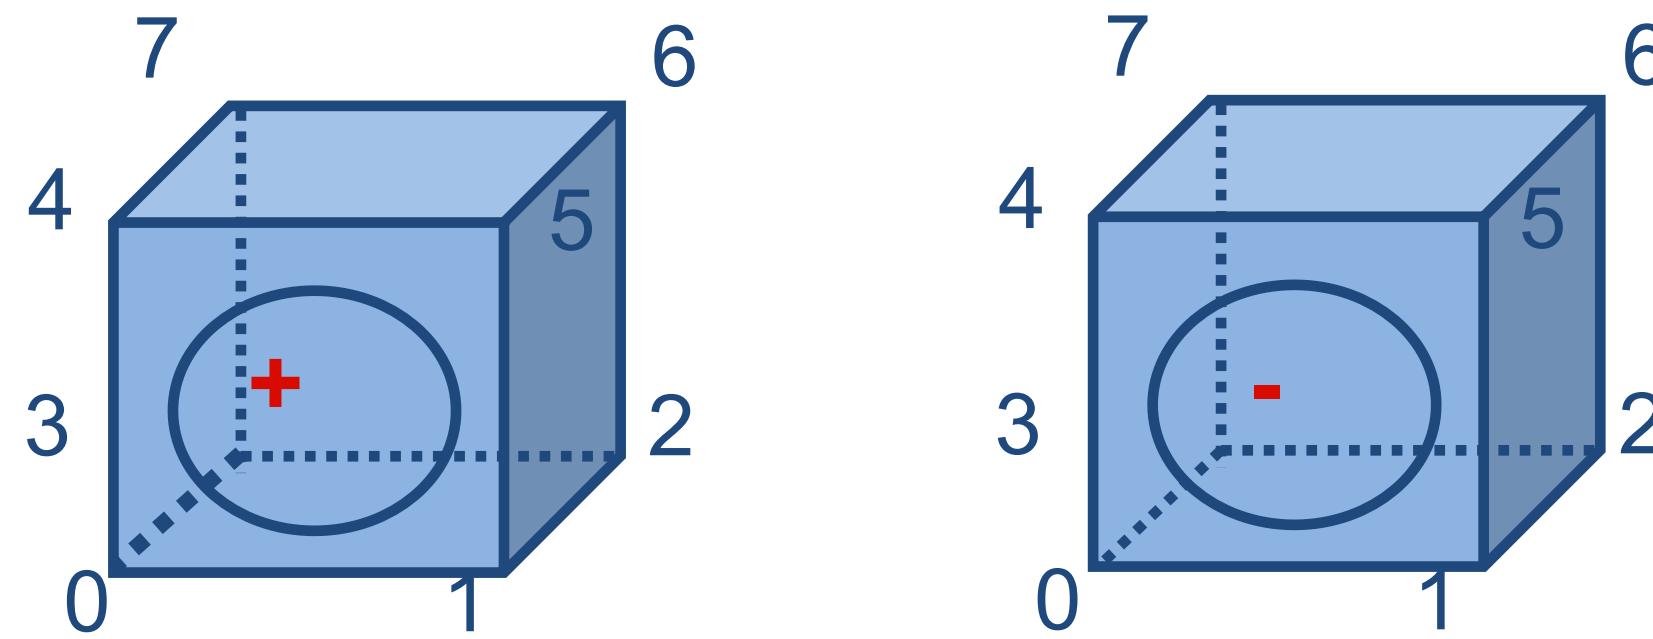
Manifolds

- If closed and not intersecting, a manifold divides the space into inside and outside
- A closed manifold polygonal mesh is called **polyhedron**



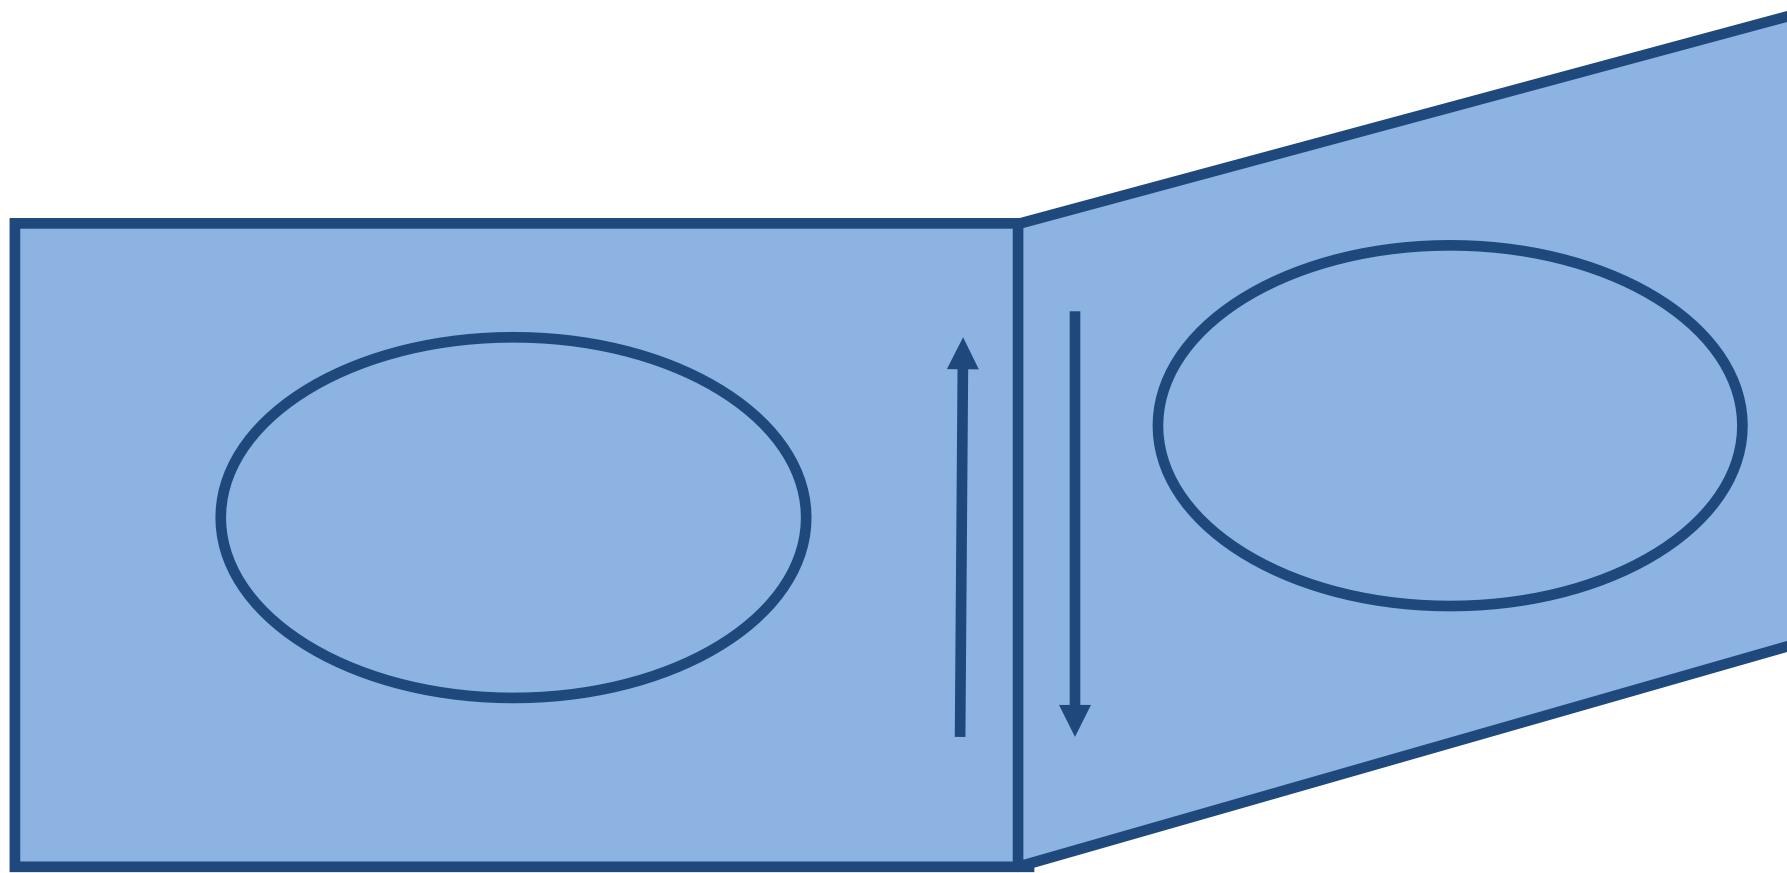
Orientation

- Every face of a polygonal mesh is orientable
 - Clockwise vs. counterclockwise order of face vertices
 - Defines sign/direction of the surface normal



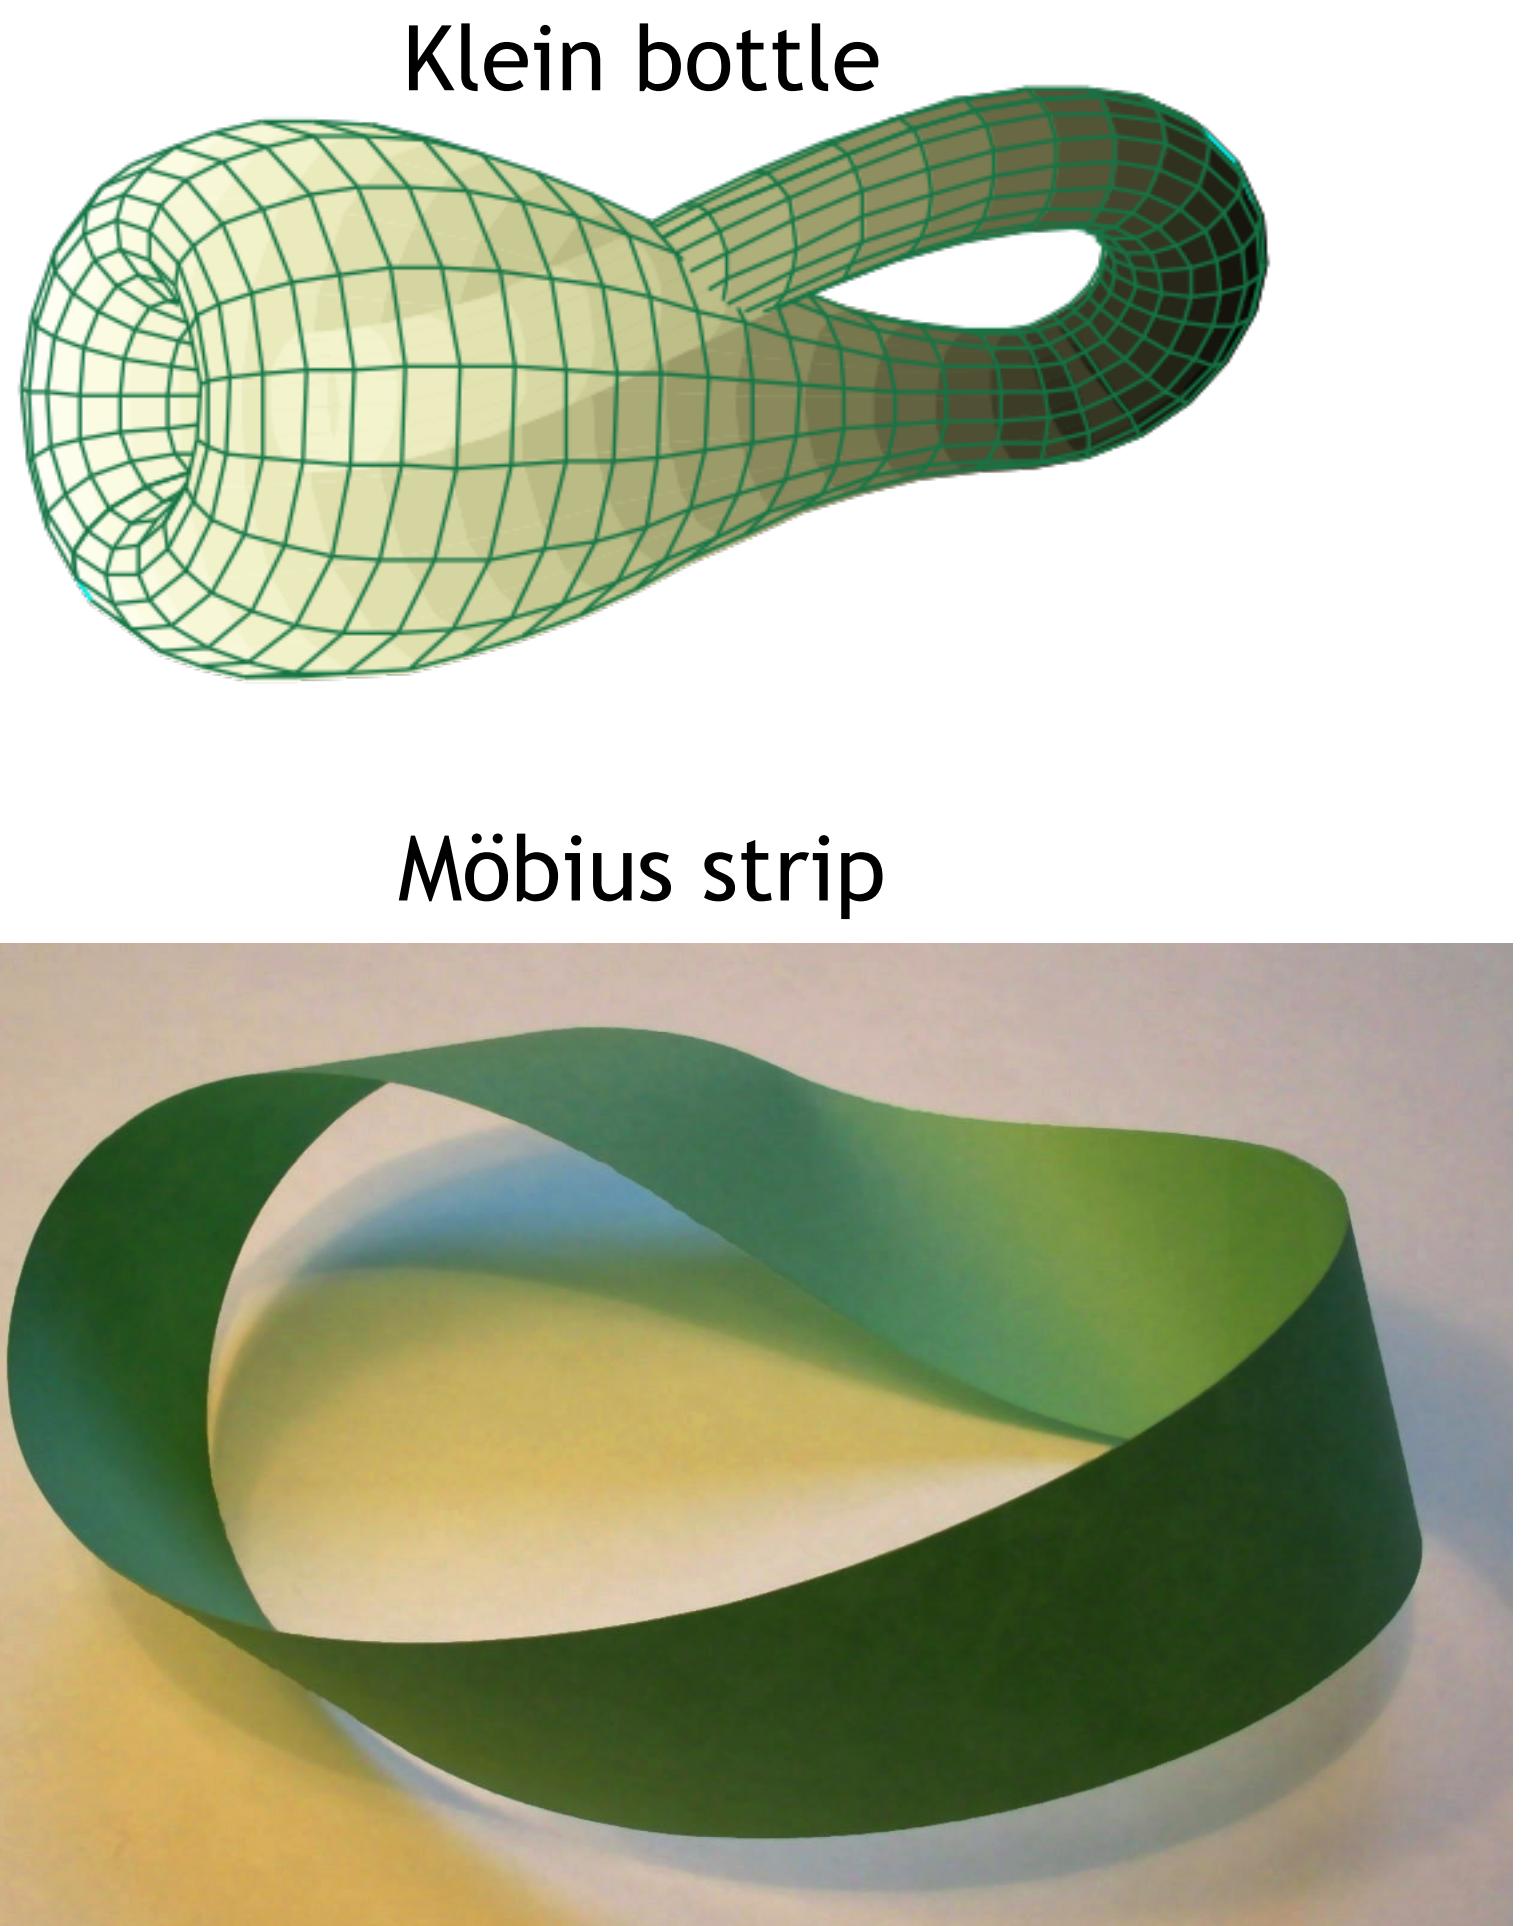
Orientation

- Consistent orientation of neighboring faces:



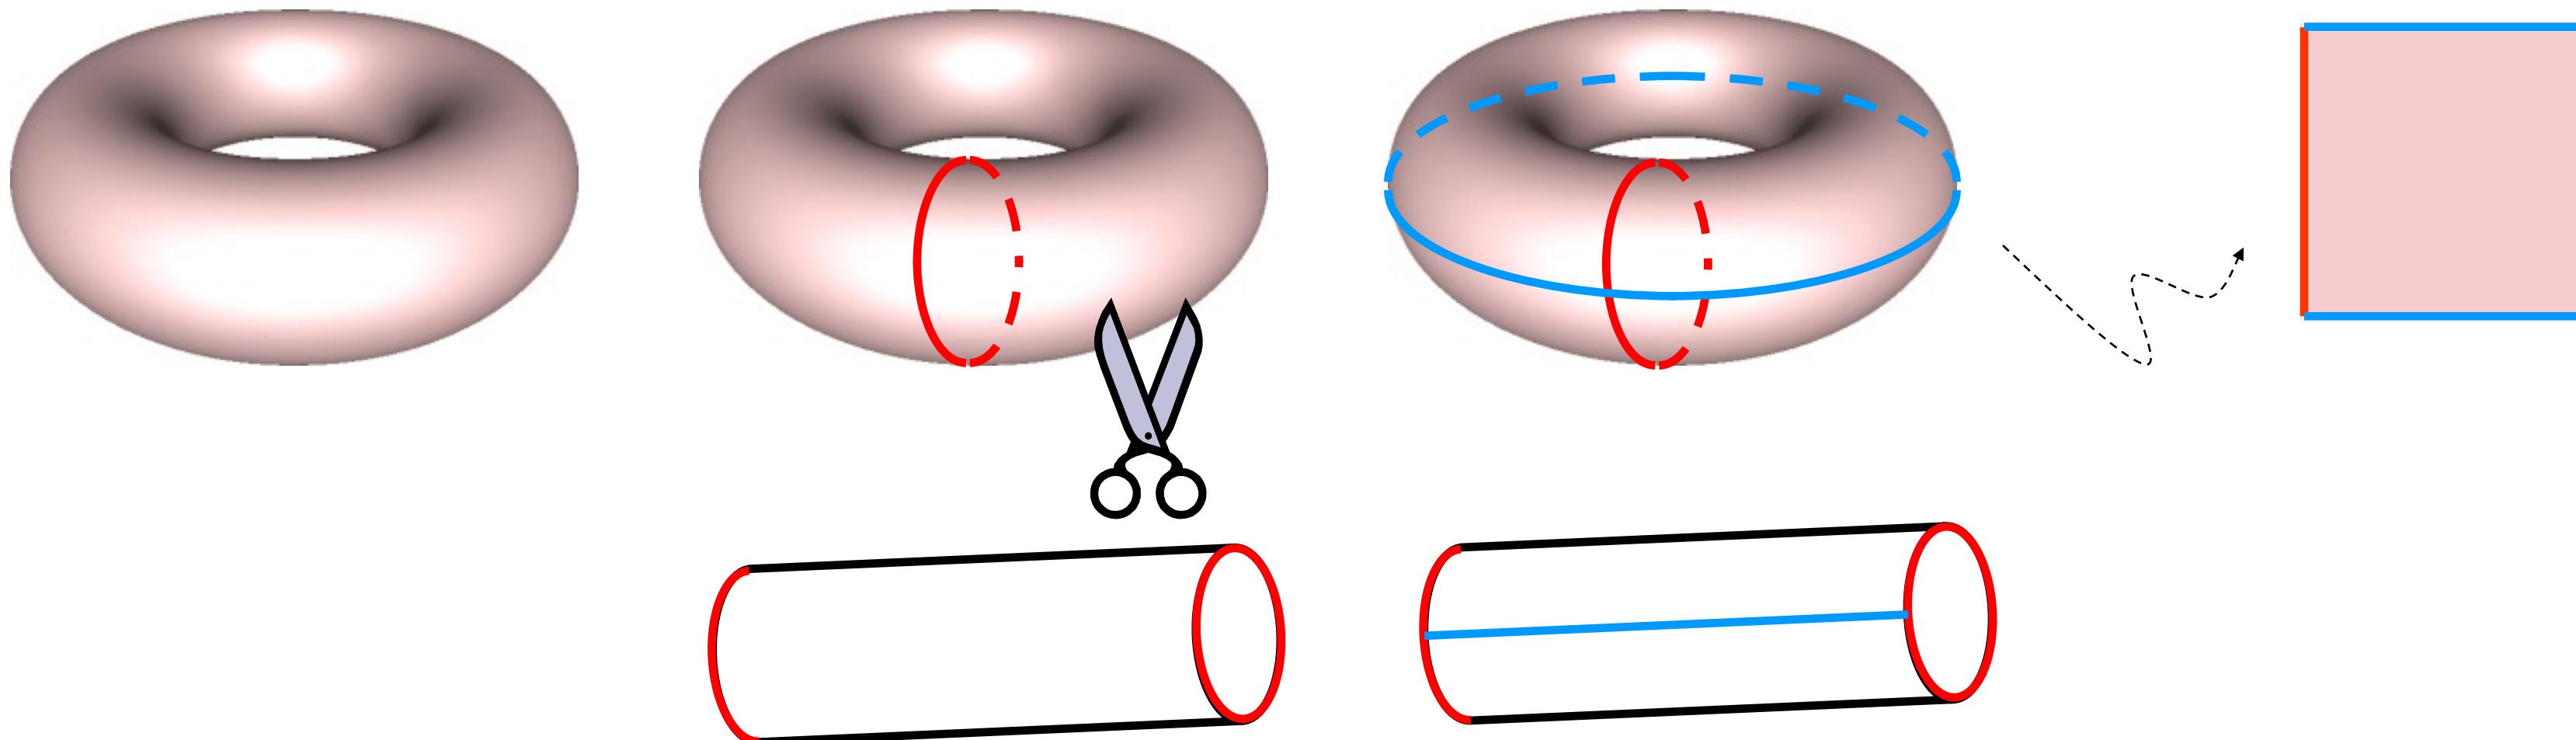
Orientability

- A polygonal mesh is orientable, if the incident faces to every edge can be consistently oriented
 - If the faces are consistently oriented for every edge, the mesh is oriented
- Notes
 - Every non-orientable closed mesh embedded in \mathbb{R}^3 intersects itself
 - The surface of a polyhedron is always orientable



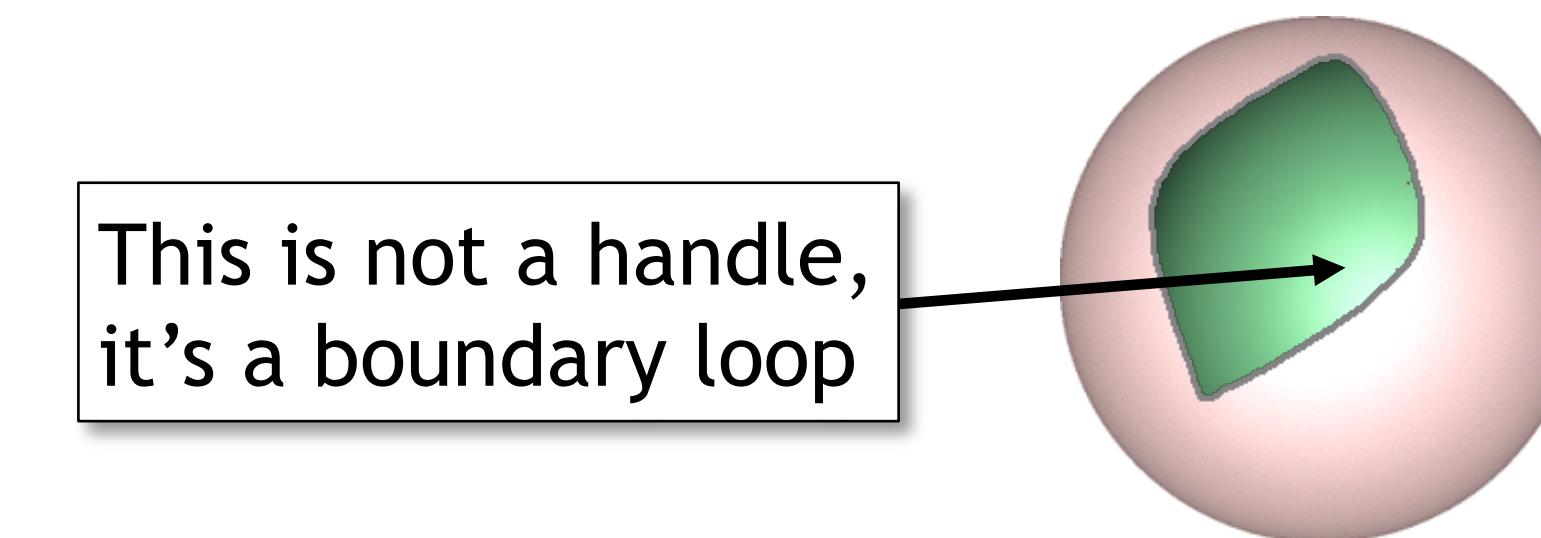
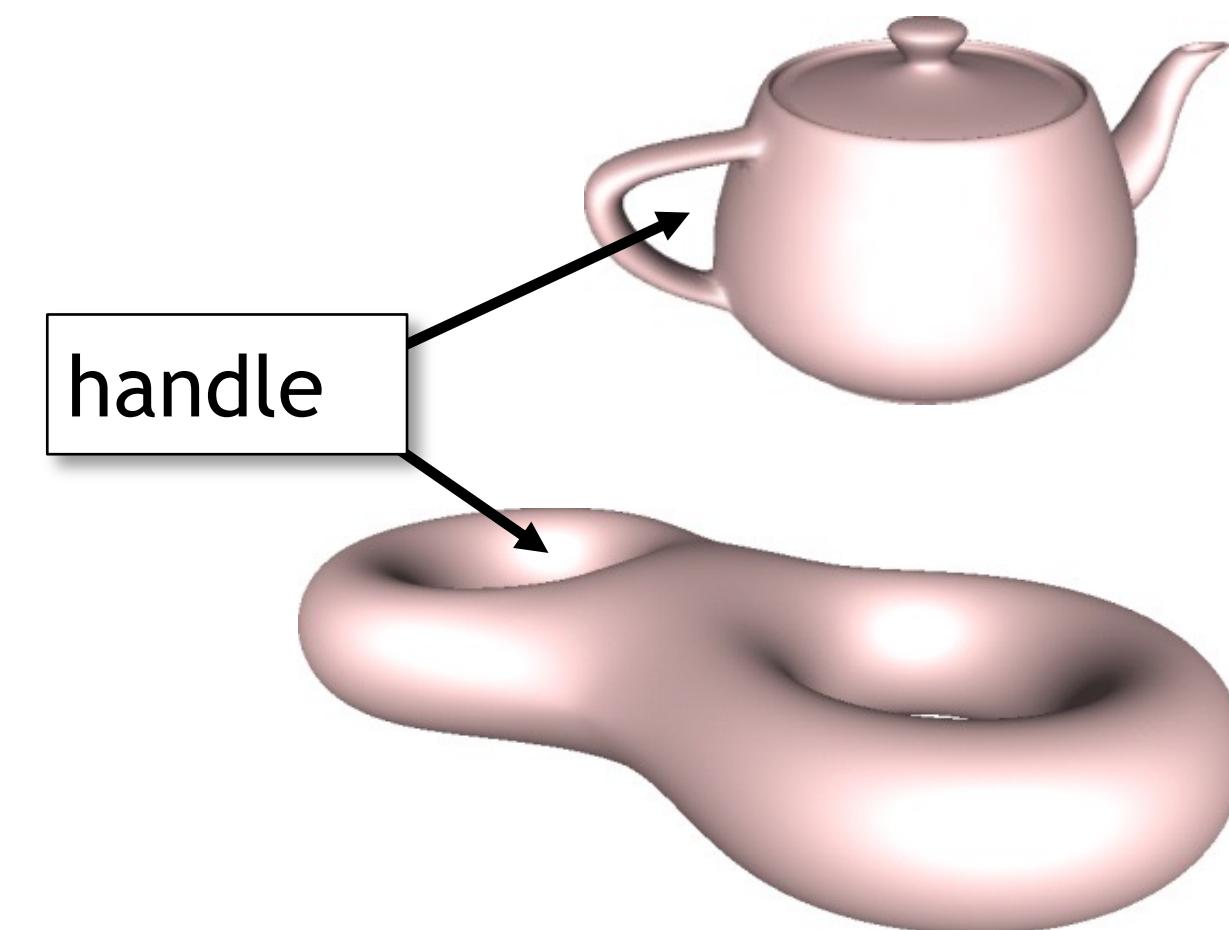
Global Topology of Meshes

- **Genus:** $\frac{1}{2} \times$ the maximal number of closed paths that do not disconnect the graph.
 - Informally, the number of handles (“donut holes”).



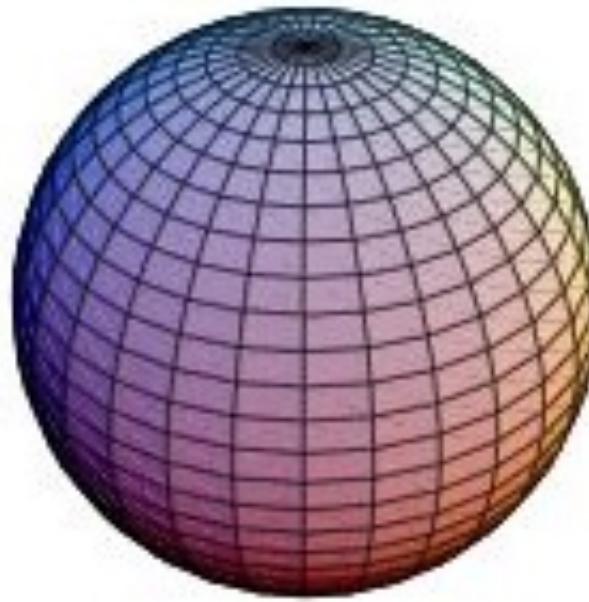
Global Topology of Meshes

- **Genus:** $\frac{1}{2} \times$ the maximal number of closed paths that do not disconnect the graph.
 - Informally, the number of handles (“donut holes”).

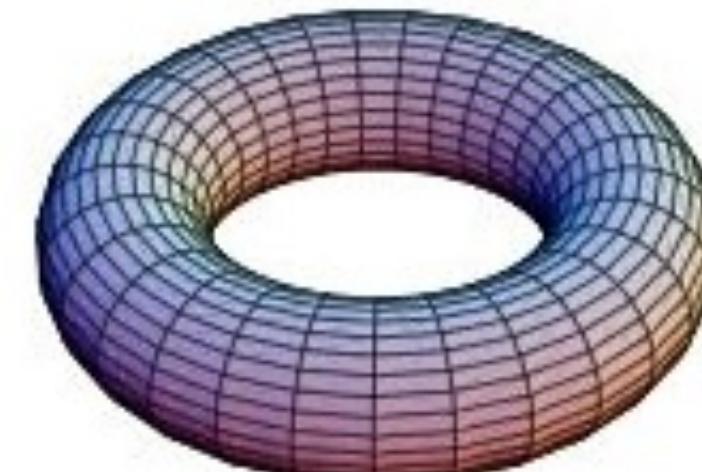


Global Topology of Meshes

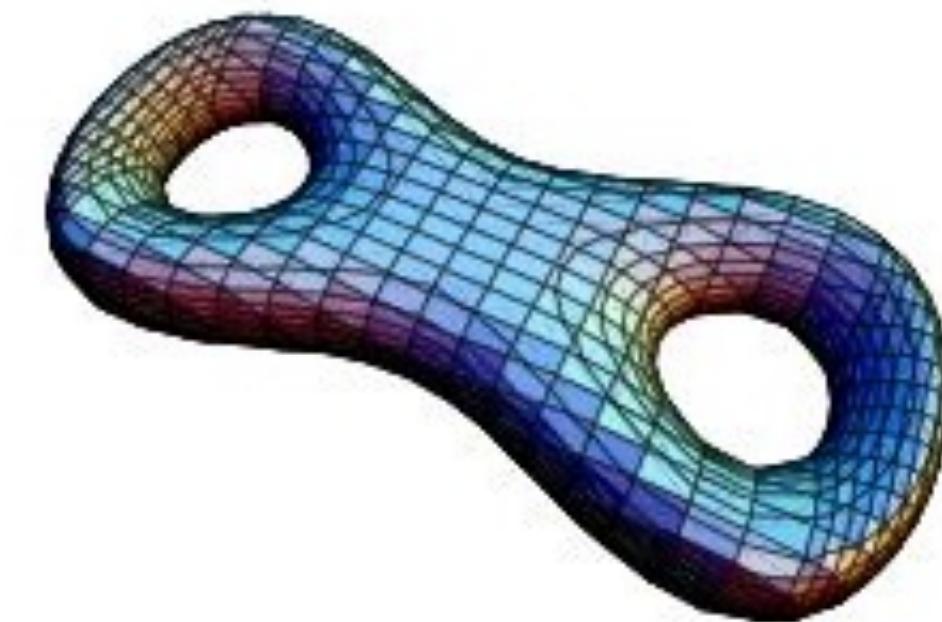
- **Genus:** $\frac{1}{2} \times$ the maximal number of closed paths that do not disconnect the graph.
 - Informally, the number of handles (“donut holes”).



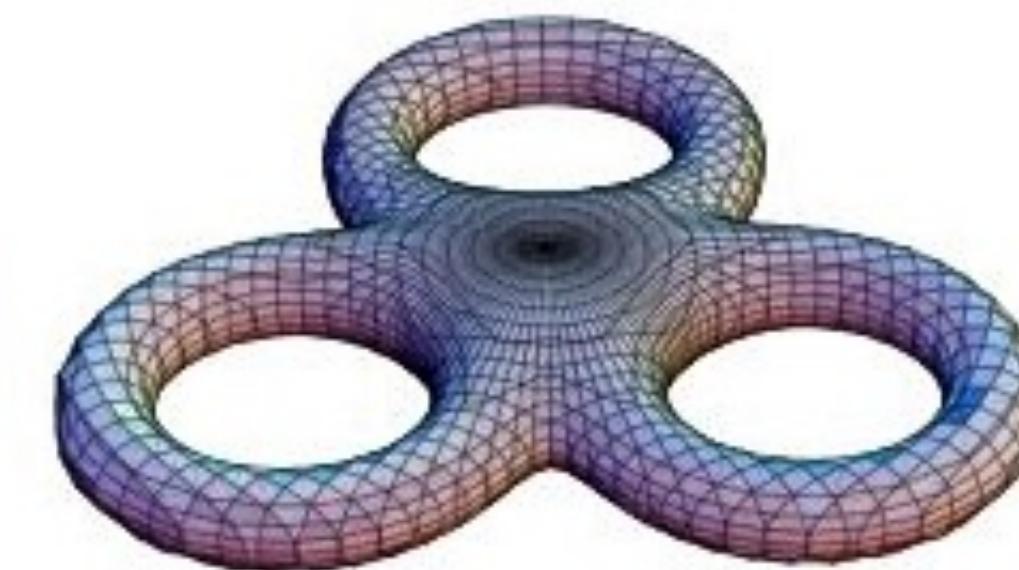
Genus 0



Genus 1



Genus 2



Genus 3

Euler-Poincaré Formula

- Theorem (Euler): The sum

$$\chi(M) = v - e + f$$

is **constant** for a given surface topology, no matter which (manifold) mesh we choose.

- v = number of vertices
- e = number of edges
- f = number of faces

Euler-Poincaré Formula

- For orientable meshes:

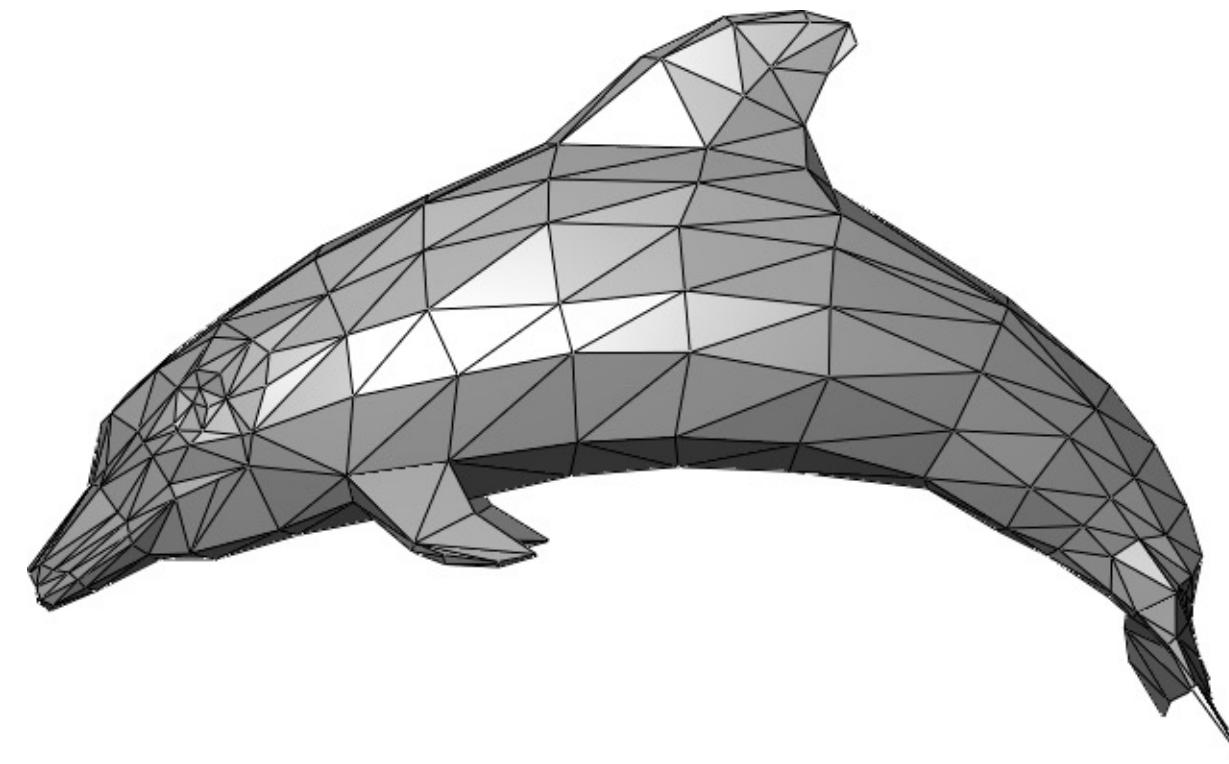
$$v - e + f = 2(c - g) - b = \chi(M)$$

- c = number of connected components
- g = genus
- b = number of boundary loops

$$\chi(\text{Sphere}) = 2 \quad \chi(\text{Torus}) = 0$$

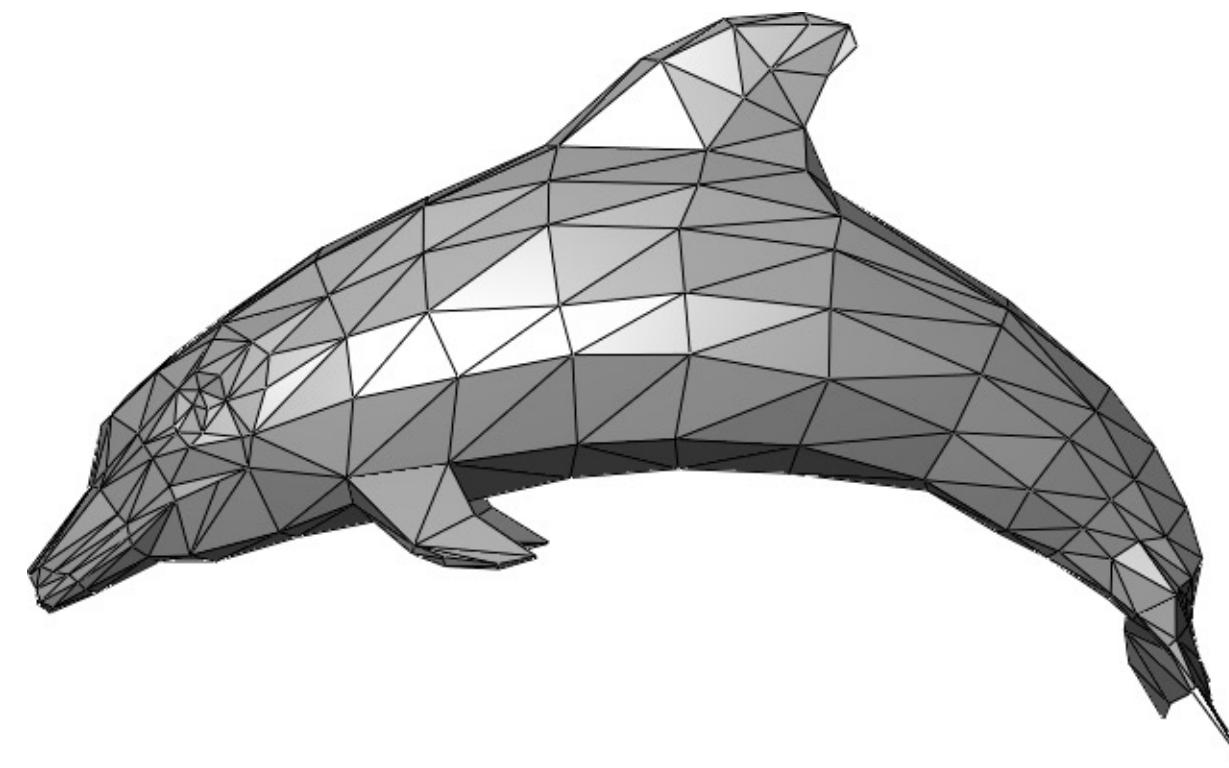
Implication for Mesh Storage

- Let's count the edges and faces in a closed **triangle mesh**:
 - Ratio of edges to faces: $e = 3/2 f$
 - each edge belongs to exactly 2 triangles
 - each triangle has exactly 3 edges



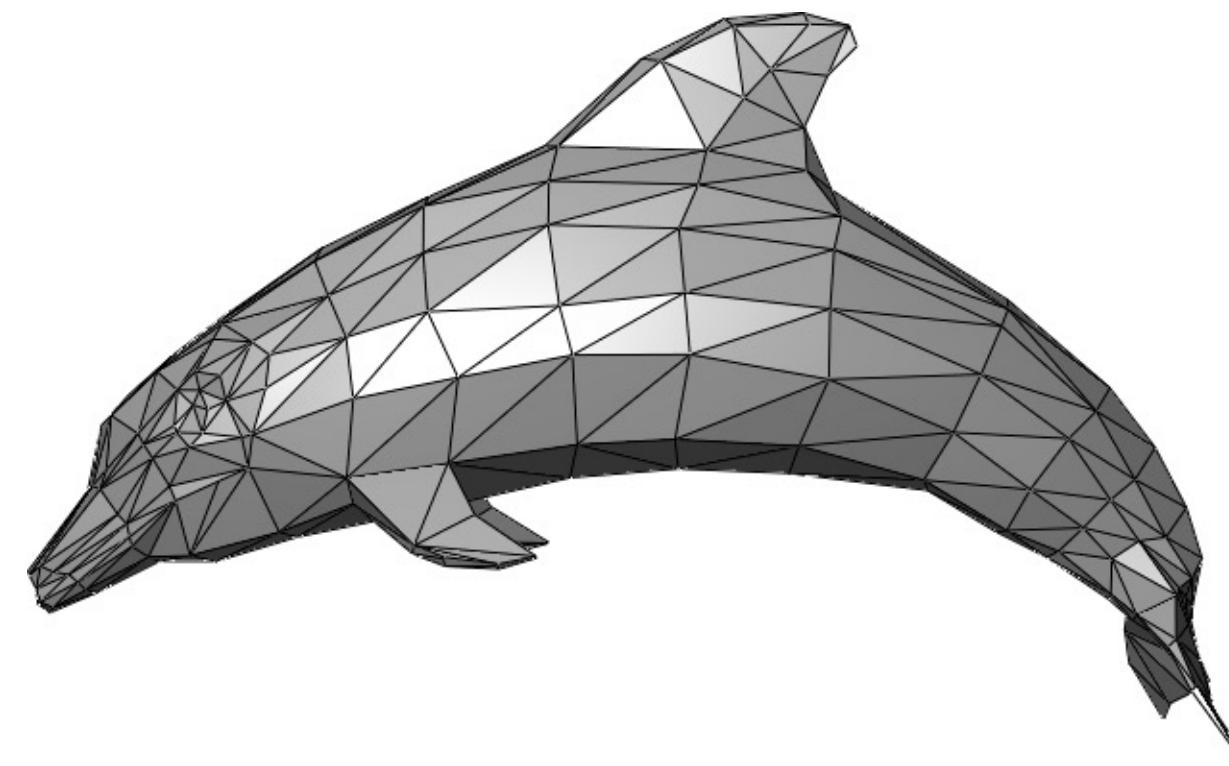
Implication for Mesh Storage

- Let's count the edges and faces in a closed **triangle mesh**:
 - Ratio of edges to faces: $e = 3/2 f$
 - each edge belongs to exactly 2 triangles
 - each triangle has exactly 3 edges
 - Ratio of vertices to faces: $f \sim 2v$
 - $2 = v - e + f = v - 3/2 f + f$
 - $2 + f / 2 = v$



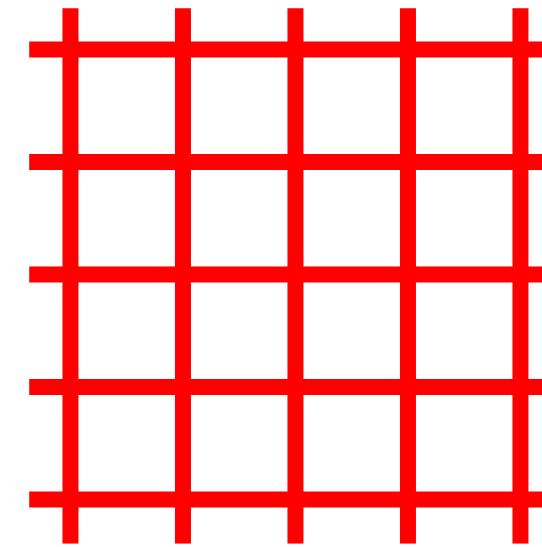
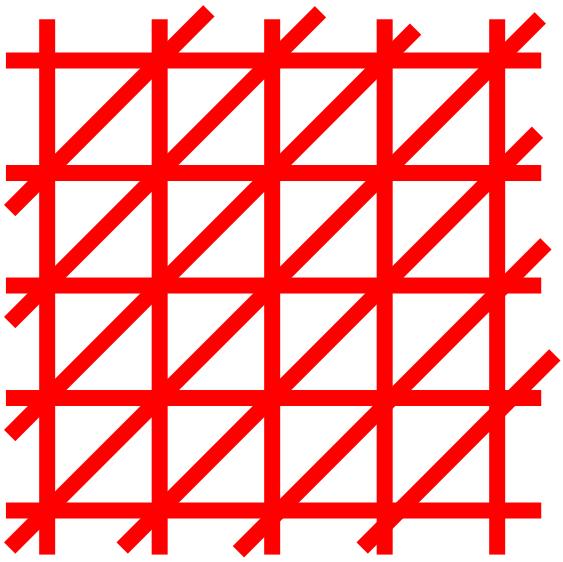
Implication for Mesh Storage

- Let's count the edges and faces in a closed **triangle mesh**:
 - Ratio of edges to faces: $e = 3/2 f$
 - each edge belongs to exactly 2 triangles
 - each triangle has exactly 3 edges
 - Ratio of vertices to faces: $f \sim 2v$
 - $2 = v - e + f = v - 3/2 f + f$
 - $2 + f / 2 = v$
 - Ratio of edges to vertices: $e \sim 3v$
 - **Average degree of a vertex:** 6



Regularity

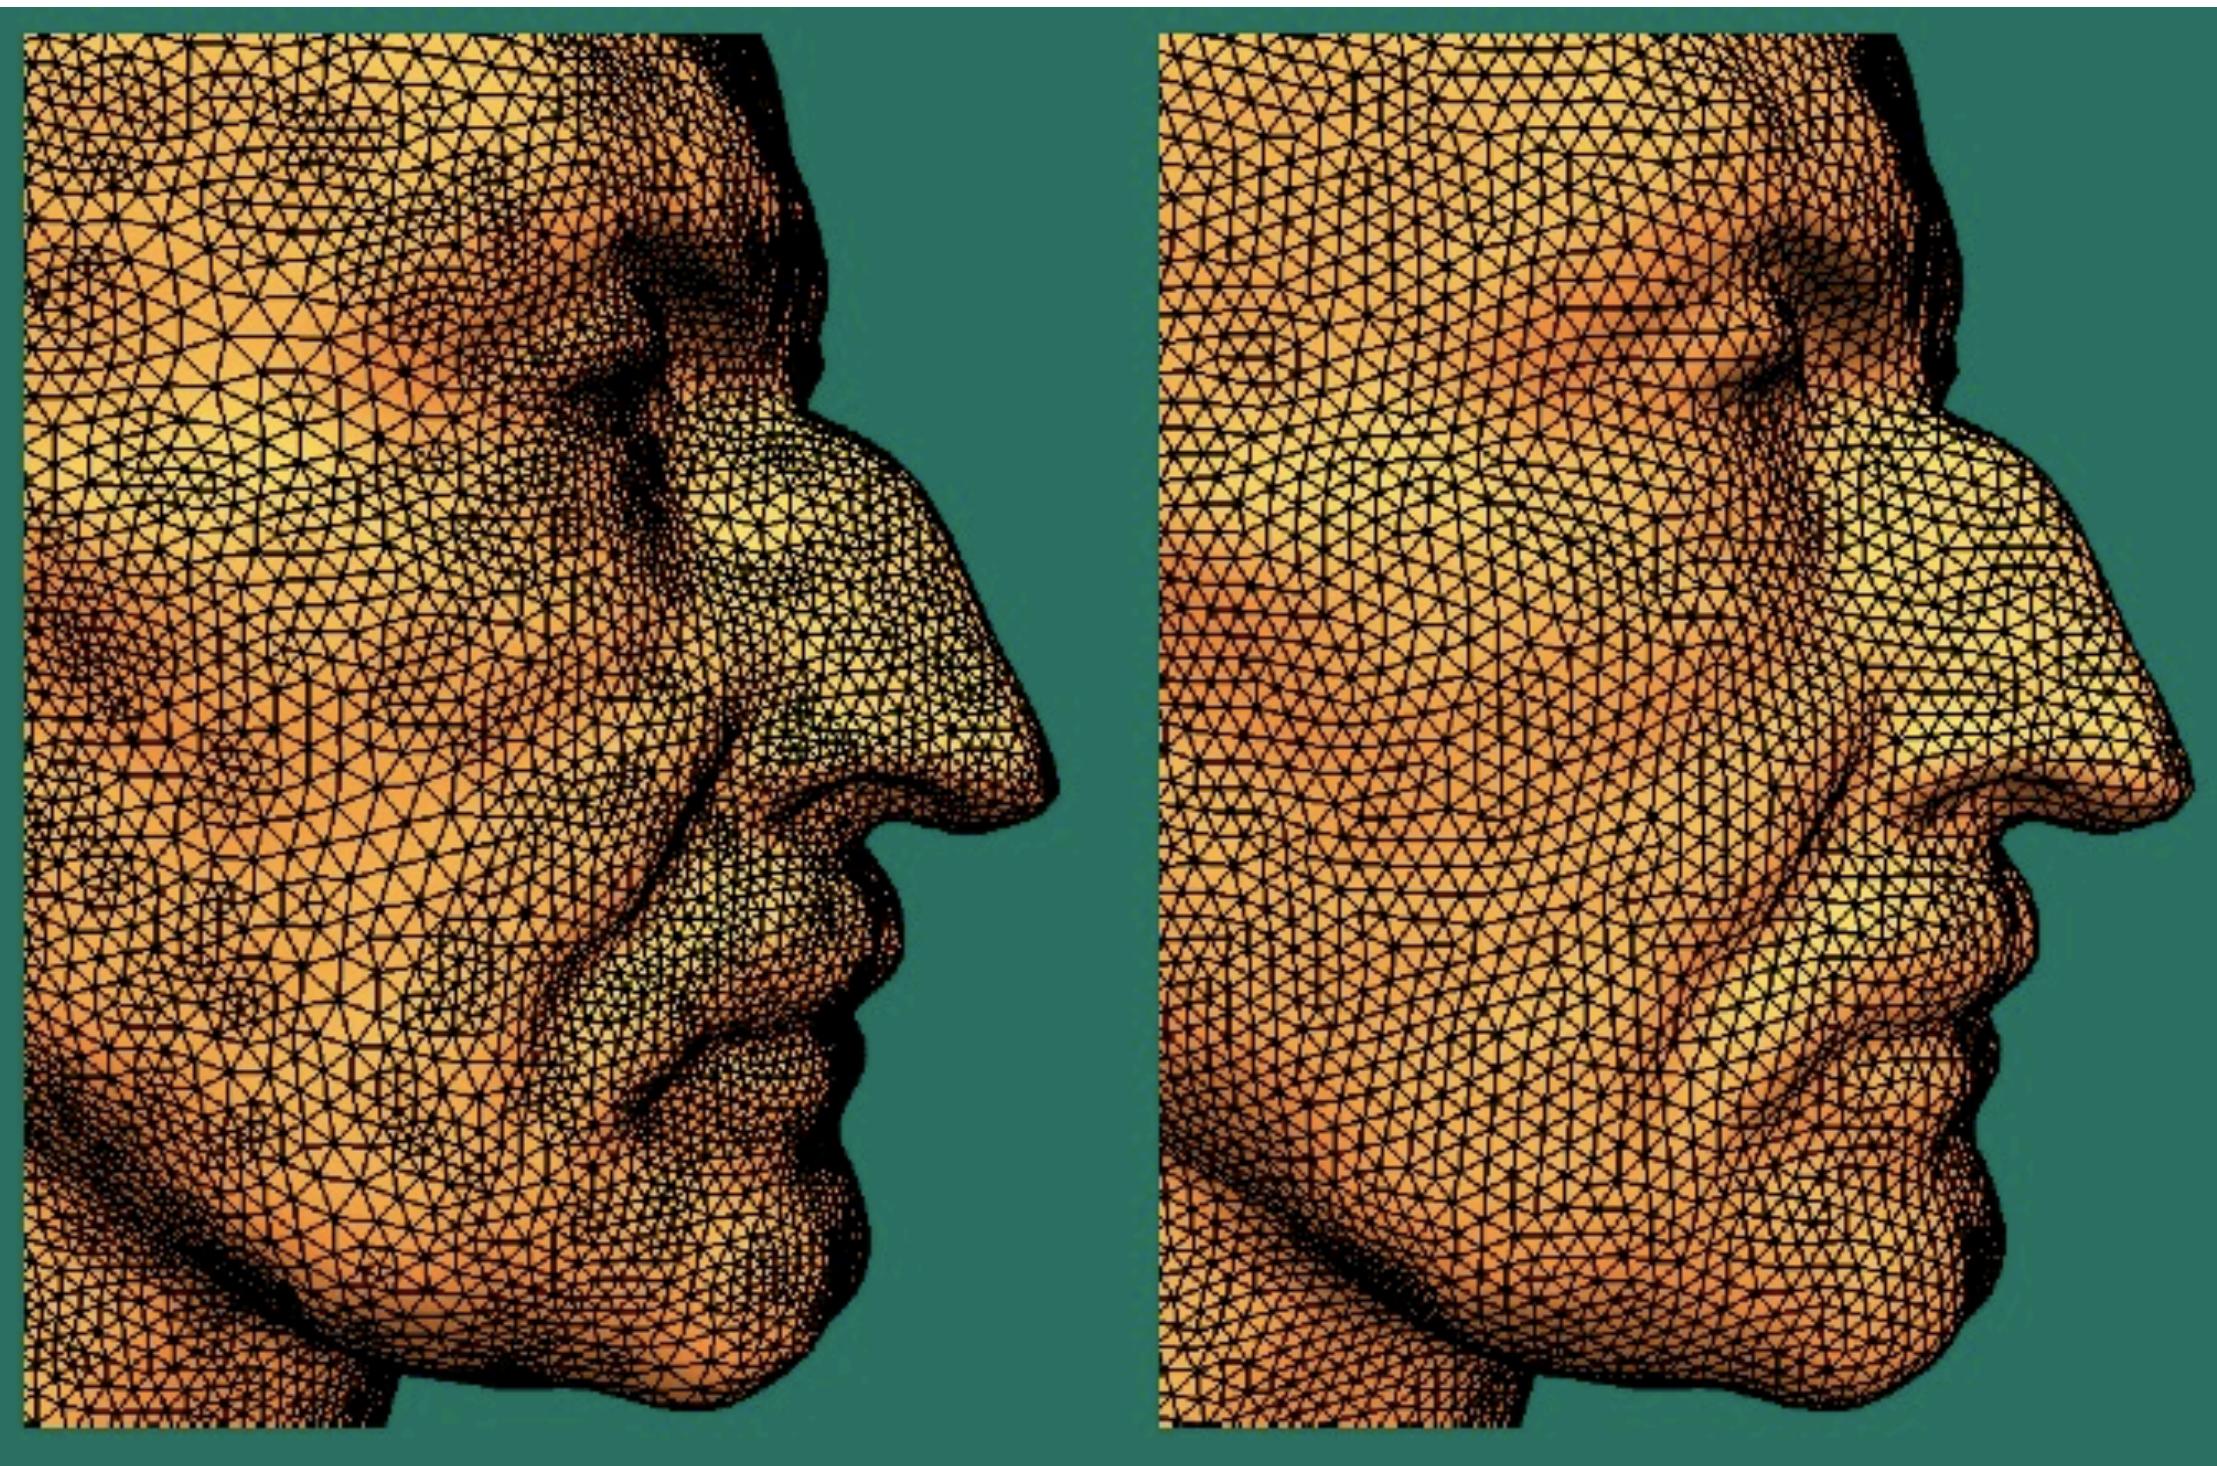
- Triangle mesh: average valence = 6
- Quad mesh: average valence = 4



- **Regular mesh:** all faces have the same number of edges and all vertex degrees are equal
- **Quasi-regular mesh:**
 - a lot of vertices have degree 6 (4). Sometimes also refers to mostly equilateral faces.

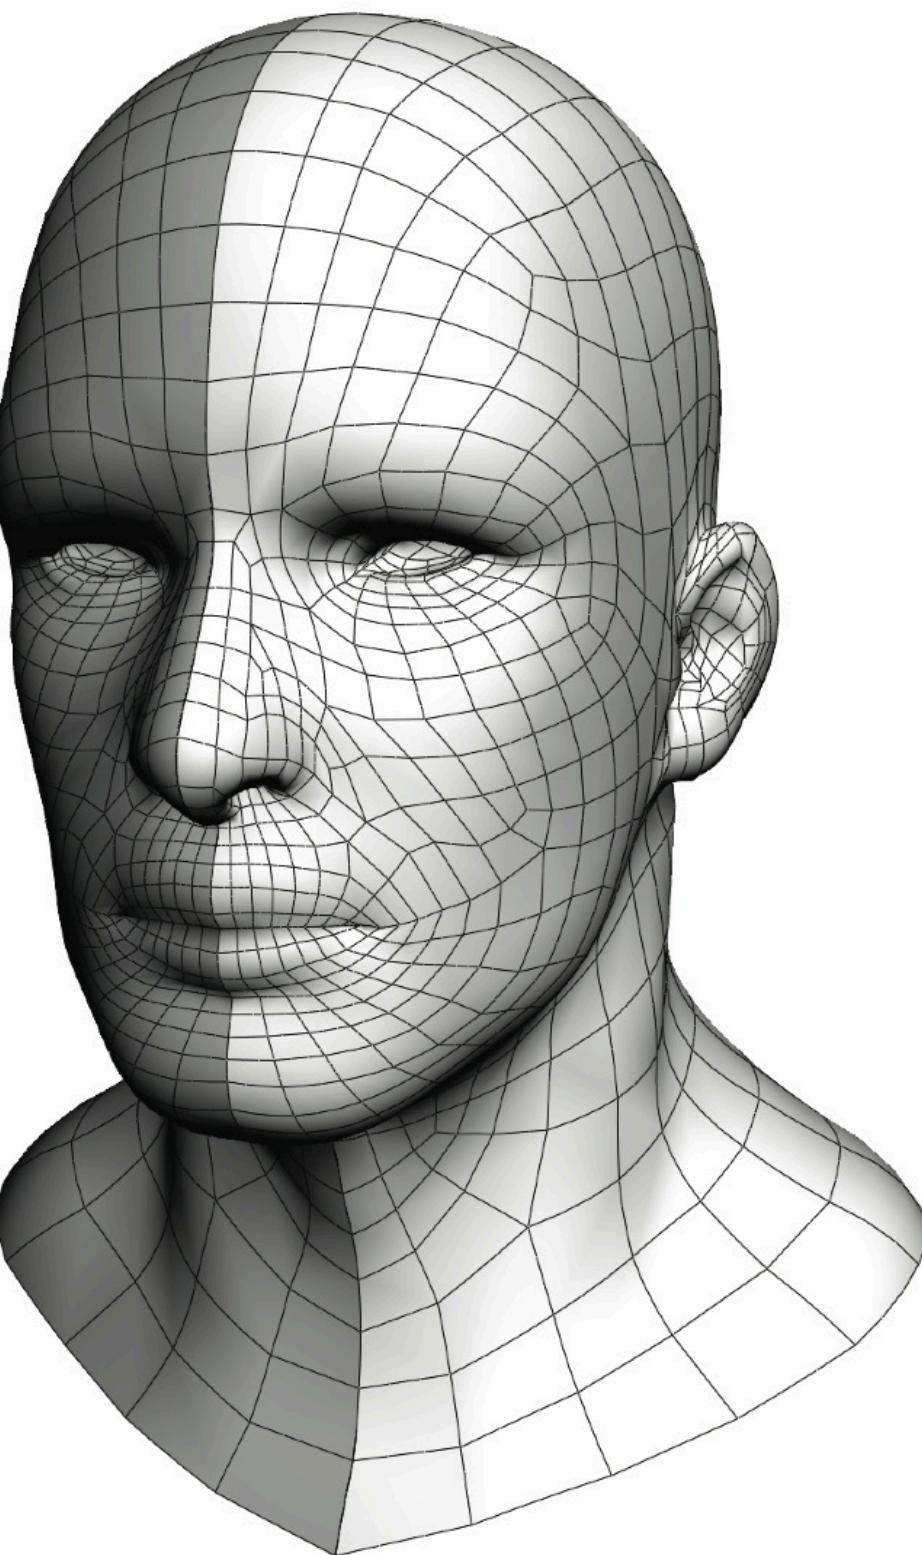
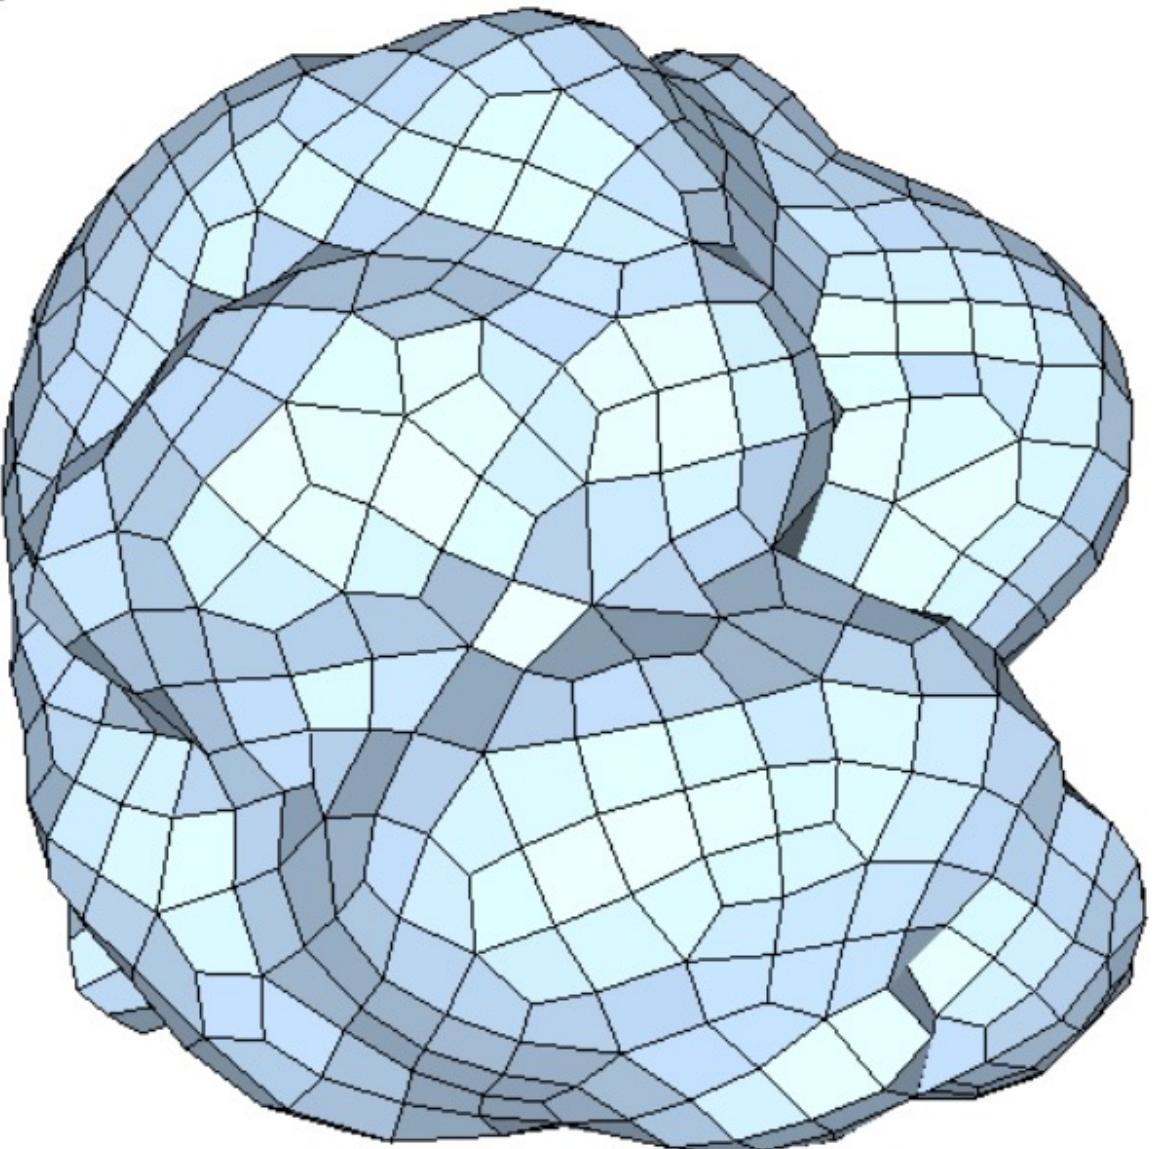
Regularity

- Quasi-regular



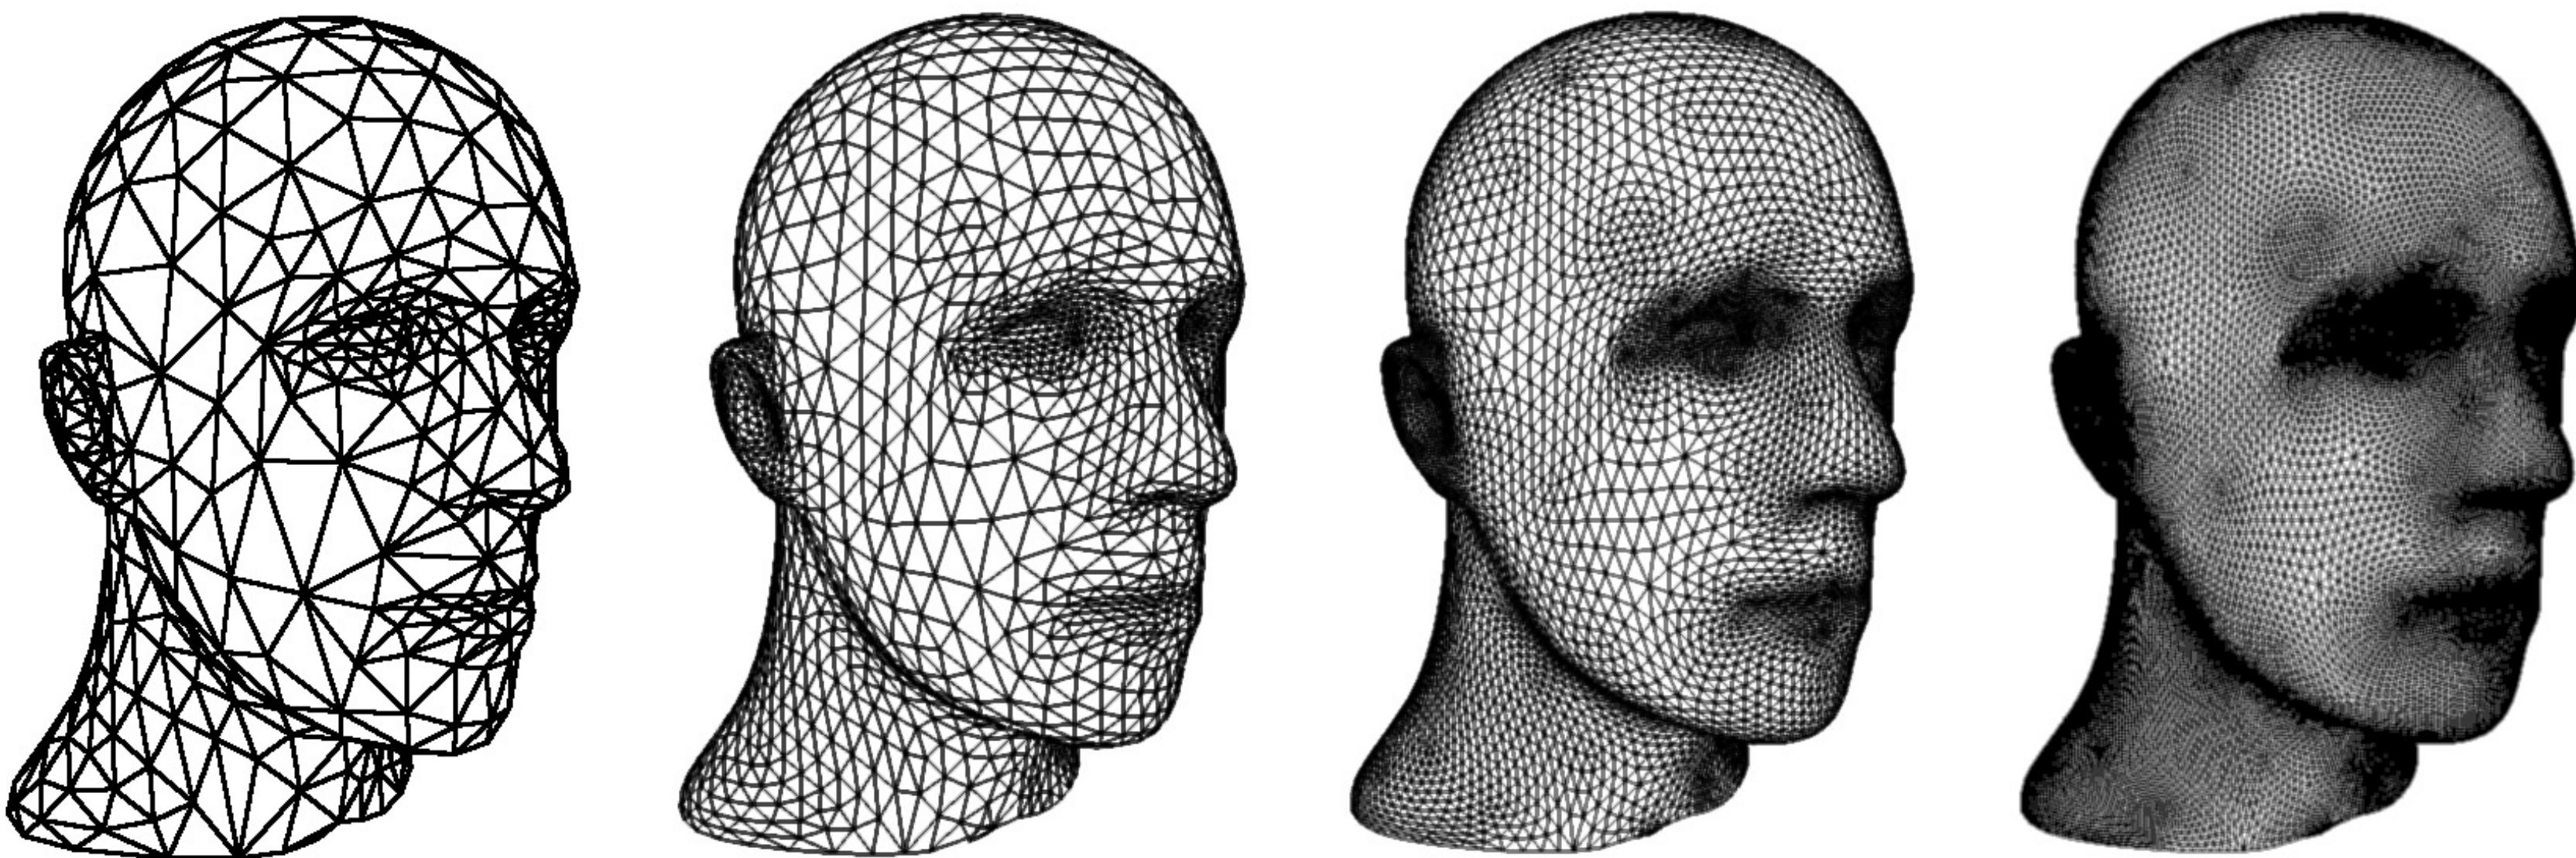
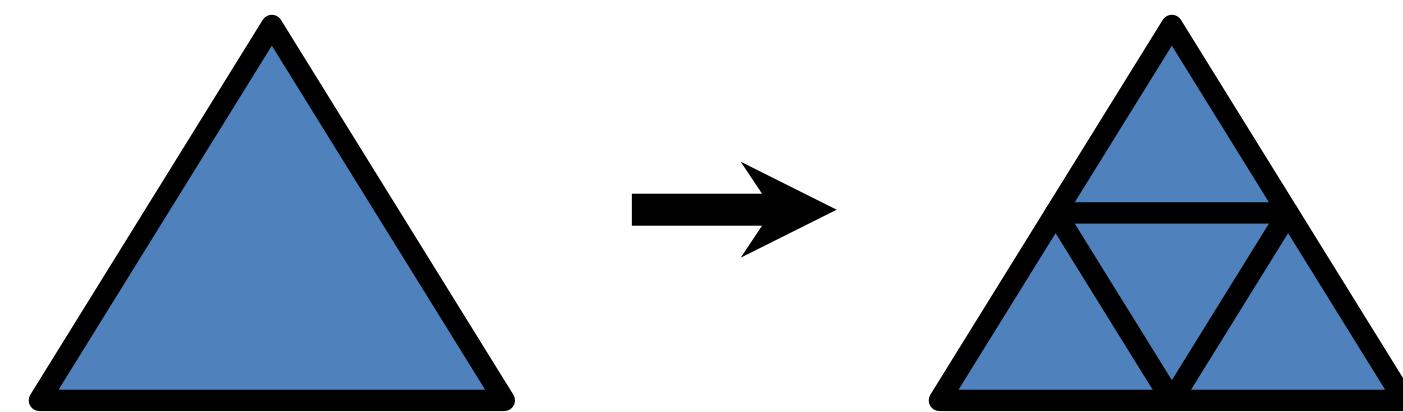
Regularity

- Quasi-regular



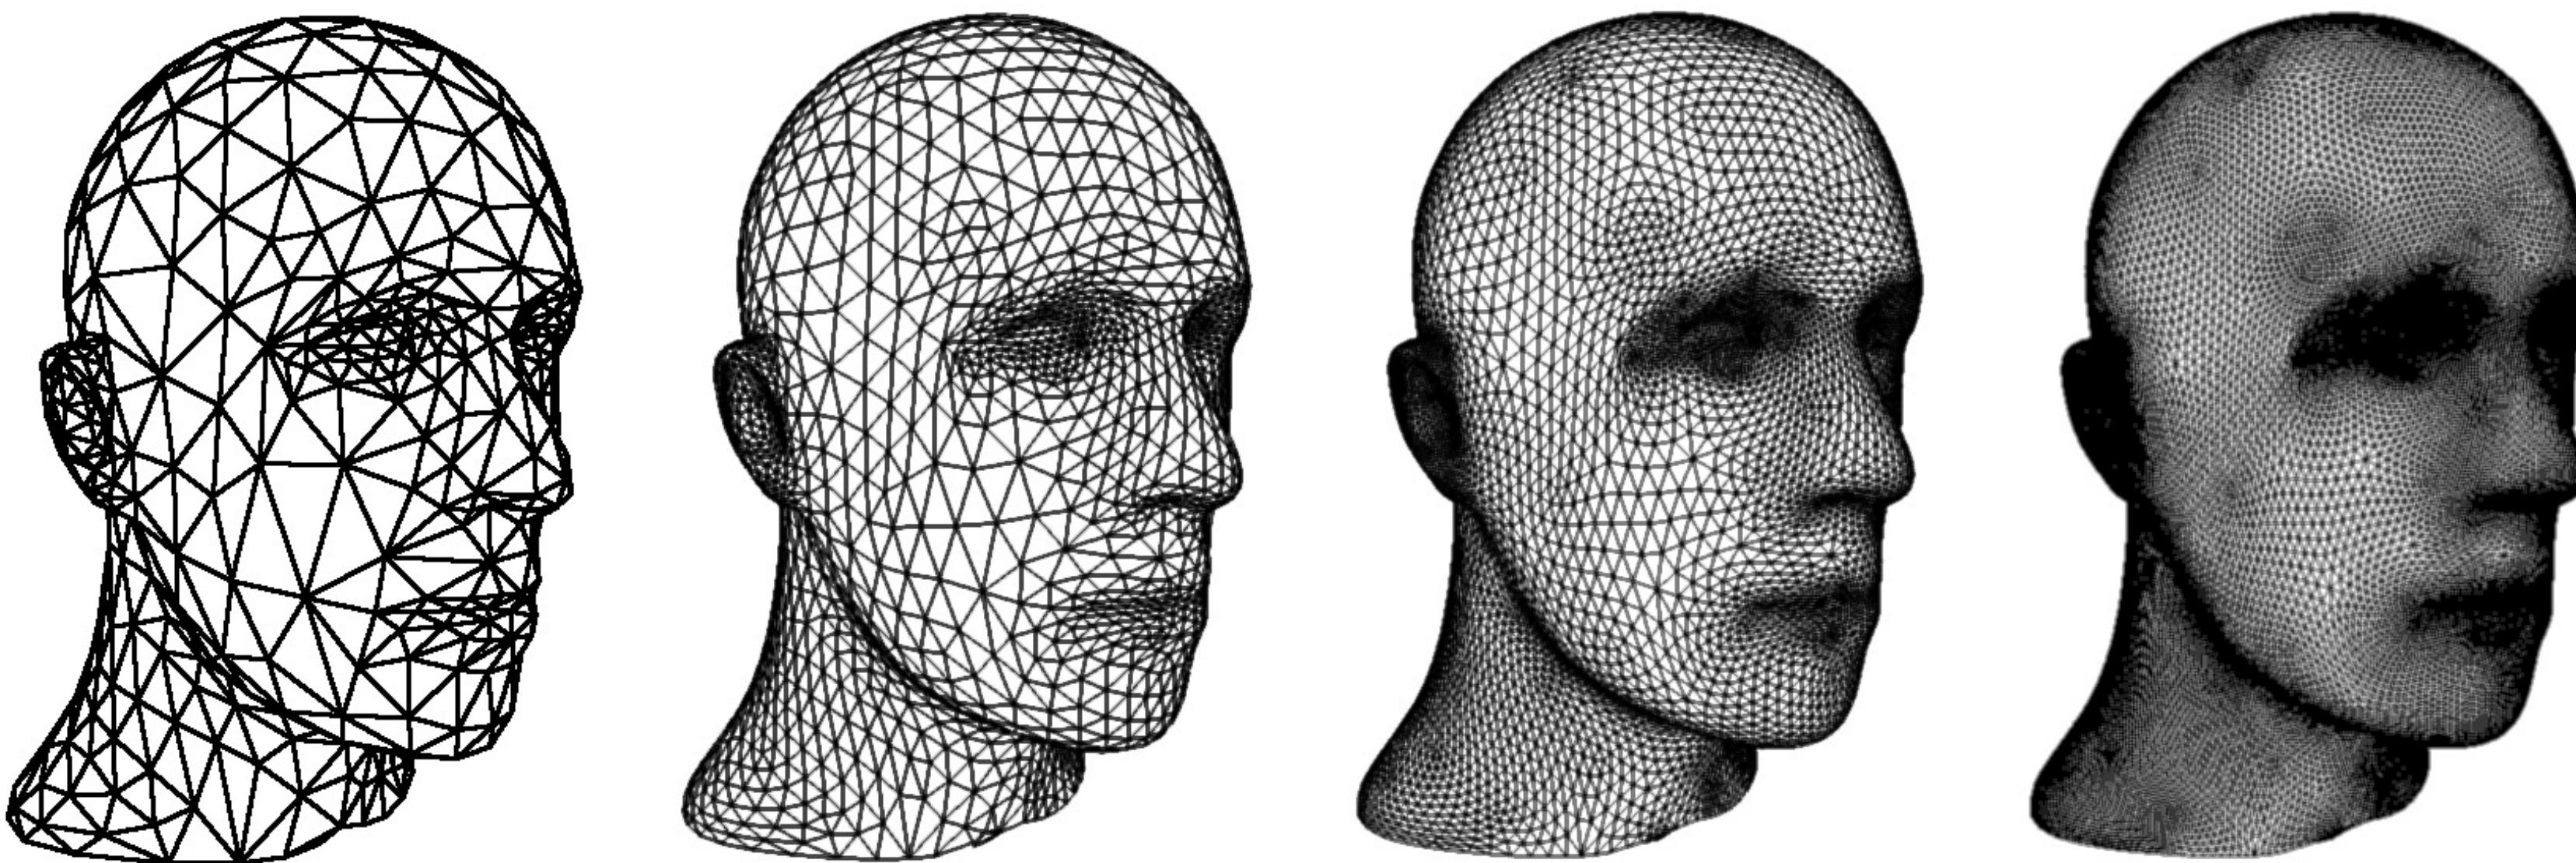
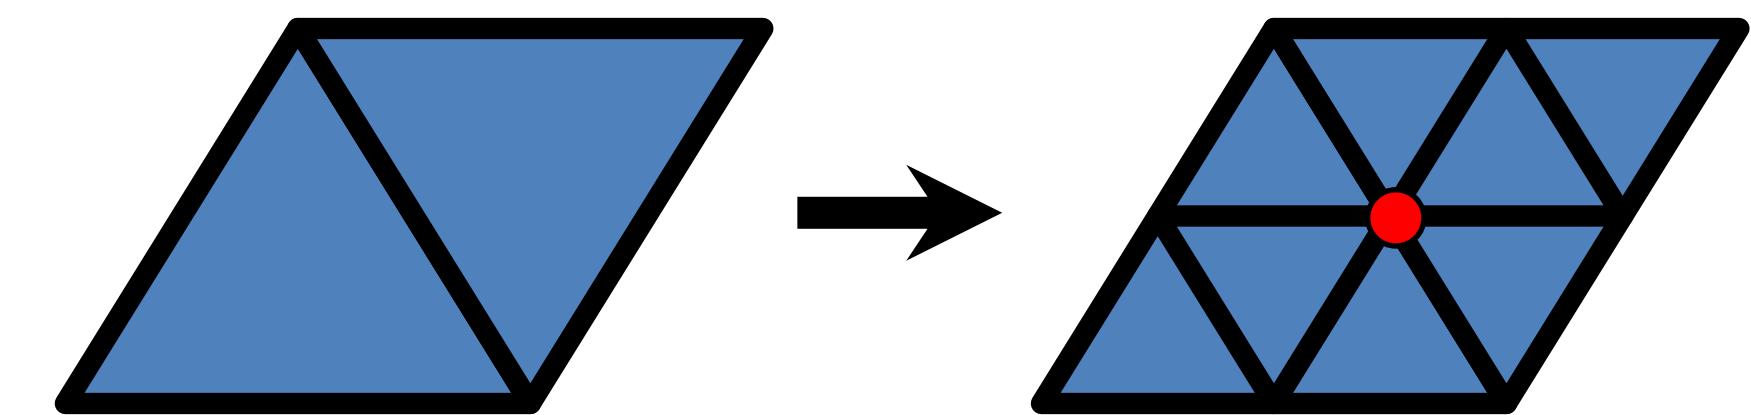
Regularity

- **Semi-regular mesh:**
connectivity is a result
of $N > 0$ subdivision steps

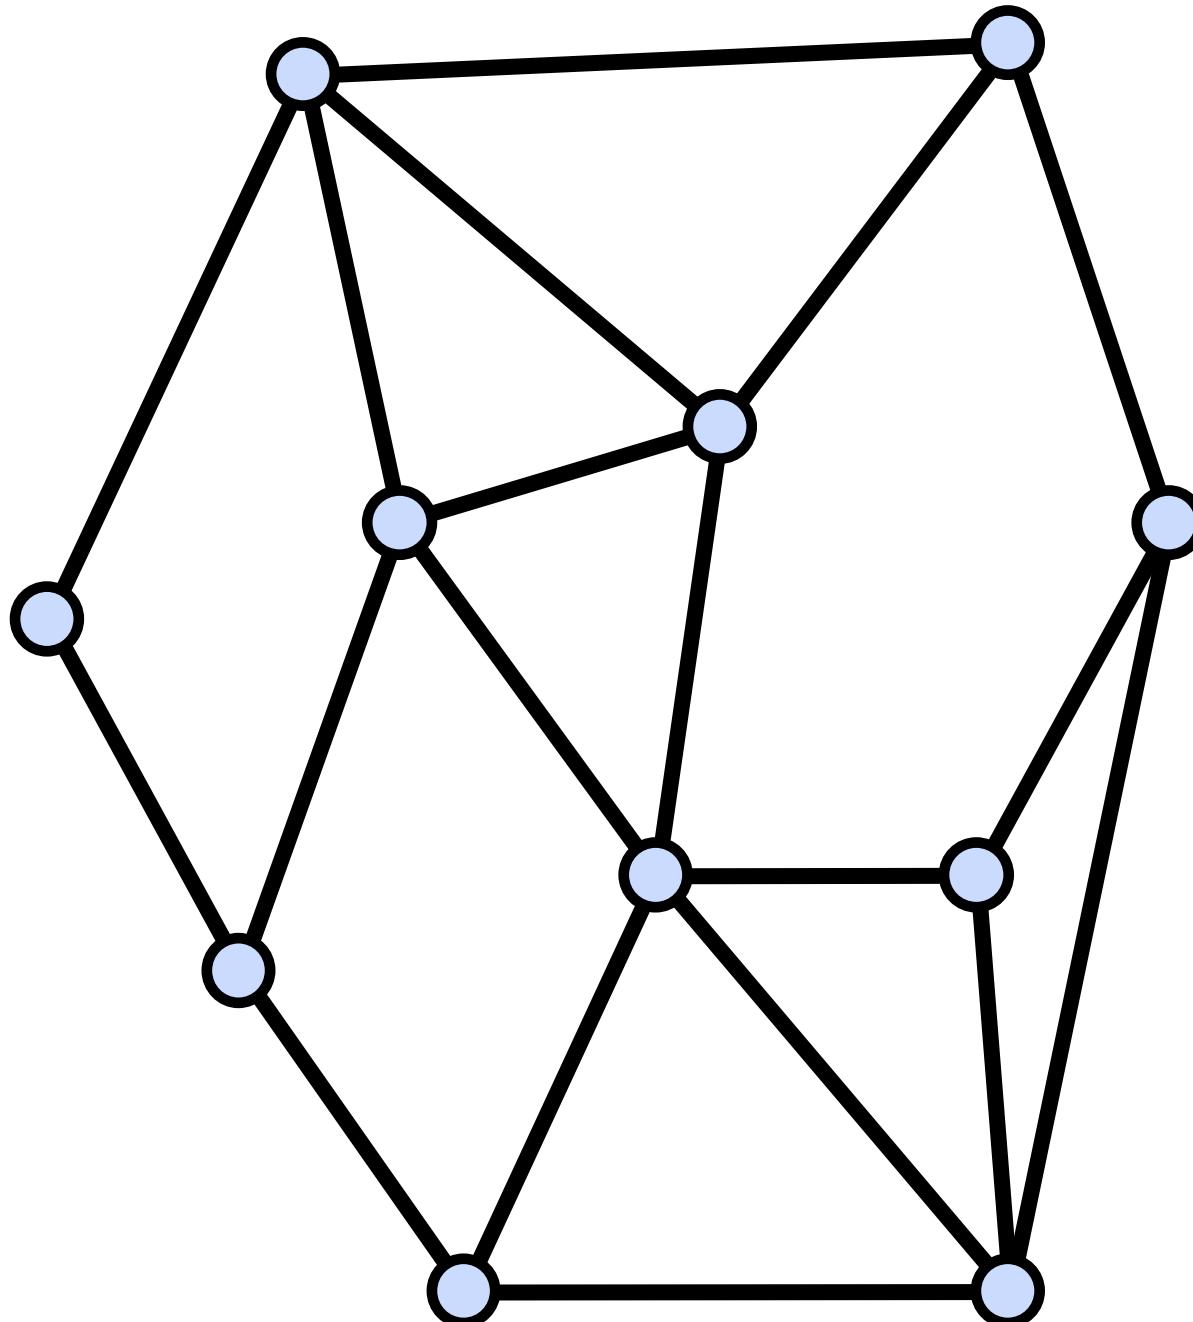


Regularity

- **Semi-regular mesh:**
connectivity is a result
of $N > 0$ subdivision steps

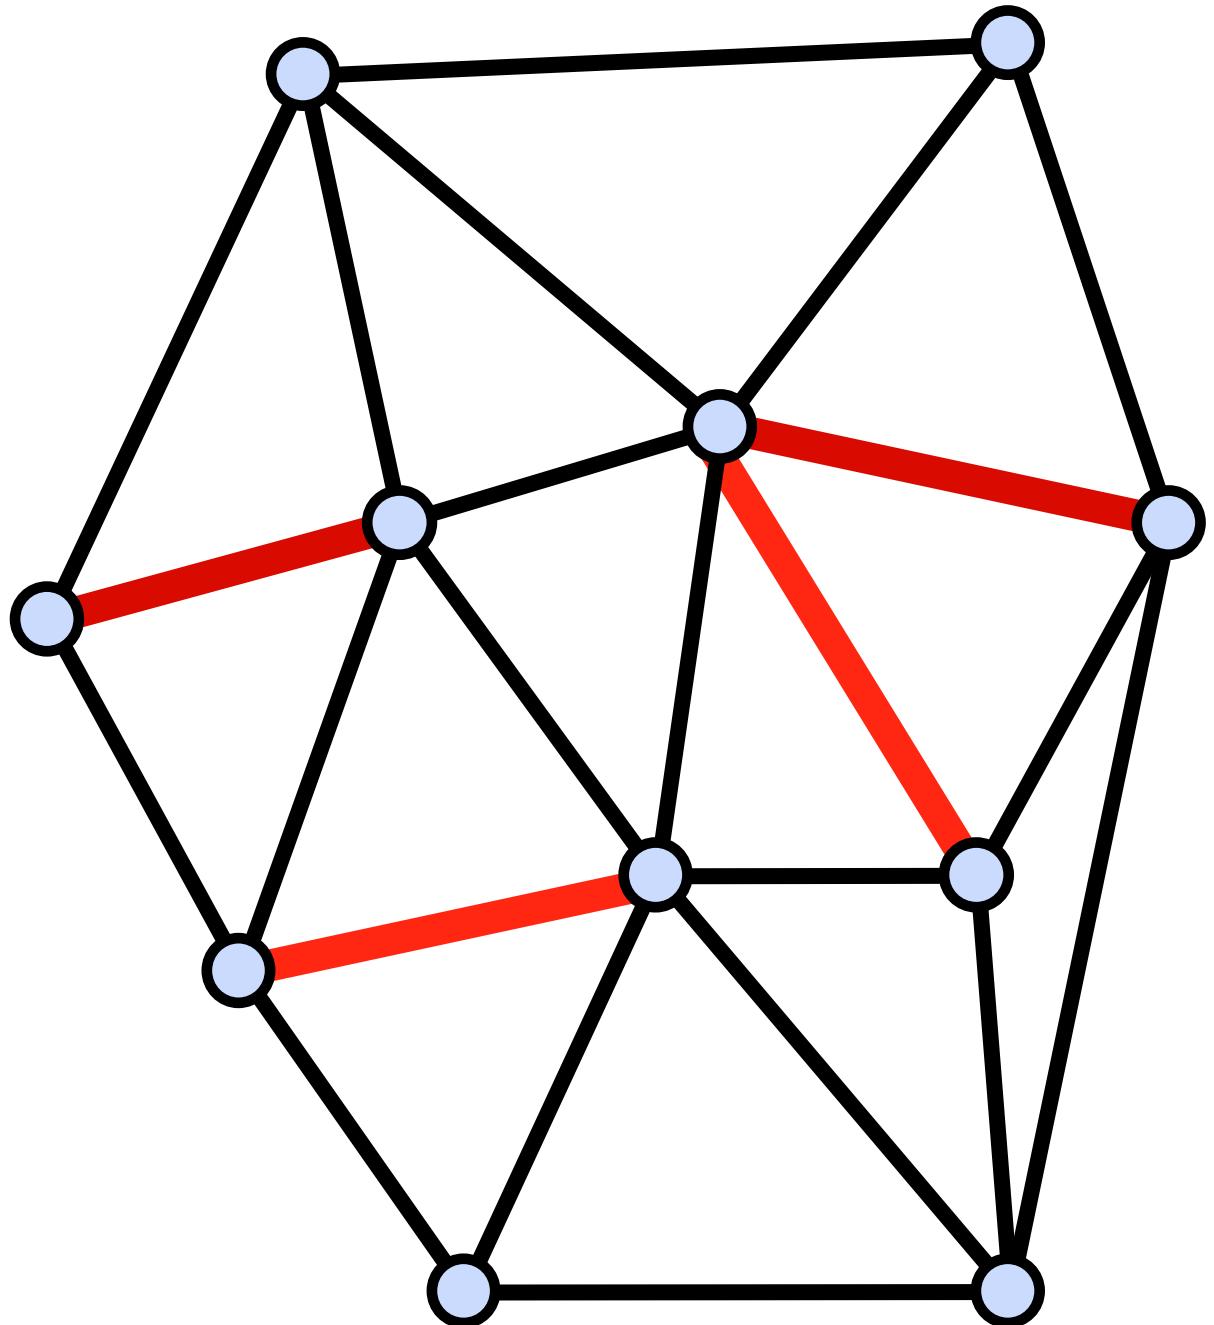


Triangulation



- Polygonal mesh where every face is a triangle
- Simplifies data structures
- Simplifies rendering
- Simplifies algorithms
- Each face planar and convex
- Any polygon can be triangulated

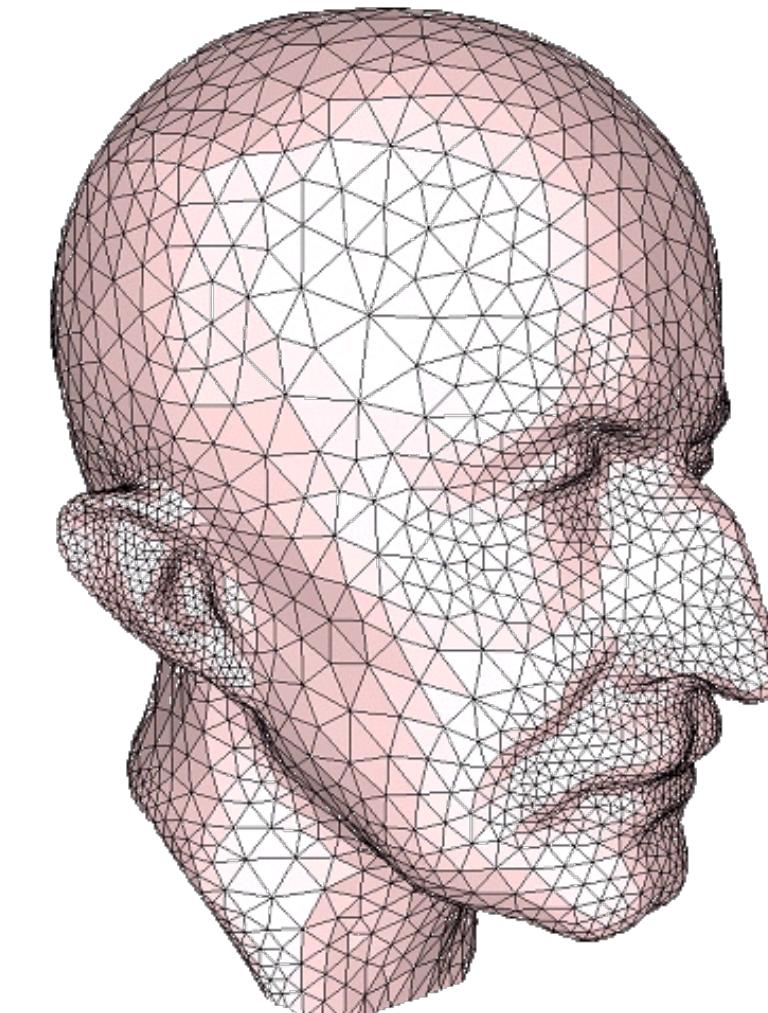
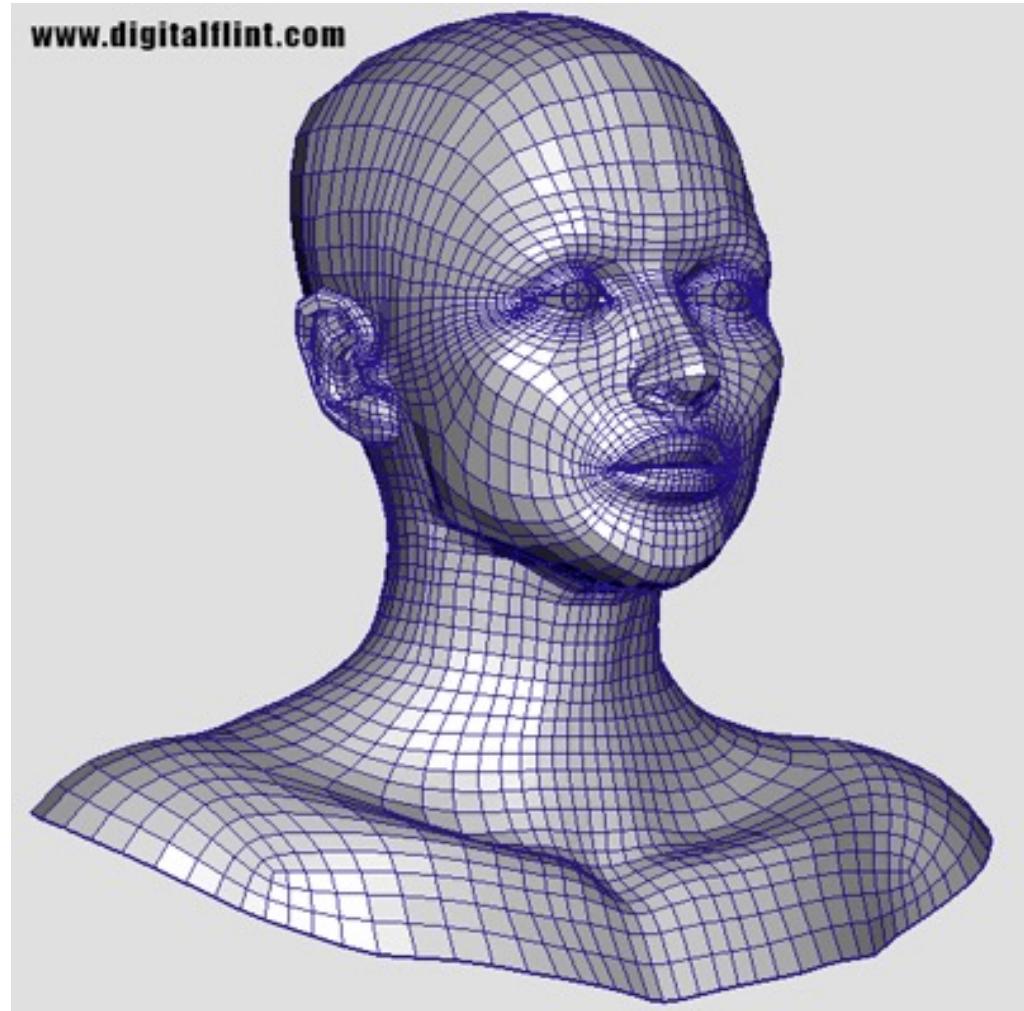
Triangulation



- Polygonal mesh where every face is a triangle
- Simplifies data structures
- Simplifies rendering
- Simplifies algorithms
- Each face planar and convex
- Any polygon can be triangulated

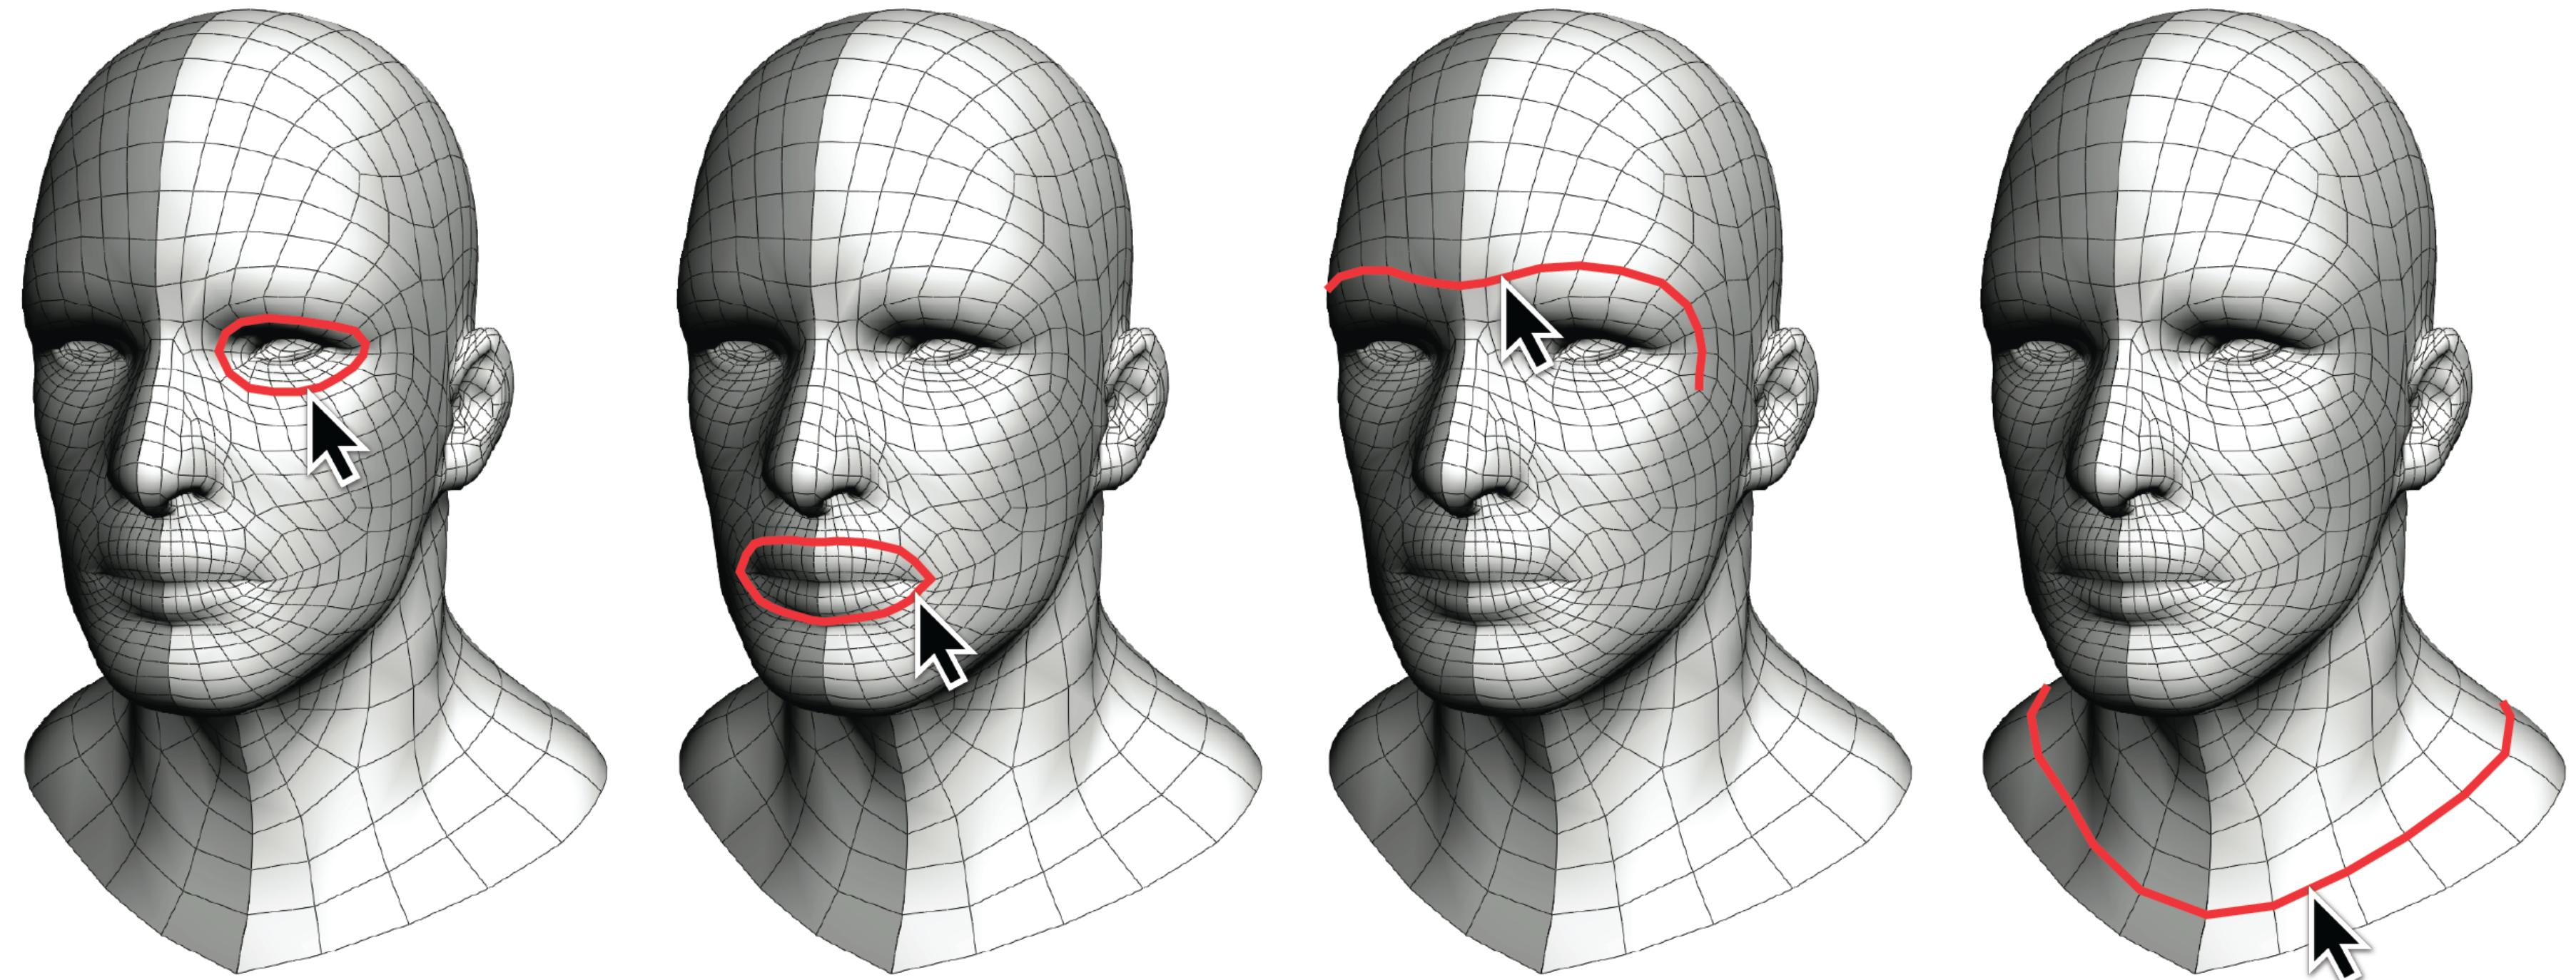
Polygonal vs. Triangle Meshes

- Triangles are flat and convex
 - Easy rasterization, normals
 - Uniformity (same # of vertices)
- General polygons are flexible
- Can be non-planar, non-convex
 - Difficult for graphics hardware
- Varying number of vertices



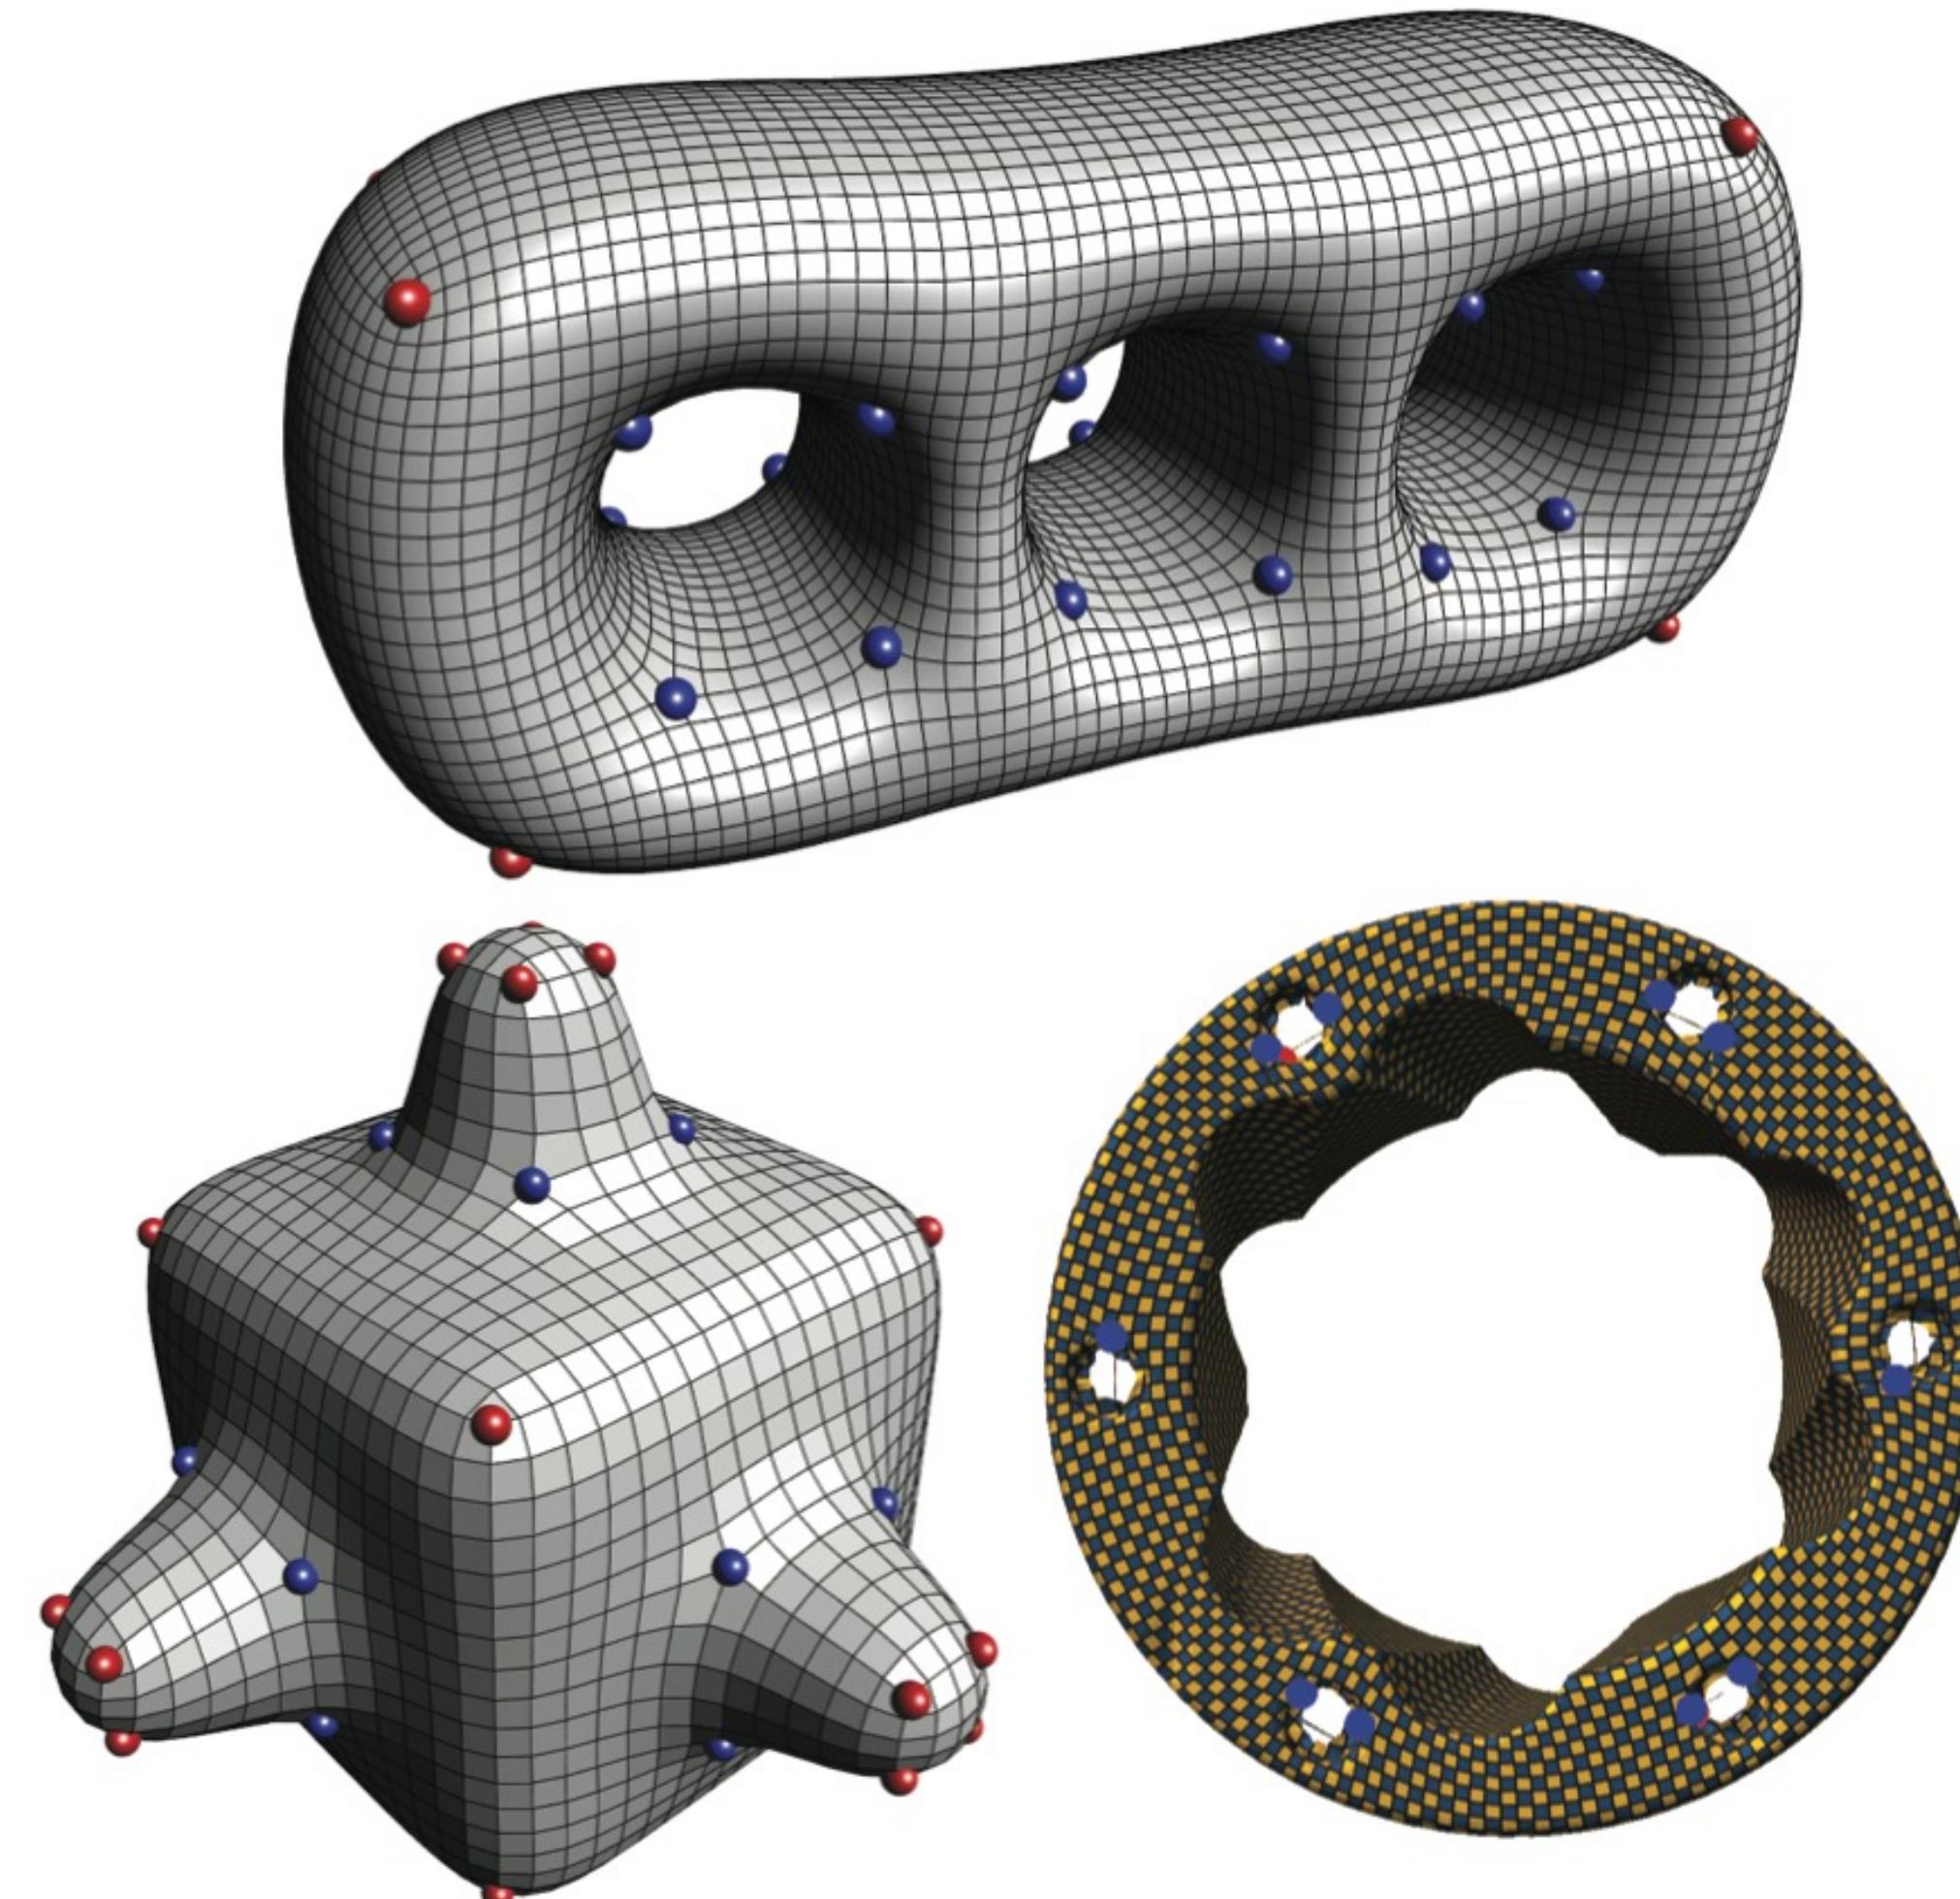
Polygonal vs. Triangle Meshes

- Edge loops are ideal for editing

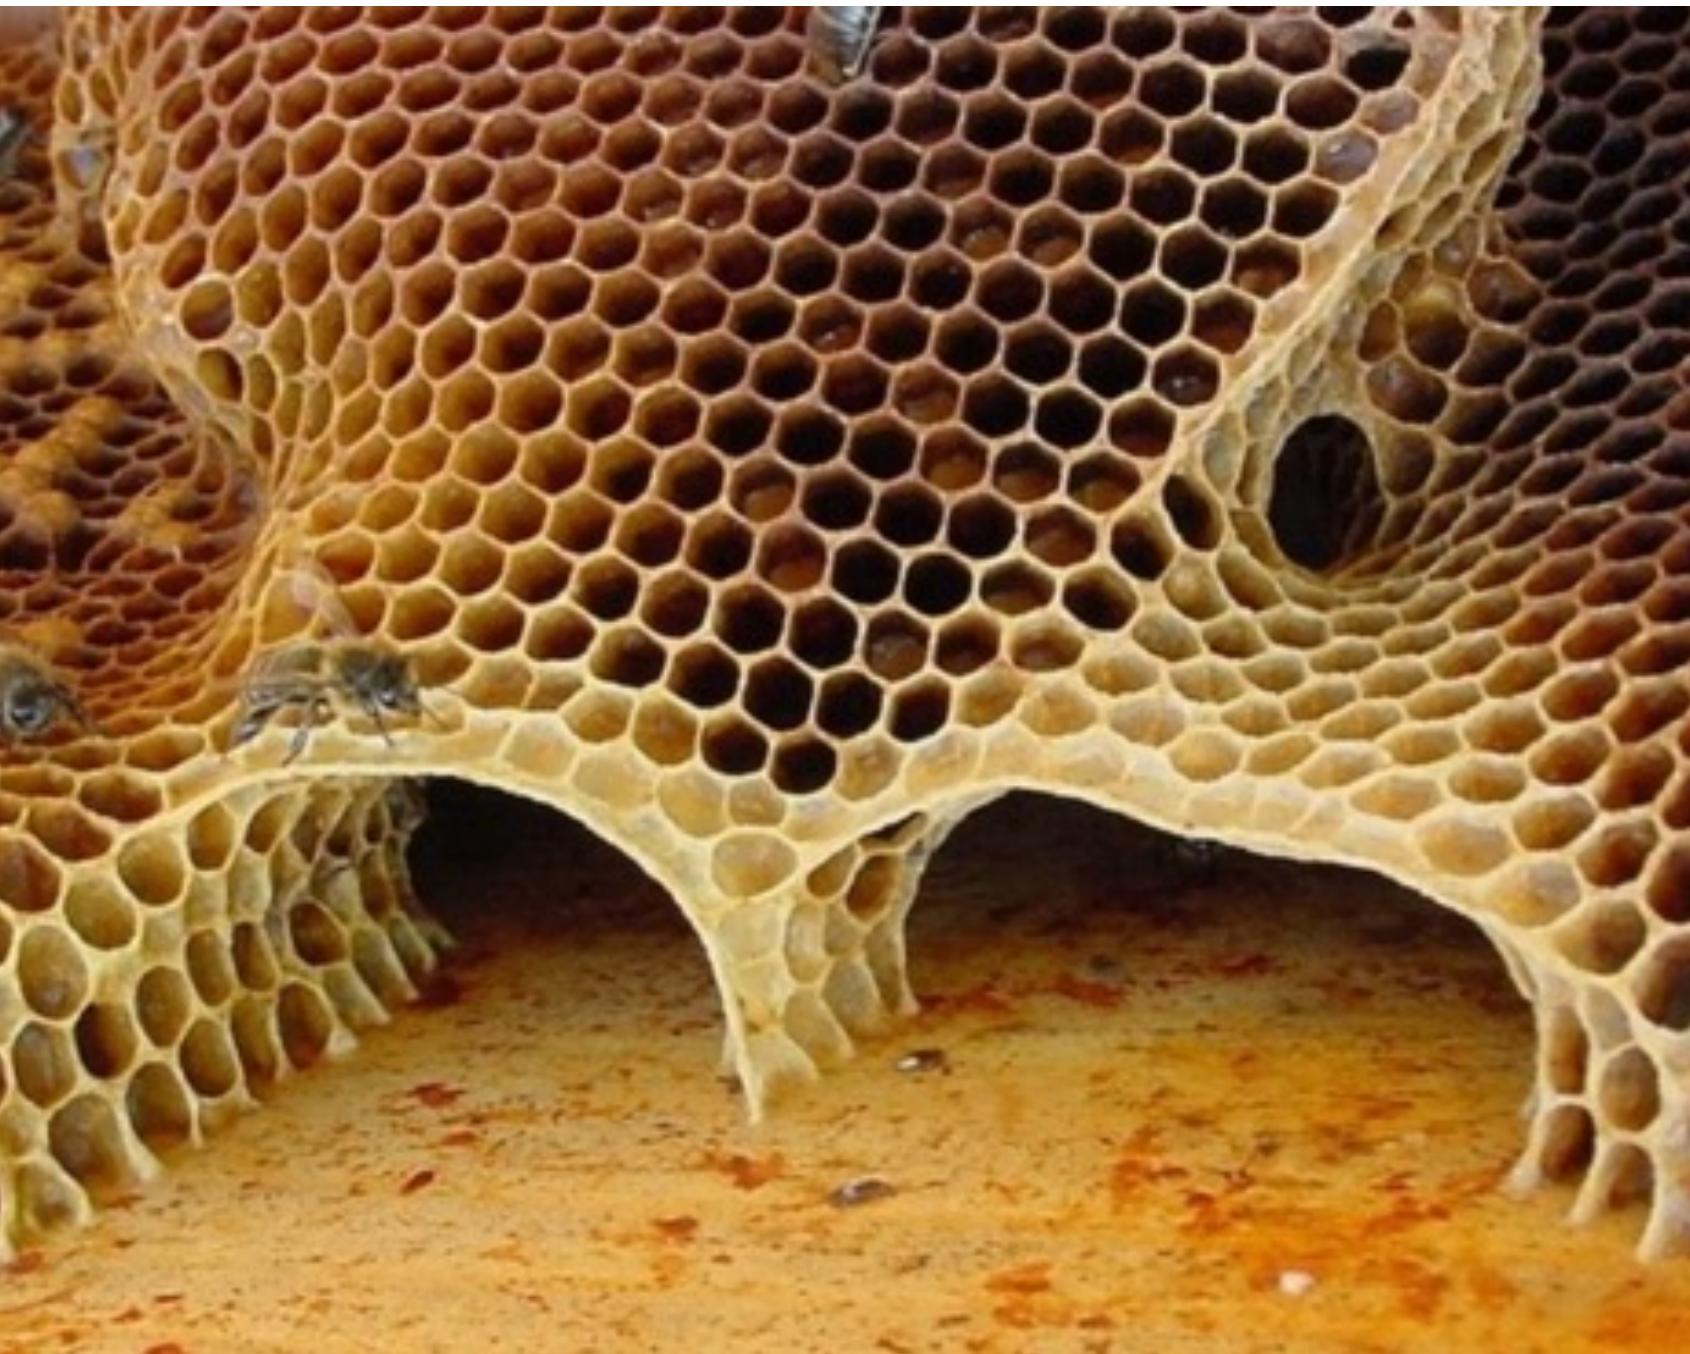


Polygonal vs. Triangle Meshes

- Quality of triangle meshes
 - Uniform Area
 - Angles close to 60
- Quality of quadrilateral meshes
 - Number of irregular vertices
 - Angles close to 90
 - Good edge flow



Polygonal vs. Triangle Meshes



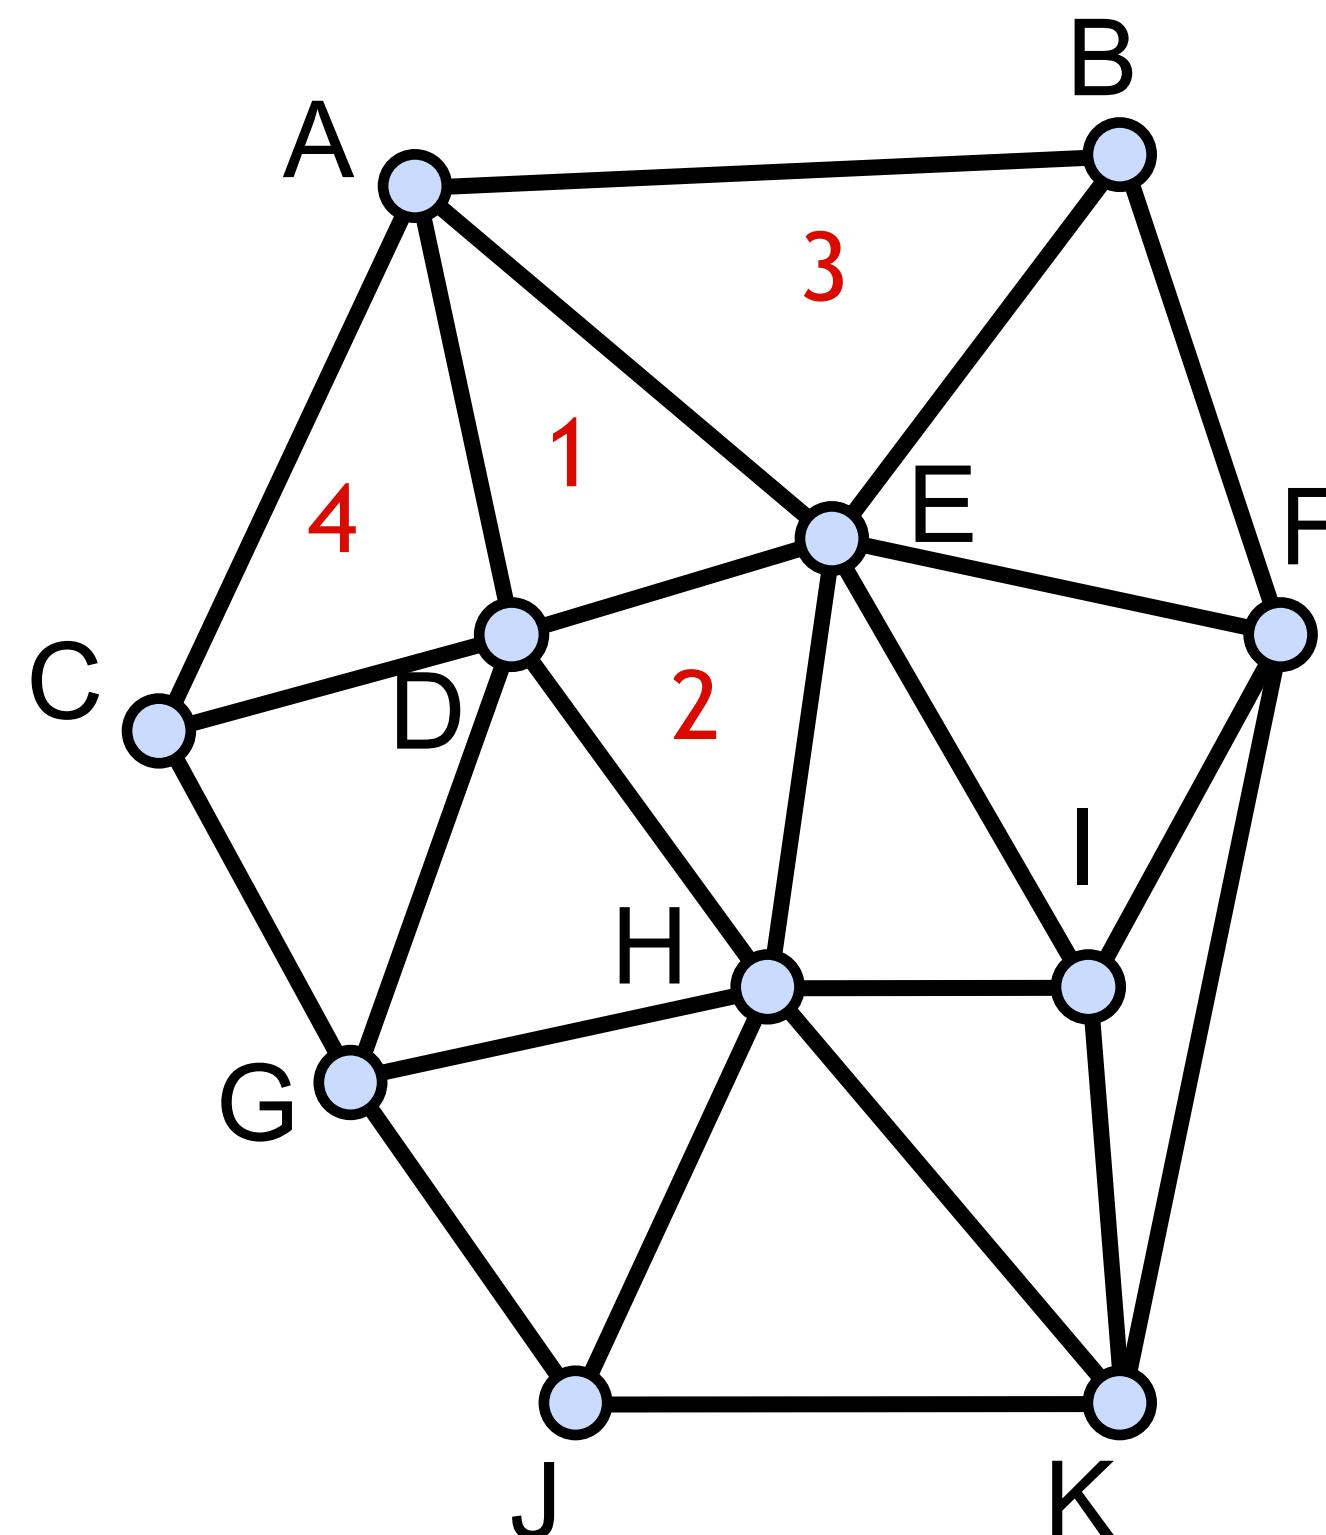
E. Van Egeraat

Data Structures

- What should be stored?
 - Geometry: 3D coordinates
 - Connectivity
 - Adjacency relationships
 - Attributes
 - Normal, color, texture coordinates
 - Per vertex, face, edge

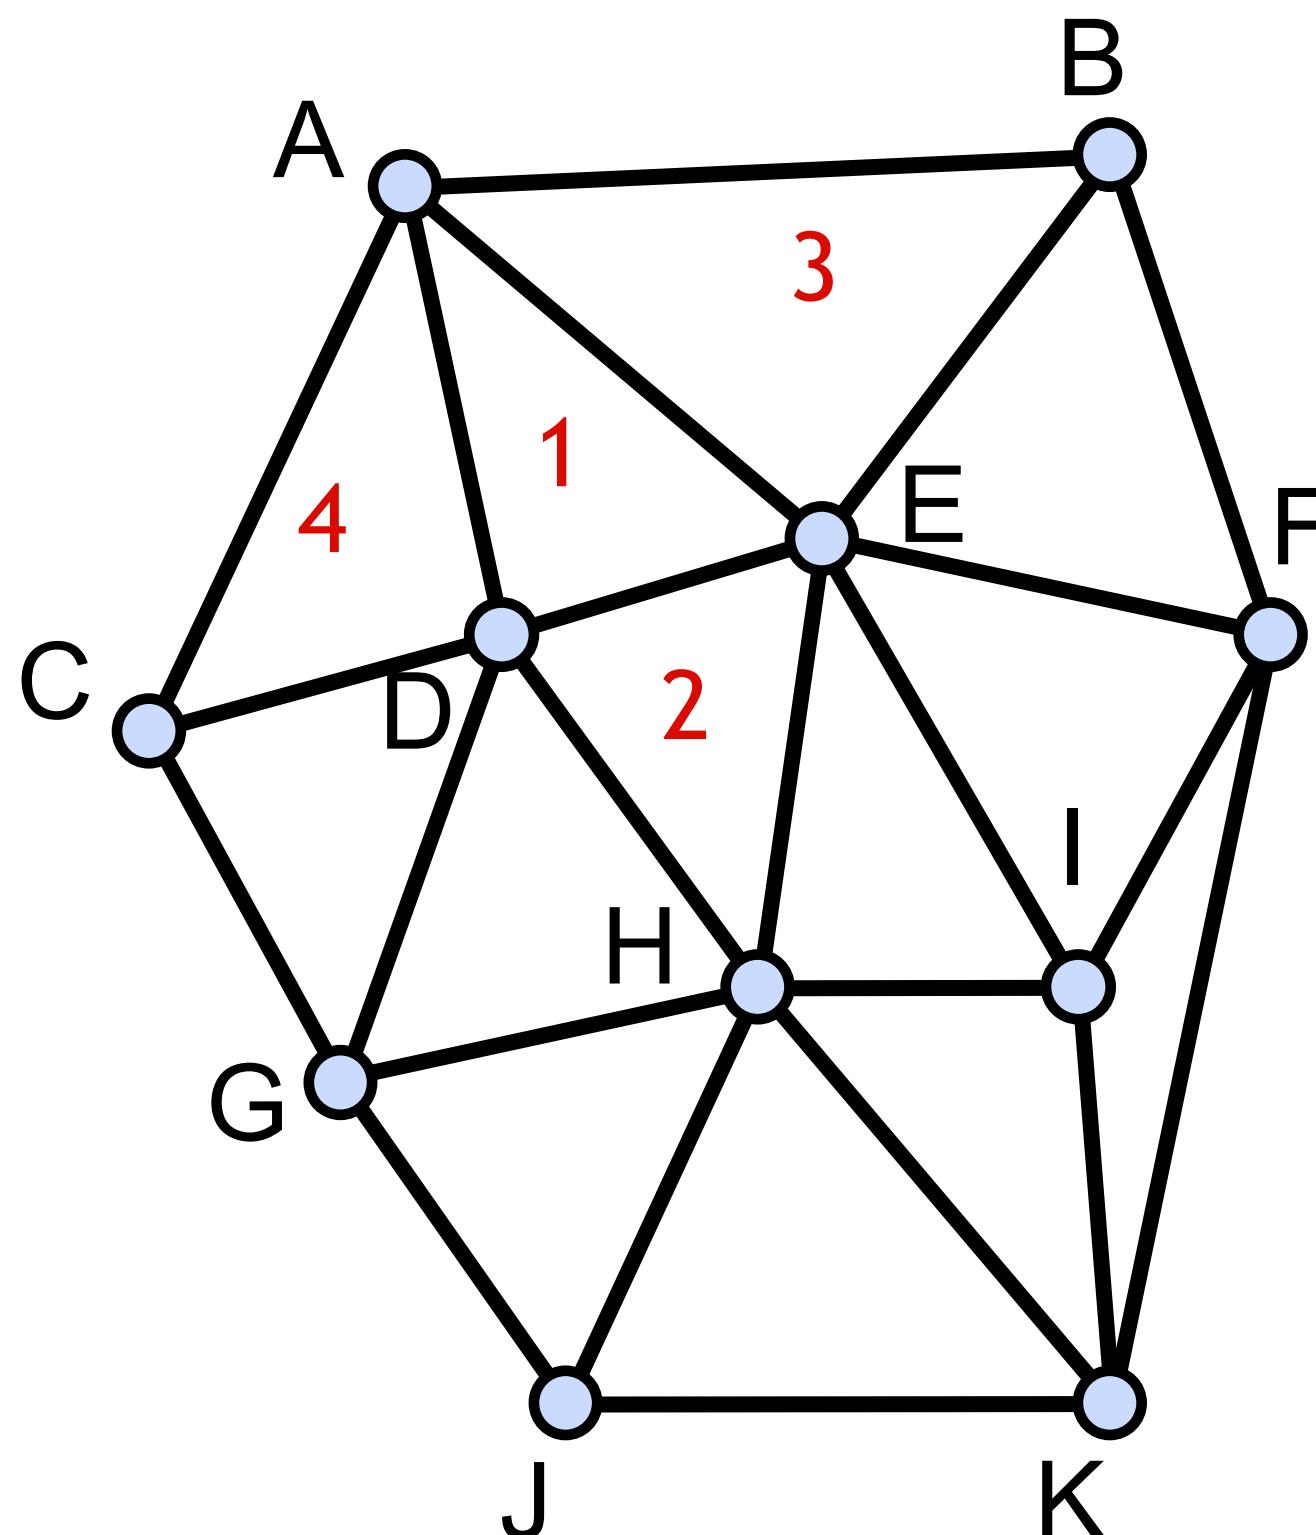


Data Structures



- What should be supported?
 - Rendering
 - Geometry queries
 - What are the vertices of face #2?
 - Is vertex A adjacent to vertex H ?
 - Which faces are adjacent to face #1?
 - Modifications
 - Remove/add a vertex/face
 - Vertex split, edge collapse

Data Structures



- How good is a data structure?
 - Time to construct
 - Time to answer a query
 - Time to perform an operation
 - Space complexity
 - Redundancy
- Criteria for design
 - Expected number of vertices
 - Available memory
 - Required operations
 - Distribution of operations

Triangle List

- STL format (used in CAD)
- Storage
 - Face: 3 positions
 - 4 bytes per coordinate
 - 36 bytes per face
 - Euler: $f = 2v$
 - $72*v$ bytes for a mesh with v vertices
- No connectivity information

Triangles			
0	x0	y0	z0
1	x1	x1	z1
2	x2	y2	z2
3	x3	y3	z3
4	x4	y4	z4
5	x5	y5	z5
6	x6	y6	z6
...

Indexed Face Set

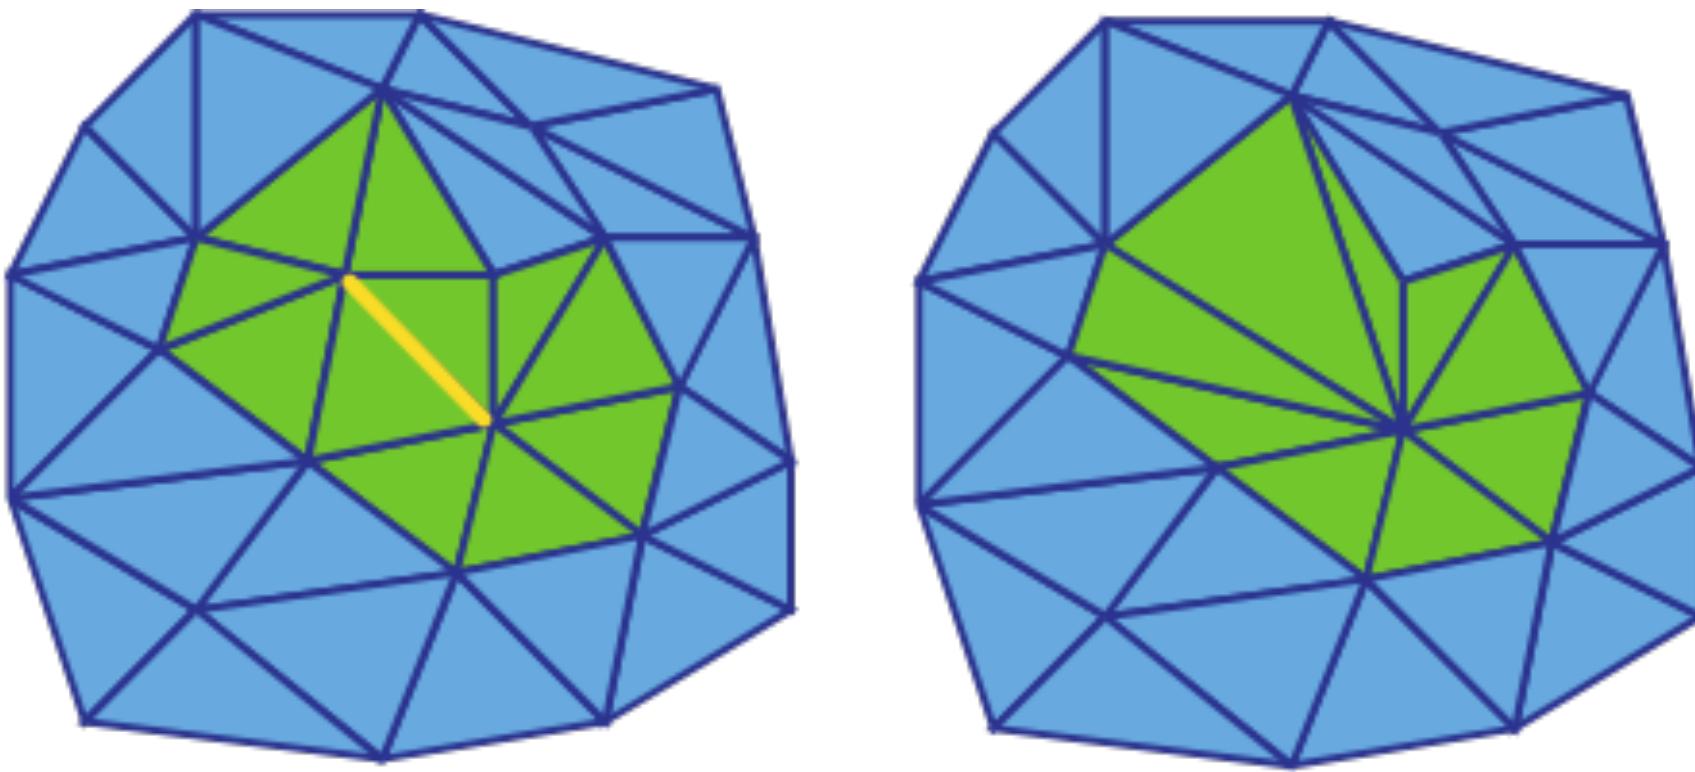
- Used in formats
OBJ, OFF, WRL
- Storage
 - Vertex: position
 - Face: vertex indices
 - 12 bytes per vertex
 - 12 bytes per face
 - $36*\nu$ bytes for the mesh
- No *explicit* neighborhood info

Vertices			
v0	x0	y0	z0
v1	x1	x1	z1
v2	x2	y2	z2
v3	x3	y3	z3
v4	x4	y4	z4
v5	x5	y5	z5
v6	x6	y6	z6
...

Triangles			
t0	v0	v1	v2
t1	v0	v1	v3
t2	v2	v4	v3
t3	v5	v2	v6
...

Indexed Face Set: Problems

- Information about neighbors is not explicit
 - Finding neighboring vertices/edges/faces costs $O(\#V)$ time!
 - Local mesh modifications cost $O(V)$



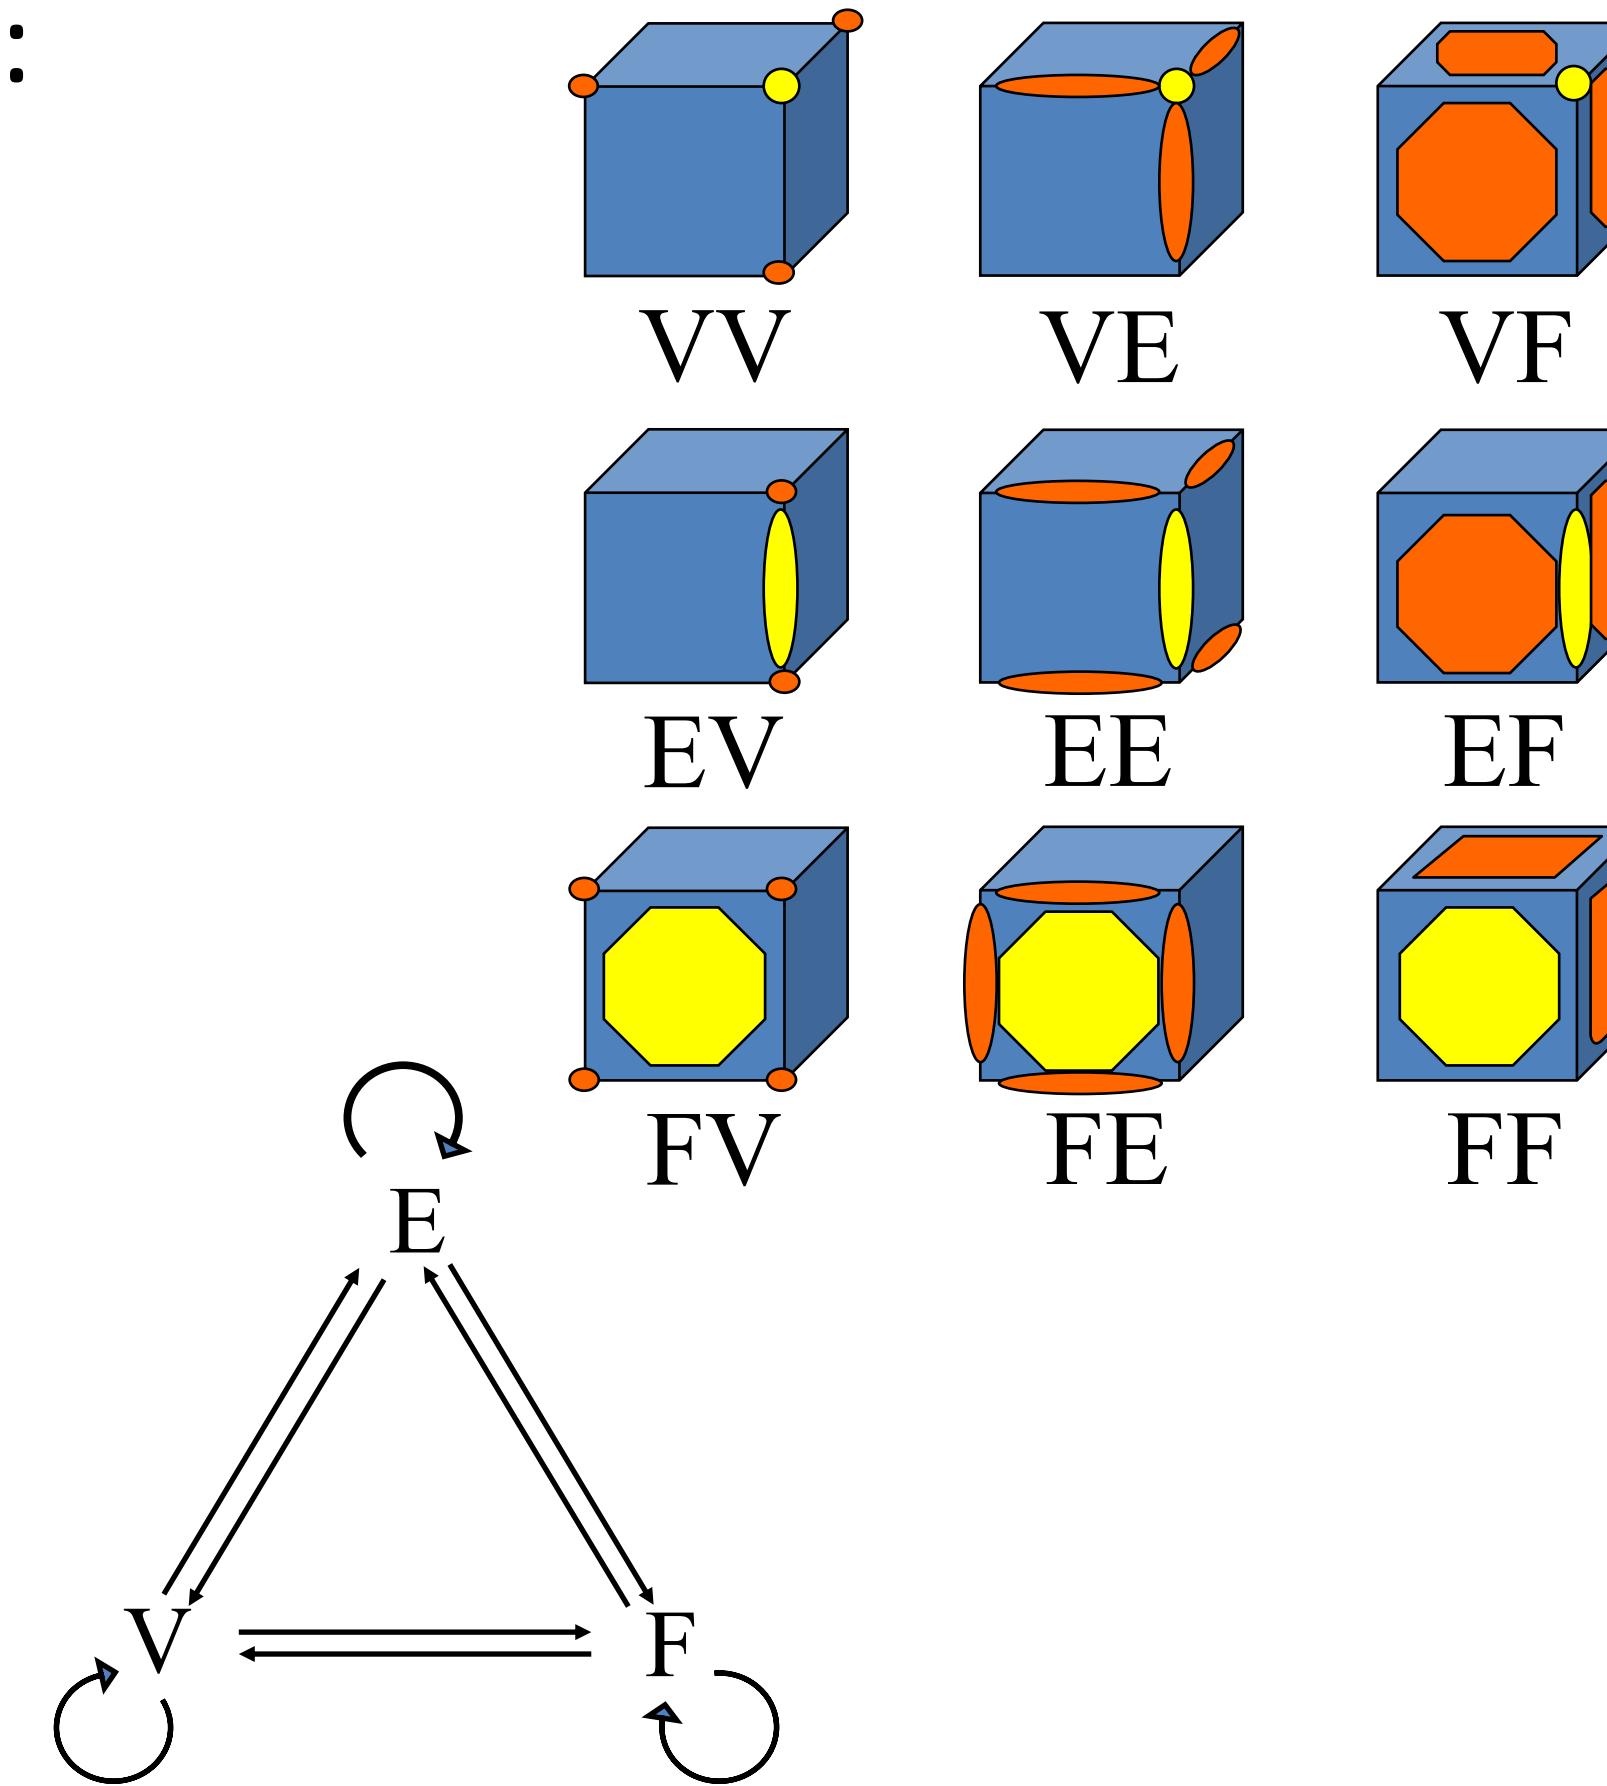
- Breadth-first search costs $O(k * \#V)$ where $k = \#$ found vertices

Neighborhood Relations

- All possible neighborhood relationships:

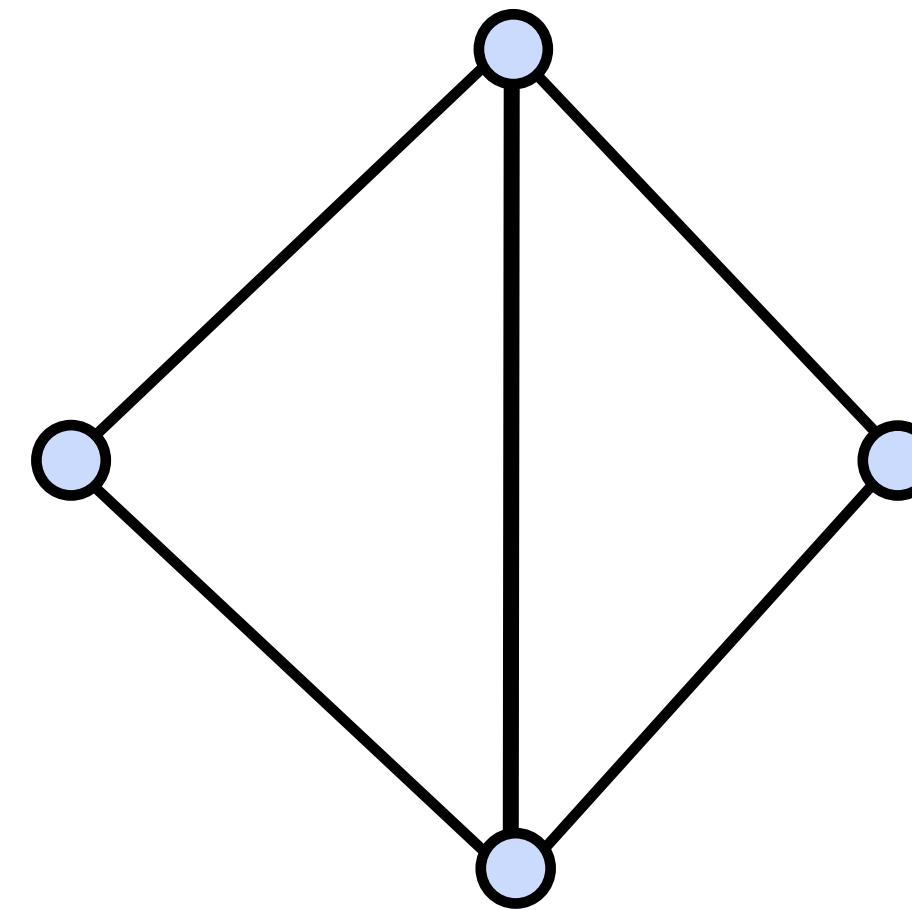
1. Vertex	– Vertex	VV
2. Vertex	– Edge	VE
3. Vertex	– Face	VF
4. Edge	– Vertex	EV
5. Edge	– Edge	EE
6. Edge	– Face	EF
7. Face	– Vertex	FV
8. Face	– Edge	FE
9. Face	– Face	FF

We'd like $O(1)$ time for queries and local updates of these relationships



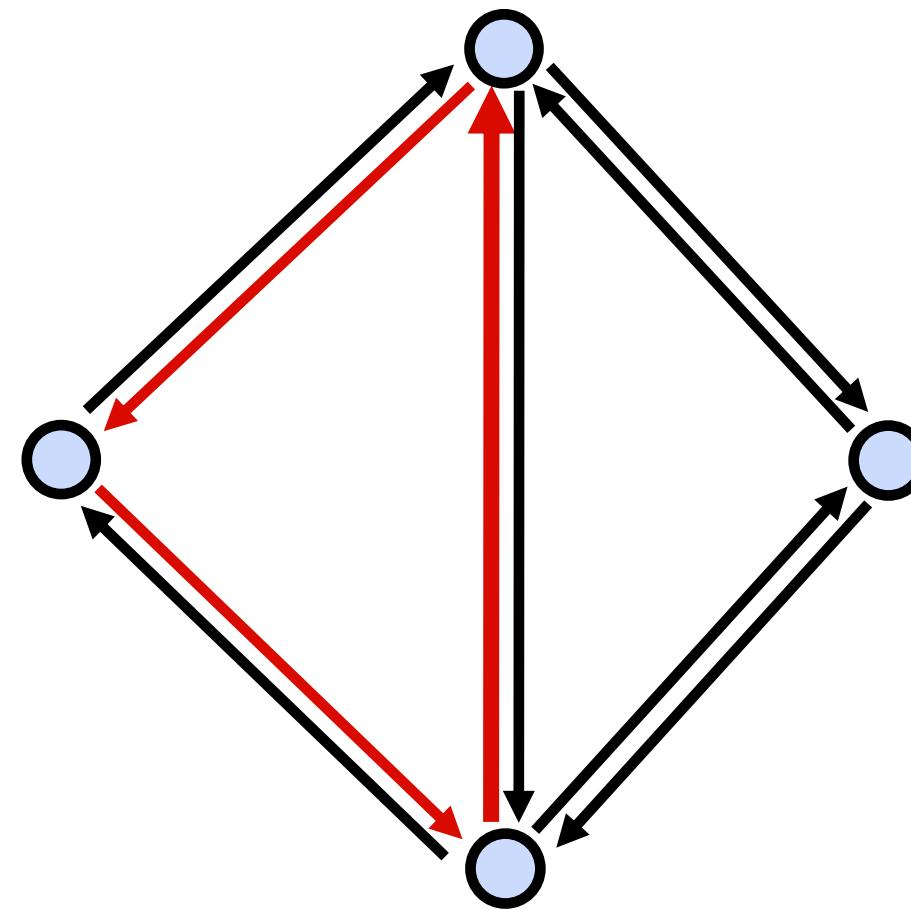
Halfedge data structure

- Introduce orientation into data structure
 - Oriented edges



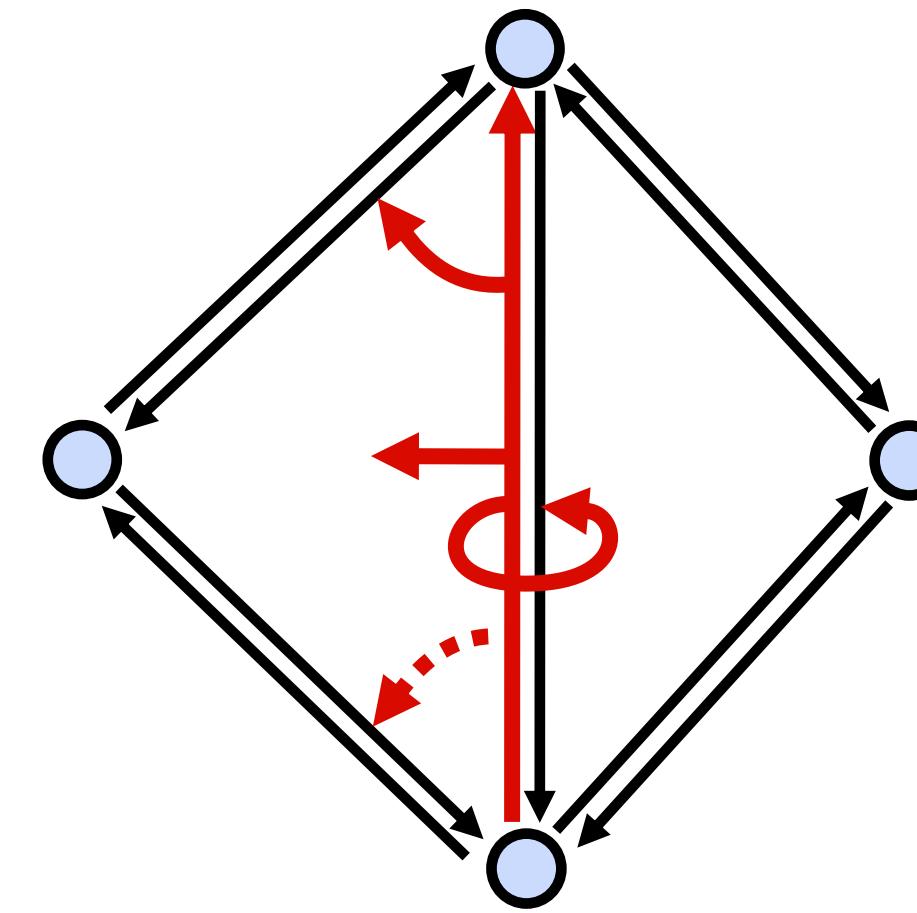
Halfedge data structure

- Introduce orientation into data structure
 - Oriented edges



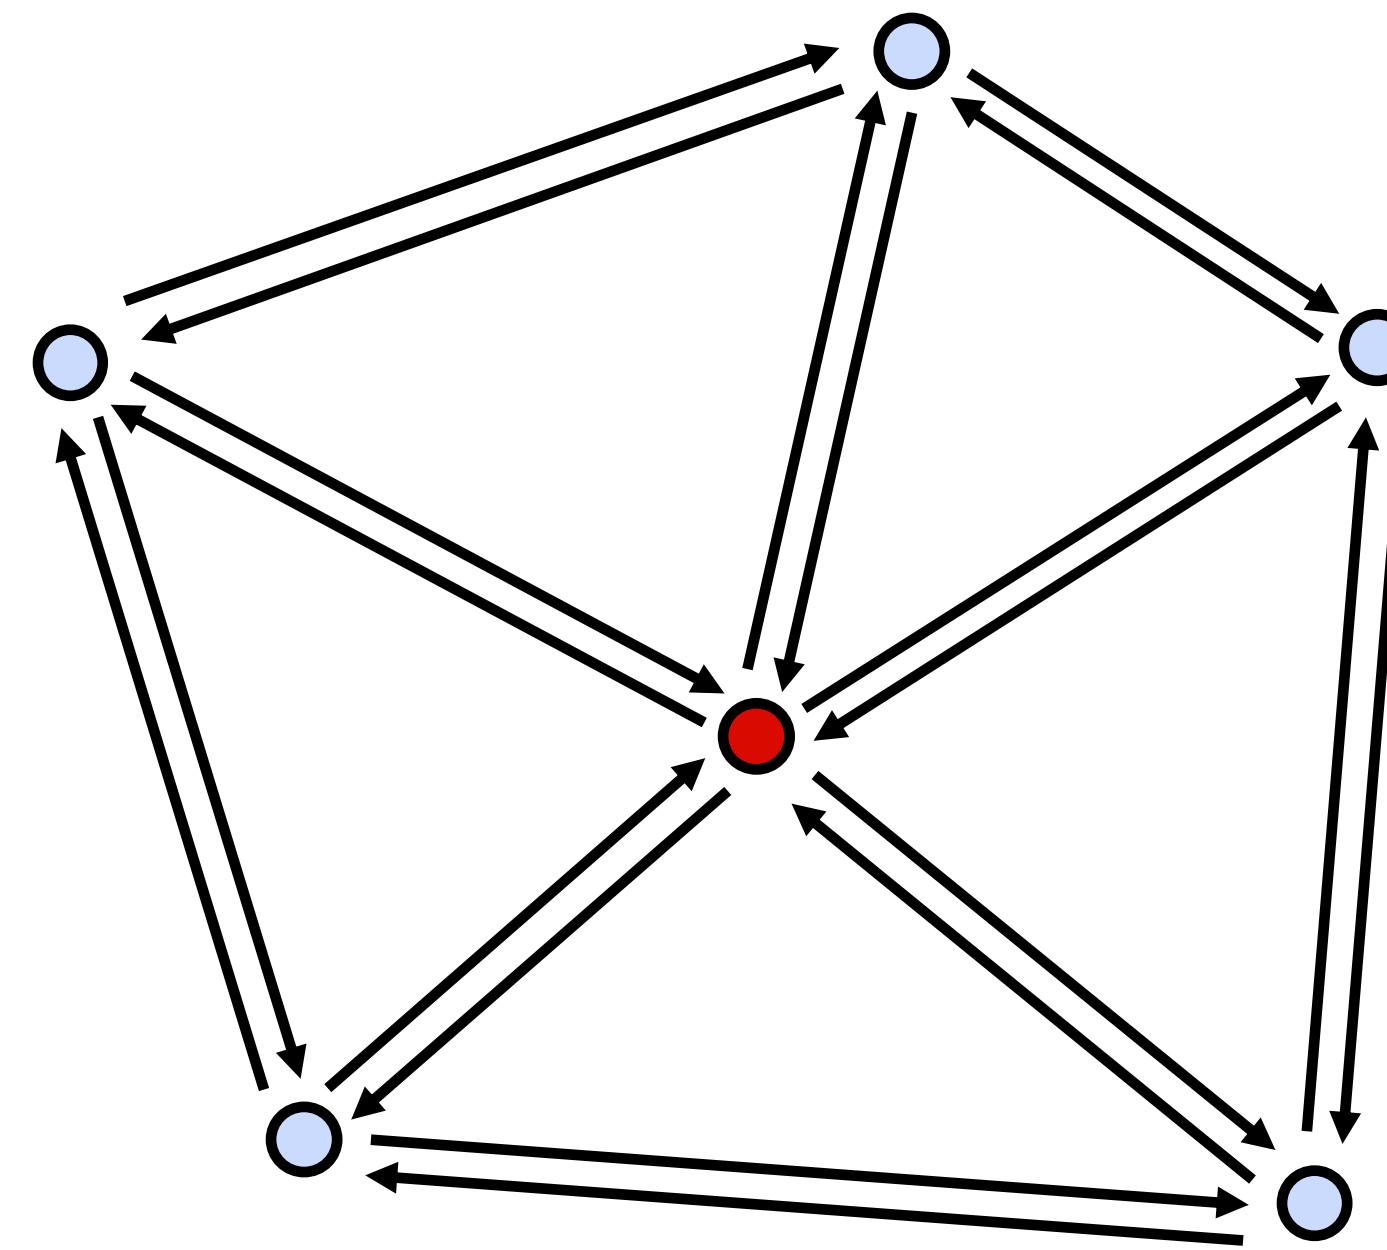
Halfedge data structure

- Introduce orientation into data structure
 - Oriented edges
- Vertex
 - Position
 - 1 outgoing halfedge index
- Halfedge
 - 1 origin vertex index
 - 1 incident face index
 - 3 next, prev, twin halfedge indices
- Face
 - 1 adjacent halfedge index
- Easy traversal, full connectivity



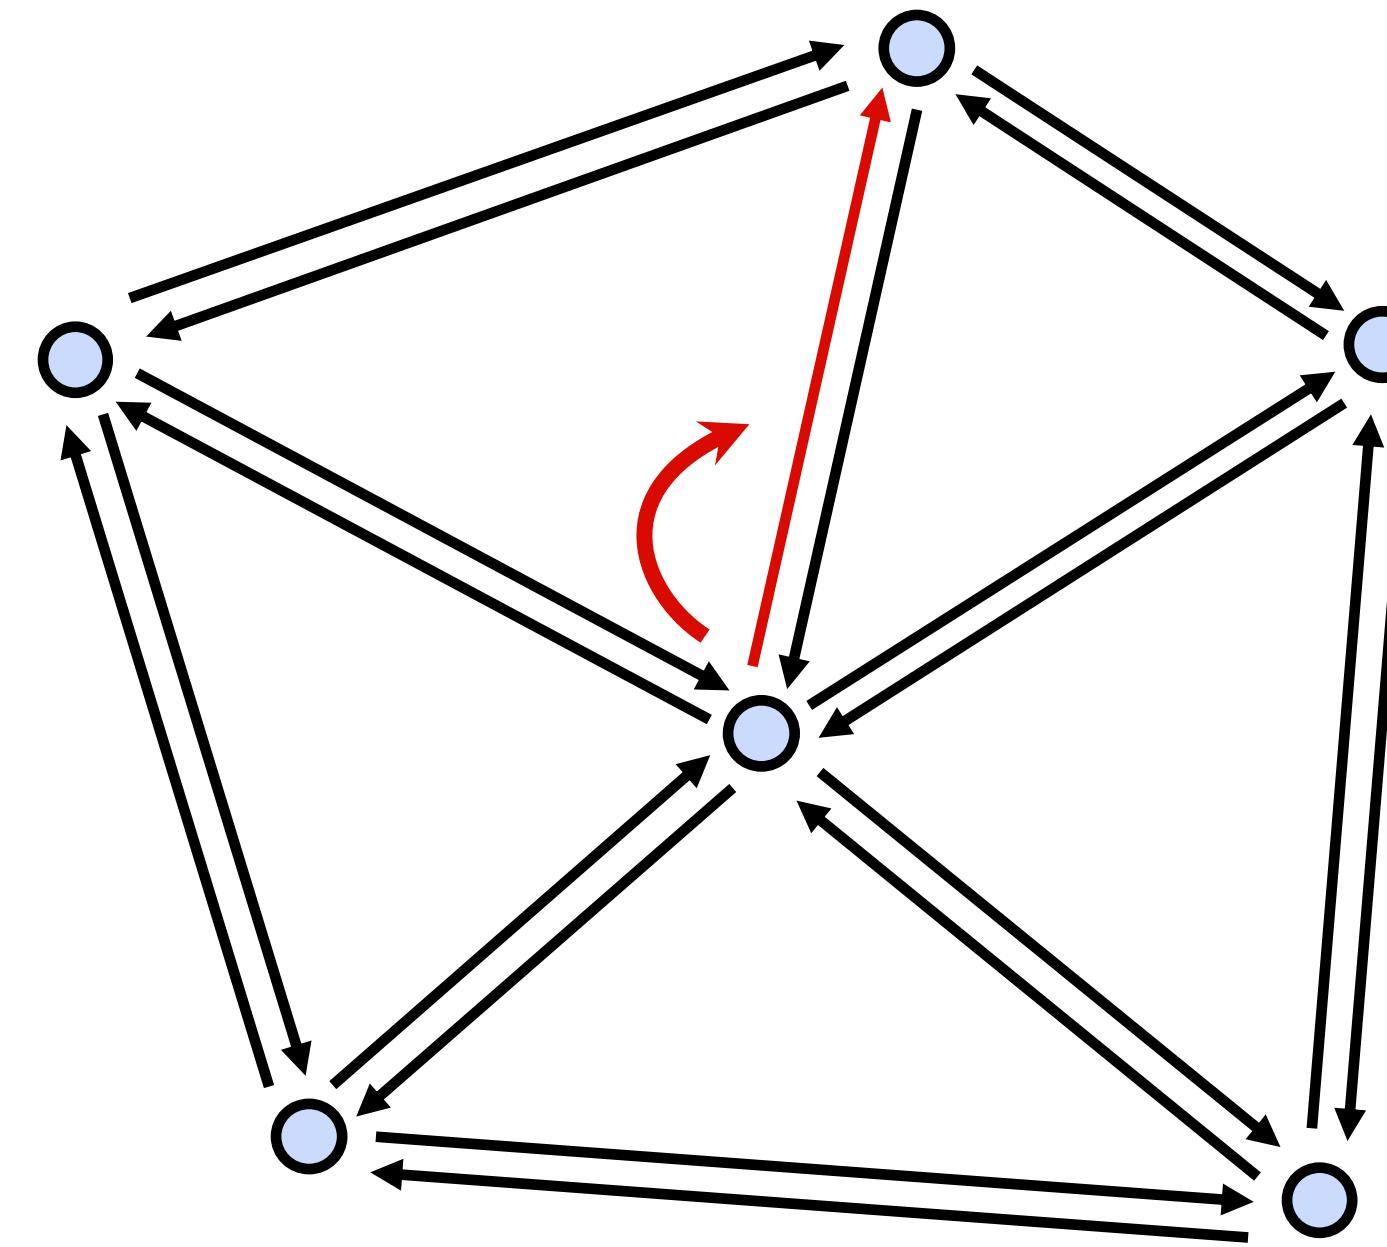
Halfedge data structure

- One-ring traversal
 - Start at vertex



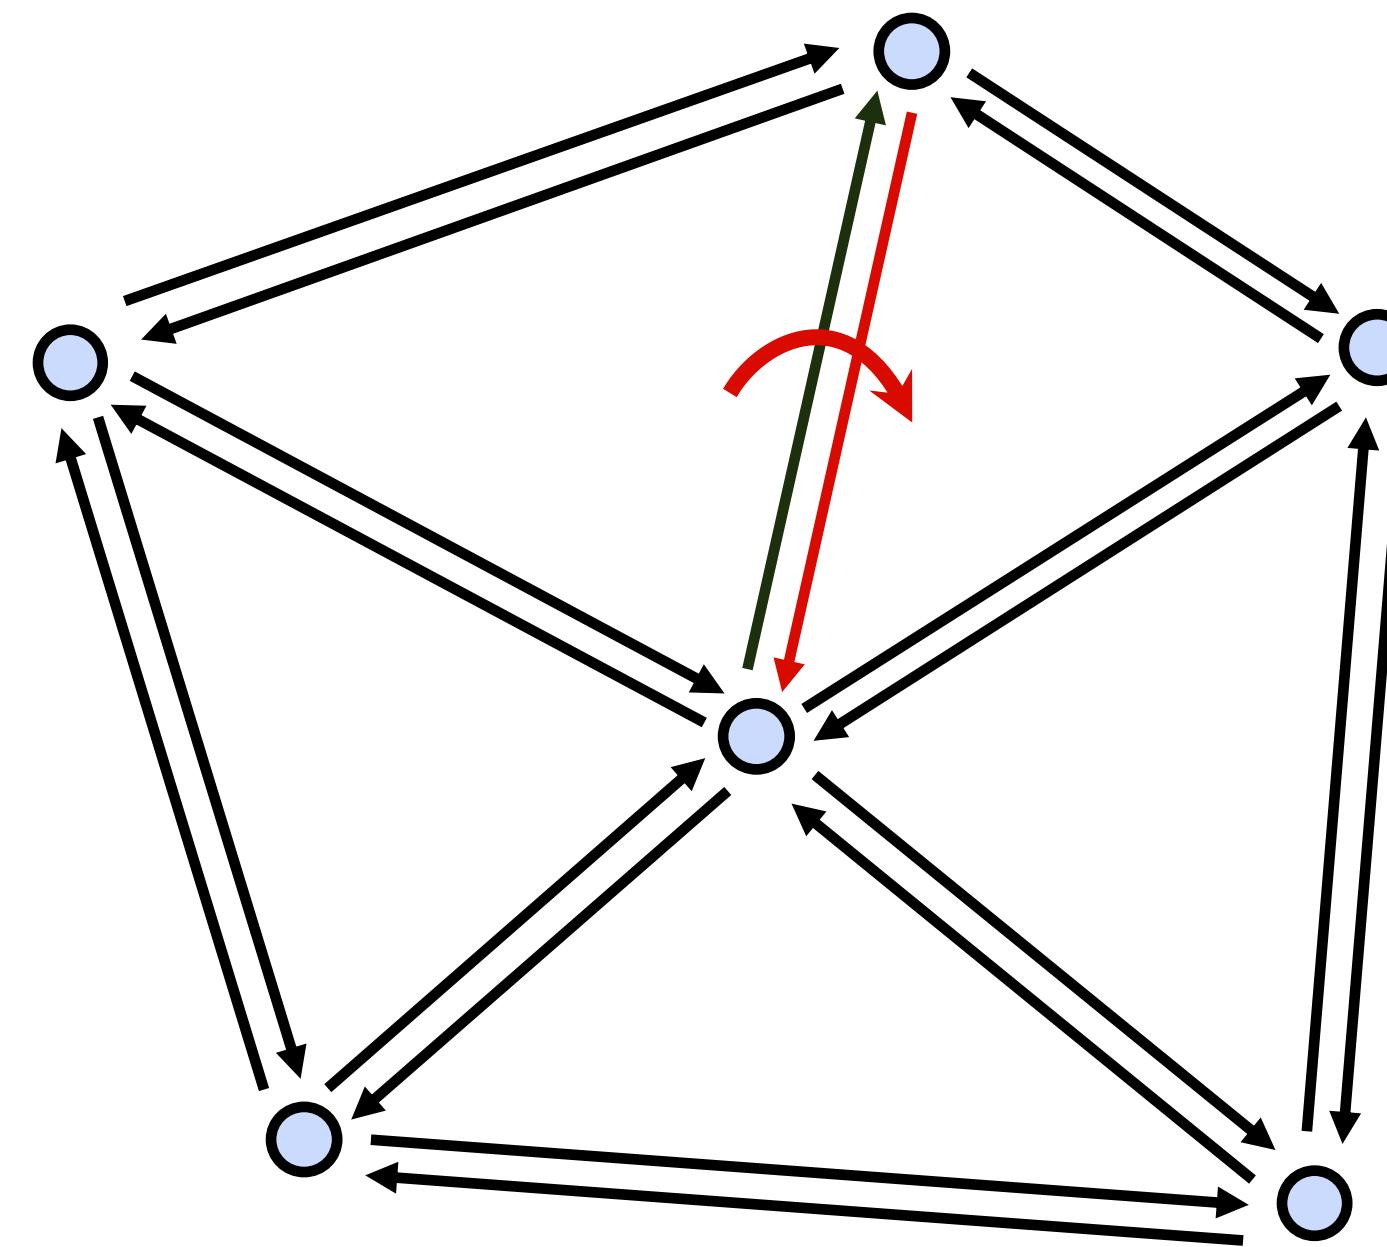
Halfedge data structure

- One-ring traversal
 - Start at vertex
 - Outgoing halfedge



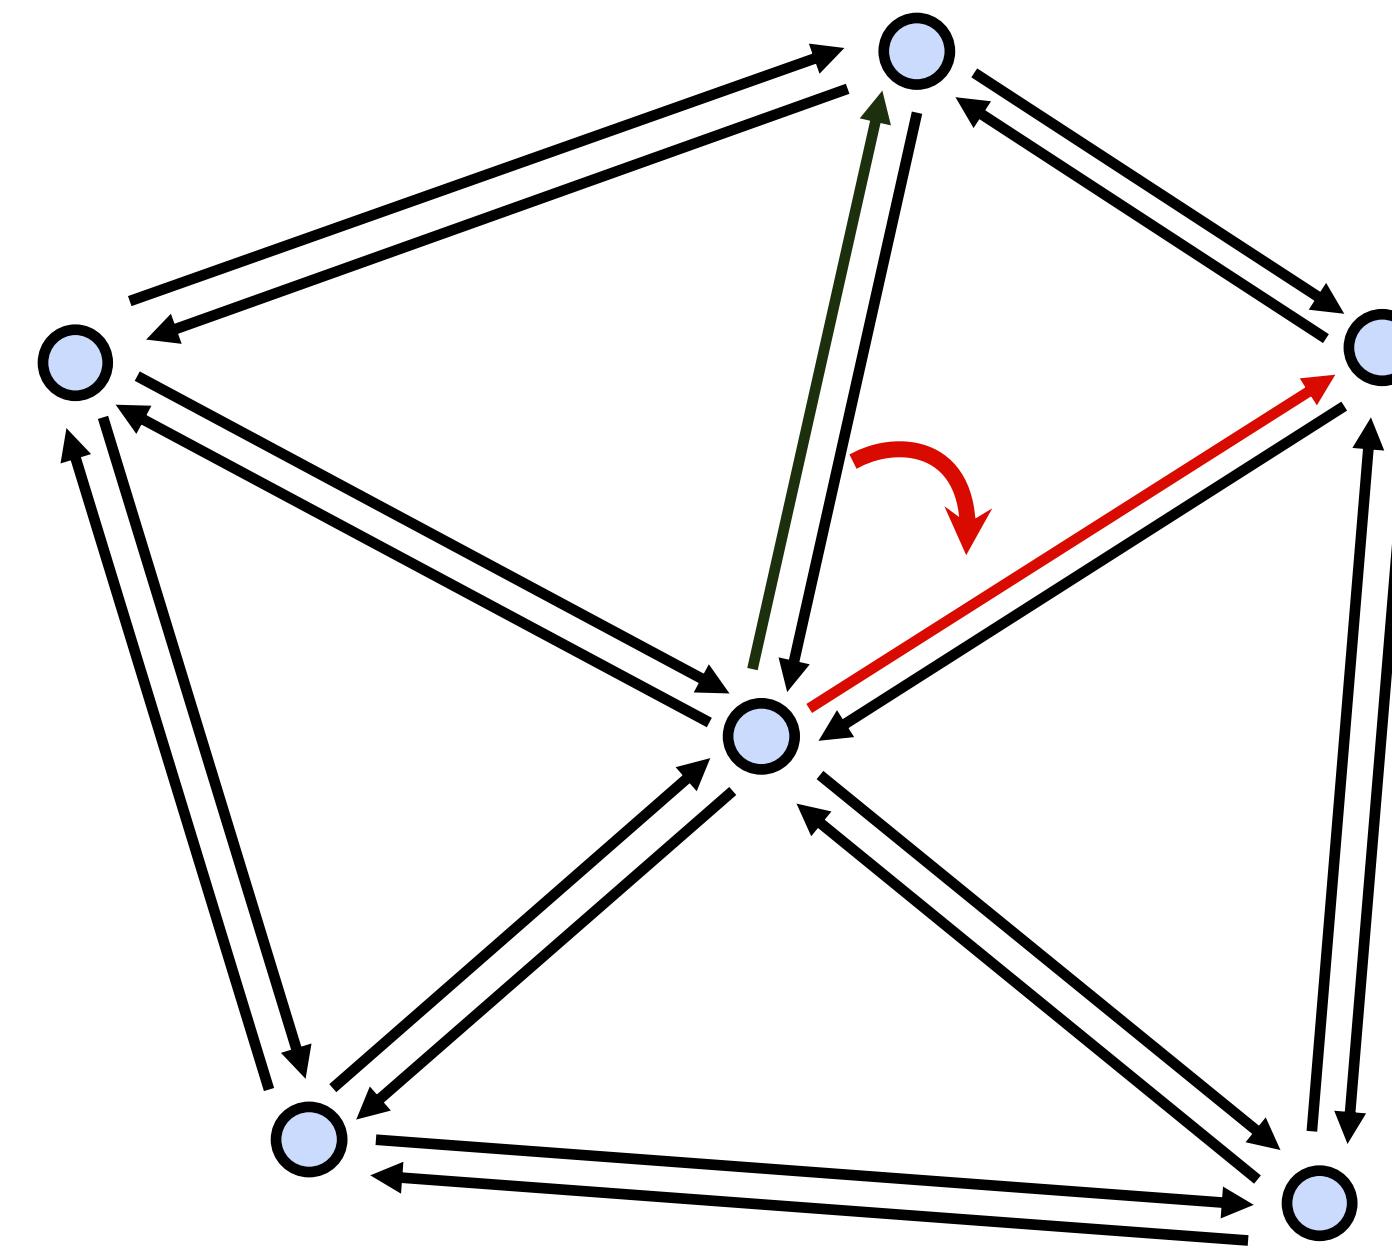
Halfedge data structure

- One-ring traversal
 - Start at vertex
 - Outgoing halfedge
 - Twin halfedge



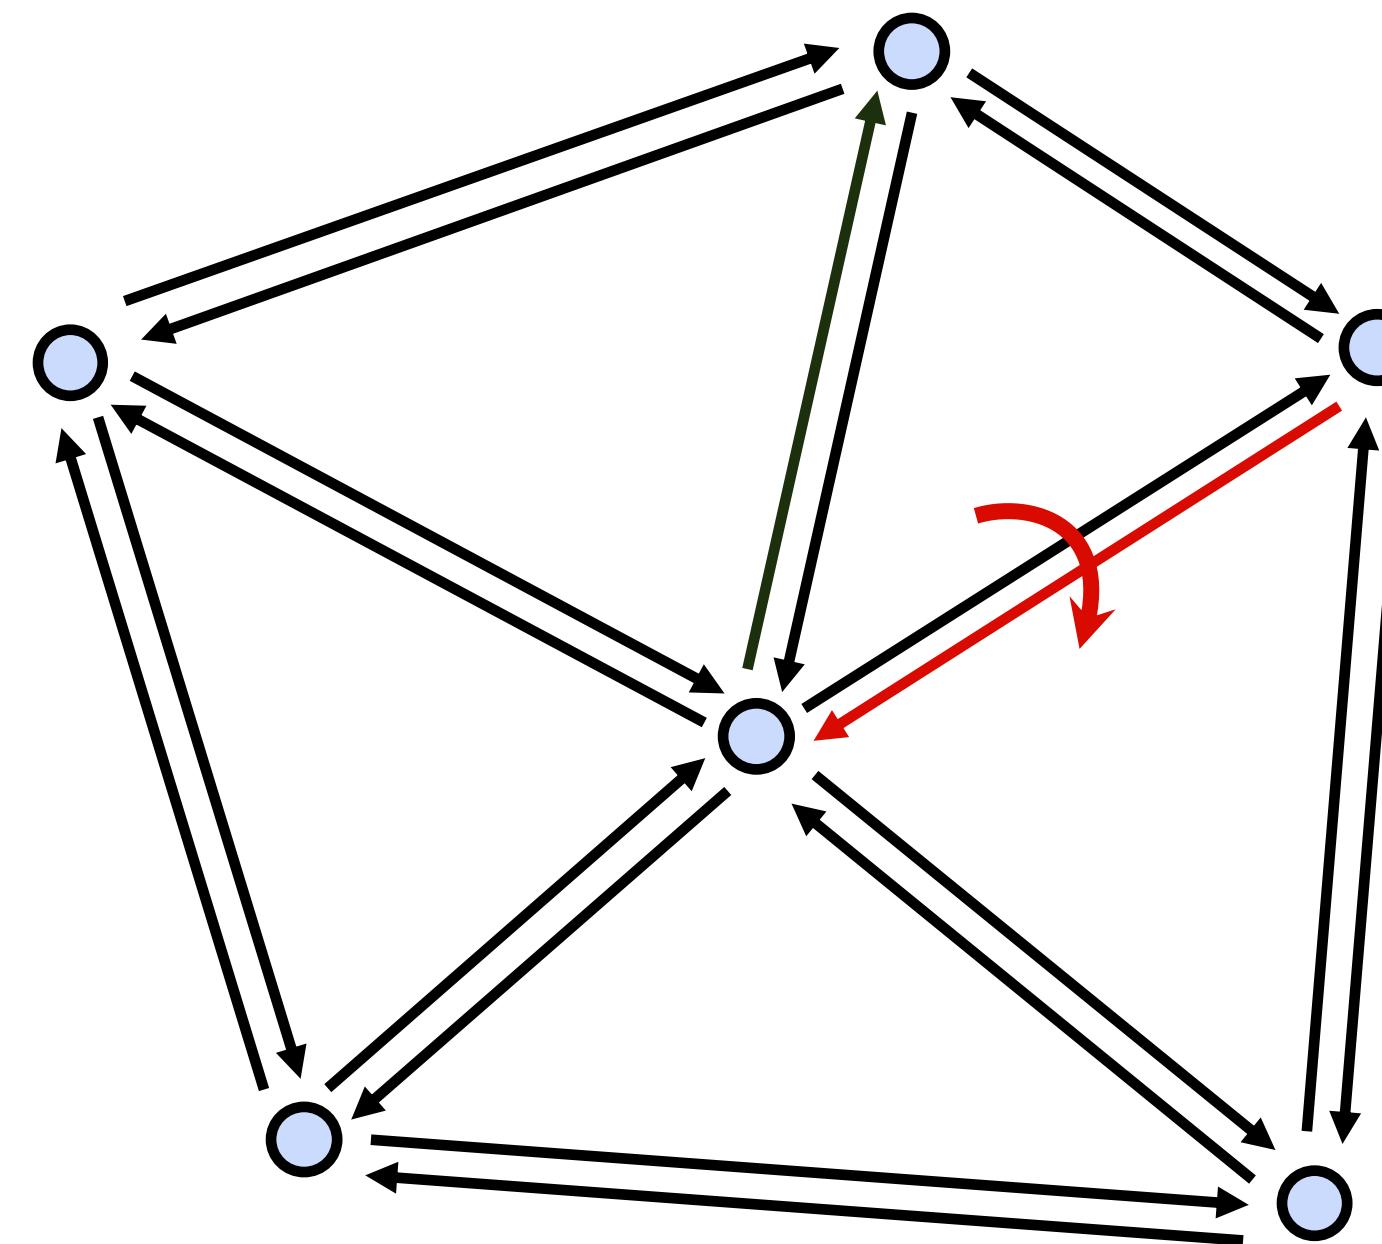
Halfedge data structure

- One-ring traversal
 - Start at vertex
 - Outgoing halfedge
 - Twin halfedge
 - Next halfedge



Halfedge data structure

- One-ring traversal
 - Start at vertex
 - Outgoing halfedge
 - Twin halfedge
 - Next halfedge
 - Twin ...



Halfedge data structure

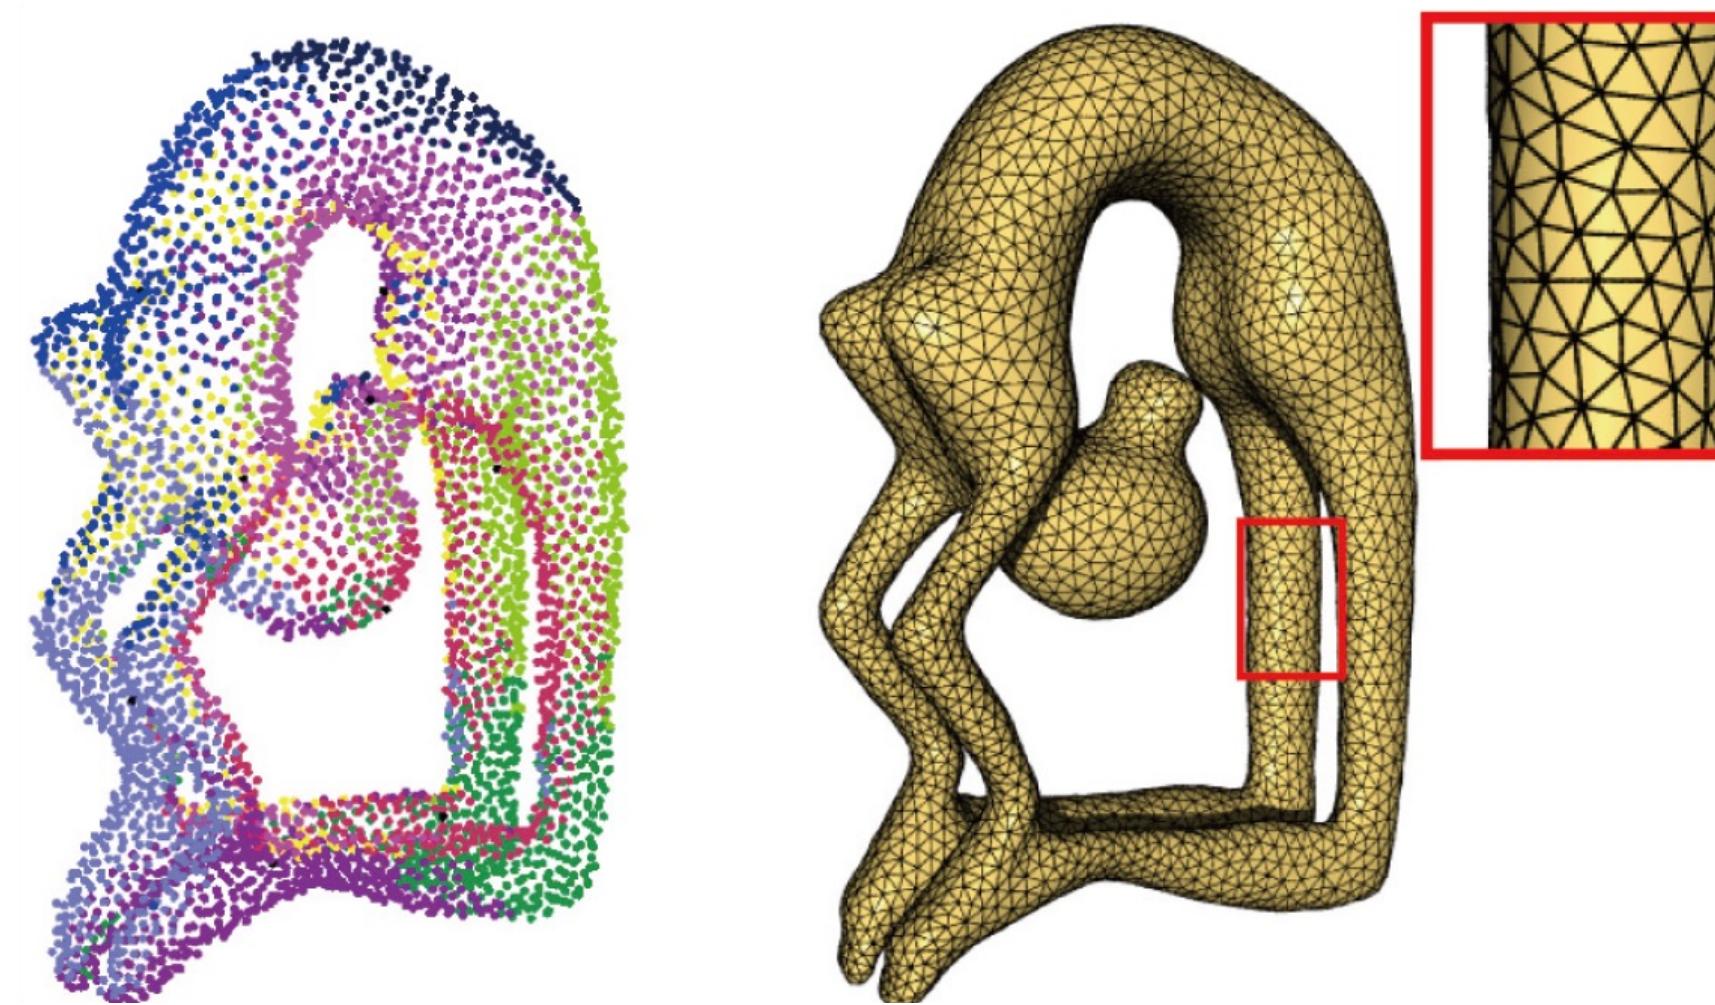
- Pros: *(assuming bounded vertex valence)*
 - $O(1)$ time for neighborhood relationship queries
 - $O(1)$ time and space for local modifications (edge collapse, vertex insertion...)
- Cons:
 - Heavy – requires storing and managing extra pointers
 - Not as trivial as Indexed Face Set for rendering with OpenGL / Vertex Buffer Objects

Halfedge Libraries

- CGAL
 - www.cgal.org
 - Computational geometry
- OpenMesh
 - www.openmesh.org
 - Mesh processing
- We will not implement a half-edge data structure in the class. Instead we will work with Indexed Face Set and augment it to have fast queries.

Next Lecture

- How to get a nice, watertight surface mesh from a sampled point set



- The most popular way:
points → implicit function → surface mesh

Thank you